

CS 475

Operating Systems



Department of Mathematics
and Computer Science

Lecture 1
History and Trends in
Computer Systems

Today's Topics

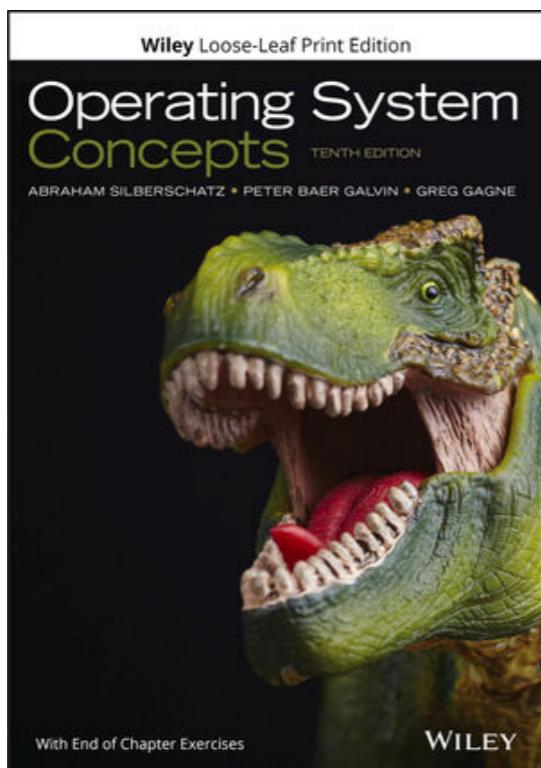
- ▶ Course Overview
- ▶ Technological Trends
- ▶ What Is an OS, Anyway?
- ▶ A Brief History of Computer and Operating Systems
- ▶ Conclusion

About David

- ▶ David Chiu (pronounced "chew")
 - "Professor Chiu", "Dr. Chiu", or just "David"
- ▶ Background
 - PhD, Computer Science & Engineering, Ohio State, 2005-2010.
 - (Majored in systems, high performance computing, and databases)
 - Assistant Professor at Washington State University, 2010-2014.
 - Professor at Puget Sound since 2014.

Course Requirements

- ▶ Required textbooks (both required)
 - **Dive into Systems (free e-book)**
 - <https://diveintosystems.org/book/>
 - **Silberschatz, Galvin, and Gagne. Operating System Concepts. 8th ed. or higher.**



How to Reach Me

- ▶ How to reach me
 - Email: dchiu@pugetsound.edu
 - Office: TH 303 (near south elevator)

- ▶ Drop in when door is open/cracked
 - Or, schedule an appointment here:



Course Webpages

- ▶ Two important webpages to bookmark.
 - Course Calendar - <https://davidtchiu.github.io/teaching/cs475/>
 - Canvas - <https://canvas.pugetsound.edu>
 - Assignment submission, code examples, study guides found here
 - [Let's take a quick tour]
- ▶ Scan below to take you to the course calendar



Course Policies

▶ Class Disruption

- Put phone on silent
- *No laptops out, unless permitted for note-taking by SAA or permitted on lab days*



▶ Cheating: Zero-Tolerance

- I compare current and past assignments
- OK to brainstorm, and ask for help, but
you must produce your own code



ChatGPT and Other Generative AI

- ▶ **Do** use it to be your personalized tutor, and that is as far as it should go.
Use it to explain C code you don't fully grasp.
 - ▶ Prompt: “**Without giving me any code, explain ...**”
- ▶ **Do** use it to explain code to you. Is there a piece of code we went over in class that's hard to grasp? Paste it, and have it explain line-by-line as well as holistically.
- ▶ **Do** use it to explain errors to you. Paste your code and the errors you get when compiling or running it. Give hints on what might be the issue.
- ▶ **Do** use it to explain any OS concepts that you need to know.

ChatGPT and Other Generative AI (2)

- ▶ **Don't turn in** anything that was generated by these tools. Copying-and-pasting AI output is considered plagiarism and will be treated as such.
- ▶ **Don't underestimate** how easy it is for us to detect cases where students are turning in code written by generative AI tools.
- ▶ **Don't forget** that you *still* need to demonstrate proficiency on all your exams to pass the course.

Evaluation

► Breakdown of Evaluation

- 40% - Homework Assignments
- 15% - Midterm Exam I
- 18% - Midterm Exam II
- 22% - Final Exam (Comprehensive)
- 5% - Participation
 - Starts at 10. Unexcused absence (-2), late (-1), non-class related activities (-1)



Homework

- ▶ You can expect 7-8 homework assignments
 - Programs must compile and run on the remote server
 - $-3^d\%$ point late penalty, including weekends
 - All assignments must be submitted using **git/github**
- ▶ **Budget more time than usual for coding assignments**
 - Assignments *seem* easy at first
 - But C is very finicky and low-level
 - The **gcc** compiler doesn't do a great job verbalizing runtime errors
 - Carefully doing and learning from the early "labs" is crucial



Comprehensive Exams

- ▶ Two **1-hr midterms** and a **2-hr final** will be based on
 - Lectures, notes, homework
 - Comprehensive, but weighted heavily on new material
- ▶ I don't test anything I did not specifically cover in class
 - But things covered in class, but isn't found in the readings, is fair game
 - So don't miss class!
- ▶ Calculator and 1-page notes (front/back) allowed



How to Succeed in This Course

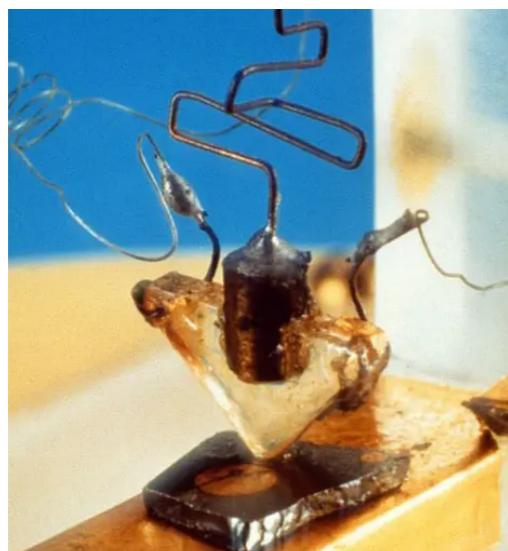
- ▶ Dedicate time to learning C
 - C is not intrinsically *hard*, but it's nuanced and debugging is frustrating!
 - There's no "stack trace" when a program crashes.
 - Trivial tasks will take you longer to write in C (vs. other languages)
 - Imperative that you *master* strings and pointers early
 - (They *will* drive you nuts)
 - Meanwhile...
 - You must also develop a solid grounding with Linux and its command line shell

Today's Topics

- ▶ Course Overview
- ▶ Technological Trends
- ▶ What Is an OS, Anyway?
- ▶ A Brief History of Computer and Operating Systems
- ▶ Conclusion

Bell Labs and the Transistor

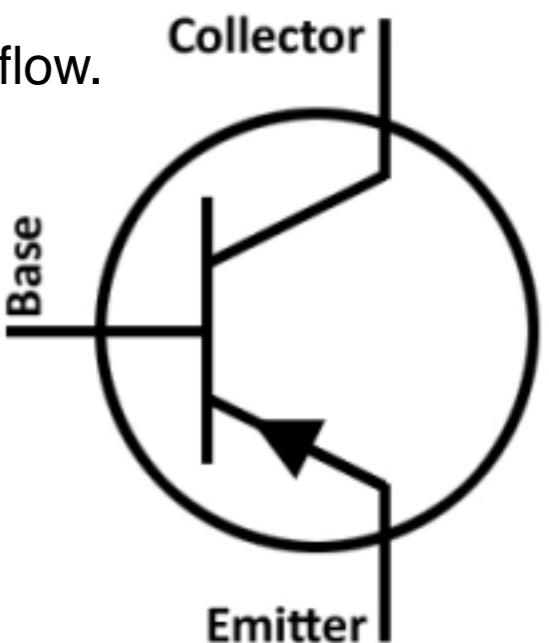
- ▶ The "*transistor*" is one the most important inventions in the 20th century.
 - Basic building blocks of logic gates, adders, flip-flops, and so on.
- ▶ Invented in Bell Labs in 1947
 - John Bardeen, Walter Brattain, and William Shockley
 - Earned the 1956 Nobel Prize in Physics



The original Bell Labs transistor

It acts like a faucet (or *switch!*) for electric flow.

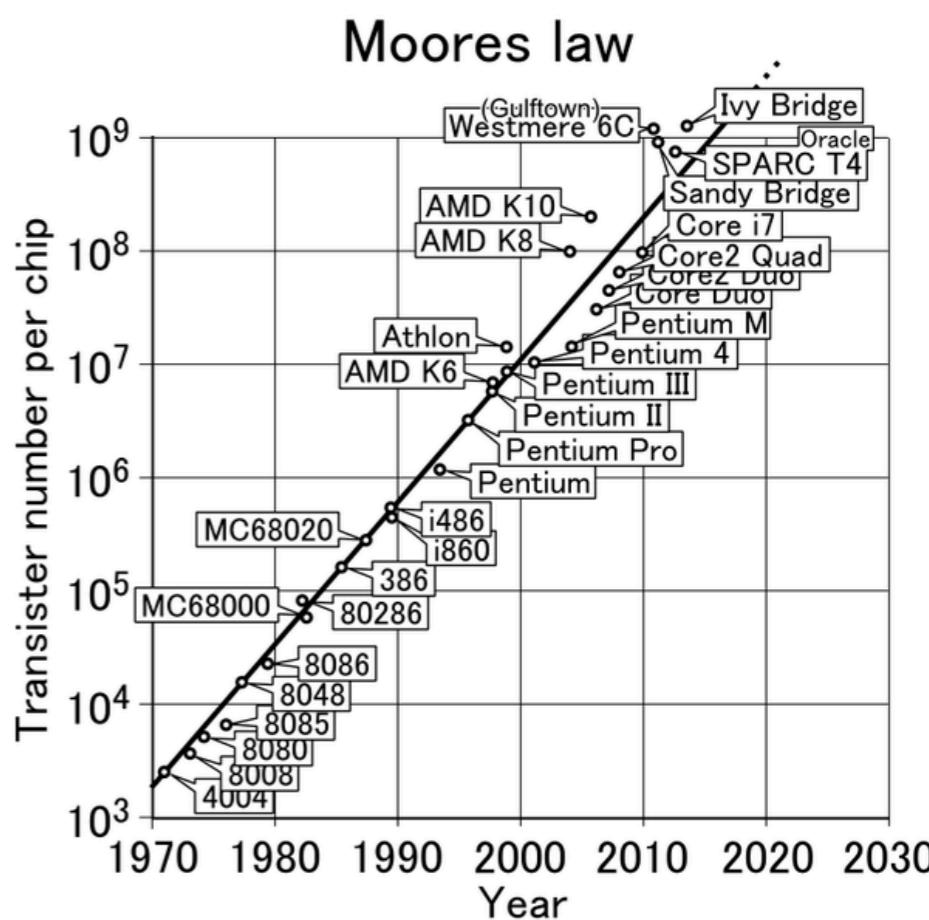
1. **Emitter** (electricity flows in)
2. **Base** (faucet handle)
3. **Collector** (electricity flows out)



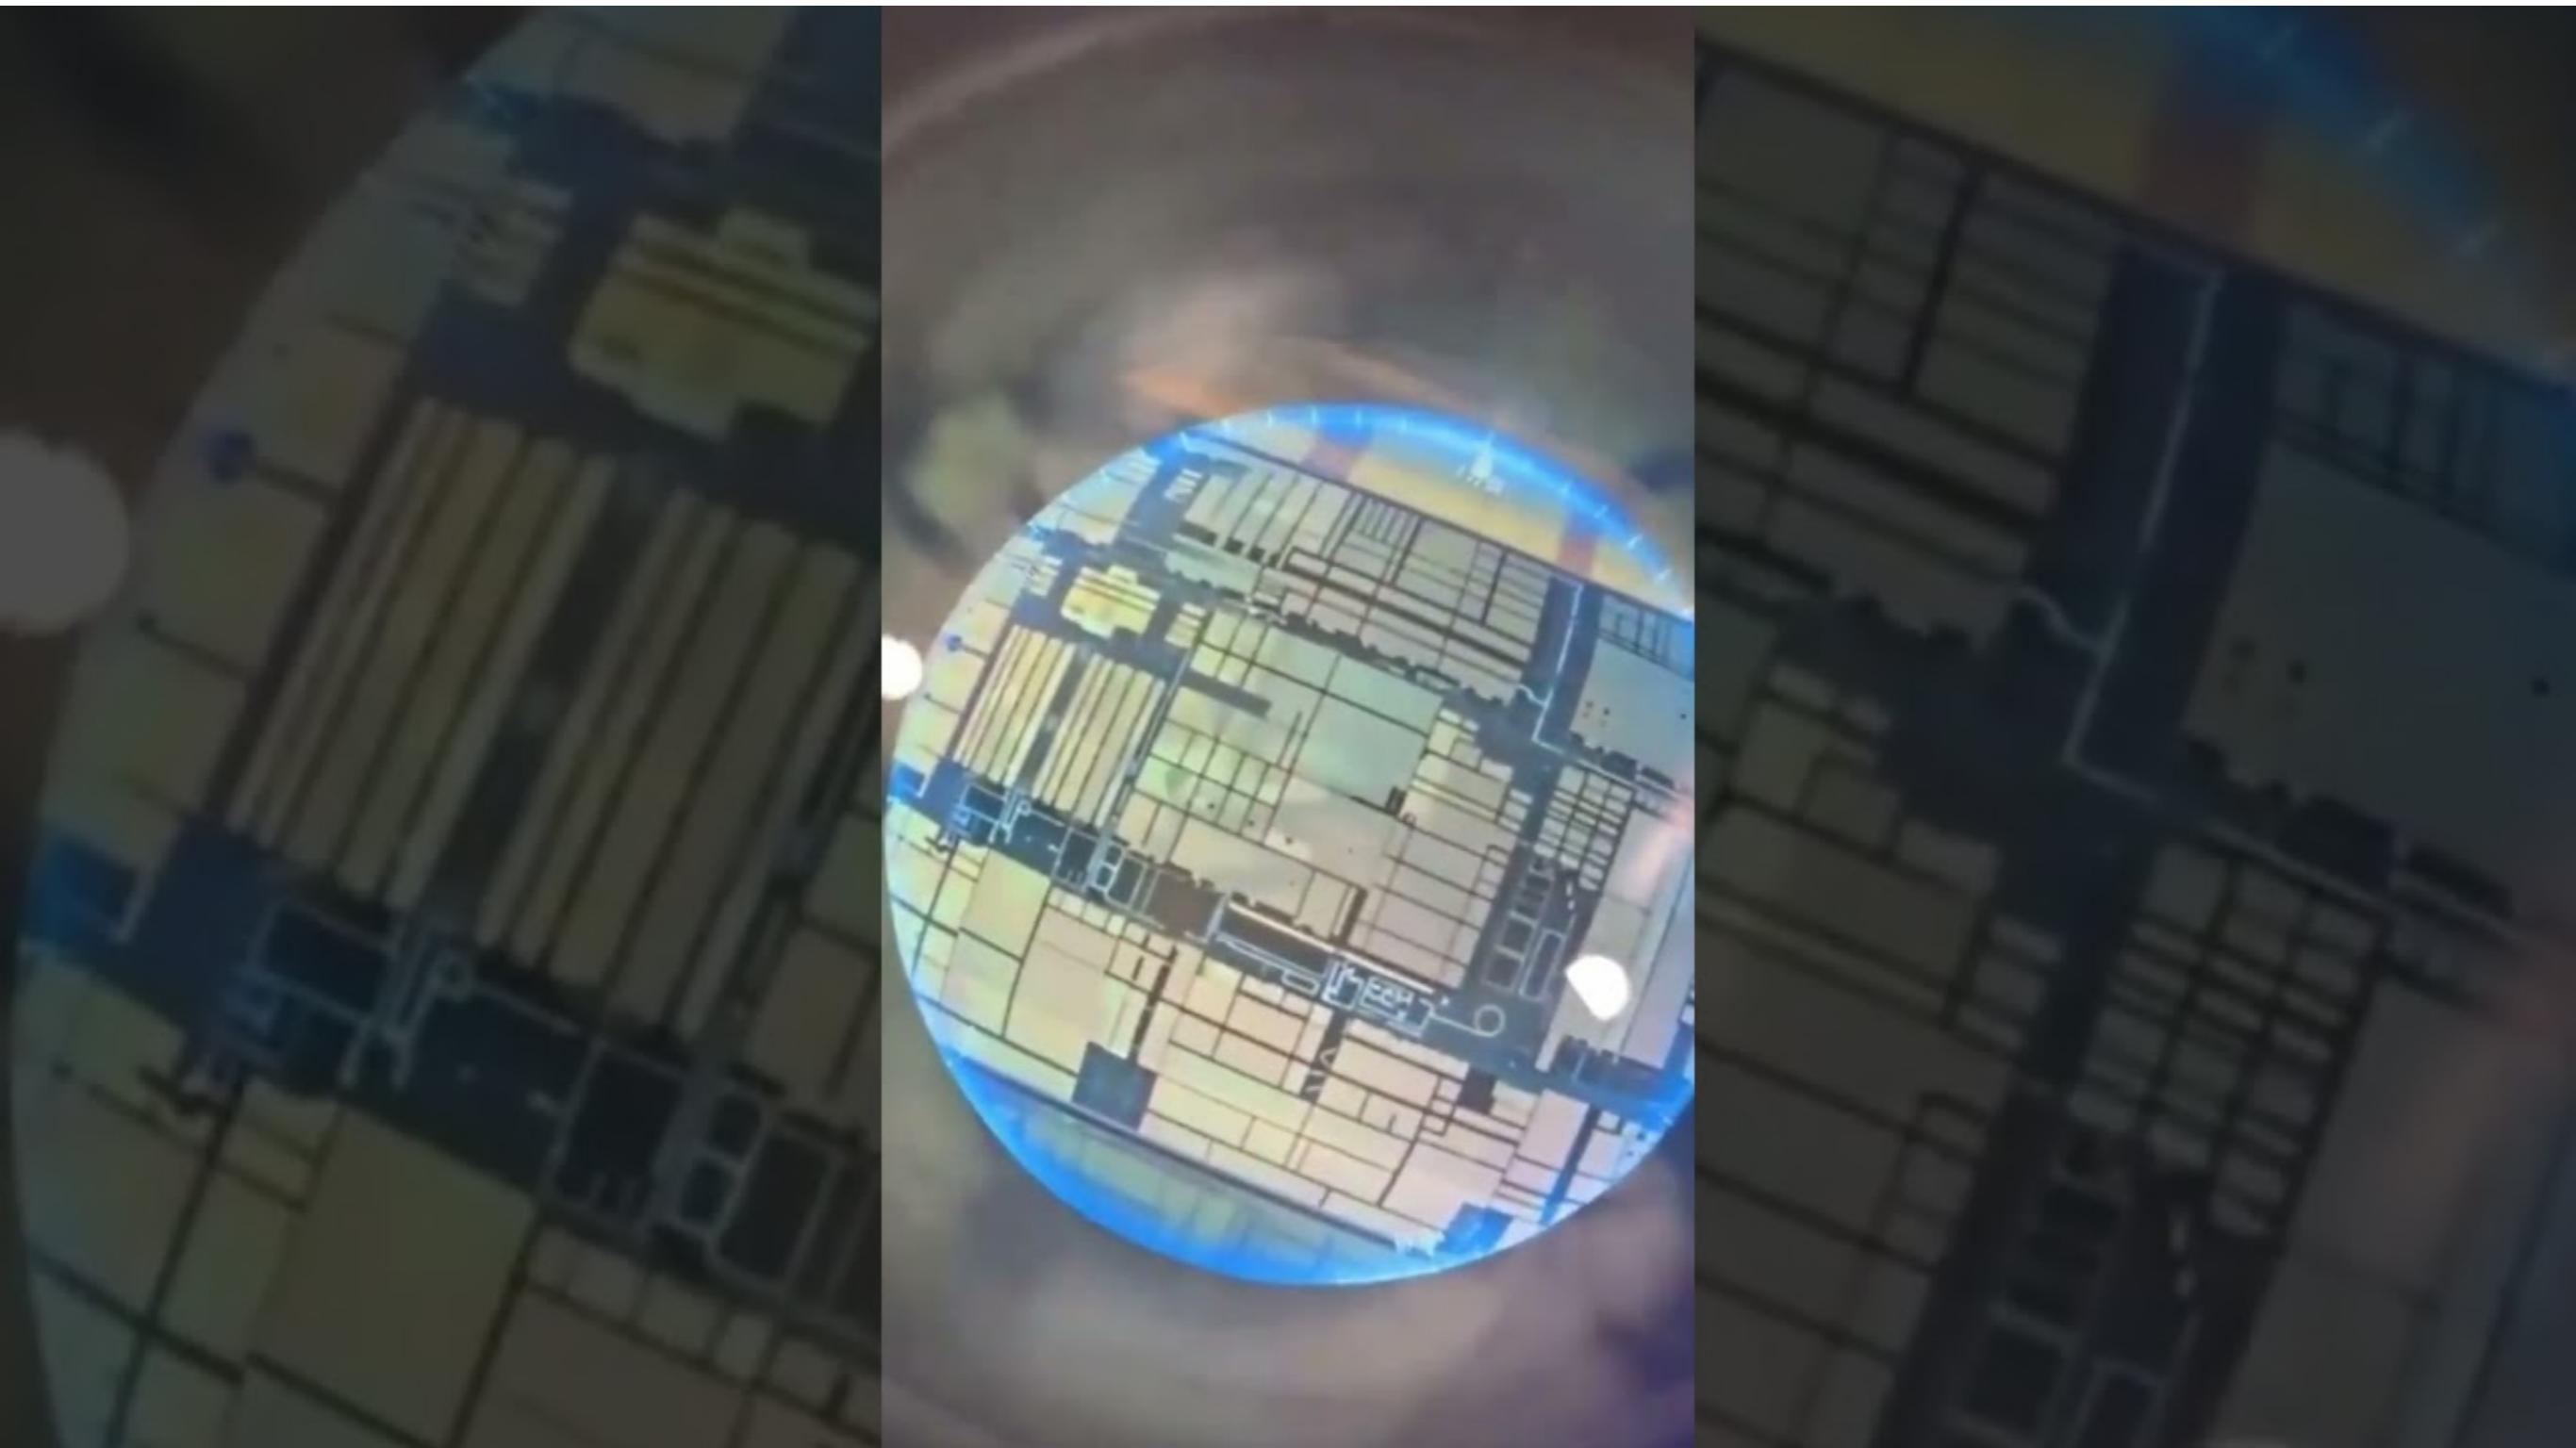
Transistors Enabled Microchips

- ▶ Gordon Moore (Co-Founder of Intel)
 - Former employee of Shockley Semiconductor
 - Part of the infamous "*Traitorous 8*"

- ▶ Known for his "law" or rather, prediction:
 - *"The number of transistors on a chip doubles every 18-24 months."*
 - Reaching physical limits: if transistor width is a few **nanometers**, quantum tunneling can occur, making them unreliable.

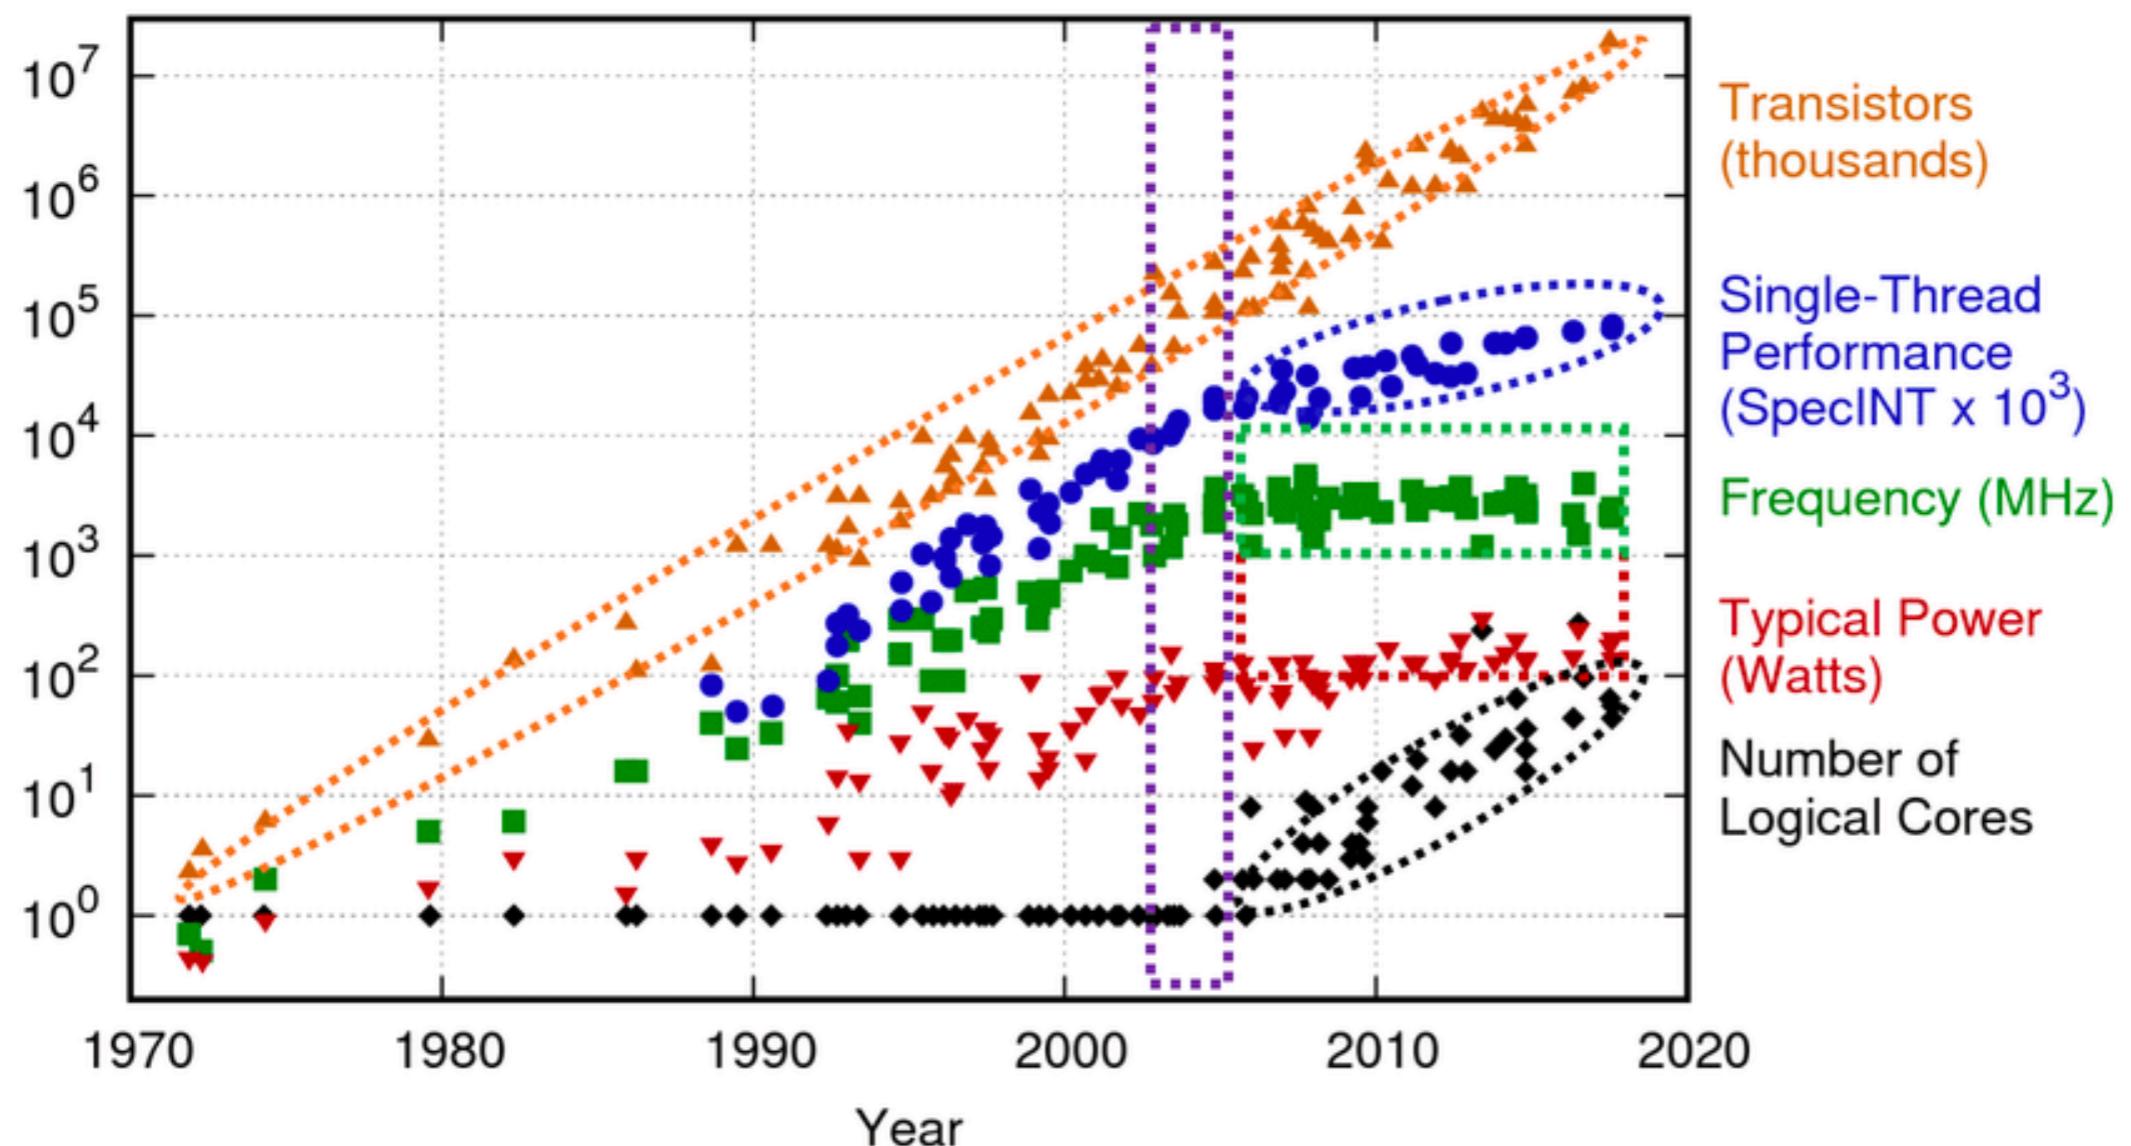


Scale of Chip Components



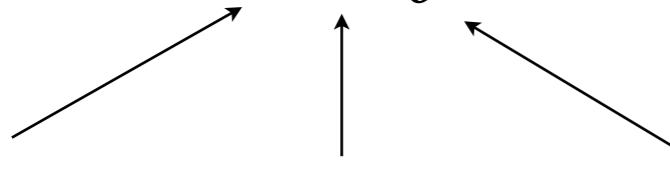
The Slowdown of Joy's Law

- ▶ *Bill Joy's Law:* CPU Performance doubles every ~18 months
 - But notice that the **CPU Performance** suddenly dropped off circa 2001.



Dynamic Power Consumption of CPUs

- ▶ P approximates the power consumed by the CPU:

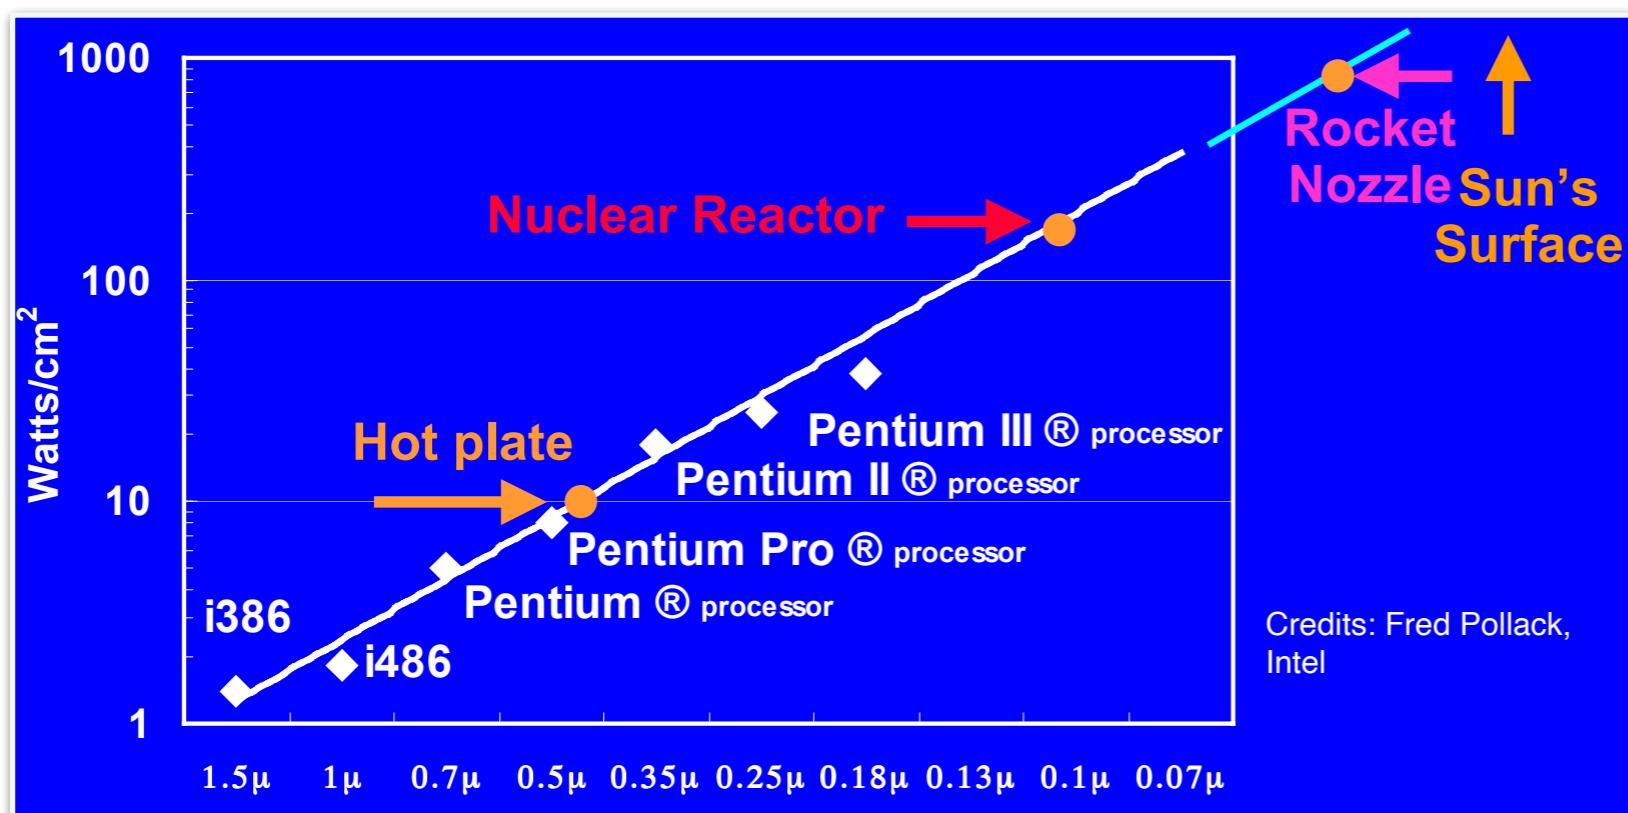
$$P = CV^2f$$


capacitance (chip area) voltage CPU frequency (i.e., clock rate)

- ▶ Before 2001, chip manufacturers had always found ways to increase clock rate f to increase CPU performance.
 - (The CPU used to be just one massive, monolithic core!)

Problem: Power (Heat) Wall

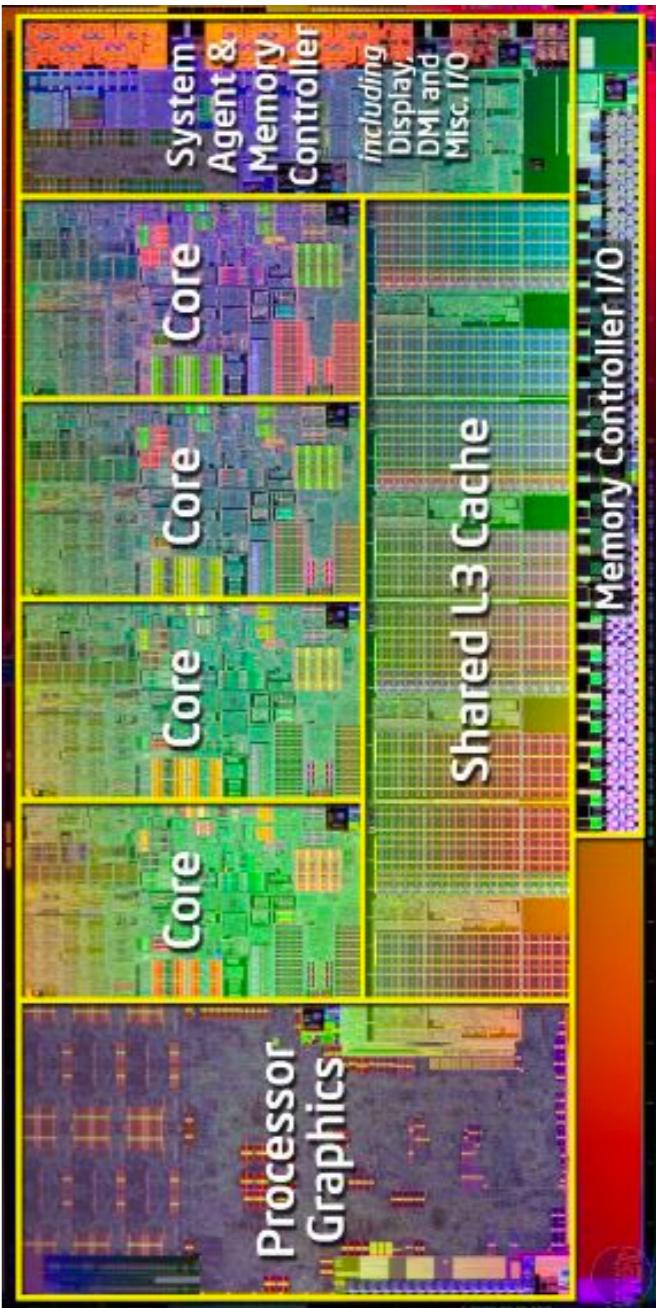
- ▶ "Constraint [of chip making] is power, not manufacturability."
- Dr. Fred Pollack, Intel Fellow
- ▶ Chip manufacturers were met with the "*power wall*" in the early 2000s.
 - Profound amount of heat generated in a very small area.



Solution: Multicore CPUs

- ▶ Reduce *f* (clock rate) on the single CPU to make it puny
 - But Moore's Law permitted more transistors per unit area.
 - Ah ha! Add more puny CPUs ("cores") to each chip

- ▶ Dynamic Power *P* can now be held constant, but what about *performance*?
 - Instead of one massive 4 GHz CPU, make a CPU with two puny cores running at 2 GHz a piece. $2+2 = 4$ right?
 - Can we get close to 4 GHz performance when both cores are utilized simultaneously? (*The need for parallelism*)



Big Picture Today

- ▶ Life used to be easy. Slow code? No problem, wait for Moore's Law.
- ▶ Today
 - Multi-core CPUs are now commonplace
 - Need support from OS
 - Increased need for parallel programmers
- ▶ What's Next?
 - The decline of Moore's Law reveals new opportunities.
 - Co-Processors like Nvidia GPUs
 - 3D stacking of CPU planes
 - Advanced materials, like graphene, instead of silicon to overcome physical limits
 - Quantum computers

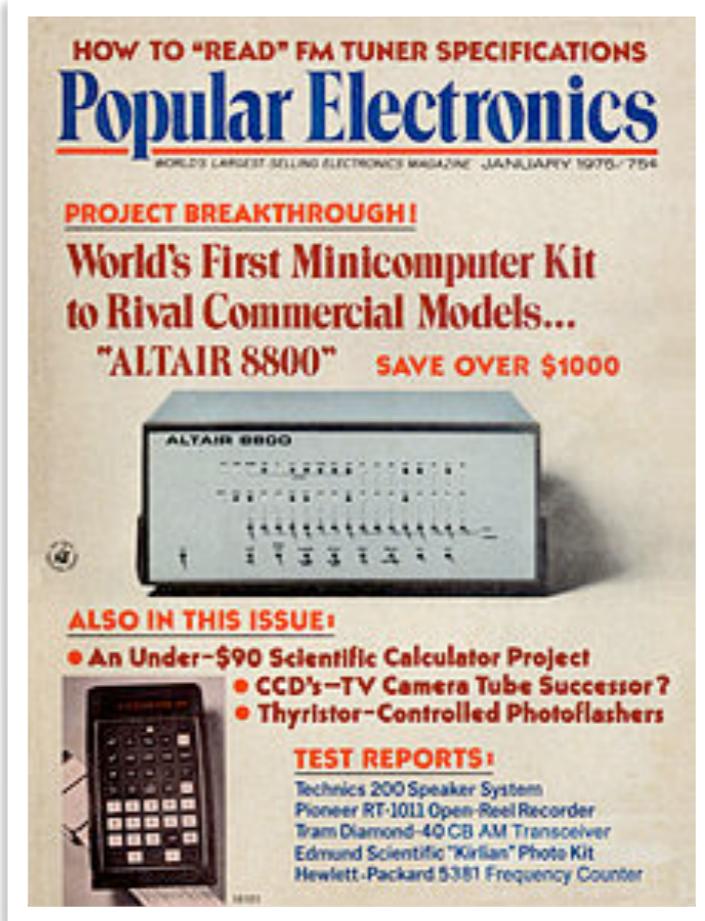
Today's Topics

- ▶ Course Overview
- ▶ **What Is an Operating System, Anyway?**
- ▶ A Brief History of Computer and Operating Systems
- ▶ Conclusion

What's Life Like without OS?

► MITS Altair 8800 (1975)

- World's first affordable PC
 - Intel 8080 (8-bits, 2MHz), 256B RAM, \$397
- No keyboard, no monitor display, no OS
 - Program by toggling binary code, hit run, read results off the LED display.
 - Sparked many startups (Microsoft was born!)
- For the curious
 - <https://youtu.be/suyiMfzmZKs>



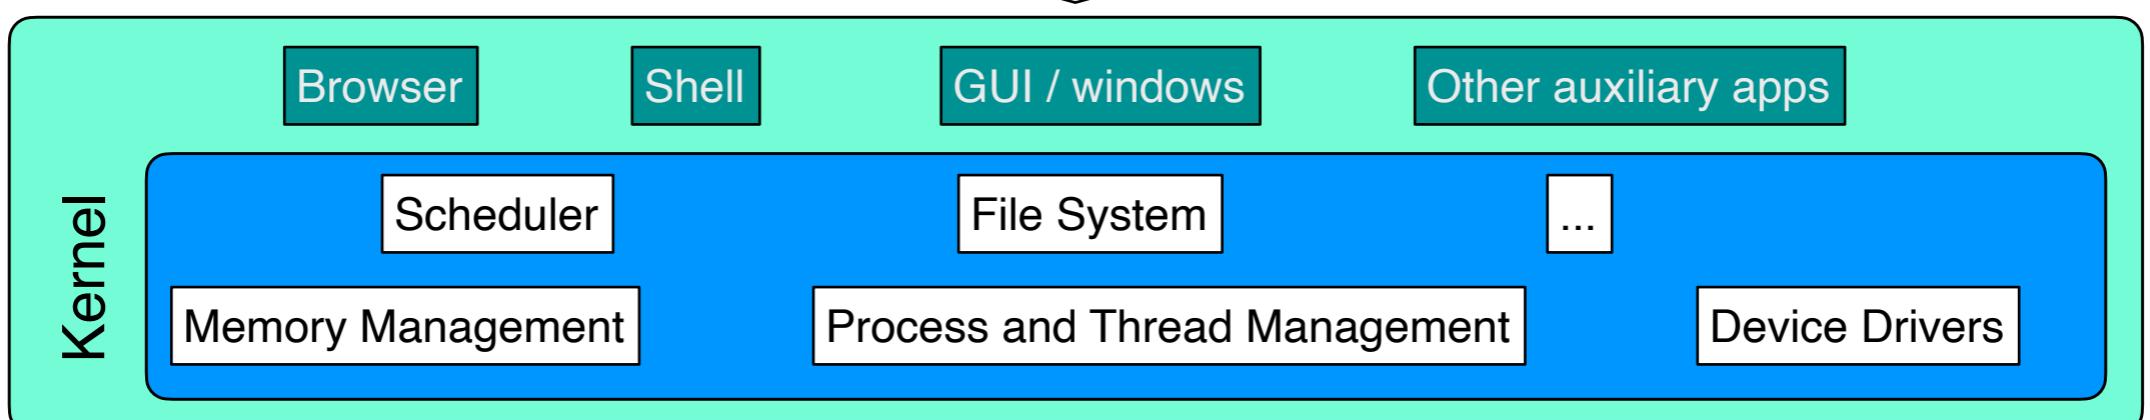
What Is an Operating System (Kernel)?

User Applications



*OS Services
(via System Call API)*

Operating System



Hardware



*Hardware Interface
(called "device drivers")*

What Is an Operating System?

- ▶ The OS is an illusionist
- ▶ Illusion:
 - Programs (processes) think they get unrestricted access to all hardware units.
- ▶ Reality: OS must establish controlled sharing among multiple processes.
 - This is called "*hardware virtualization*"



Illusionist: Virtualizing the CPU

- ▶ Each running program *thinks* it gets to run on its own CPU.

Keynote

Technological Trends

► Gordon Moore (Co-Founder, Intel)

► Known for "Moore's Law"

- Prediction that transistor-density on an IC chip would double every 2 years*

► Link to 1965 paper

- http://download.intel.com/museum/Moores_Law/Articles-Press_Releases/Gordon_Moore_1965_Article.pdf

*Original prediction was even more ambitious: doubling every year.

CS 475: Operating Systems - 1 - Introduction

I think I have my
own CPU

Spotify



So do I!



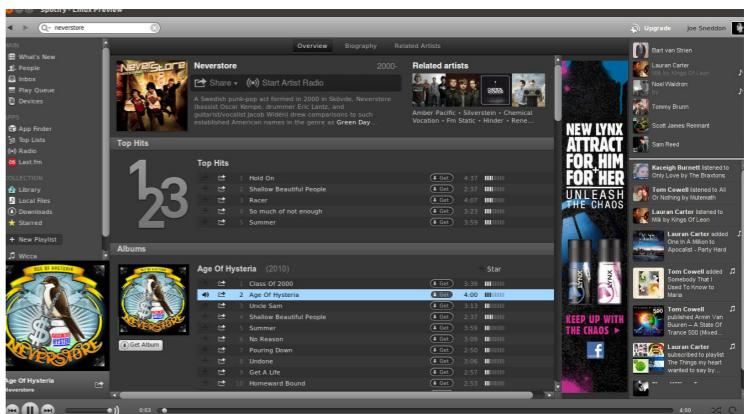
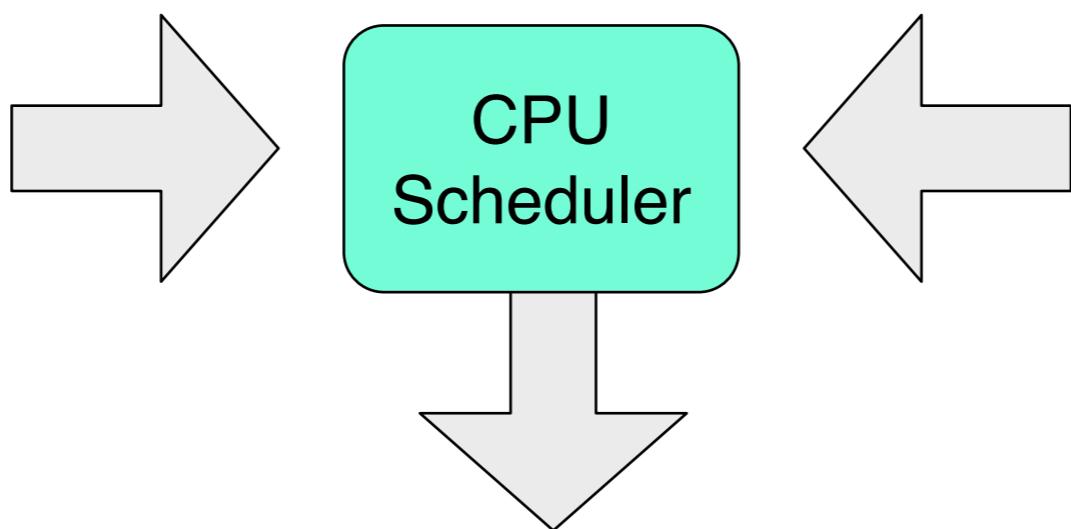
Illusionist: Virtualizing the CPU (Cont.)

- ▶ Reality: number of processes > number of CPUs
 - *CPU scheduler* establishes time-sharing via "context switching."

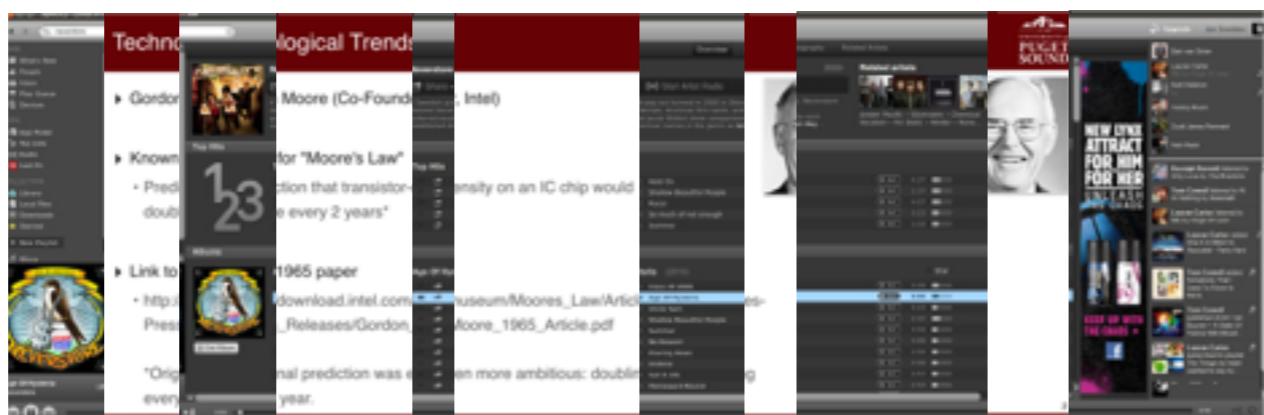
Technological Trends

- ▶ Gordon Moore (Co-Founder, Intel)
 - 
- ▶ Known for "Moore's Law"
 - Prediction that transistor-density on an IC chip would double every 2 years*
- ▶ Link to 1965 paper
 - http://download.intel.com/museum/Moores_Law/Articles-Press_Releases/Gordon_Moore_1965_Article.pdf
- *Original prediction was even more ambitious: doubling every year.

CR 47% Operating Systems - 1 - Introduction

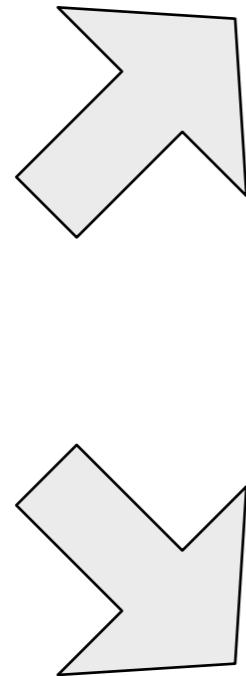


Timeshare a single CPU core



Illusionist: Virtualizing Memory

- ▶ Provide an illusion of having infinite memory dedicated to each program

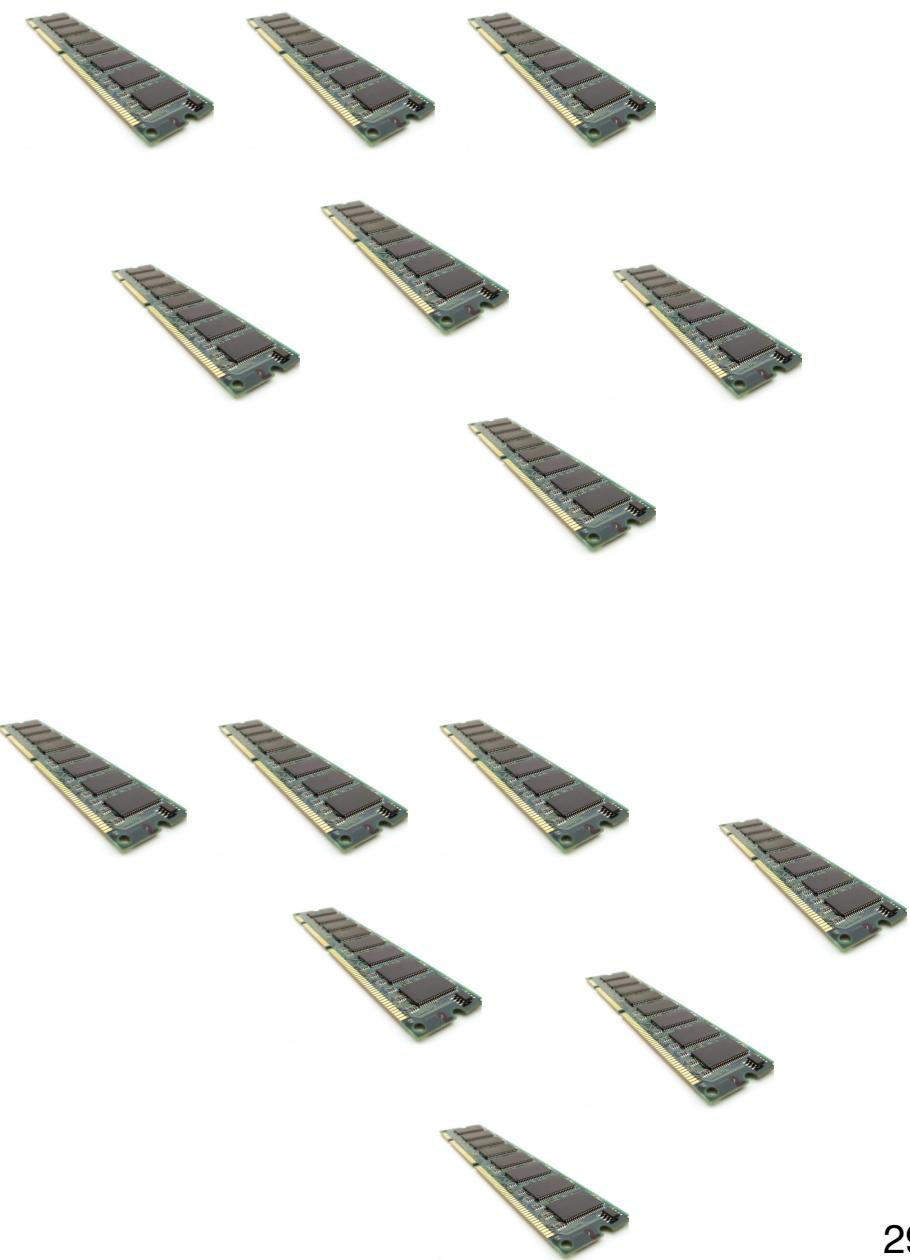


Technological Trends

- ▶ Gordon Moore (Co-Founder, Intel)
- ▶ Known for "Moore's Law"
 - Prediction that transistor-density on an IC chip would double every 2 years*
- ▶ Link to 1965 paper
 - http://download.intel.com/museum/Moores_Law/Articles-Press_Releases/Gordon_Moore_1965_Article.pdf



*Original prediction was even more ambitious: doubling every year.



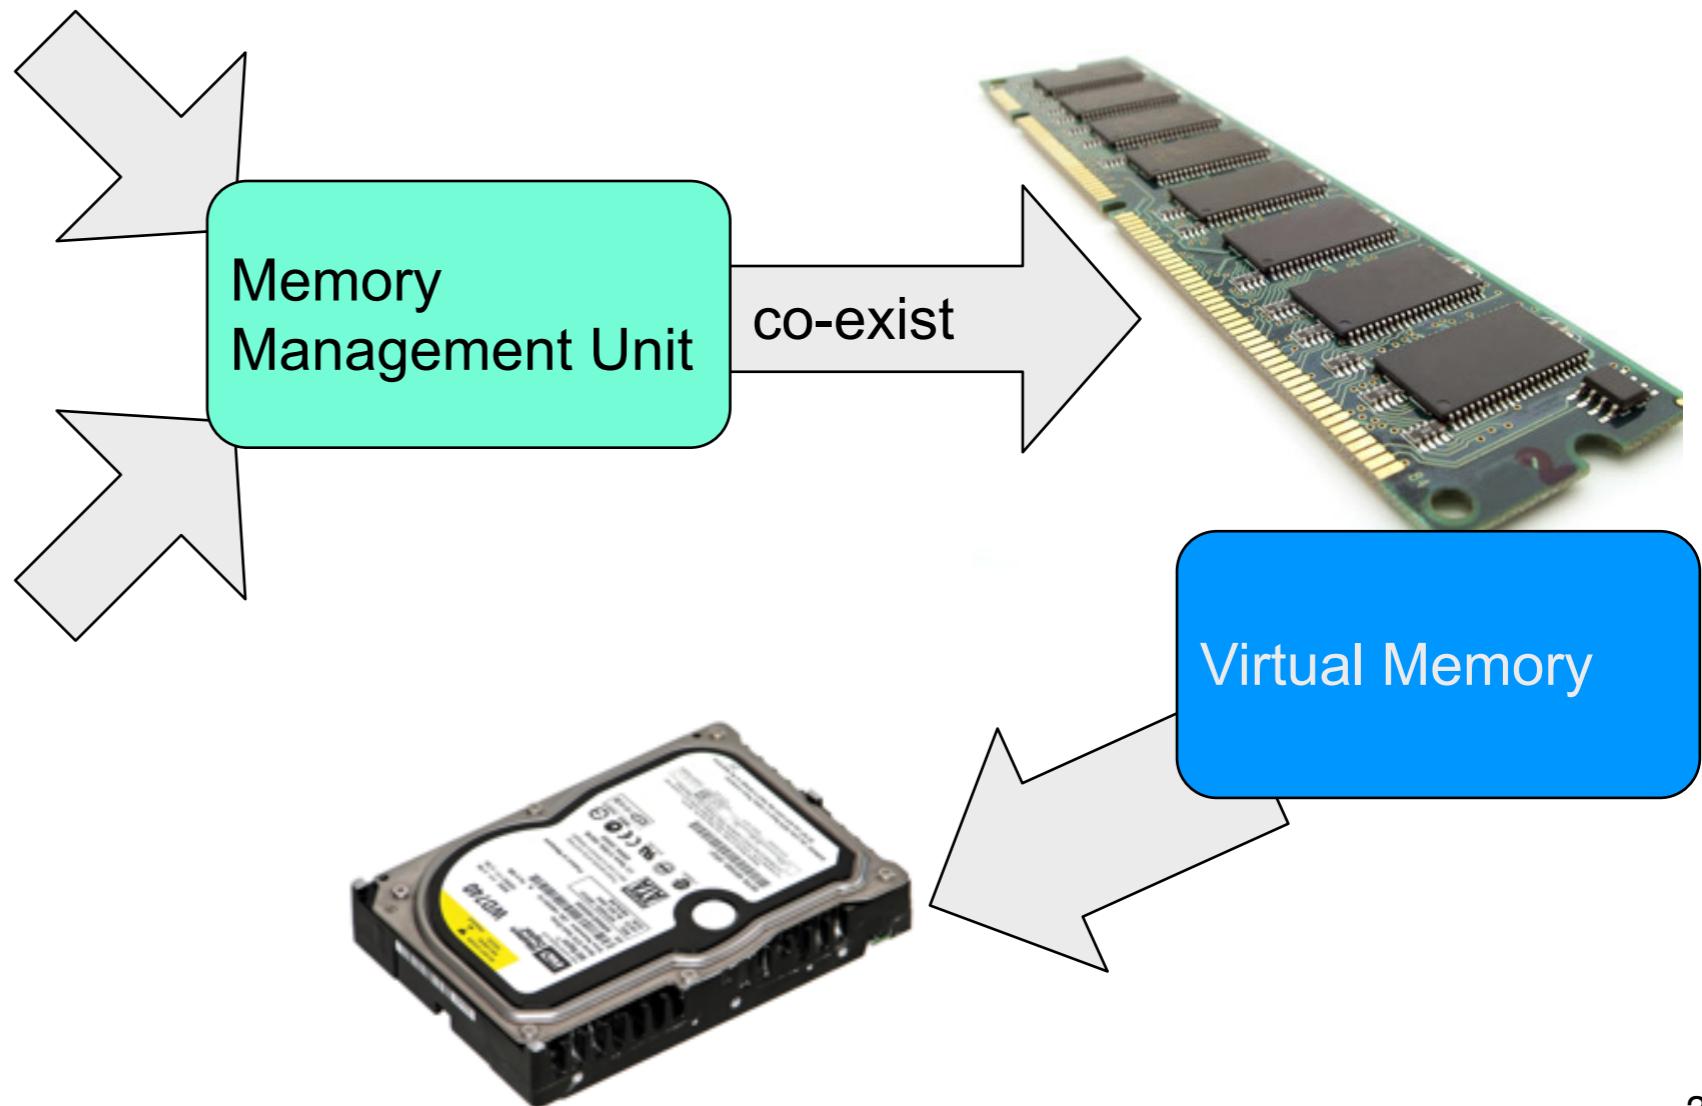
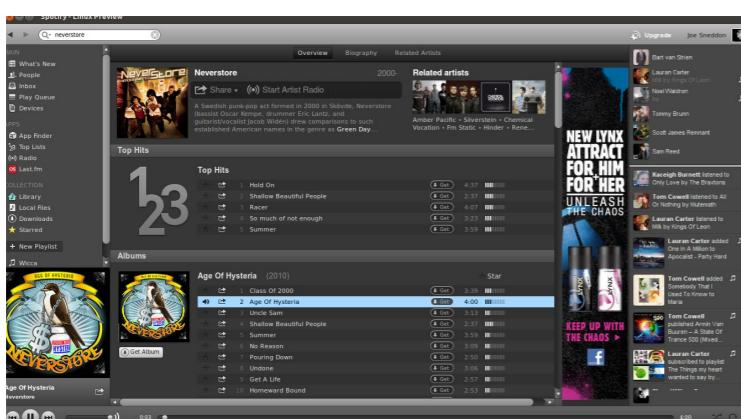
Illusionist: Virtualizing Memory

- ▶ **Reality:** Limited amount of memory. OS's memory management unit allows for sharing (and protection).

Technological Trends

- ▶ Gordon Moore (Co-Founder, Intel)
- ▶ Known for "Moore's Law"
 - Prediction that transistor-density on an IC chip would double every 2 years*
- ▶ Link to 1965 paper
 - http://download.intel.com/museum/Moores_Law/Articles_Press_Releases/Gordon_Moore_1965_Article.pdf
- Original prediction was even more ambitious: doubling every year.

CS 475: Operating Systems - 1 - Introduction

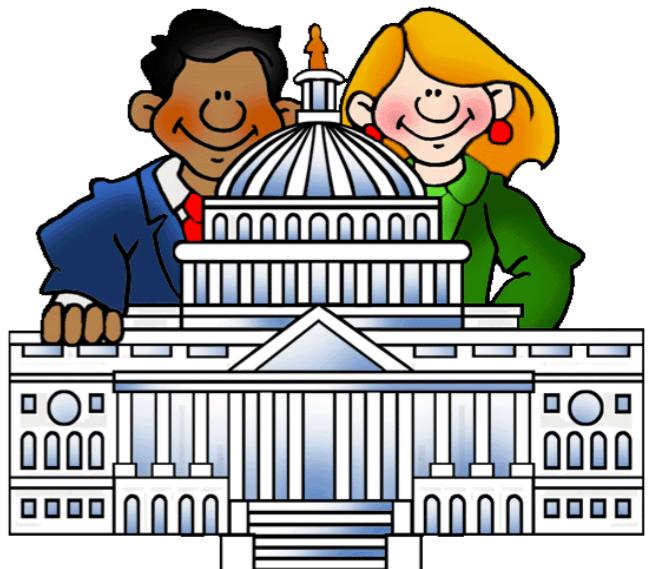


What Else Does an OS Do for Us?

► The OS is also like a government

- It provides a set of regulated services that everyone needs but are too risky to let users use without oversight. Examples:

- Create, terminate, suspend a process
- Open/read/write a file
- Printing to the screen
- Requesting more memory during runtime



- Users can request the OS to take action made through a "*system call*"
 - Each OS provides a standard library of system functions
 - (In Linux, calls made through a C library called **unistd.h**)

Sample System Call APIs

- ▶ *System calls:* These actions are privileged, so you must request the OS to do them on your behalf. (Defined in **unistd.h**)

Category	Linux
Process Control	fork()
	exit()
	wait()
File Manipulation	open()
	read()
	write()
	close()
File Protection	chmod()
	chown()
...	...

Today's Topics

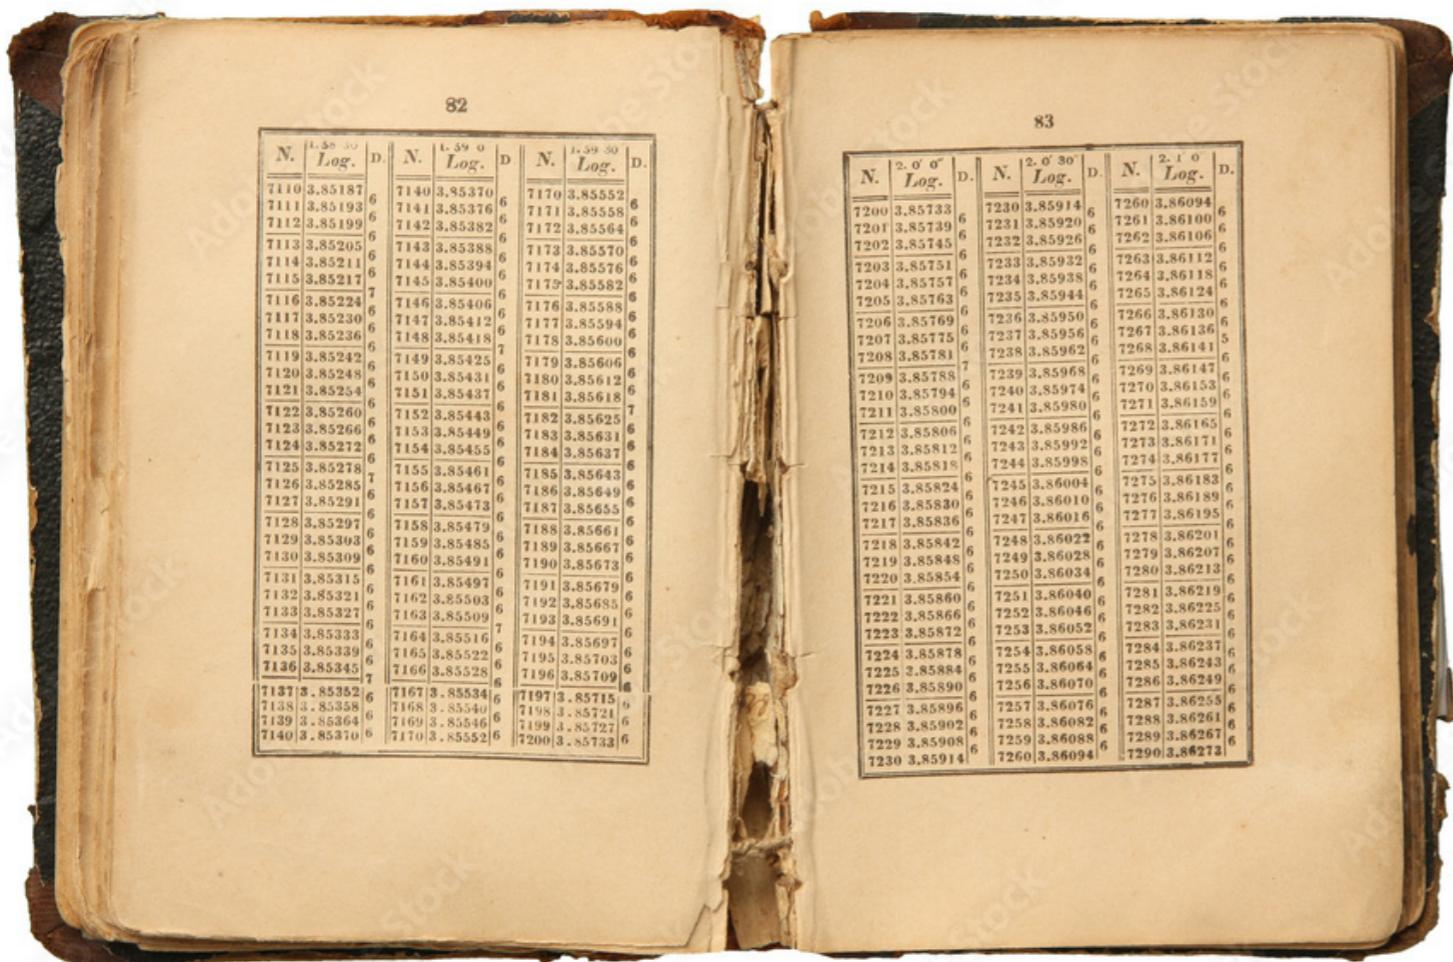
- ▶ Course Overview
- ▶ What Is an OS, Anyway?
- ▶ A Brief History of Computer and Operating Systems
- ▶ Conclusion

1700s - 1800s: Origins of Computing

► Importance of data tables

- Used for engineering, studying astronomy, ...
- Calculations done by hand. Lacked precision and accuracy

Example: Table of Logarithms

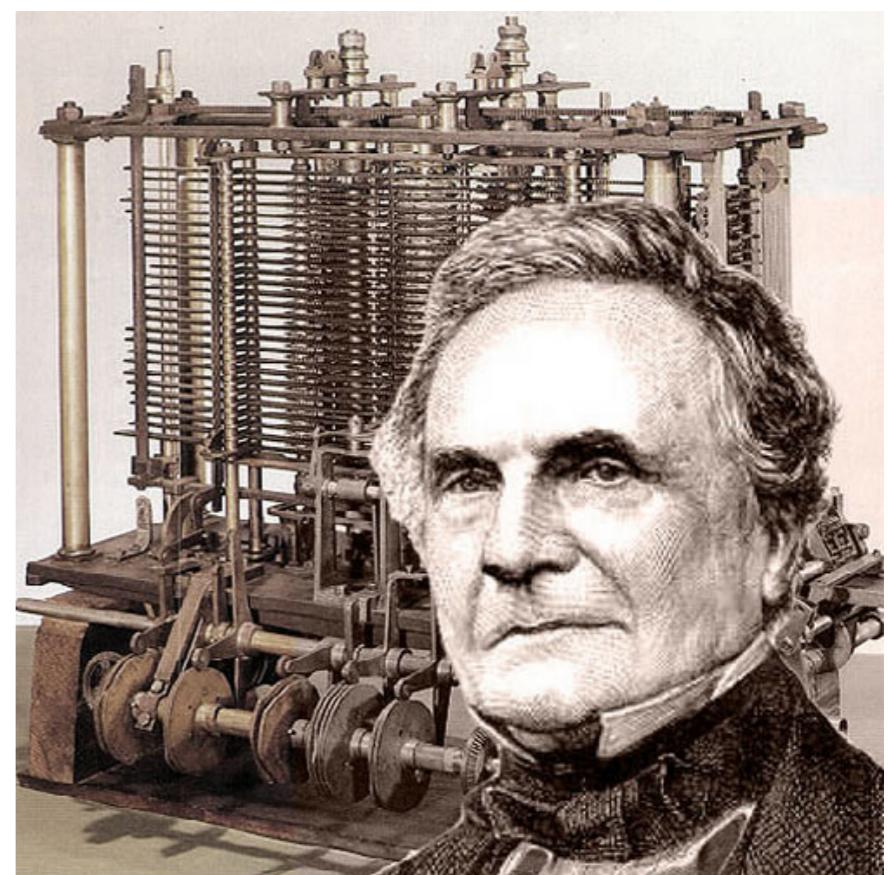


The image shows an open book with two pages of a logarithm table. The left page is numbered 82 and the right page is numbered 83. Both pages contain tables of logarithmic values for various numbers.

N.	Log.	D.	N.	Log.	D.	N.	Log.	D.
7110	3.85187	6	7140	3.85370	6	7170	3.85552	6
7111	3.85193	6	7141	3.85376	6	7171	3.85558	6
7112	3.85199	6	7142	3.85382	6	7172	3.85564	6
7113	3.85205	6	7143	3.85388	6	7173	3.85570	6
7114	3.85211	6	7144	3.85394	6	7174	3.85576	6
7115	3.85217	6	7145	3.85400	6	7175	3.85582	6
7116	3.85224	7	7146	3.85406	6	7176	3.85588	6
7117	3.85230	6	7147	3.85412	6	7177	3.85594	6
7118	3.85236	6	7148	3.85418	7	7178	3.85600	6
7119	3.85242	6	7149	3.85425	6	7179	3.85606	6
7120	3.85248	6	7150	3.85431	6	7180	3.85612	6
7121	3.85254	6	7151	3.85437	6	7181	3.85618	6
7122	3.85260	6	7152	3.85443	6	7182	3.85625	6
7123	3.85266	6	7153	3.85449	6	7183	3.85631	6
7124	3.85272	6	7154	3.85455	6	7184	3.85637	6
7125	3.85278	7	7155	3.85461	6	7185	3.85643	6
7126	3.85285	6	7156	3.85467	6	7186	3.85649	6
7127	3.85291	6	7157	3.85473	6	7187	3.85655	6
7128	3.85297	6	7158	3.85479	6	7188	3.85661	6
7129	3.85303	6	7159	3.85485	6	7189	3.85667	6
7130	3.85309	6	7160	3.85491	6	7190	3.85673	6
7131	3.85315	6	7161	3.85497	6	7191	3.85679	6
7132	3.85321	6	7162	3.85503	6	7192	3.85685	6
7133	3.85327	6	7163	3.85509	6	7193	3.85691	6
7134	3.85333	6	7164	3.85516	7	7194	3.85697	6
7135	3.85339	6	7165	3.85522	6	7195	3.85703	6
7136	3.85345	7	7166	3.85528	6	7196	3.85709	6
7137	3.85352	6	7167	3.85534	6	7197	3.85715	6
7138	3.85358	6	7168	3.85540	6	7198	3.85721	6
7139	3.85364	6	7169	3.85546	6	7199	3.85727	6
7140	3.85370	6	7170	3.85552	6	7200	3.85733	6

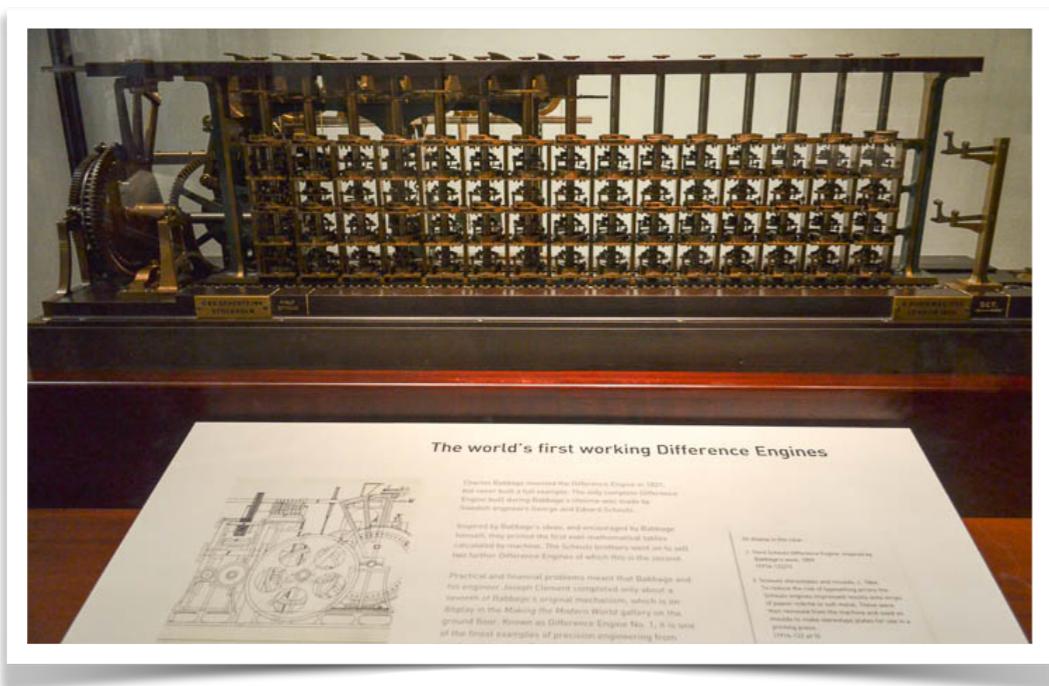
N.	2. 0' 0"	D.	N.	2. 0' 30"	D.	N.	2. 1' 0"	D.
7200	3.85733	6	7230	3.85914	6	7260	3.86094	6
7201	3.85739	6	7231	3.85920	6	7261	3.86100	6
7202	3.85745	6	7232	3.85926	6	7262	3.86106	6
7203	3.85751	6	7233	3.85932	6	7263	3.86112	6
7204	3.85757	6	7234	3.85938	6	7264	3.86118	6
7205	3.85763	6	7235	3.85944	6	7265	3.86124	6
7206	3.85769	6	7236	3.85950	6	7266	3.86130	6
7207	3.85775	6	7237	3.85956	6	7267	3.86136	5
7208	3.85781	7	7238	3.85962	6	7268	3.86141	6
7209	3.85788	6	7239	3.85968	6	7269	3.86147	6
7210	3.85794	6	7240	3.85974	6	7270	3.86153	6
7211	3.85800	6	7241	3.85980	6	7271	3.86159	6
7212	3.85806	6	7242	3.85986	6	7272	3.86165	6
7213	3.85812	6	7243	3.85992	6	7273	3.86171	6
7214	3.85818	6	7244	3.85998	6	7274	3.86177	6
7215	3.85824	6	7245	3.86004	6	7275	3.86183	6
7216	3.85830	6	7246	3.86010	6	7276	3.86189	6
7217	3.85836	6	7247	3.86016	6	7277	3.86195	6
7218	3.85842	6	7248	3.86022	6	7278	3.86201	6
7219	3.85848	6	7249	3.86028	6	7279	3.86207	6
7220	3.85854	6	7250	3.86034	6	7280	3.86213	6
7221	3.85860	6	7251	3.86040	6	7281	3.86219	6
7222	3.85866	6	7252	3.86046	6	7282	3.86225	6
7223	3.85872	6	7253	3.86052	6	7283	3.86231	6
7224	3.85878	6	7254	3.86058	6	7284	3.86237	6
7225	3.85884	6	7255	3.86064	6	7285	3.86243	6
7226	3.85890	6	7256	3.86070	6	7286	3.86249	6
7227	3.85896	6	7257	3.86076	6	7287	3.86255	6
7228	3.85902	6	7258	3.86082	6	7288	3.86261	6
7229	3.85908	6	7259	3.86088	6	7289	3.86267	6
7230	3.85914	6	7260	3.86094	6	7290	3.86273	6

Charles Babbage



Origins: Analytical Engine (1800s)

- ▶ Charles Babbage (1791-1871) and Ada Lovelace (1815-1852)
 - Babbage invented the Difference Engine, a mechanical calculator
 - Drew up plans for the Analytical Engine, but was never built
 - Watch: Calculating Ada documentary



Part of the Difference Engine, Science Museum in London



Ada Lovelace:
Published the first program (for the analytical engine)

Babbage's Difference Engine

- ▶ Print results of polynomials in a table format
 - Problem: labor-intensive, prone to human error

- ▶ Key insight:
 - Automate the *method of finite differences*
 - *Decompose task using simple additions*
 - For an degree- n polynomial, input all n differences for first few values of x .
 - Then $f(x) = f(x - 1) + \sum_{i=1}^n \Delta_i$
 - <https://bit.ly/3kF8srn>

- ▶ **"Registers"** held the current iteration's values

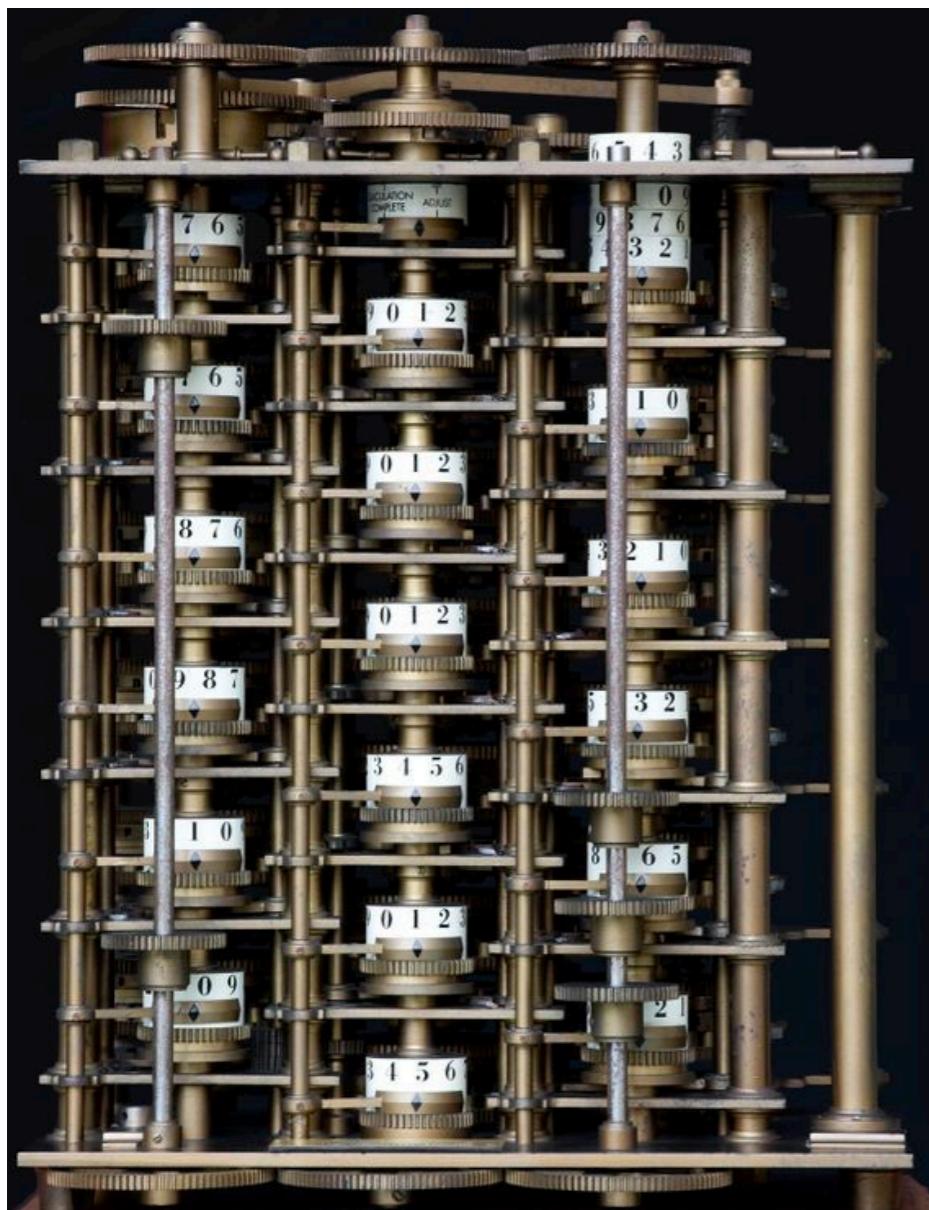
$$f(x) = 2x^2 + 4$$

x	$f(x)$	Δ_1	Δ_2
1	6		
2	12	6	
3	22	10	4
4	36	14	4
5	54	18	4
6	?

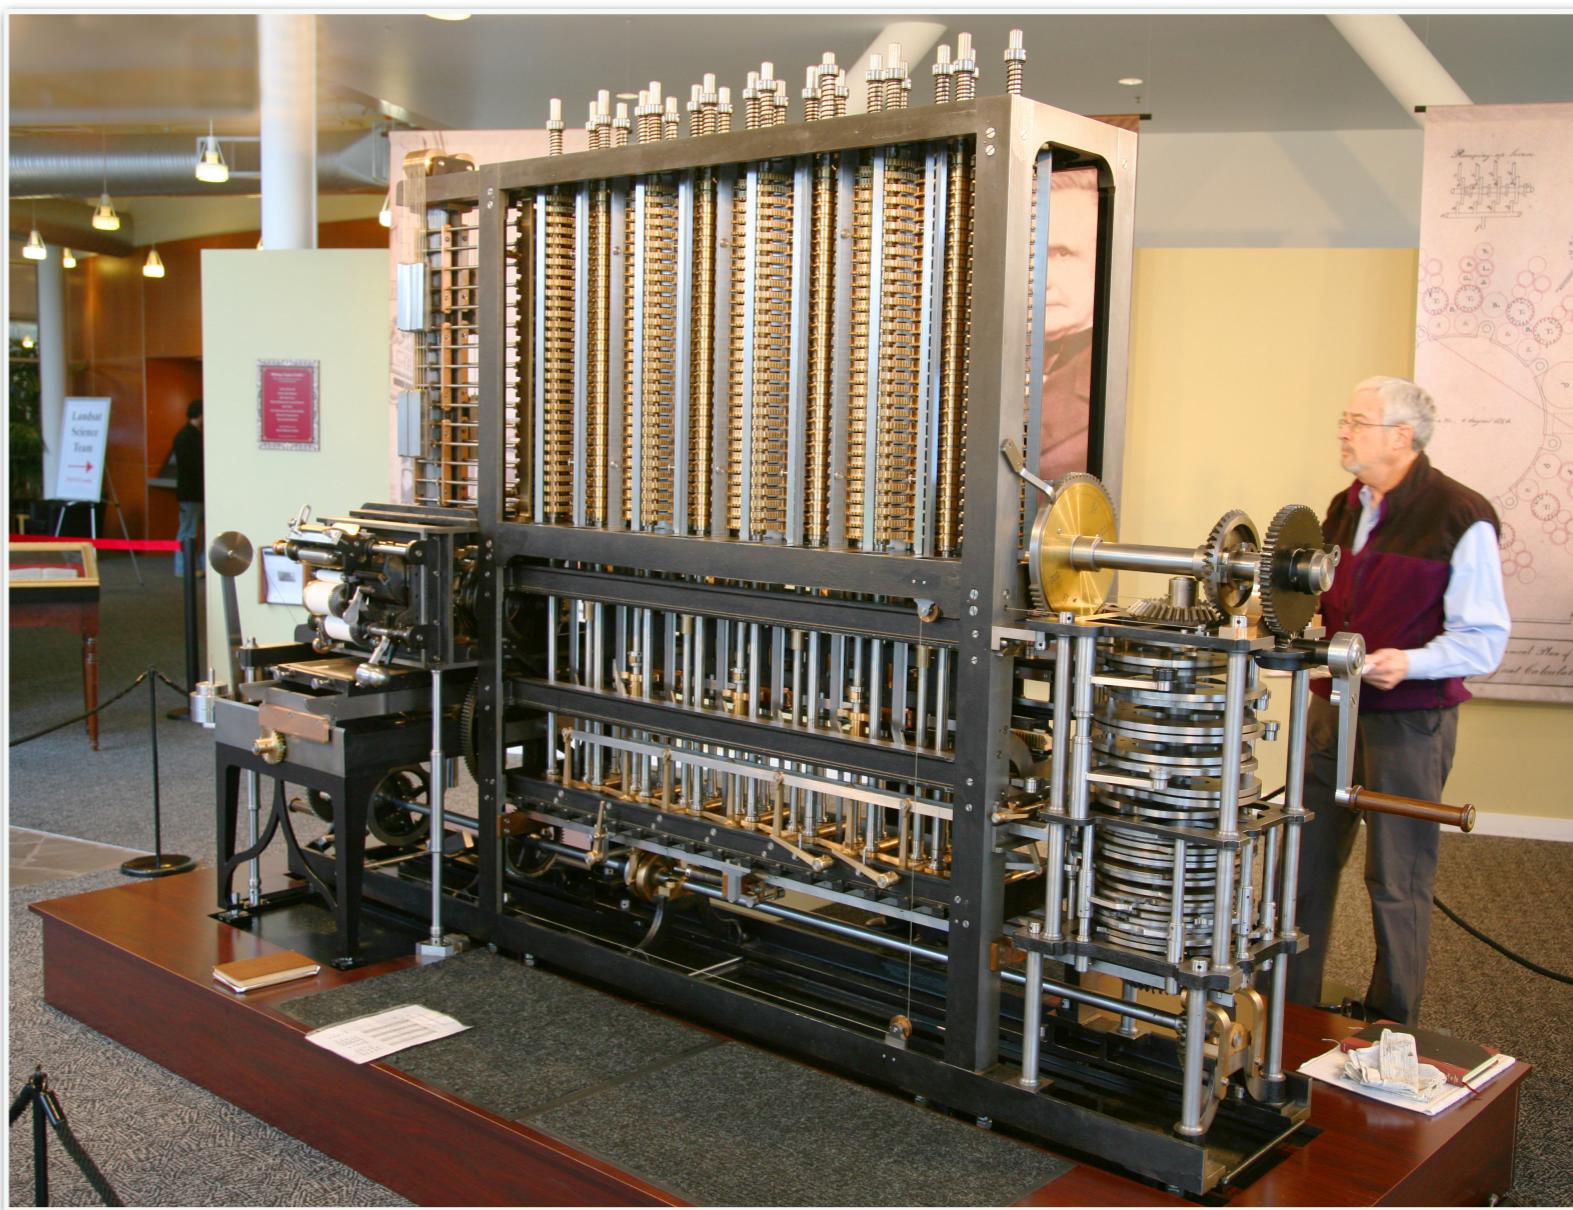
A diagram showing the computation of the 6th value in the sequence. Arrows point from the 4th value (36) and the 5th value (54) to their respective Δ_1 values (14 and 18). Another arrow points from the 5th value (54) to its Δ_2 value (4).

Babbage's Difference Engine

- ▶ A reproduction built for the Computer History Museum (CHM)



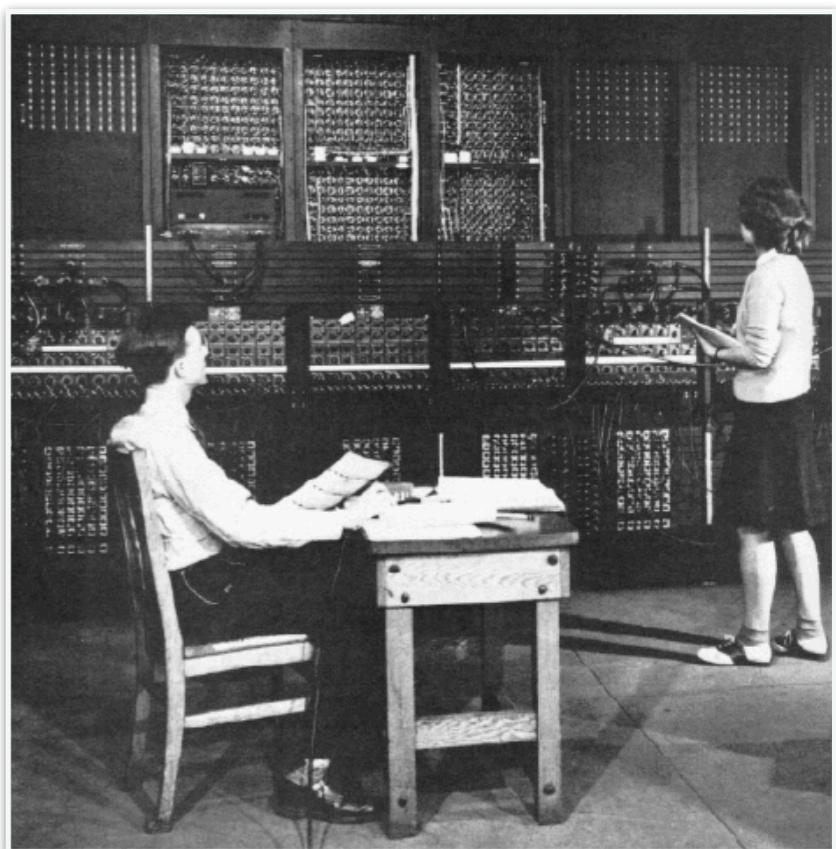
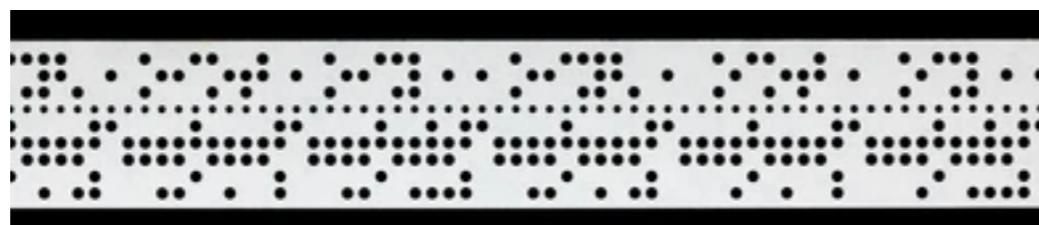
"Registers"



A Replica at the Computer History Museum

1940s: Dawn of Time

- ▶ Hardware is very expensive and experimental forms of art
 - Only universities and Federal Labs had "automatic computers"
 - First manufactured by IBM as general-purpose "calculators"
- ▶ Computers have lots of **moving parts (slow; unreliable)**
- ▶ No stored programs
- ▶ Input/output
 - Programs on punched tape
 - Printer
- ▶ No OS



1940s: Harvard IBM Mark I (1944)

- ▶ First fully-automatic, *electro-mechanical* general-purpose computer
 - 51ft x 8ft x 8ft, weighed 5 tons
 - 3 adds/subtracts per sec; 1 multiplication in 6 secs
 - Fully programmable!
 - Stopped and waited for "go" signal between operations; made for easy debugging
 - No OS, programs on tape fed by humans
- ▶ *Fun fact: "Sounded like a room full of people knitting."*



1940s: Harvard IBM Mark I (1944)

- ▶ Mark I programming team led by **Dr. Grace Hopper**
 - Initially, a Professor of Mathematics @ Vassar College
 - Then, as a US Naval Officer, became the first principle programmer of the Mark I

On the Harvard Mark I:

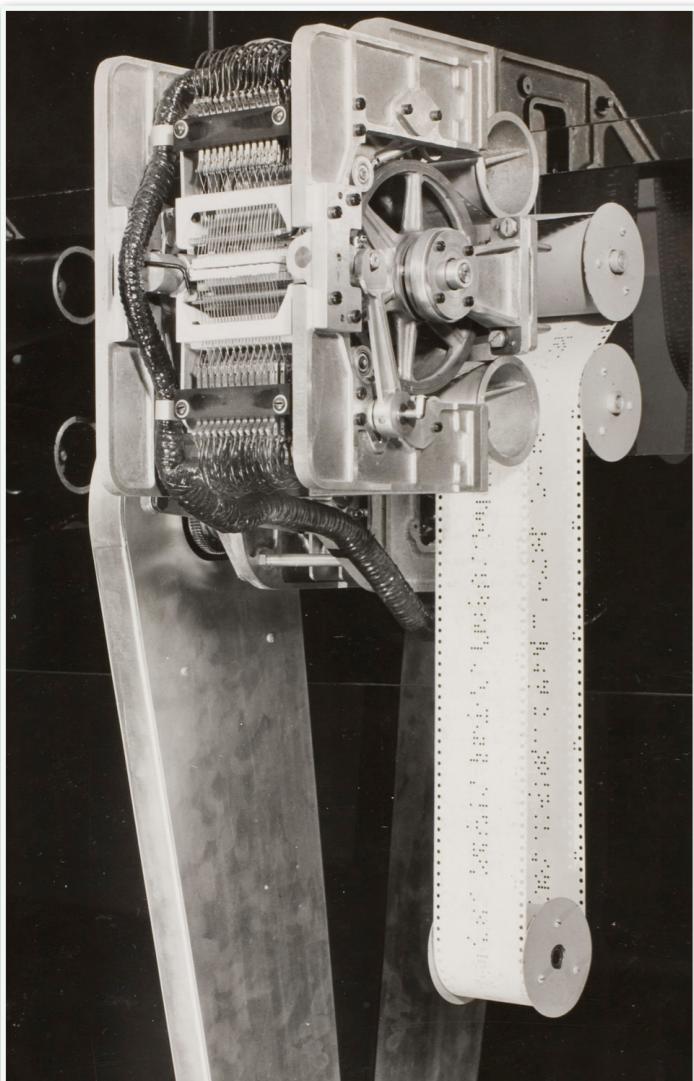
"[it was] the first machine that was built that was supposed to assist the power of man's brain instead of the strength of his arms."

Grace Murray Hopper

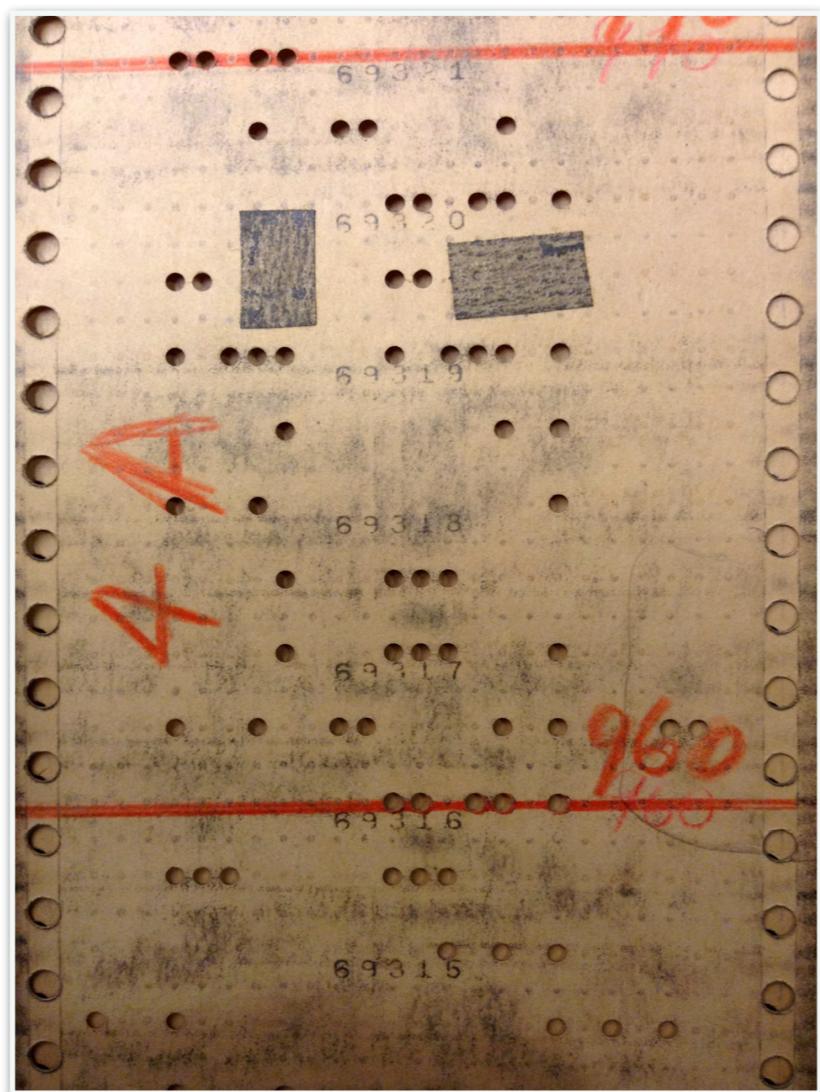


Aside: Etymology of Common Terms

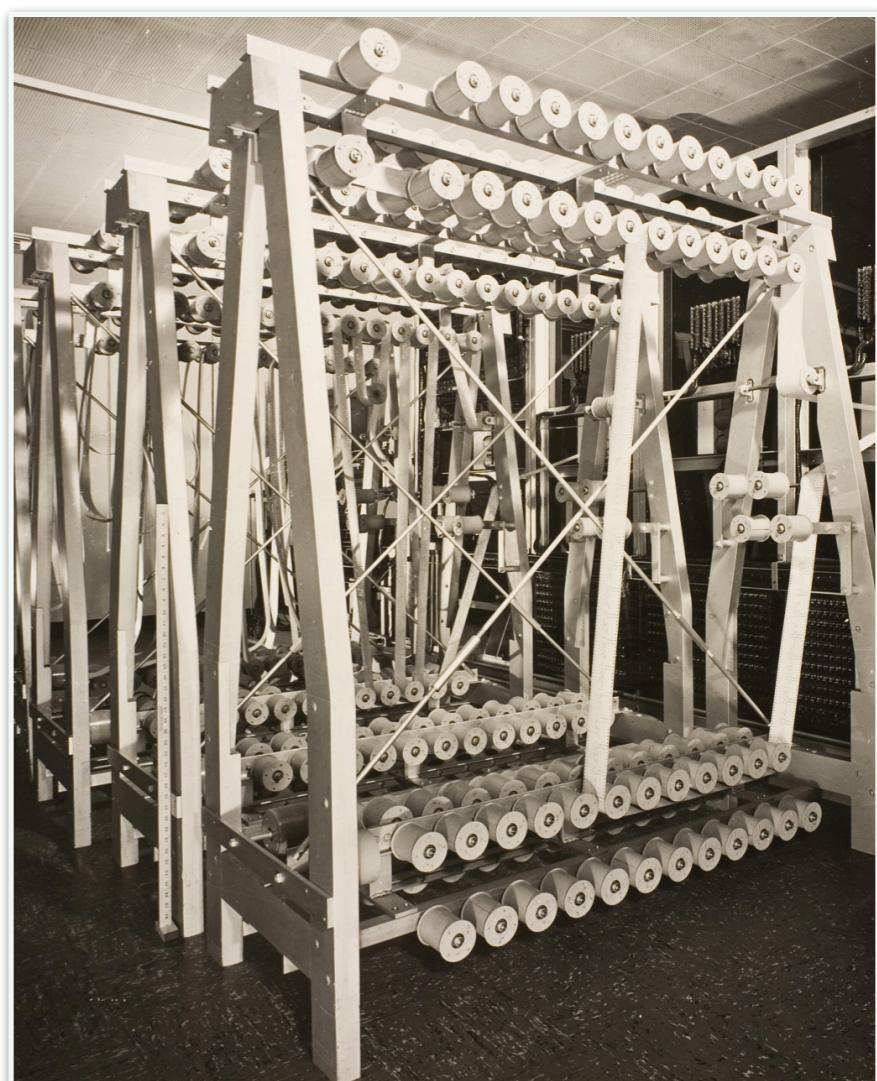
- ▶ A lot of common CS terms were coined by Mark I programmers
 - Documentary here: <https://chsi.harvard.edu/harvard-ibm-mark-1-video>



"Loop"



"Patch"

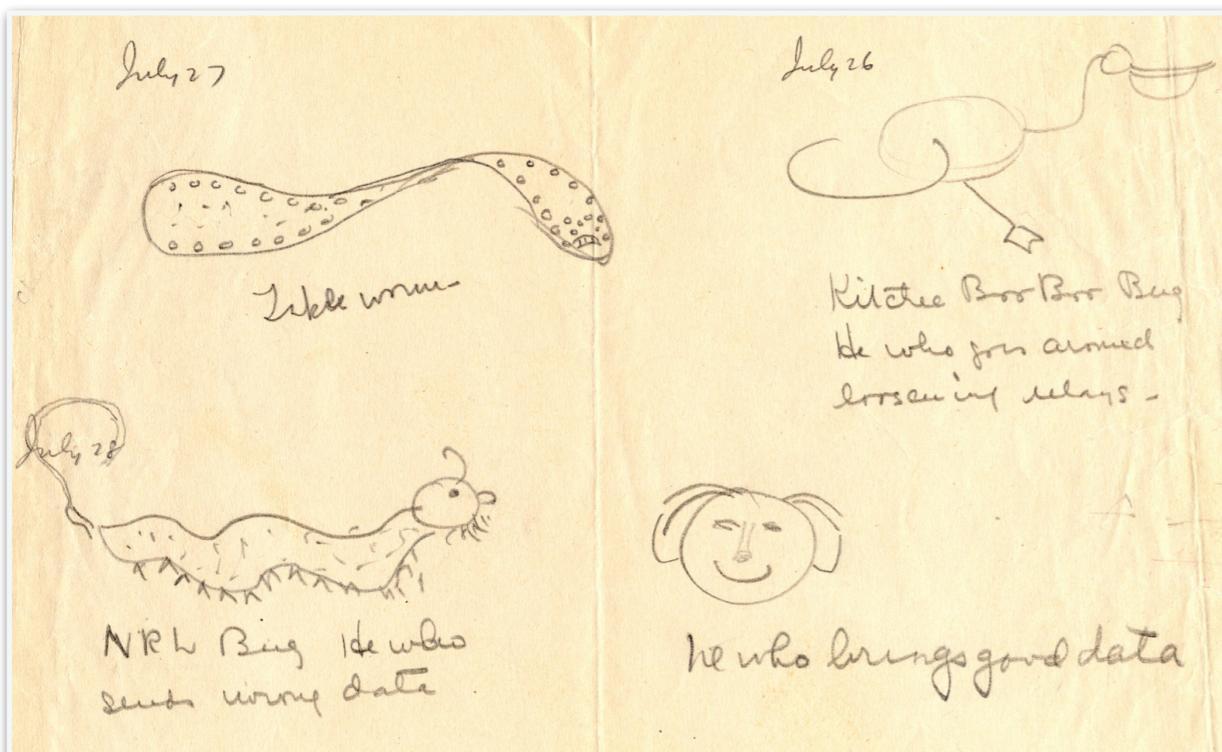


"Library"

Aside: Etymology of Common Terms

- ▶ The term "bug" had been used in the 1800s to describe flaws
 - Moth discovered in Harvard Mark II
 - "First actual case of **'bug'** being found." (1947)
 - Mark II was still electromechanical, but had special units for: sqrt, 1/x, ...

Drawings of bugs by Grace Hopper



"NRW Bug He who sends wrong data"

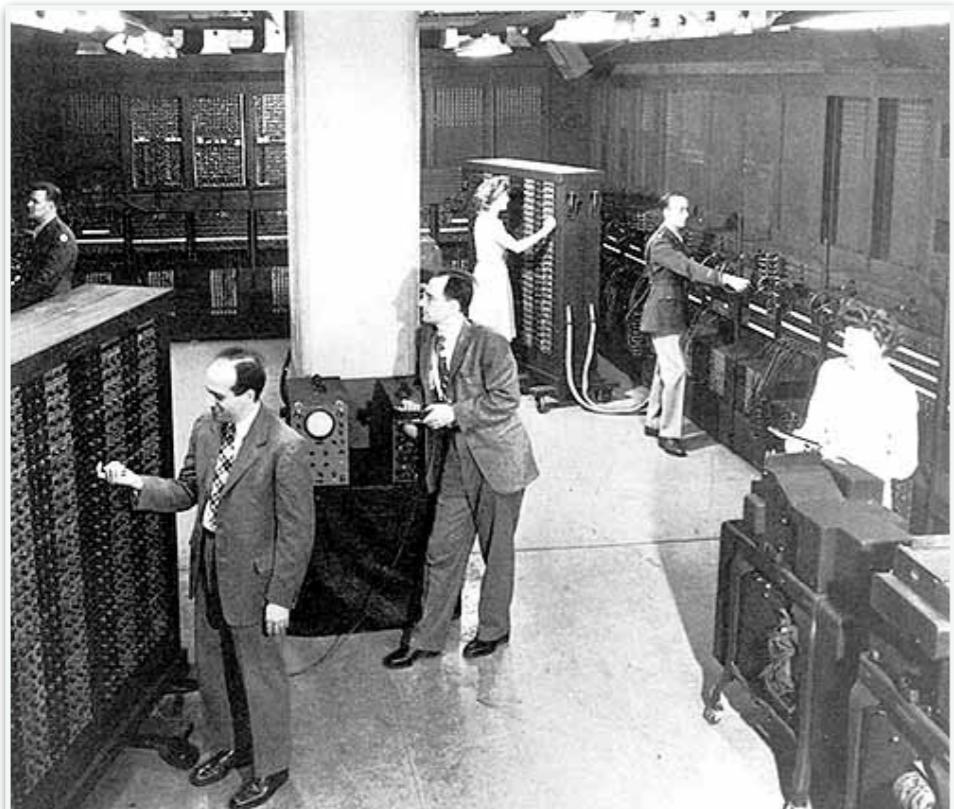
"He who brings good data"



"First actual case of bug being found"

1940s: ENIAC (1946)

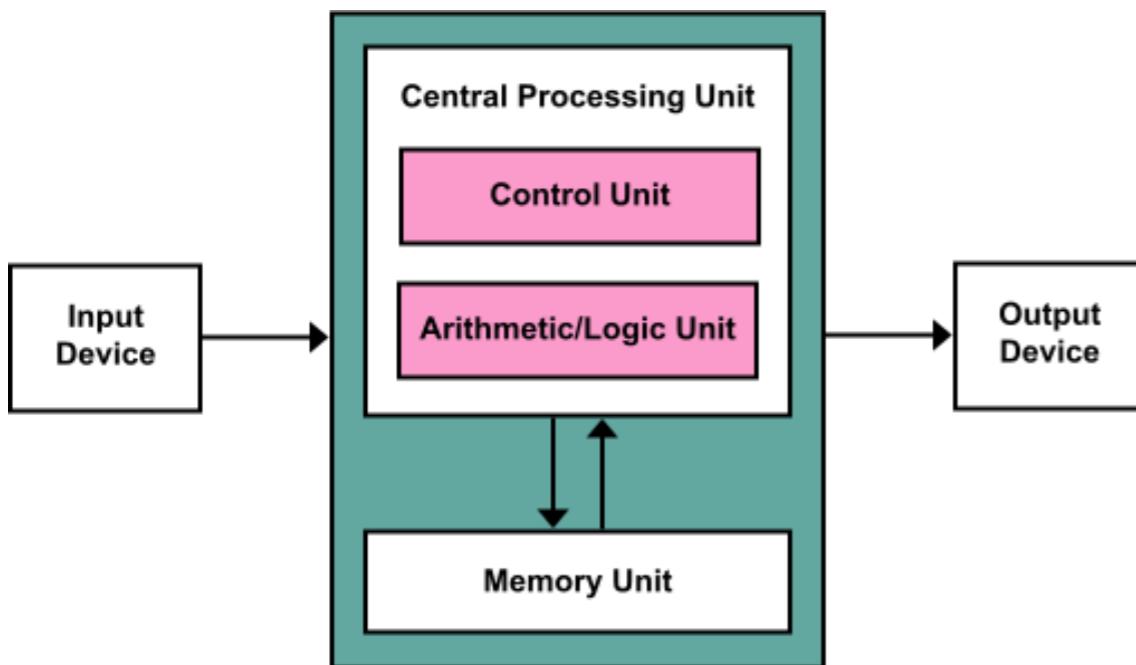
► ENIAC: Electronic Numerical Integrator and Computer



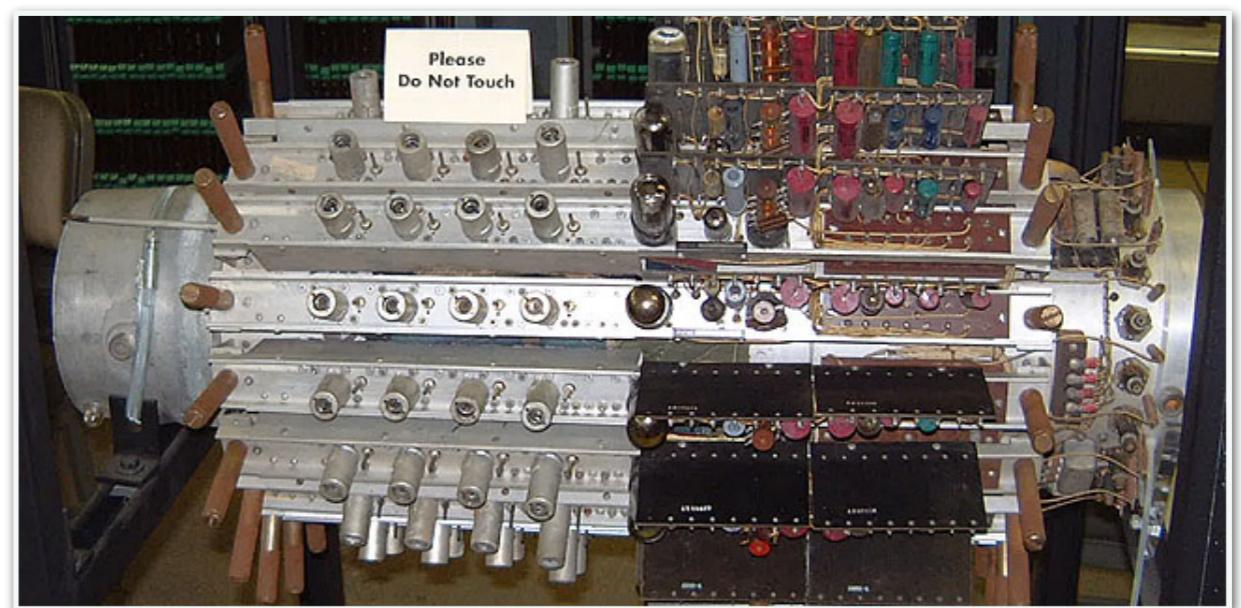
- First general-purpose fully electronic computer
 - Moore School of Engineering @ UPenn
 - J. Presper Eckert and John Mauchly
- Still no OS
 - More limiting than Mark I: Can't program ahead of time with tape/cards
 - Used switches and connections of cables (laborious) -- programmed "live" like Altair 8800

Mid-late '40s: EDVAC Changes Everything

- ▶ Mathematician John von Neumann teams up with Eckert and Mauchly
 - *Published: "First Draft of a Report on the EDVAC"*
 - The *Von Neumann model* of computing is born!
 - The *Stored-program model* (*Hmm.. how to "store" code/data in a machine?*)



The von Neumann model
(Still in use today)



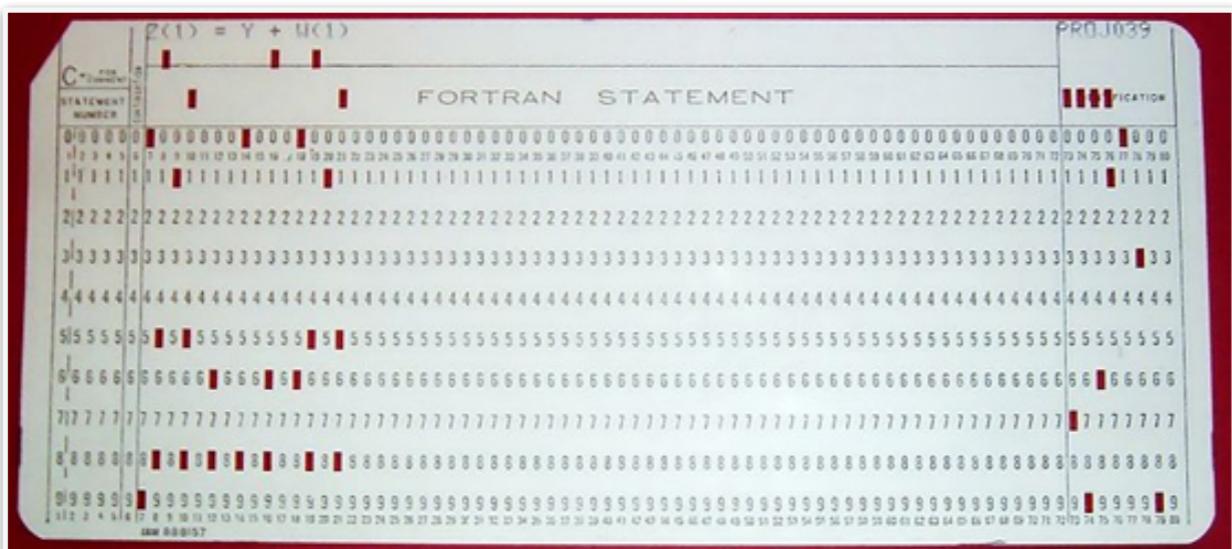
Mercury Delay-Line Storage
(https://www.youtube.com/watch?v=kignGE77I_I)

History Era 1: 1950-70's

- ▶ Hardware is expensive and humans are cheap!

- ## ► The OS begins to take form

- Stored programs now commonplace
 - Programs on punch cards or tape
 - Load program, run, print results, dump, repeat.
 - Removed humans from the loop, mostly



A-0 Simplifies Programming

- ▶ 1951-1952:
 - Eckert-Mauchly's UNIVAC becomes the first commercially-viable computer!
 - Go see a working UNIVAC @ Living Computer Museum in Seattle! (If it opens again)
- ▶ In 1952 Hopper writes the **first-ever compiler (A-0)** for the UNIVAC.
 - People belittled and resisted using the compiler for years!

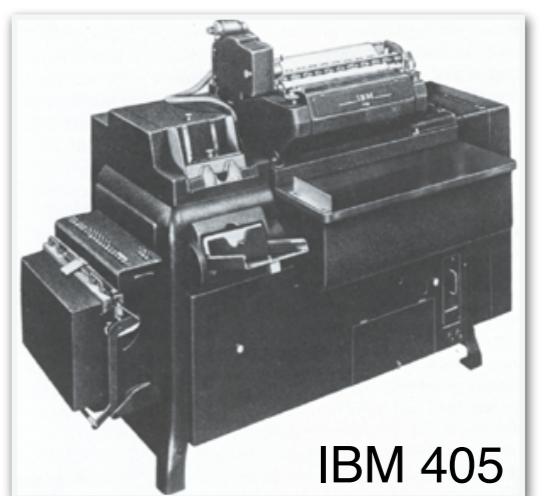
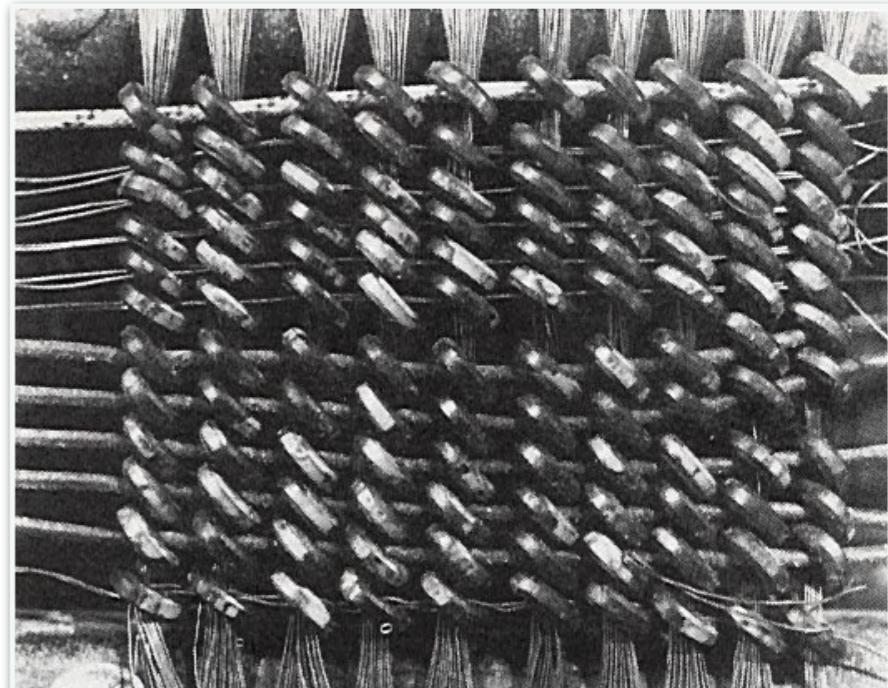
"I had a running compiler," said Hopper, "And nobody would touch it because, they *carefully* told me, computers could only do arithmetic; they could not do programs. It was a selling job to get people to try it. I think with any new idea, because people are allergic to change, you have to get out and sell the idea."

Grace Hopper: Navy Admiral and Computer Pioneer by Charlene W. Billings. 1989.



1950's: Solid-State Core Memory

- ▶ Race for solid-state storage
 - Also in 1952: The first magnetic *core memory* tested on the IBM 405 Accounting Machine
- ▶ Bits stored as magnetization in iron rings.
 - *Non-volatile!* Magnetic direction is persistent.
- ▶ A "Core Dump" still refers to the contents in memory (for investigation) after a process fails.



1950's (Batch Processing Model)

► 1952: IBM 701 Mainframe

- This machine transitions IBM from a tabulation company to the computer business.



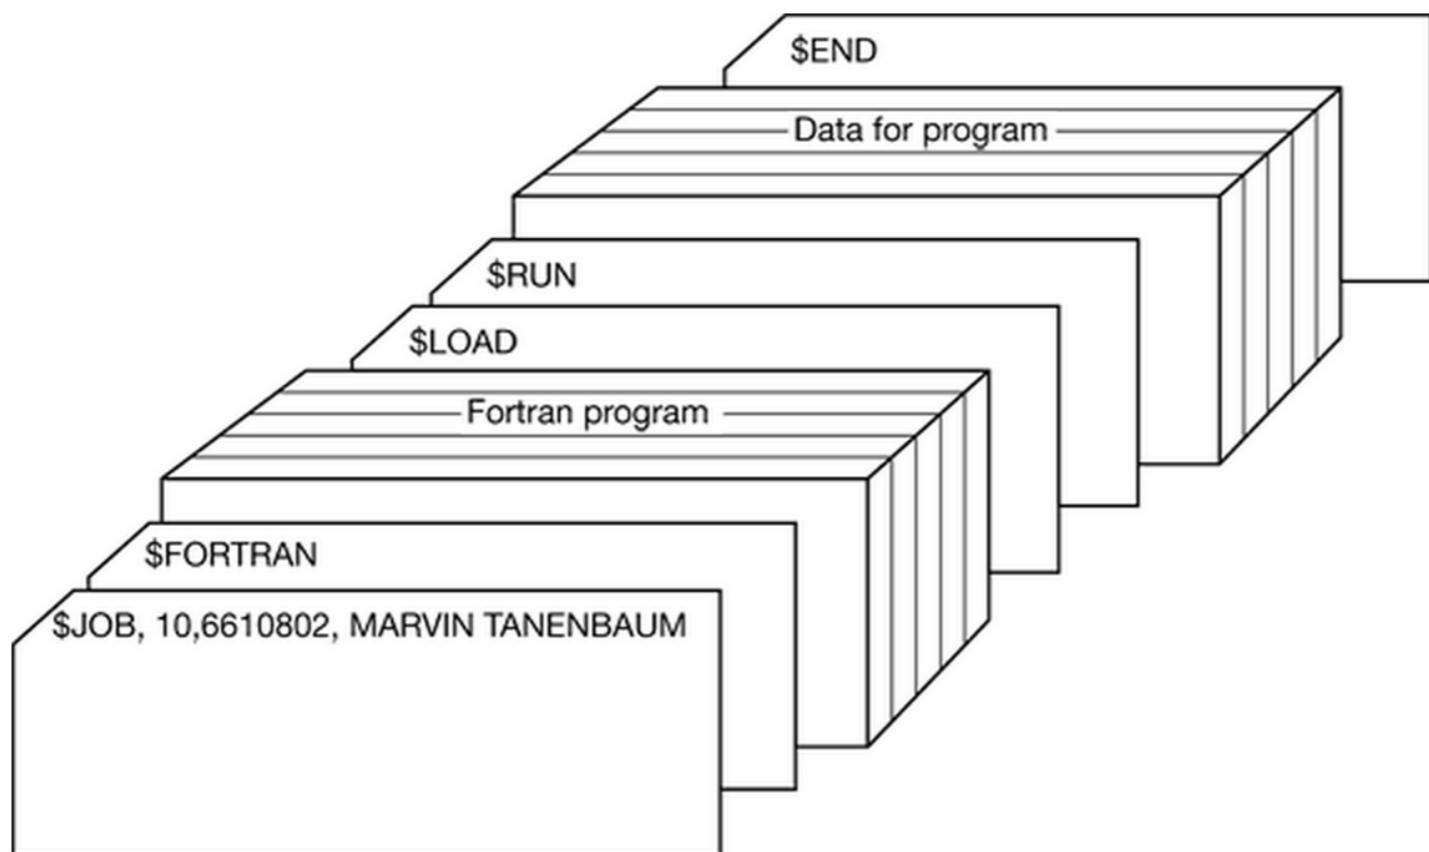
► Finally, the **First OS** called the "*resident monitor*"

- Established the *Batch Processing Model*
 - Load many programs to run.
 - Program cards separated by *control cards*.
 - Run one program (job) at a time.



1950s: Communicating with the First OS

- ▶ *Control cards* marked the beginning and end of *segments*:
 - \$JOB - First card of a job
 - \$END - Final card of a job
 - \$FORTRAN - run the fortran compiler
 - \$LOAD - run the loader
 - ...



IBM 711 Punchcard Reader

History Era 2: 1960's

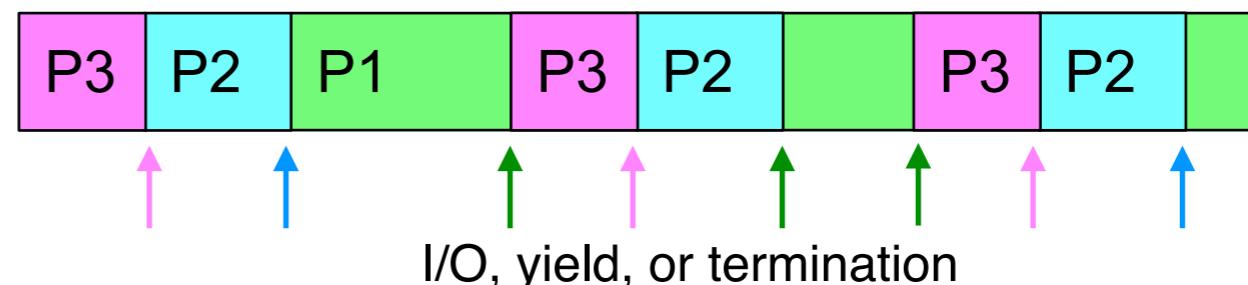
- ▶ Computers were adding many **slow** input/output (I/O) devices
 - Time is precious, CPU is expensive, I/O is slow
 - **Breakthrough:** context switching
- ▶ *Multiprogramming Model*
 - We want better CPU utilization over batch processing.
 - Make use of stored-programs: Several jobs loaded into memory at once
 - Run a program until it..
 - (1) Completes, (2) yields CPU to another program, or (3) uses I/O device
 - Put the current job back in memory, select a new job to run (*i.e.*, context switch)

1960's: Multiprogramming Is Born

► *Multiprogramming Model*

- Run a job until it either:
 - Completes, *or* yields the CPU to another program (play nice), or uses device I/O
 - Put the current job back into ready queue, and select a new job to run
 - (This mechanism is called as a "*context switch*")

When multiprogramming is good



**When multiprogramming doesn't help.
(same as batch processing)**

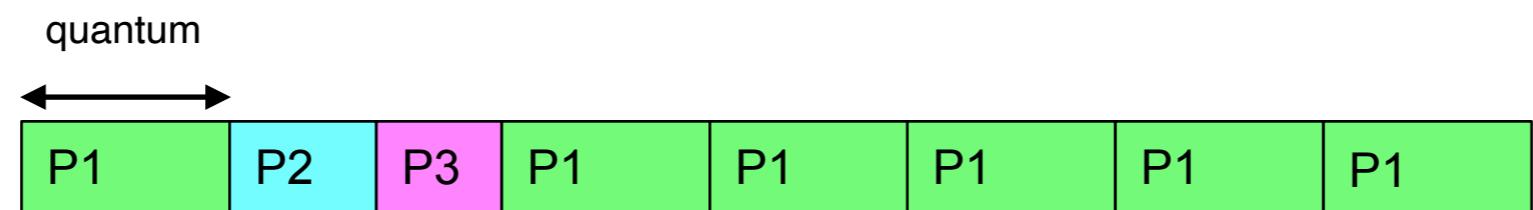


P1 never does I/O

P2 and P3 wait and suffer

1960's: The Need for Timesharing

- ▶ Multiprogramming gets complicated
 - CPU utilization increased over batch processing, but...
 - Really long jobs (that don't require I/O) can *still* monopolize computer
- ▶ *Timesharing Model:* The OS repeatedly cycles through all jobs
 - Timesharing is enhanced multiprogramming.
 - It's multiprogramming plus the following:
 - Each job is given a fixed amount of CPU time (called a time slice, or a quantum)
 - If job doesn't finish within the quantum, then it gets switched out

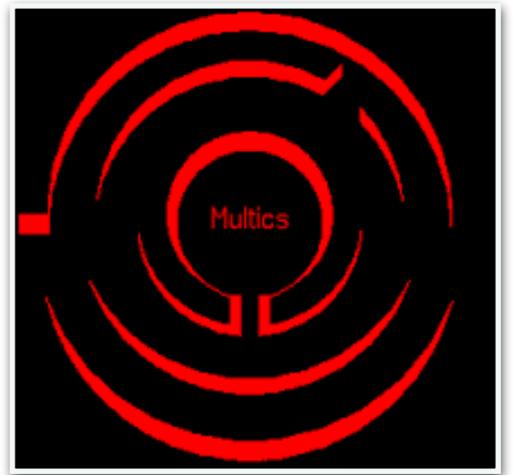


1960's: Timesharing is Hard

- ▶ Complexity gets out of hand!

- How to ensure security between running-jobs?

- Program A should not be able to access Program B's memory
 - User programs should not be able to access OS's memory



- ▶ **Multics: Multiplexed Information and Computing Service (MIT, 1969)**

- Implementing *timesharing* was the main objective
 - Formalization of a process, address space, virtual memory, shared libraries
 - Multics ultimately failed: "*Too big and slow ... an overdesigned behemoth*"
 - But concepts developed were highly influential and stood the test of time.

1970: The Success of UNIX

- ▶ UNIX operating system is developed in 1970
 - ("Unics" is a play on "Multics")
 - Developed in Bell Labs by Ken Thompson (a former Multics programmer)



Ken Thompson

- ▶ UNIX rewritten in Dennis Ritchie's new language C
 - Established C to be the *de facto* systems language
 - Lasting contribution: UNIX is portable!
 - Supported a wide-array of underlying hardware
- ▶ Co-Recipients of the ACM Turing Award in 1983



Dennis Ritchie

History Era 3 (1970's - 1980's)

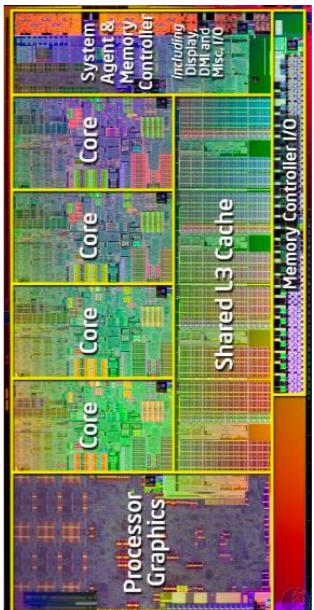
- ▶ Hardware very cheap, humans very expensive
 - Computers cost \$10,000s, but programmers now cost \$100K/yr
 - Cultural shift: Give people computers to make them more productive!
- ▶ Advocacy of Personal Computing (PC)
 - OS needed better user-friendliness and GUIs
 - A return to simplicity (due to 1-to-1 usage model)
 - Back to one program running at-a-time: MS-DOS
 - Followed by resurgence of complexity
 - People can multi-task: 1 user, many programs running



Apple Macintosh (1984)

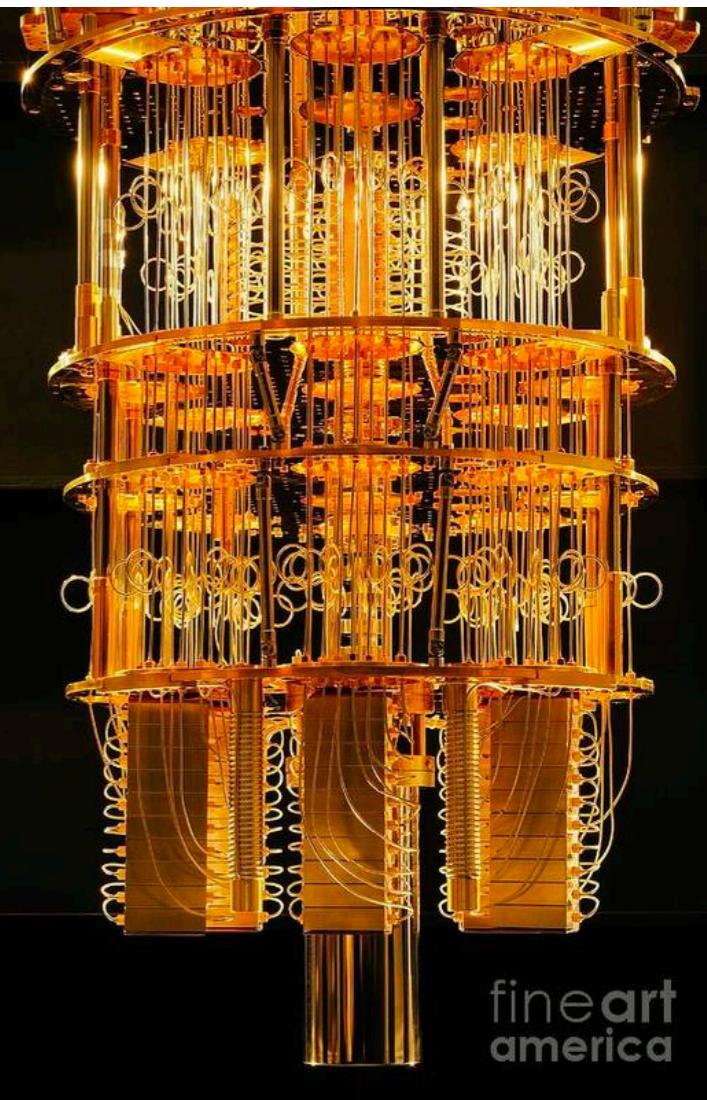
History Era 4 (Now)

- ▶ Hardware is very cheap and ubiquitous (many devices per person)
- ▶ Multicore CPUs (~ 2001)
 - Monolithic CPU approaching "*power wall*"
 - Slow them down and add more cores per CPU!
 - New problem: How to exploit parallelism?
- ▶ Mobile computing (~ 2010)
 - Power efficiency now the chief concern
 - Lack of storage
 - OS must have a small "footprint". Back to batch processing again



What's Next? Quantum Computing (Probably)

- ▶ We're entering a new era of computing!
 - The dawn of something revolutionary (Again!)
- ▶ *Quantum Computers*
 - Experimental and very expensive
 - D-WAVE quantum computer costs \$10M - \$15M
- ▶ *Qubits* harness quantum superpositioning
 - A *qubit* can represent a 0 and 1 simultaneously
 - n qubits means 2^n possibilities may be explored simultaneously, but verification of results is tricky.



fineart
america

Administrivia 1/17

► Welcome to the first day of OS!

- Reading
 - Chap 1 of Dinosaur Book
 - Chap 1-3 of Dive into Systems (on C)
- Hwk 1 due 11:59p Friday!
 - Your username is your puget sound login
 - Your password is **qwe123**