

CS 475

Operating Systems

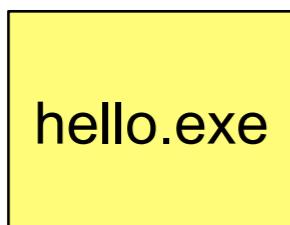


Department of Mathematics
and Computer Science

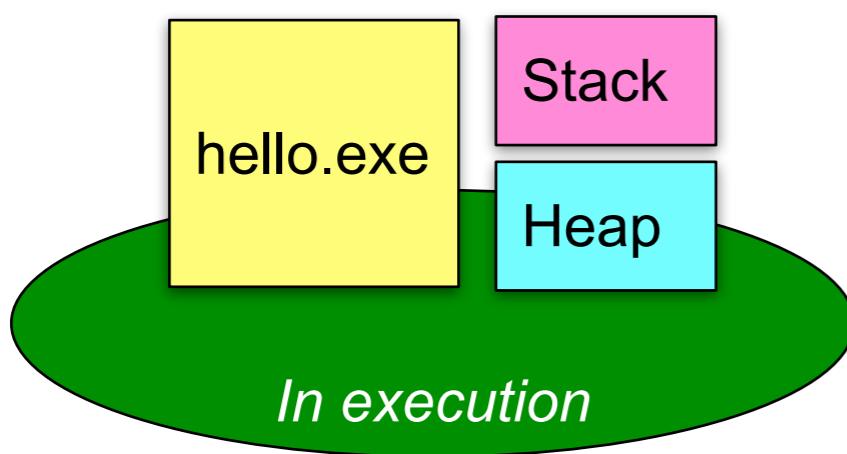
Lecture 3
Process Management

Definition: Process vs. Program

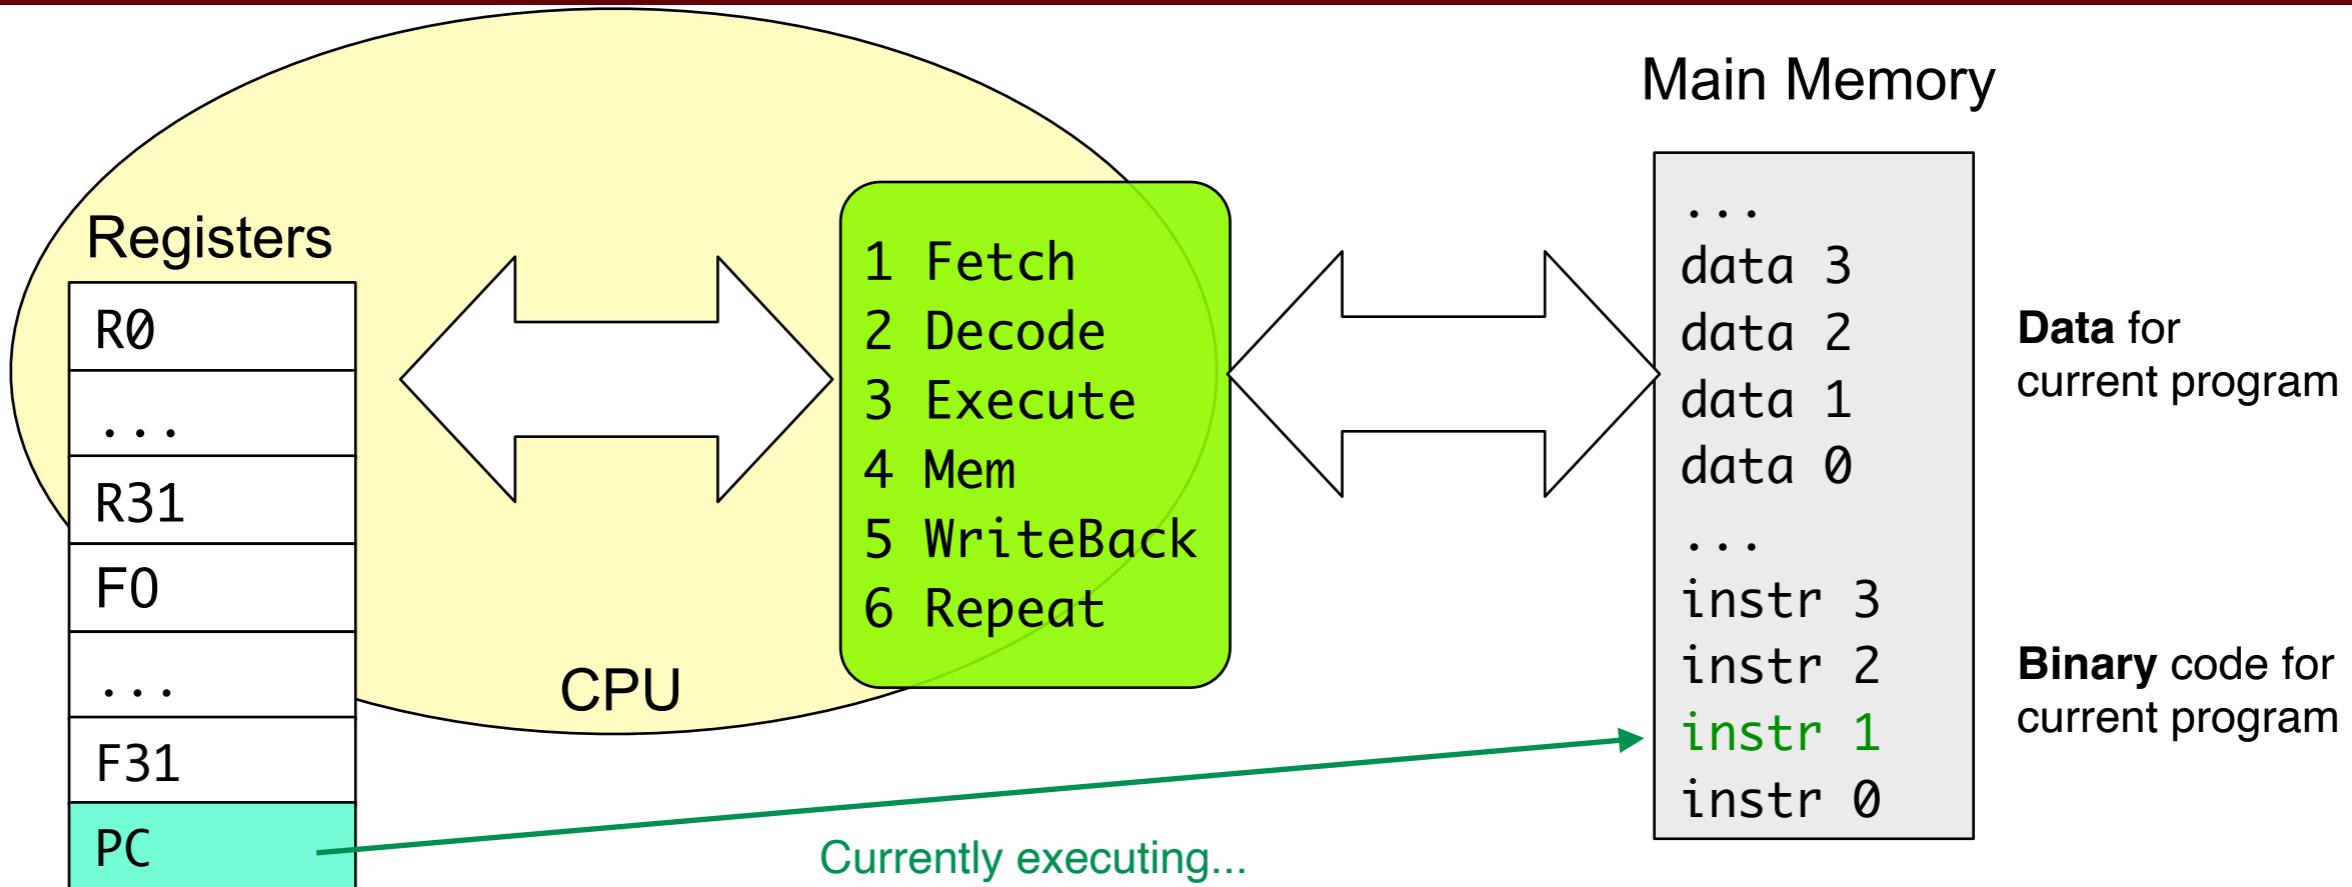
- ▶ A *program* is a *passive* entity:
 - A compiled executable (binary) file stored somewhere on disk



- ▶ A *process* is an *active* entity:
 - A process is a program that is currently running
 - It is loaded into memory and executing its instructions (or waiting to)



CPU Execution Cycle



► CPU execution sequence:

1. Fetch Instruction at address referenced in PC register ([program counter](#))
2. Decode Instruction
3. Execute (possibly using registers for intermediate storage)
4. Write results (into registers or memory)
5. $\text{PC} = \text{NextInstruction}(\text{PC})$
6. Repeat

Goals for This Lecture

- ▶ Motivation: Concurrency
- ▶ Address Space
- ▶ Processes and the PCB
- ▶ Process Management in C
 - Creation: fork()
 - Basic Synchronization: wait()
 - Running another program: exec()

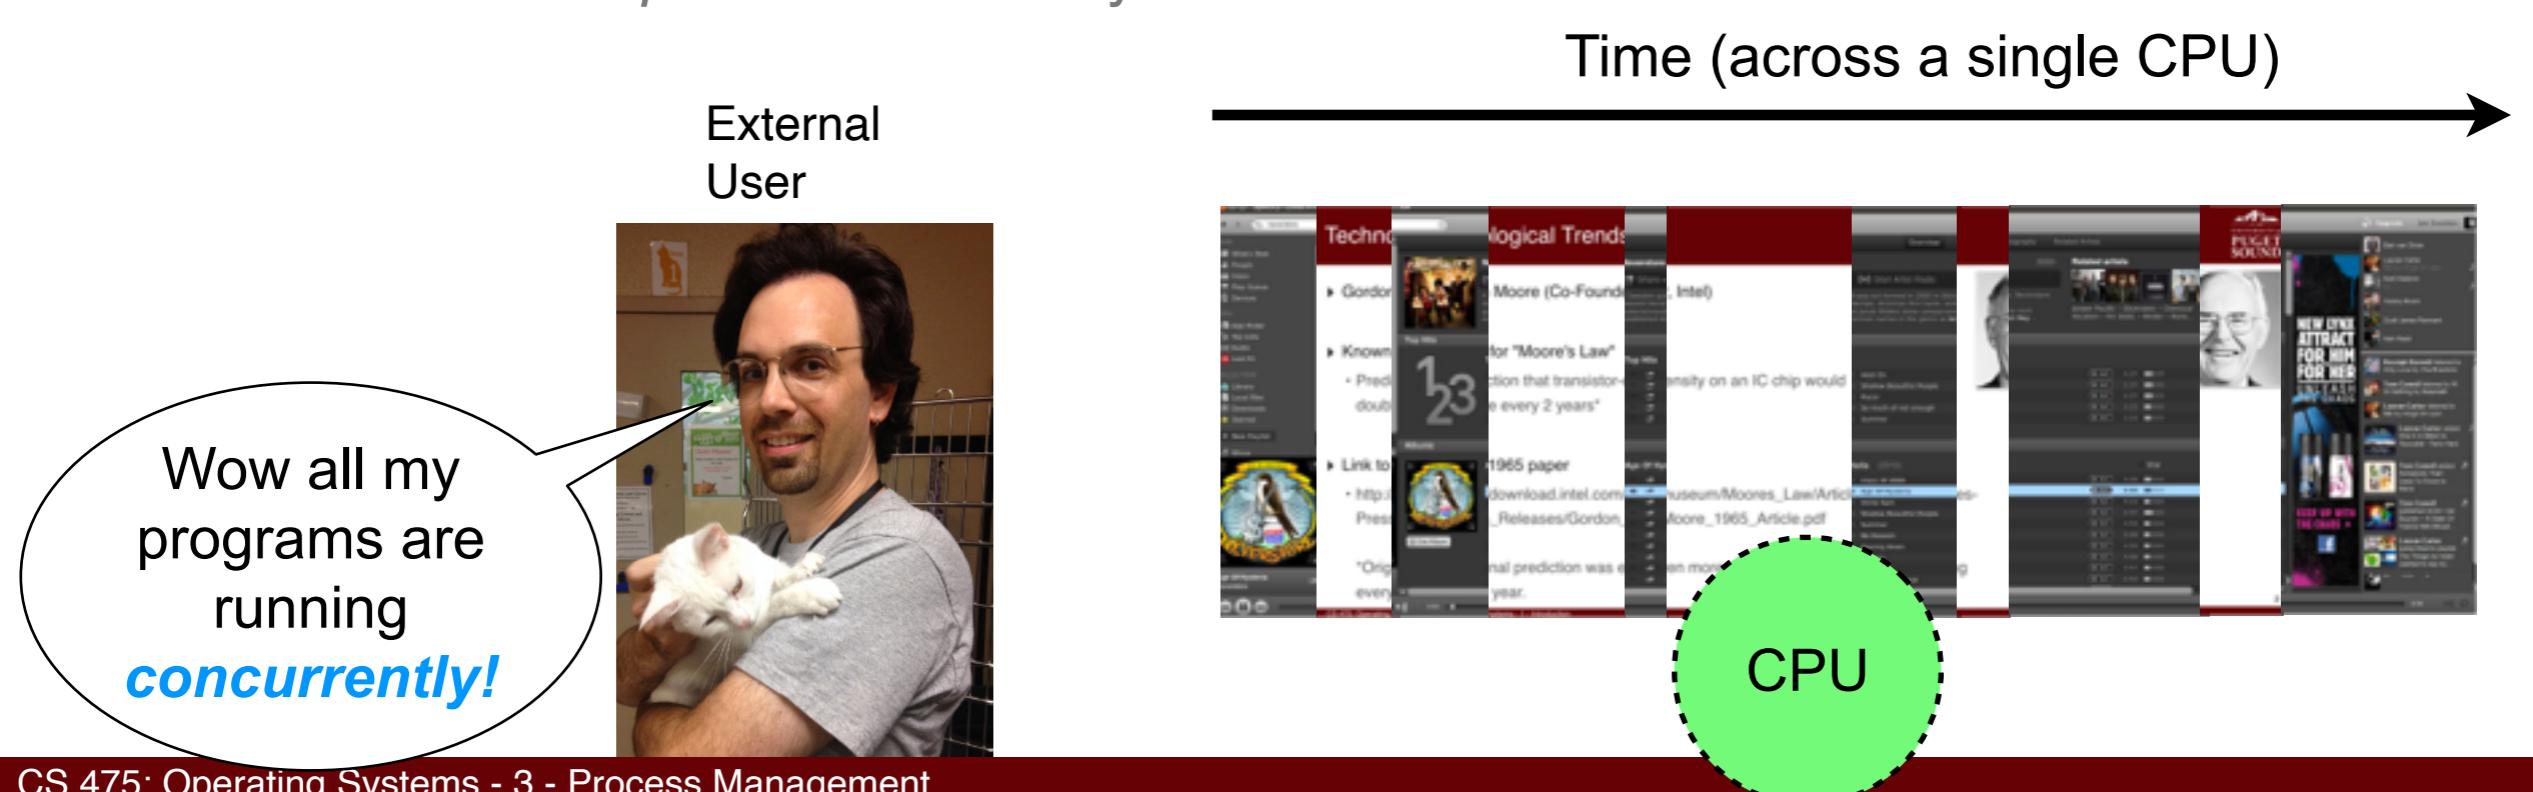
"Job Concurrency Increases CPU Utilization"

- ▶ In the 1950's, batch processing ruled the computing world!!!!
 - A computer runs **one job** at a time
 - IBM 701, MS-DOS, early Macintosh
 - Great for OS developers
 - Simplified OS design if we didn't have to worry about multiple programs running!
- ▶ 1960's-70's: Timesharing to Make User More Efficient
 - Humans are becoming more expensive. Keep them busy by multitasking!
 - Execute > 1 job at a time
 - *But 1 CPU meant that only 1 job can be really be running at any instant*
 - *(Can we fake it?)*

Definition: Concurrency vs. Parallelism

► Concurrency vs. Parallelism

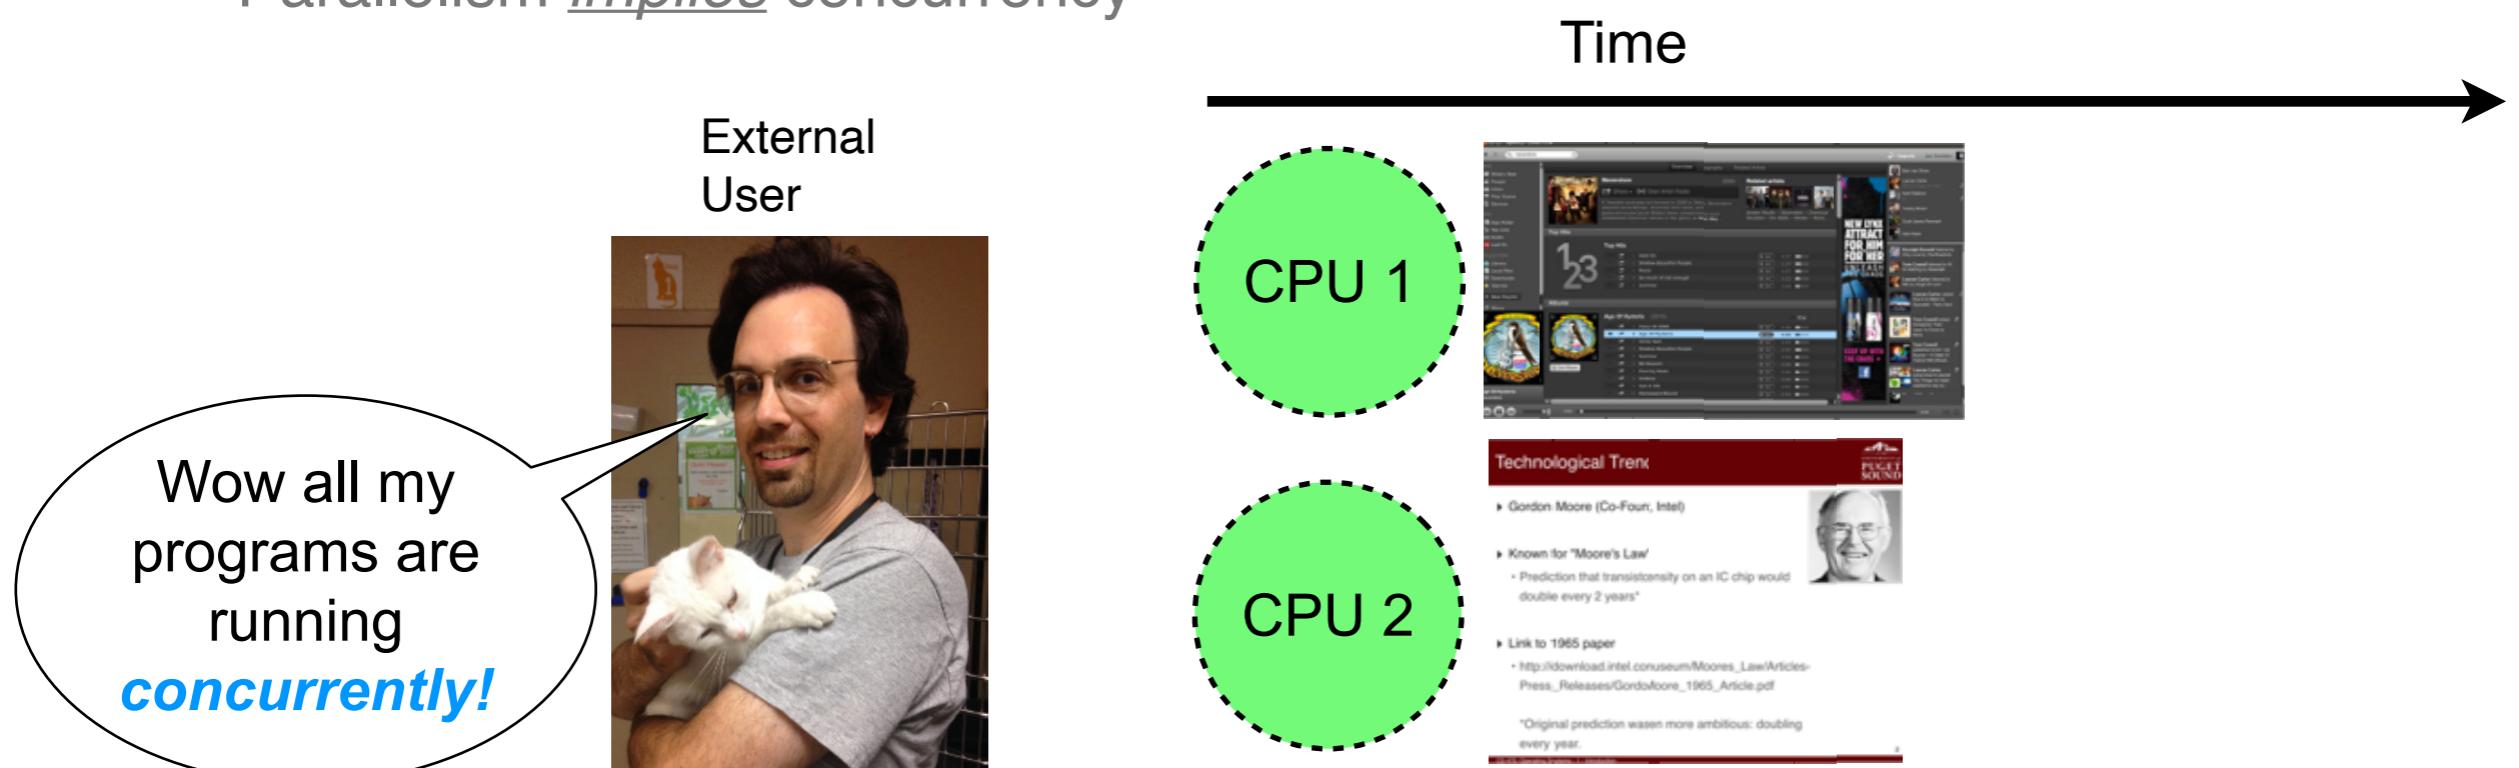
- **Concurrency** is *giving the perception* of simultaneous job execution
 - You can have one CPU multiplexing between different programs
 - Or, you can have multiple CPUs all running different programs
- **Parallelism** is *actually performing* simultaneous job execution
 - You can have multiple CPUs all running different programs
 - Parallelism *implies* concurrency



Definition: Concurrency vs. Parallelism

► Concurrency vs. Parallelism

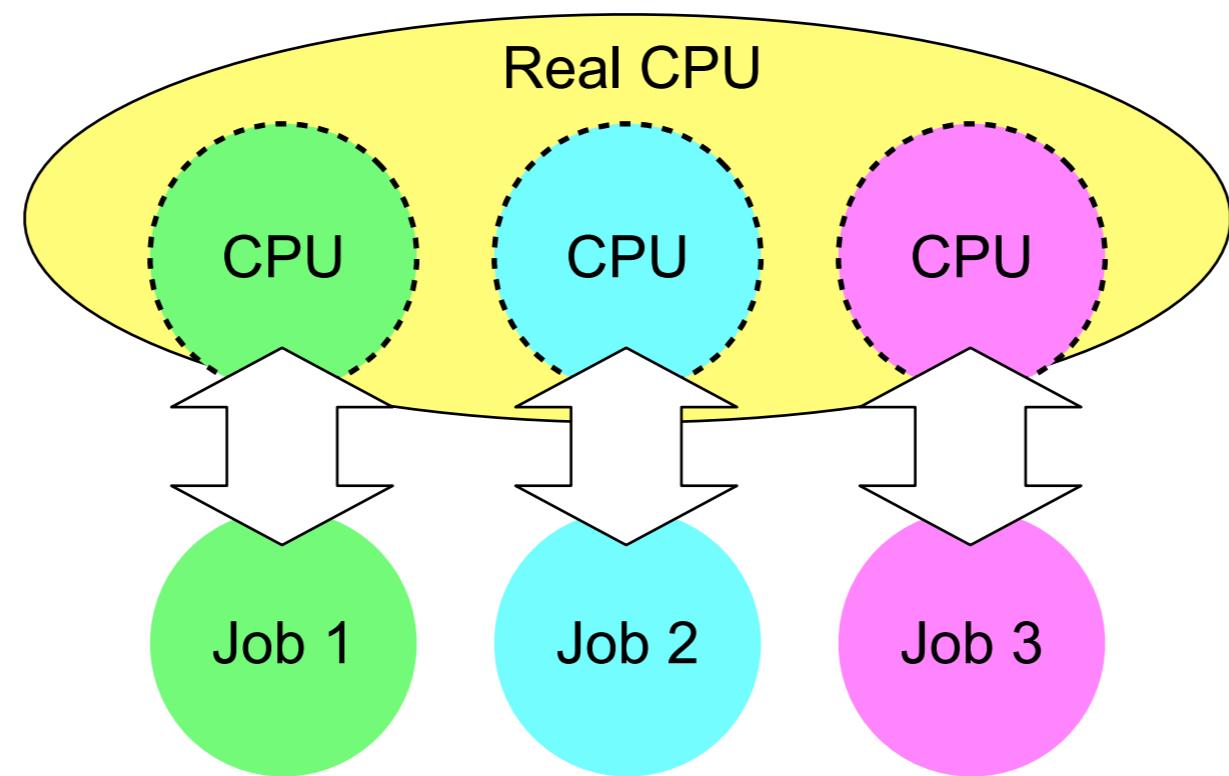
- **Concurrency** is *giving the perception* of simultaneous job execution
 - You can have one CPU multiplexing between different programs
 - Or, you can have multiple CPUs all running different programs
- **Parallelism** is *actually performing* simultaneous job execution
 - You can have multiple CPUs all running different programs
 - Parallelism *implies* concurrency



Policy: CPU Virtualization

► *Policy: We want to concurrently execute multiple jobs on 1 CPU*

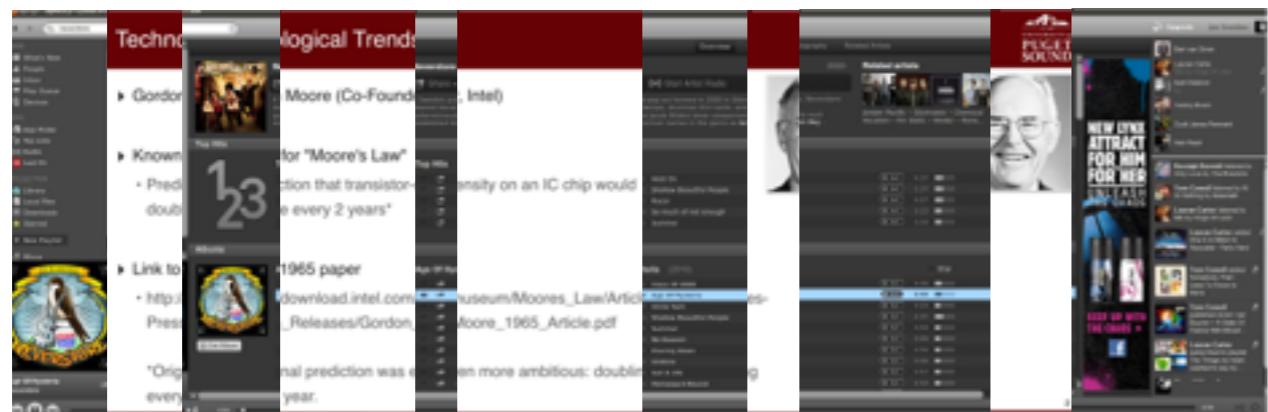
- This policy is called "*CPU Virtualization*"
- We want each job to think they have ownership of its own CPU
 - (Can you imagine the alternative? You'd have to request allocation of some usage time before you can run your program!)



Mechanism? Context Switching

► Mechanism: Context Switching

- **Save** the *current state* of a job to main memory 
- **Load** the *state* of another job state from main memory 



→

Time (across a single CPU)

- To pause and play, we need to a way to store the state of any process.

Process Control Block (PCB)

- ▶ Think games: What do we need in a "*save file*" for each process?

PCB

- ▶ The OS represents each process in a *Process Control Block (PCB)*



What's in this struct
???

- ▶ We need to capture the snapshot of a process at arbitrary points in time

- *We'll use a struct*
- *What info might a PCB struct need to store?*

Process Control Block (PCB)

► What's in a PCB Struct?

PCB Struct (One per process)

- pid (unique ID in the OS)
- uid (who created this proc?)
- Execution state (ready? running? done?)
- CPU register contents
- List of open files
- List of open sockets

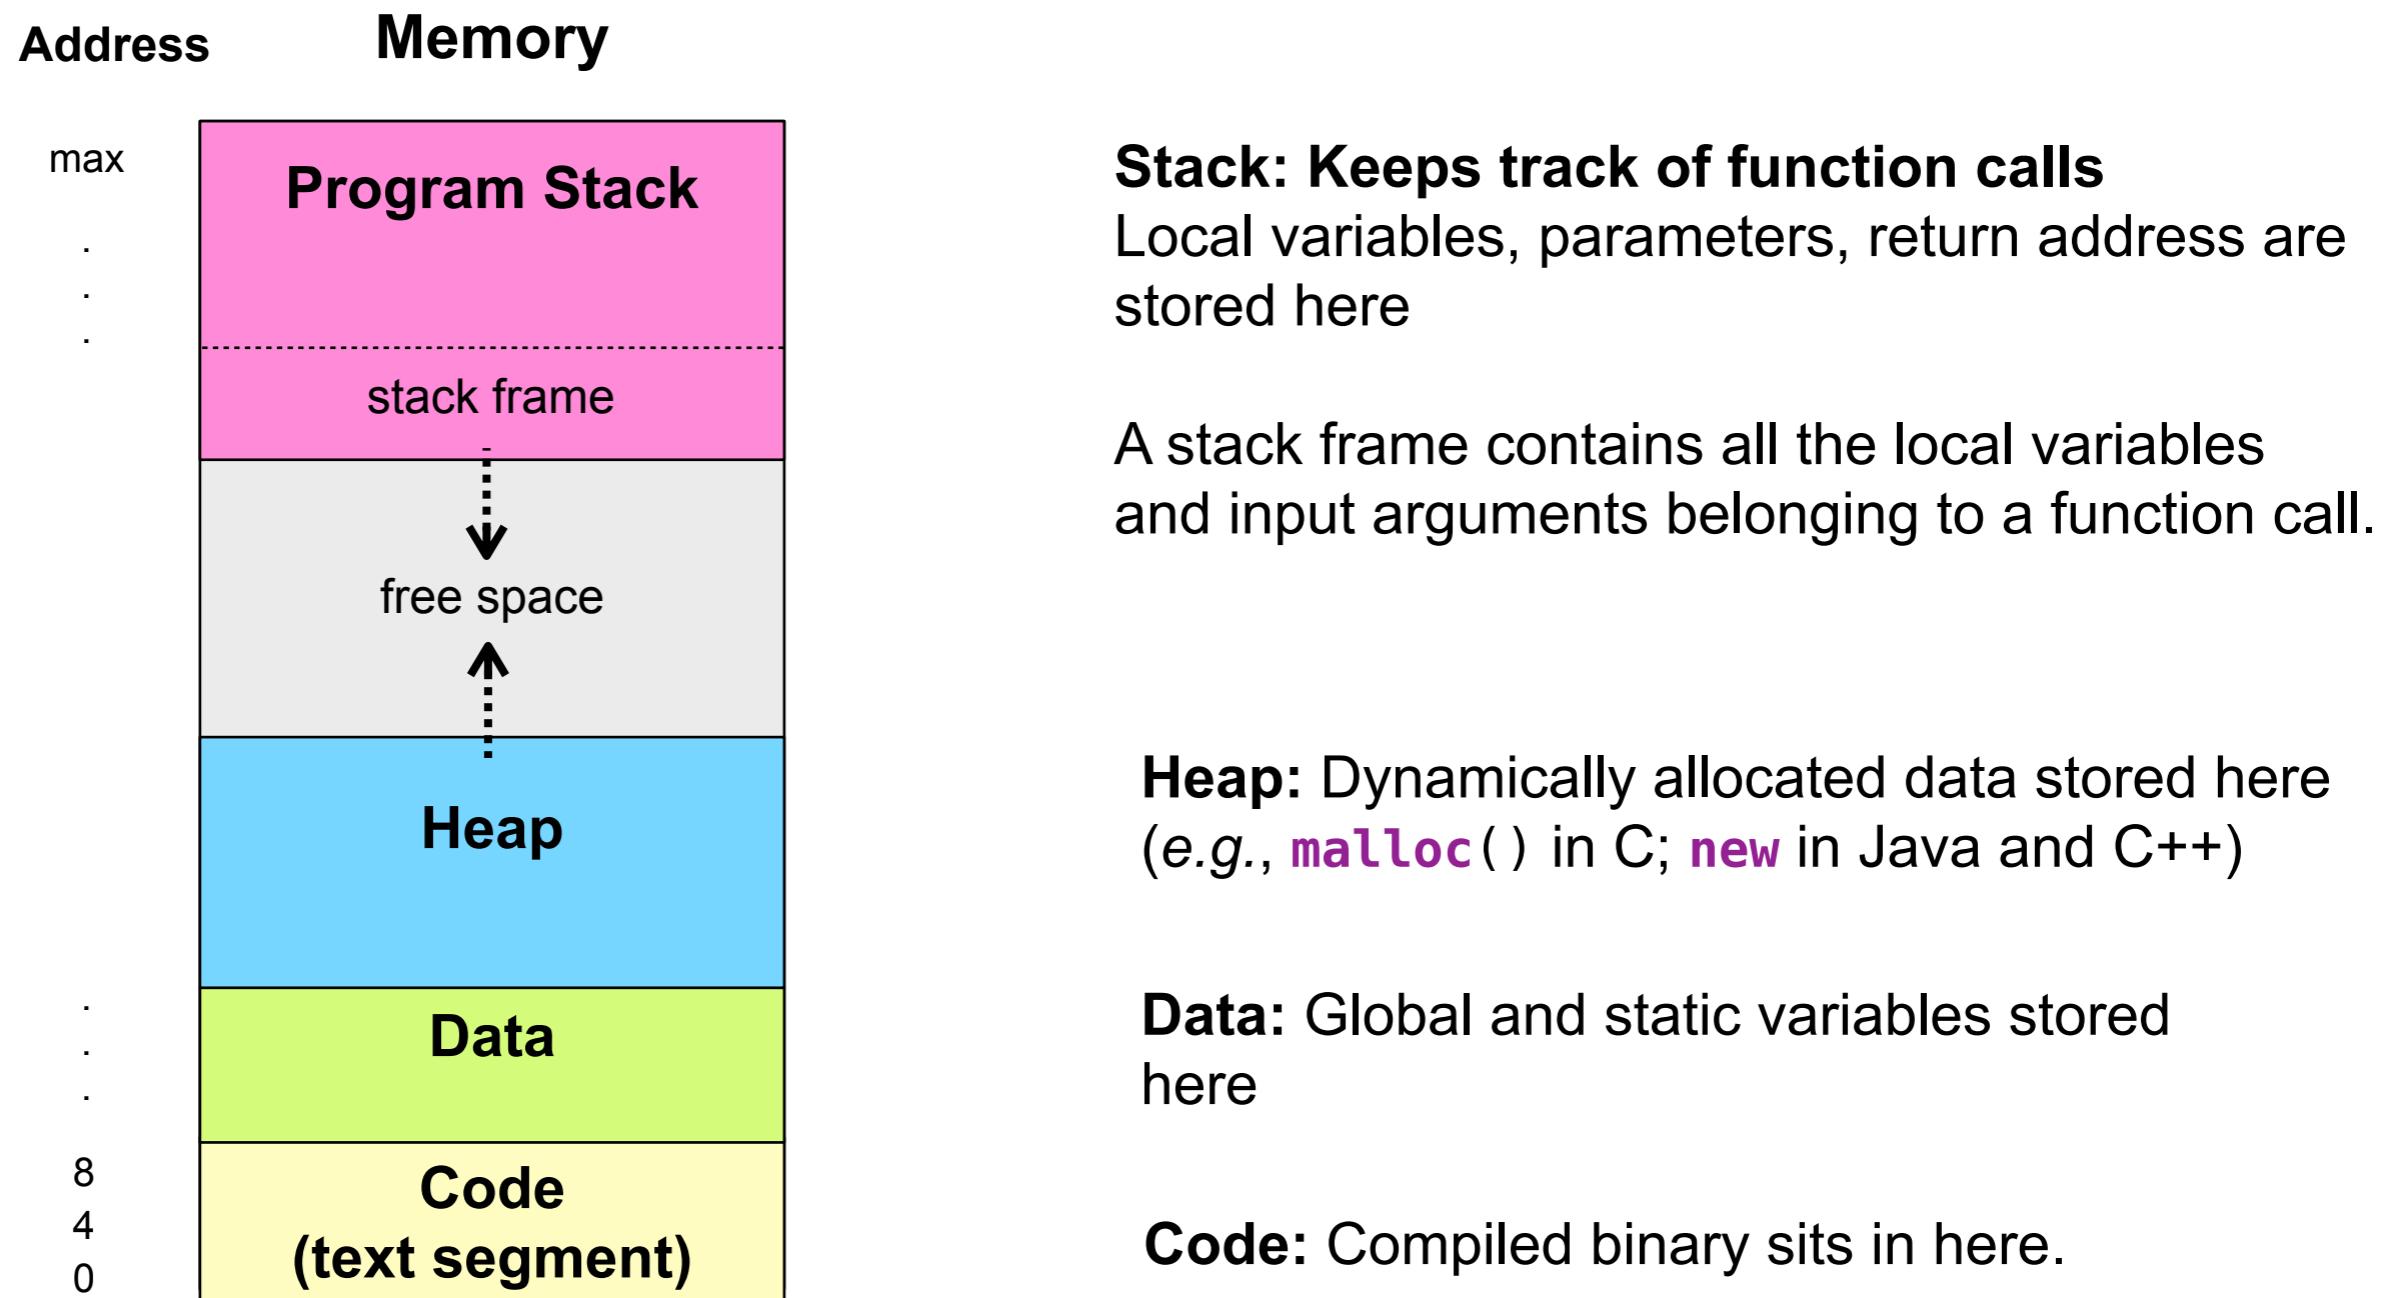
Address Space Pointers ?????

- Base Pointer (bp) ?????
- Stack Pointer (sp) ?????
- Heap Pointer (brk) ?????
- Program Counter (pc) ?????

Goals for This Lecture

- ▶ Motivation: Concurrency
- ▶ Address Space
- ▶ Processes and the PCB
- ▶ Process Management in C
 - Creation: fork()
 - Basic Synchronization: wait()
 - Running another program: exec()

Address Space (One per Process!)



Assumes a 32-bit architecture, so each CPU word is 4 bytes (= 32bits)

Process Control Block (PCB)

► What's in a PCB Struct?

PCB Struct (One per process)

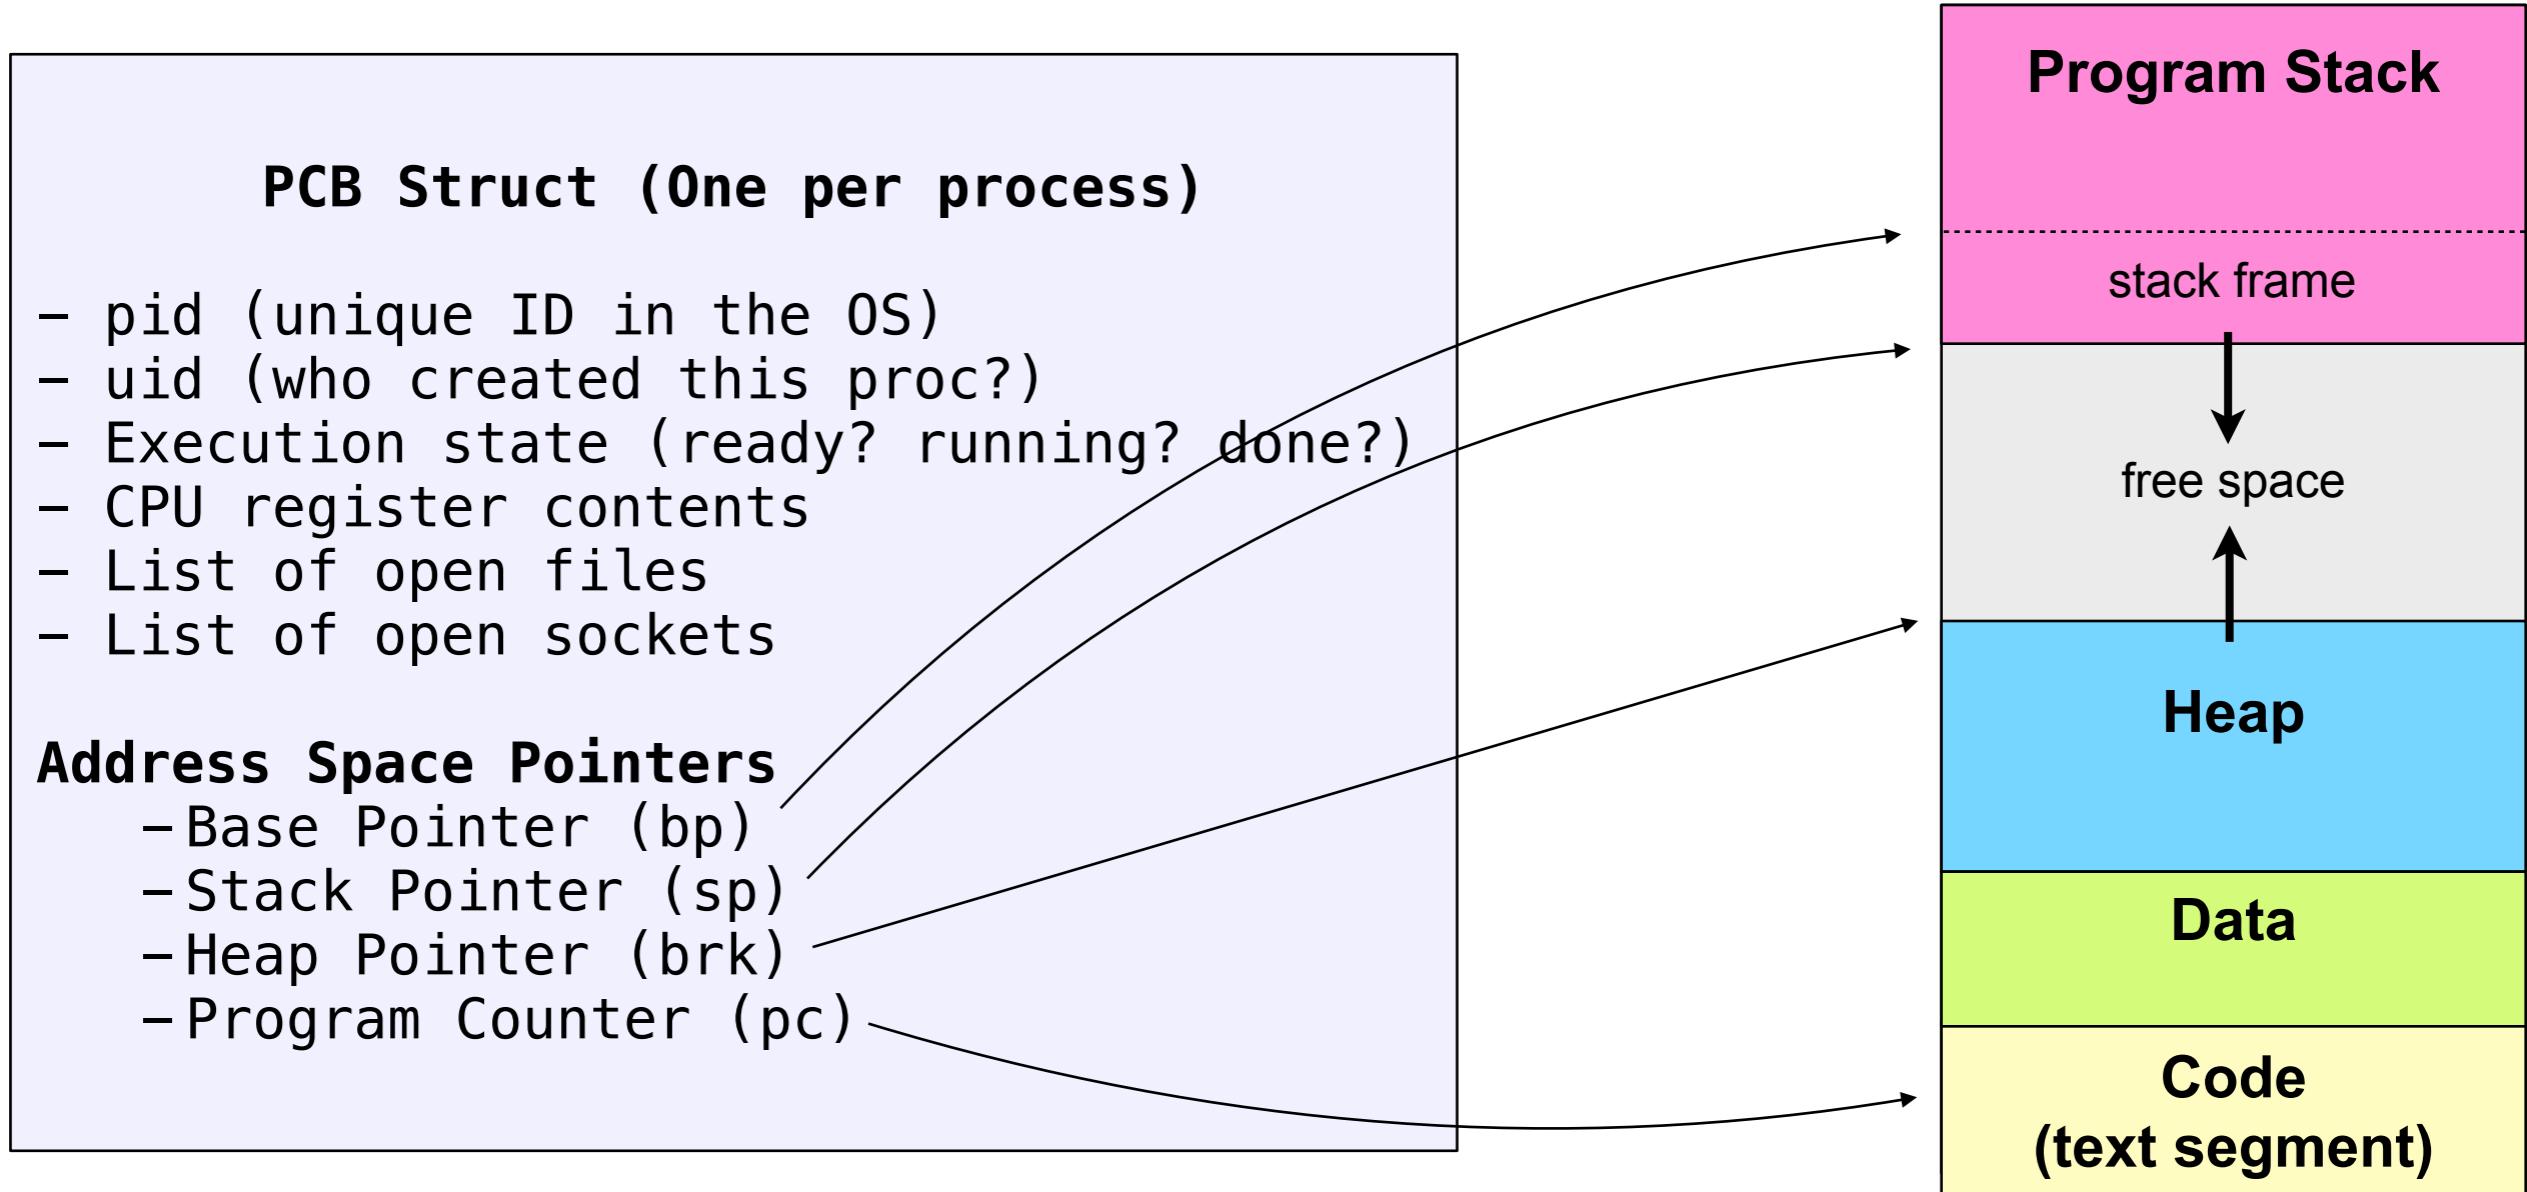
- pid (unique ID in the OS)
- uid (who created this proc?)
- Execution state (ready? running? done?)
- CPU register contents
- List of open files
- List of open sockets

Address Space Pointers ?????

- Base Pointer (bp) ?????
- Stack Pointer (sp) ?????
- Heap Pointer (brk) ?????
- Program Counter (pc) ?????

Process Control Block (PCB)

► What's in a PCB Struct?



Example: Intel x86 Architecture (32-Bit)

- ▶ x86 register names (32-bit architecture)

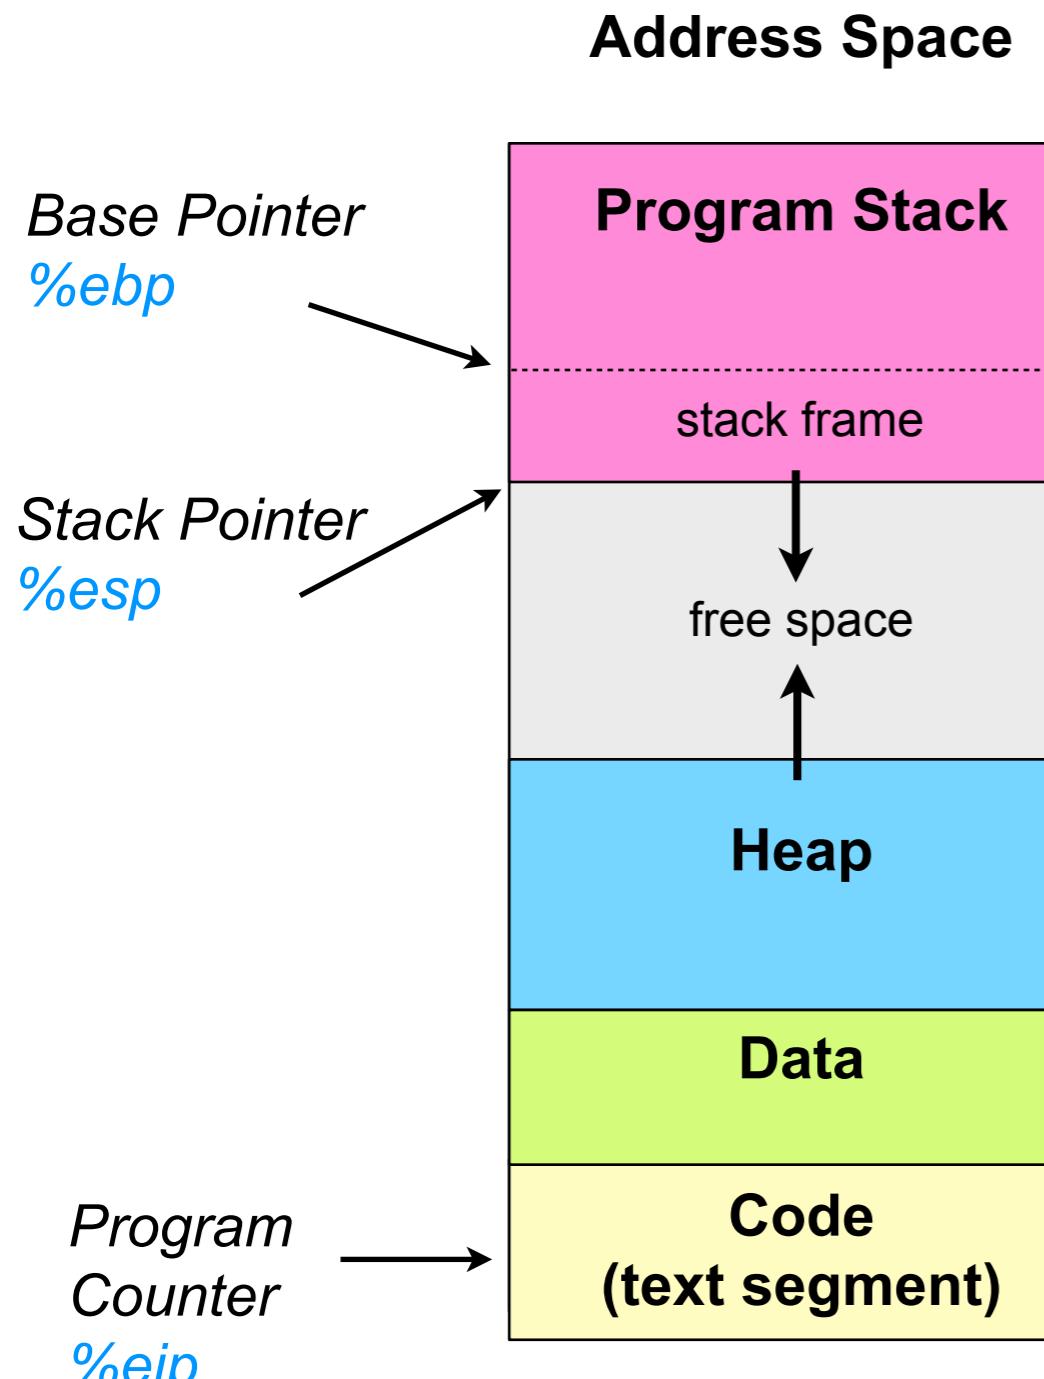
General Purpose Registers

```
%eax (usually holds a function's return value)  
%ecx  
%edx  
%ebx  
%esi  
%edi  
%esp (pointer to top of stack)  
%ebp (pointer to stack frame base address)
```

Special Purpose

```
%eip (instruction pointer or Program Counter)
```

Stack Management: x86 and C Example



```

int main() {
    int a = 5, b = 10;
    sum(a,b);
}

int sum(int x, int y) {
    int z = x + y;
    return z;
}
  
```

In code.c

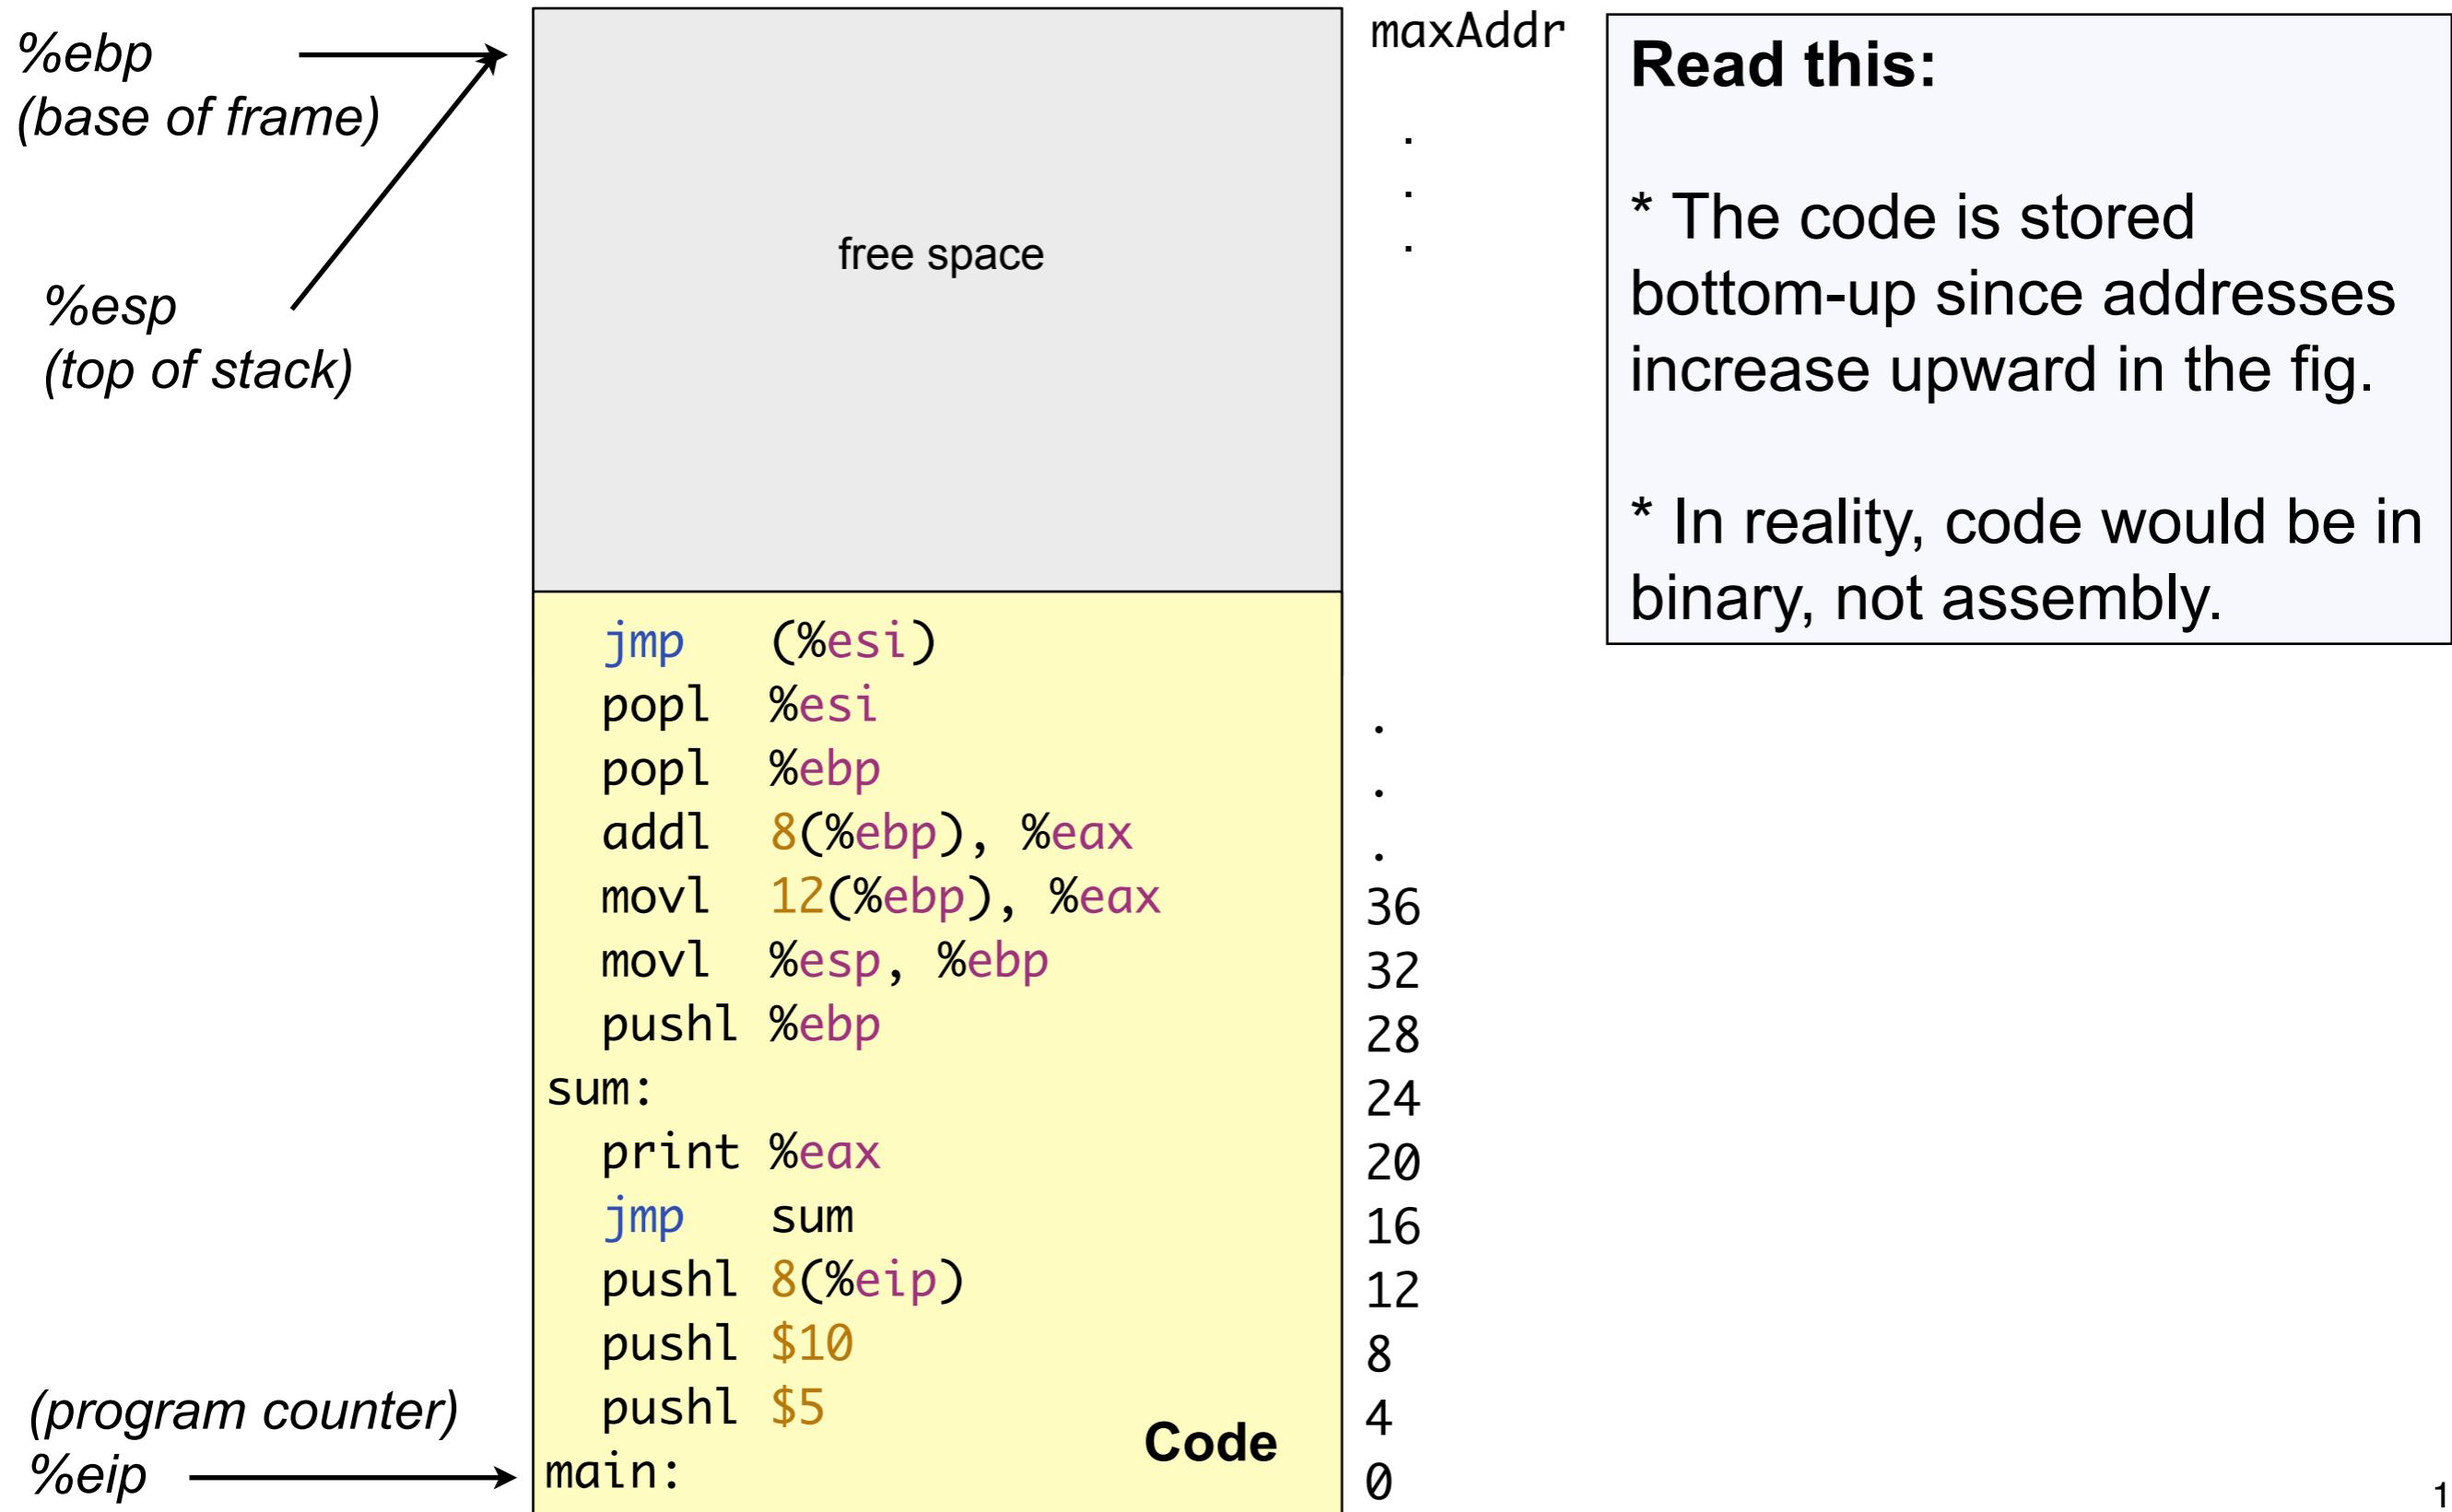
```

main:
pushl $5
pushl $10
pushl 8(%eip)
jmp  sum
print %eax

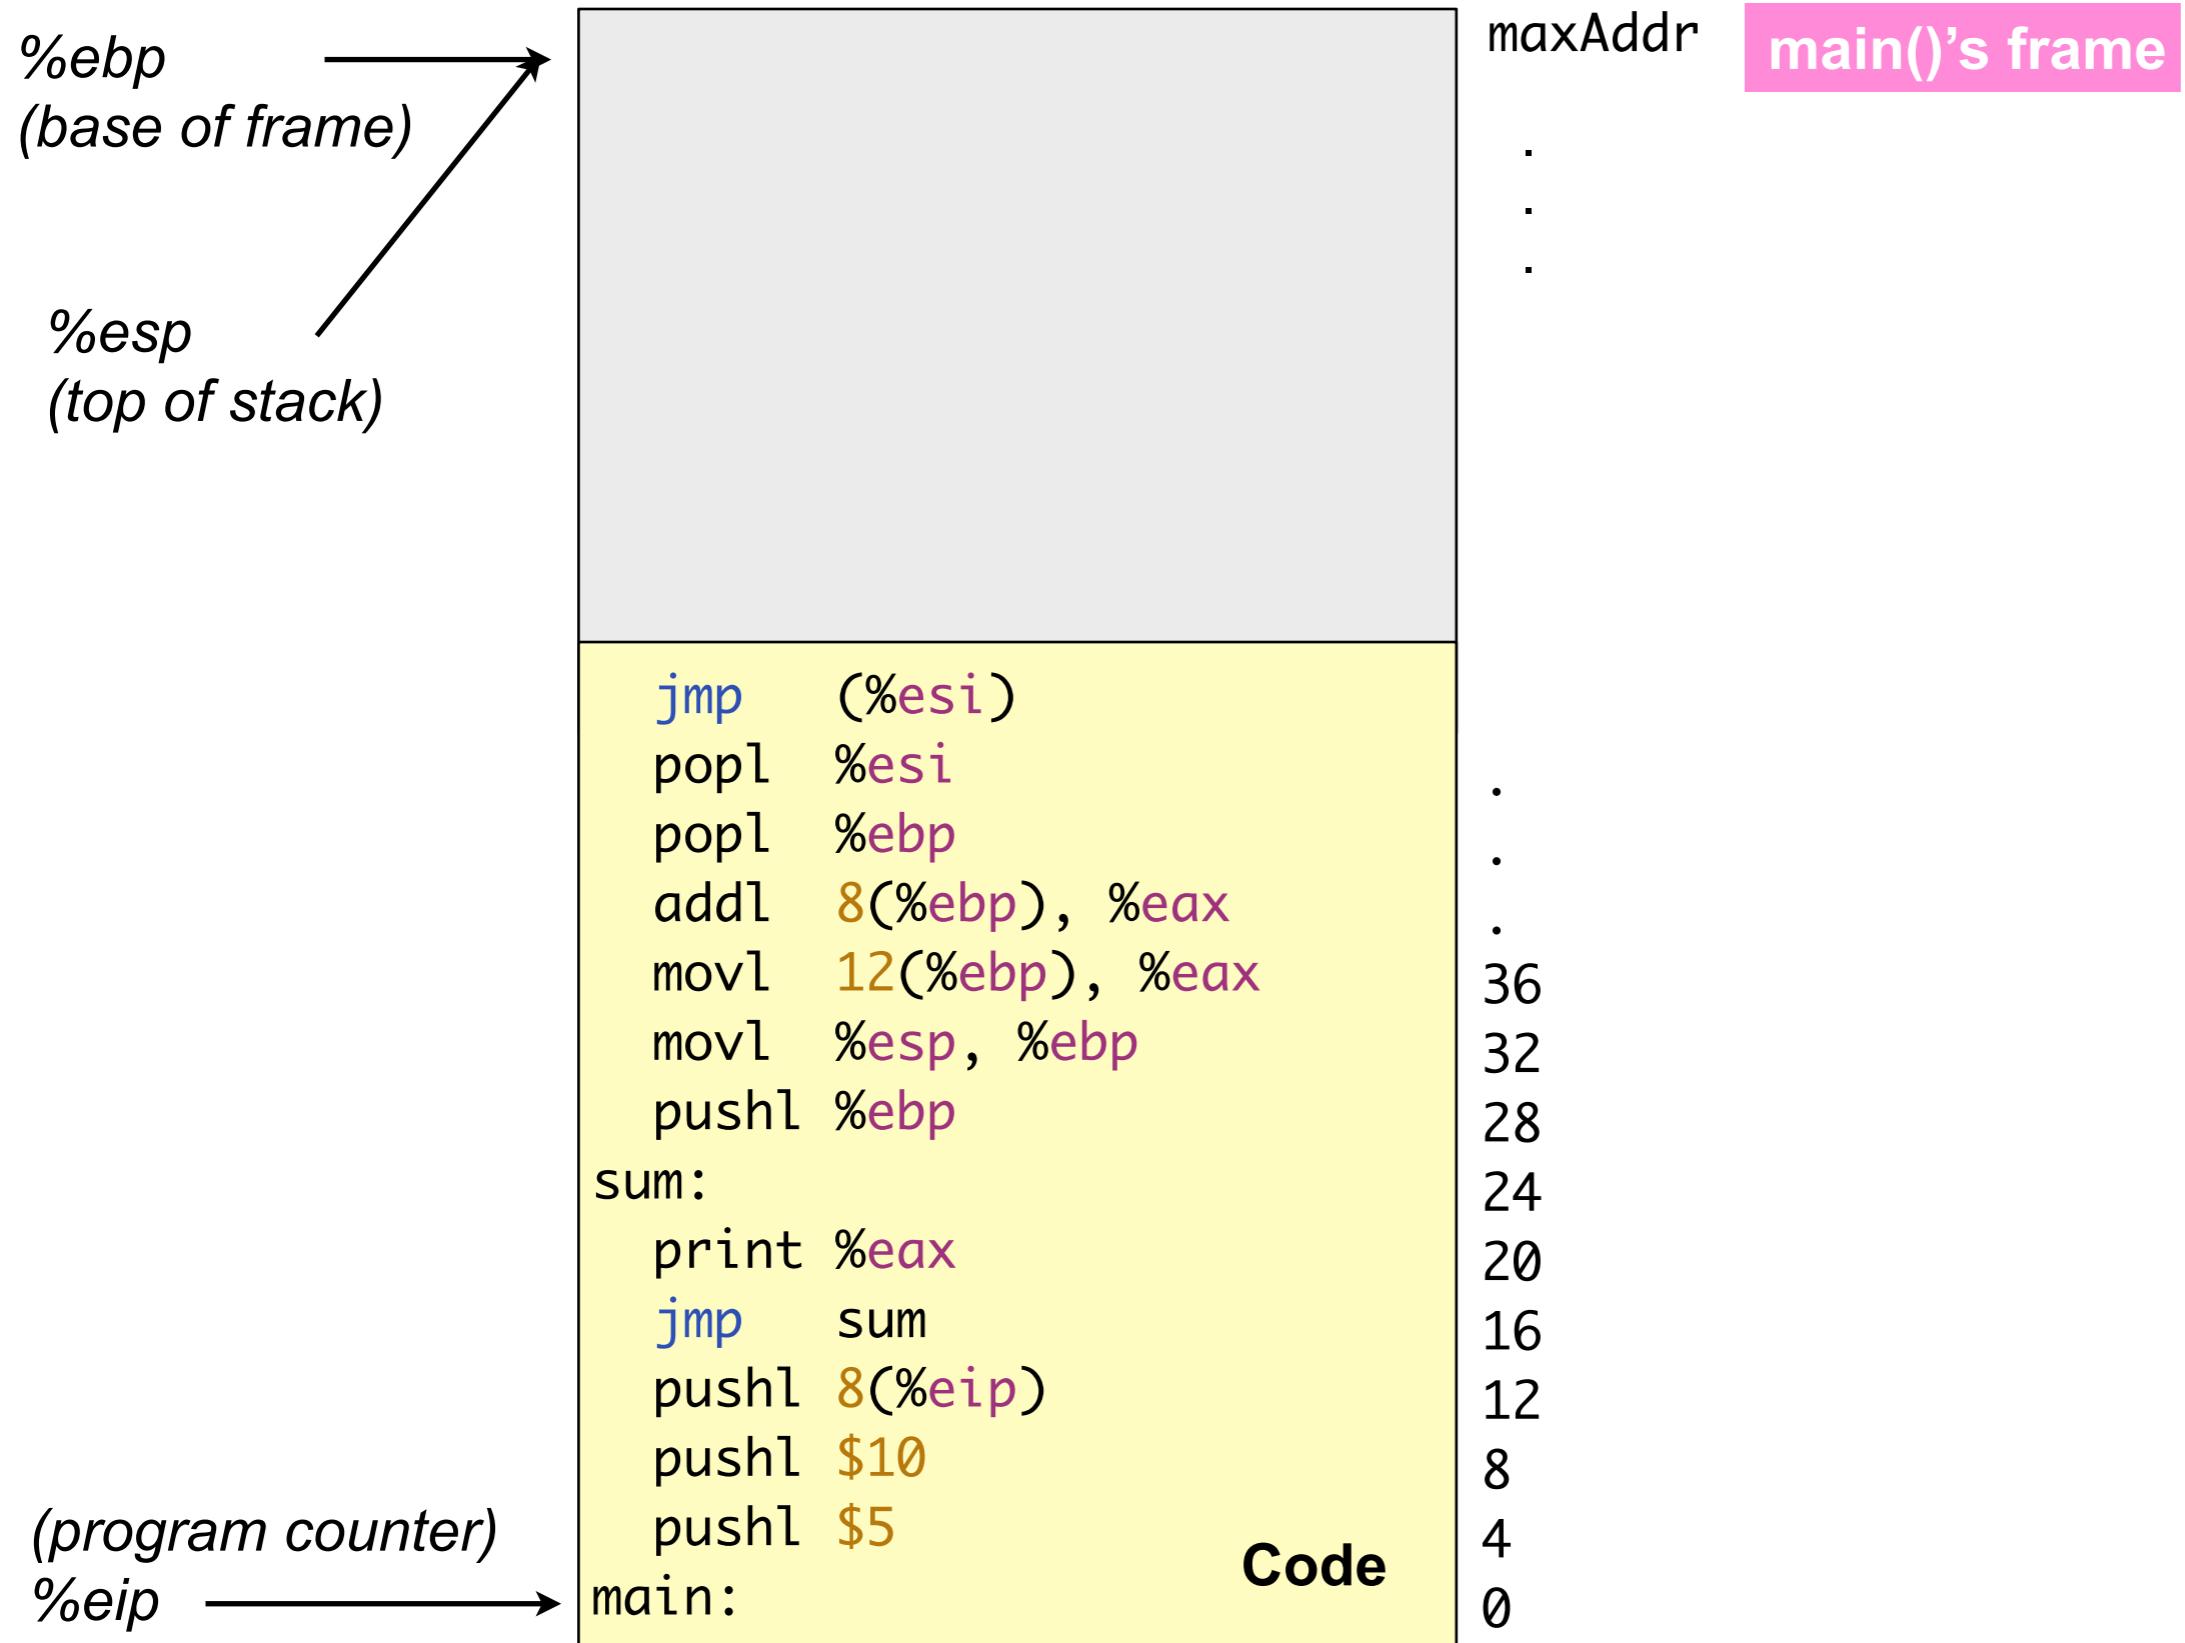
sum:
pushl %ebp
movl %esp, %ebp
movl 12(%ebp), %eax
addl 8(%ebp), %eax
popl %ebp
popl %esi
jmp (%esi)
  
```

\$ gcc -S -o sum code.c

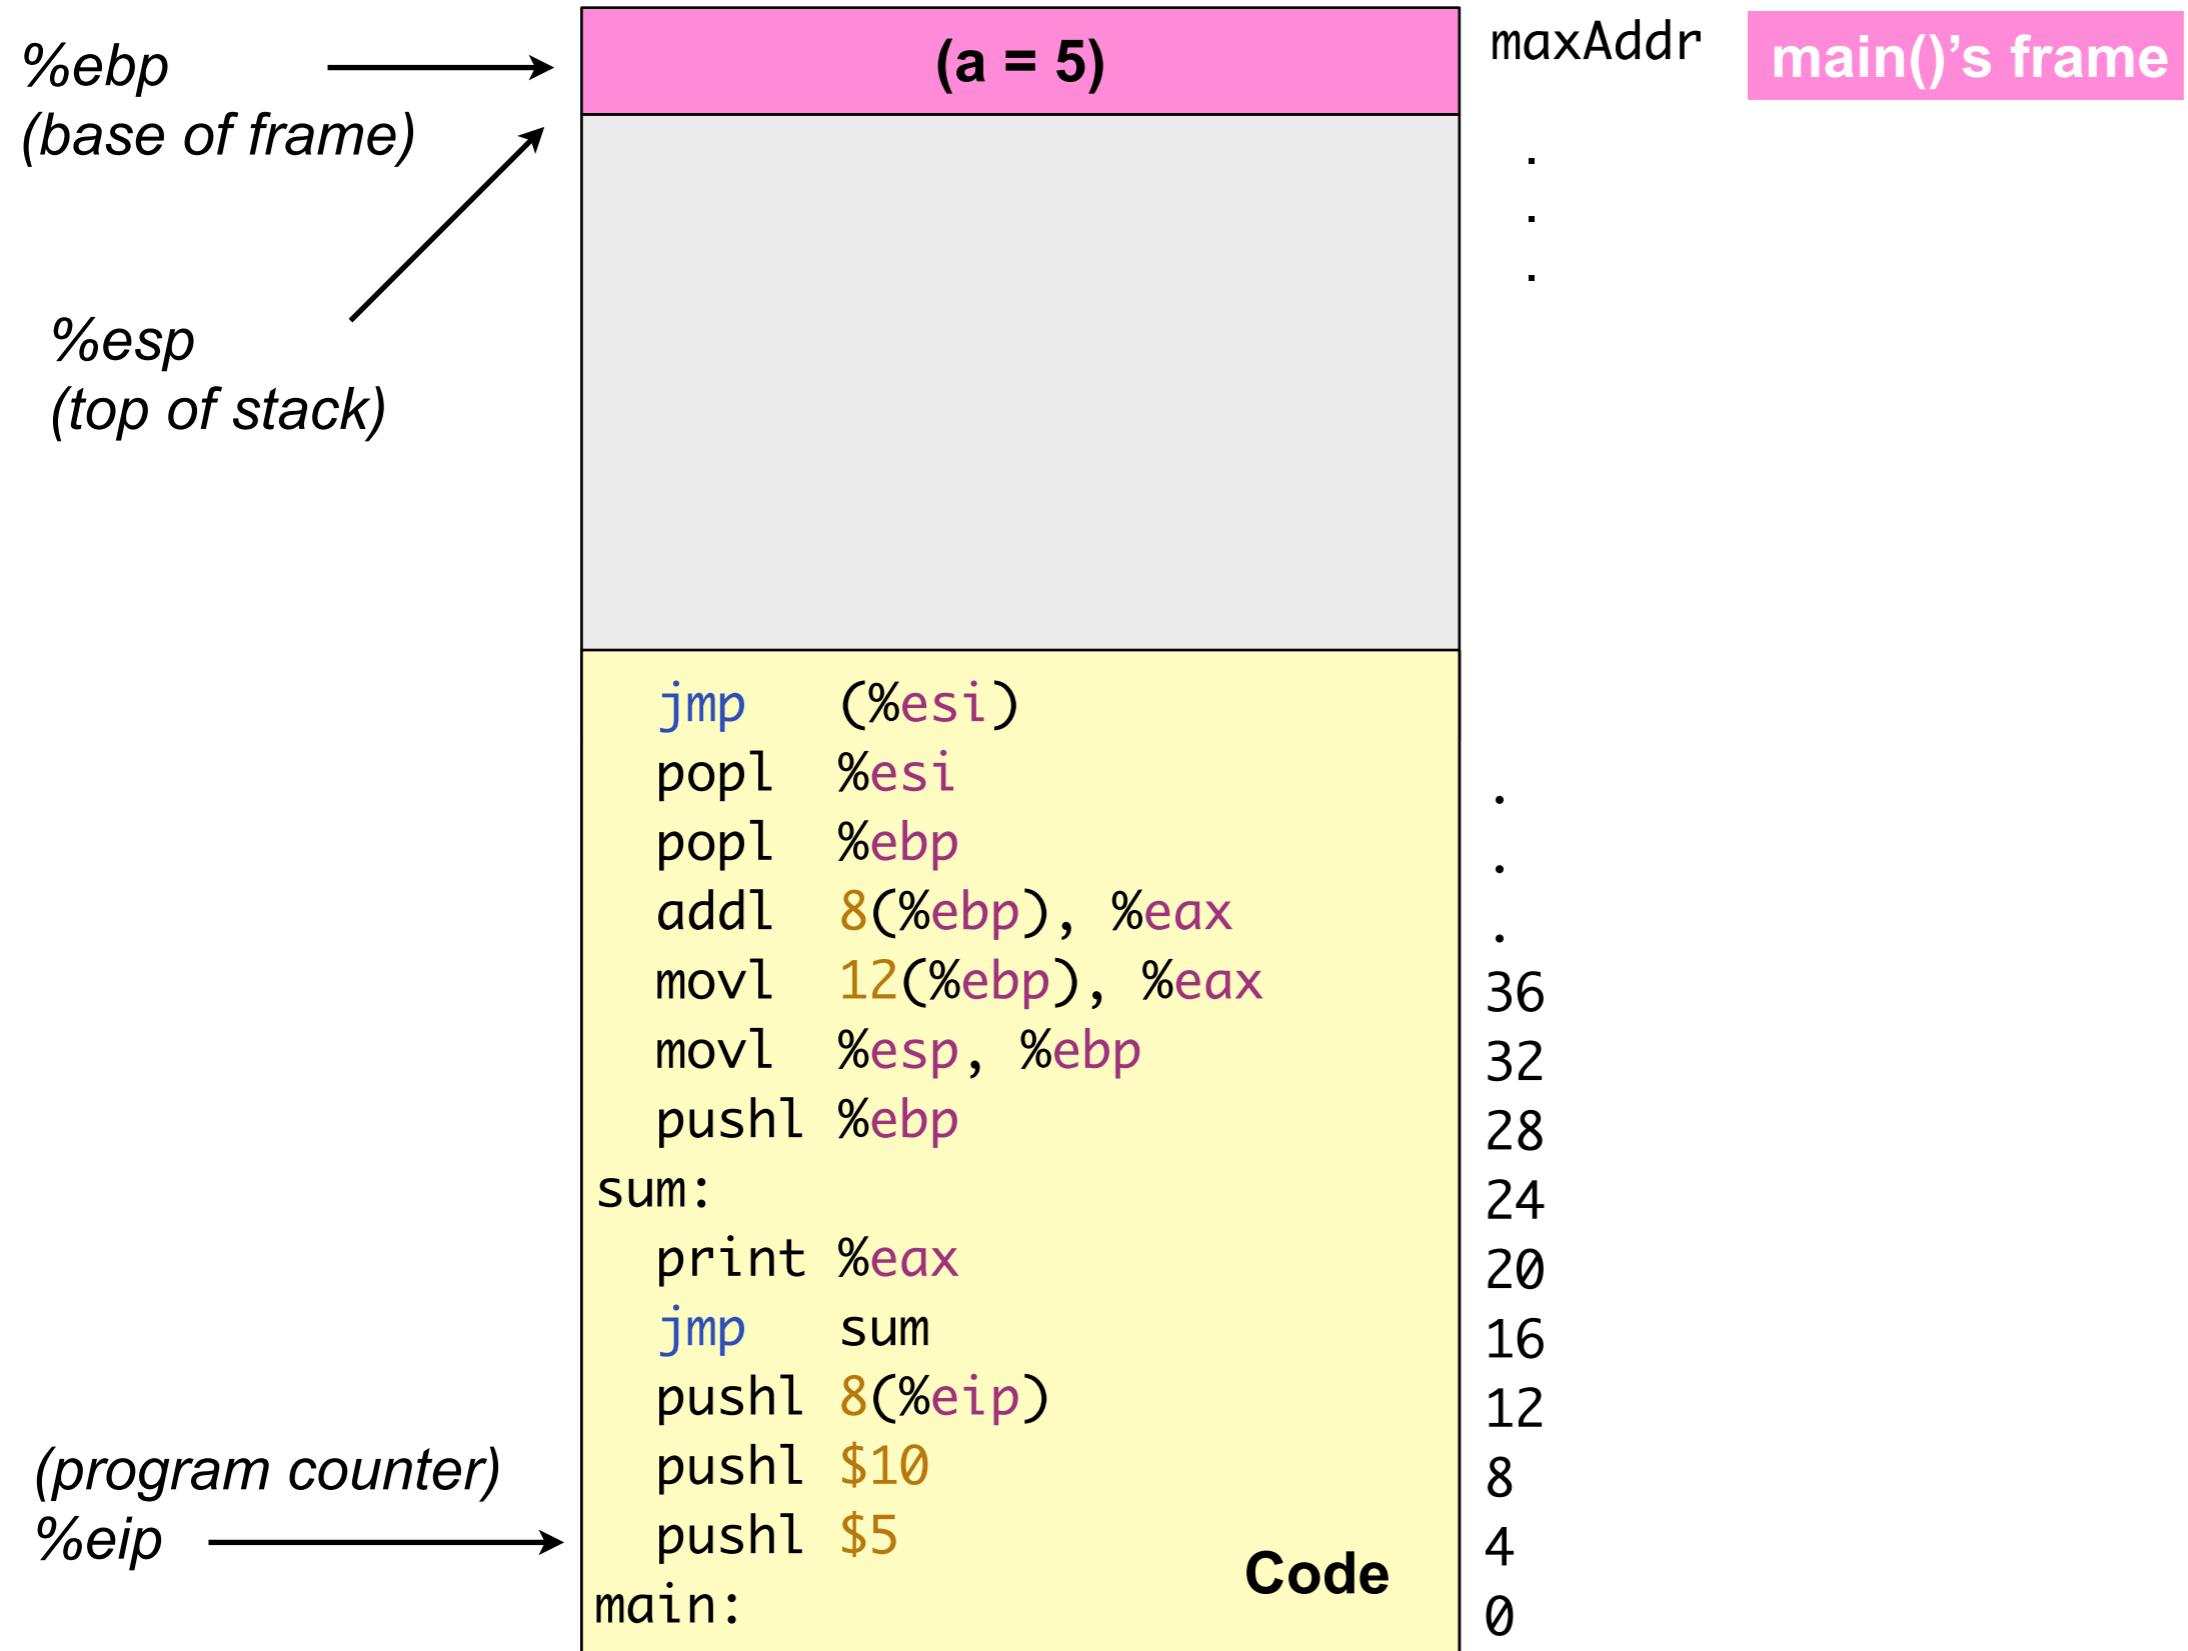
Stack Management: x86 Example



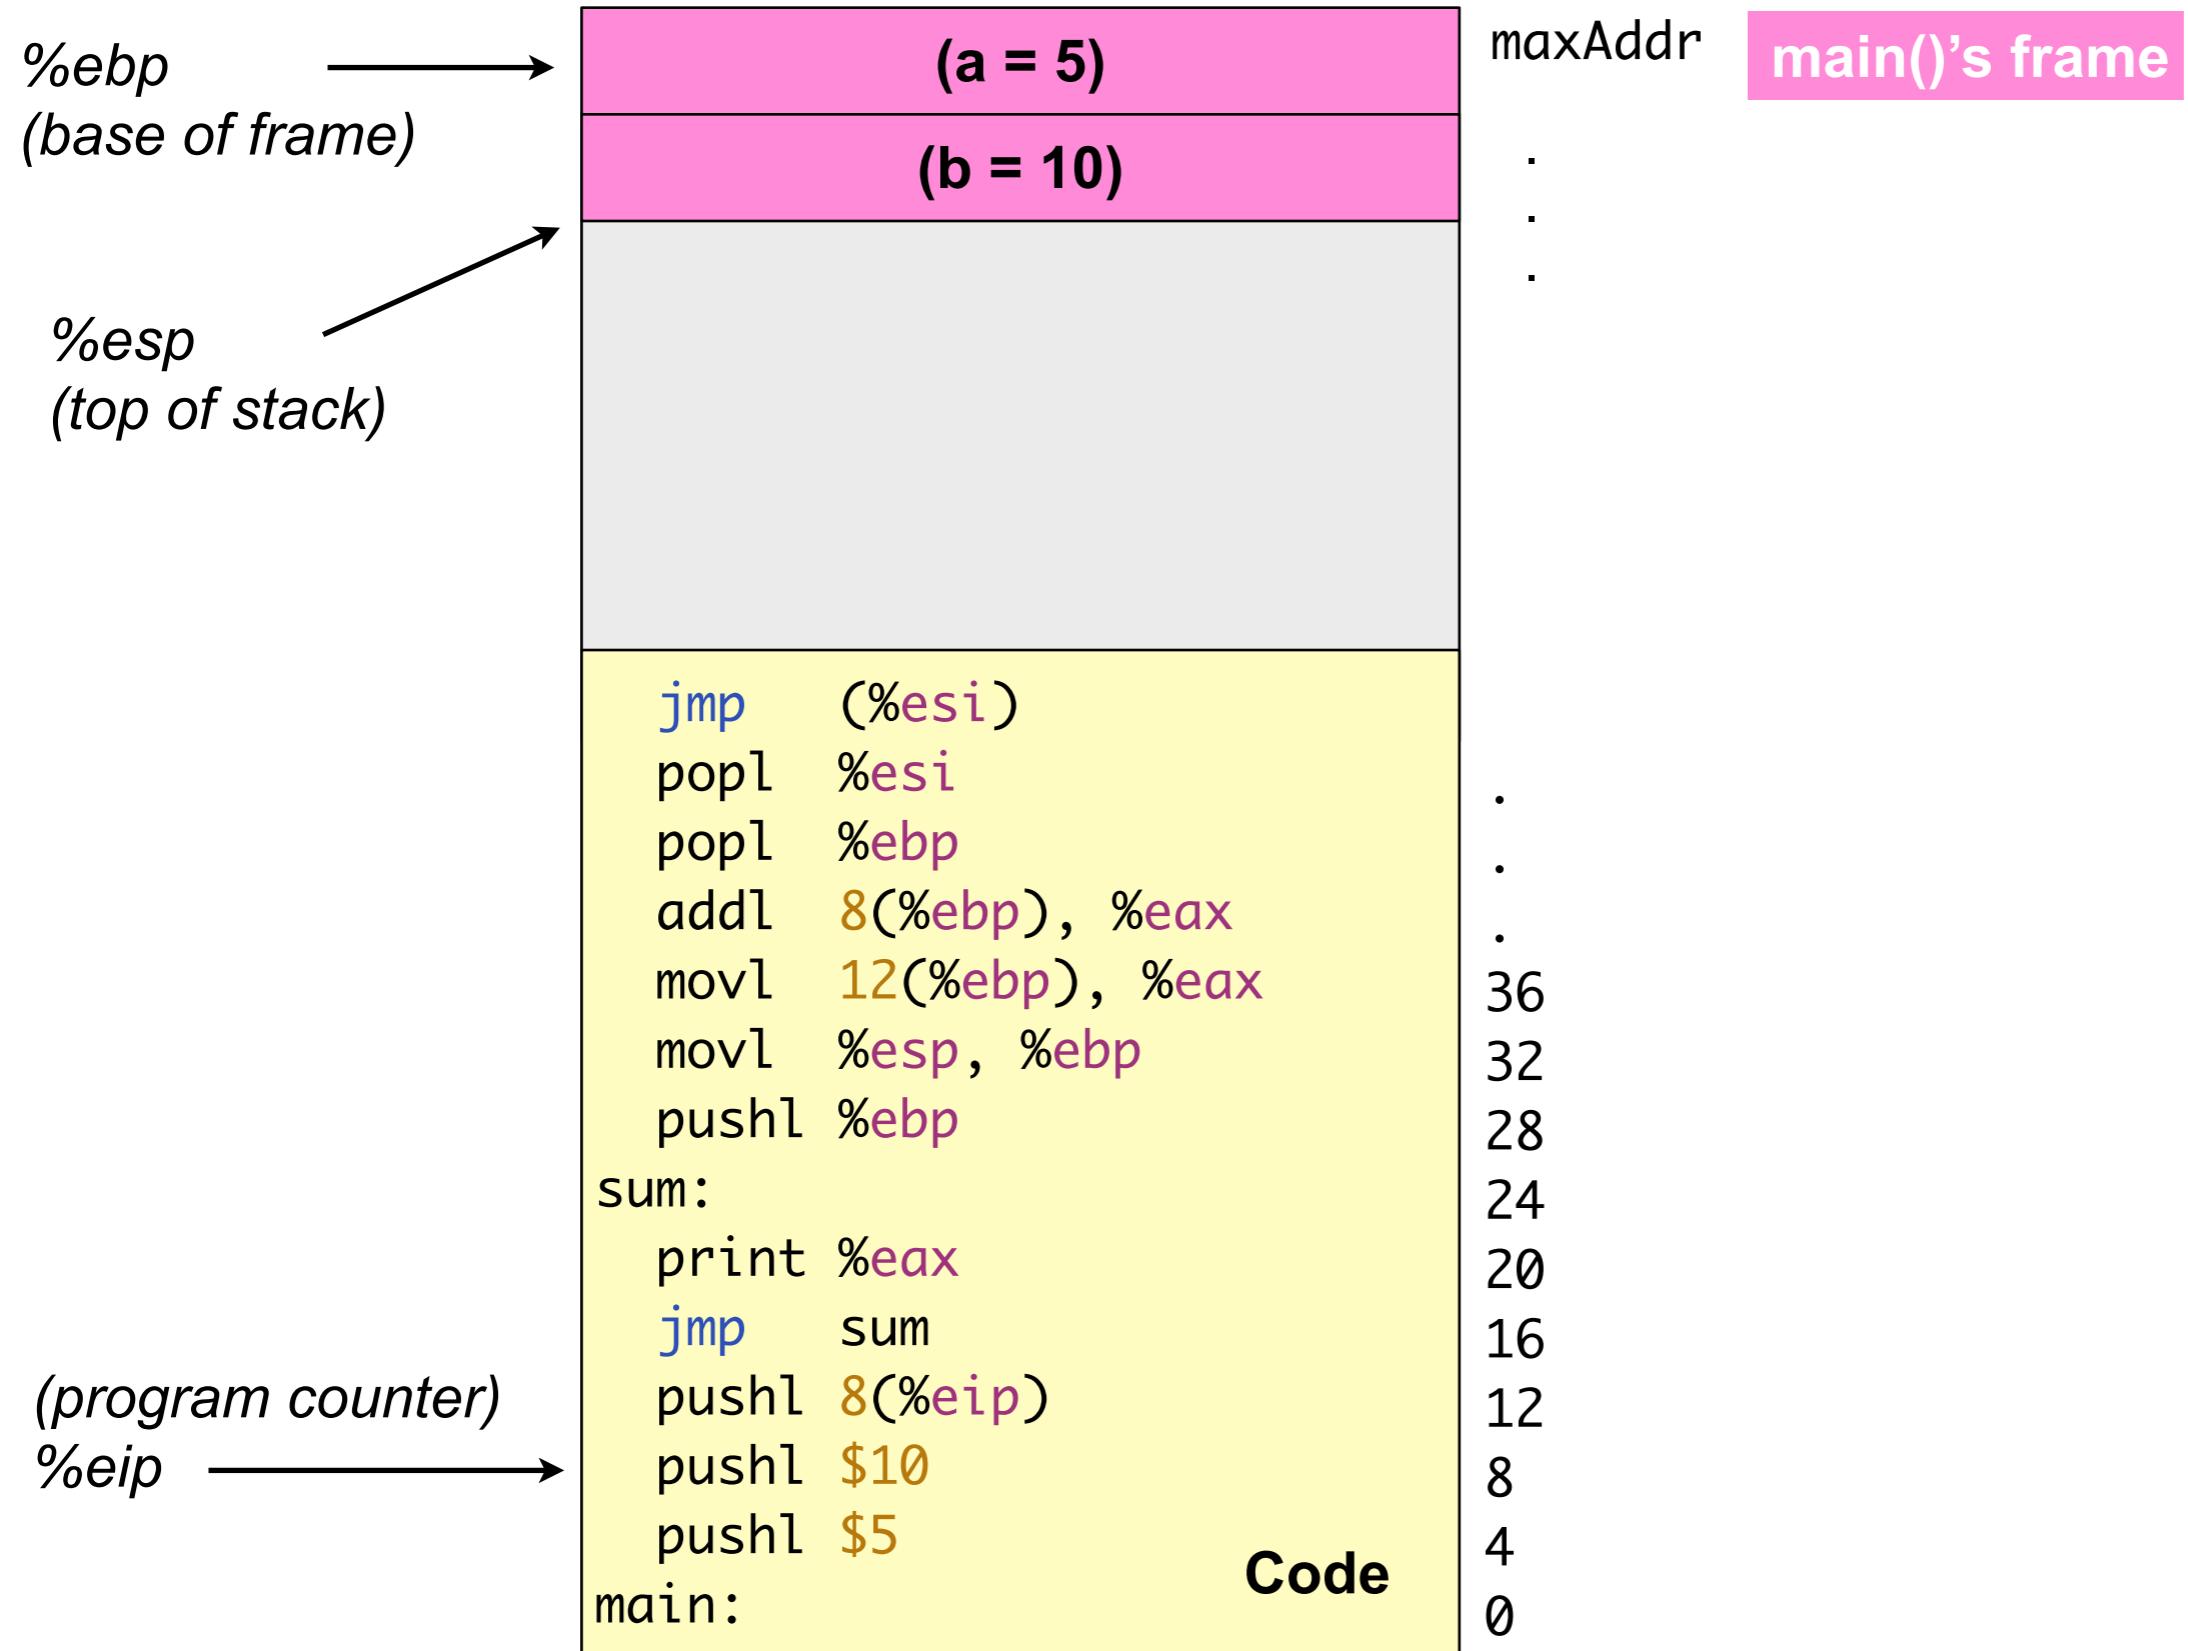
x86 Example (1)



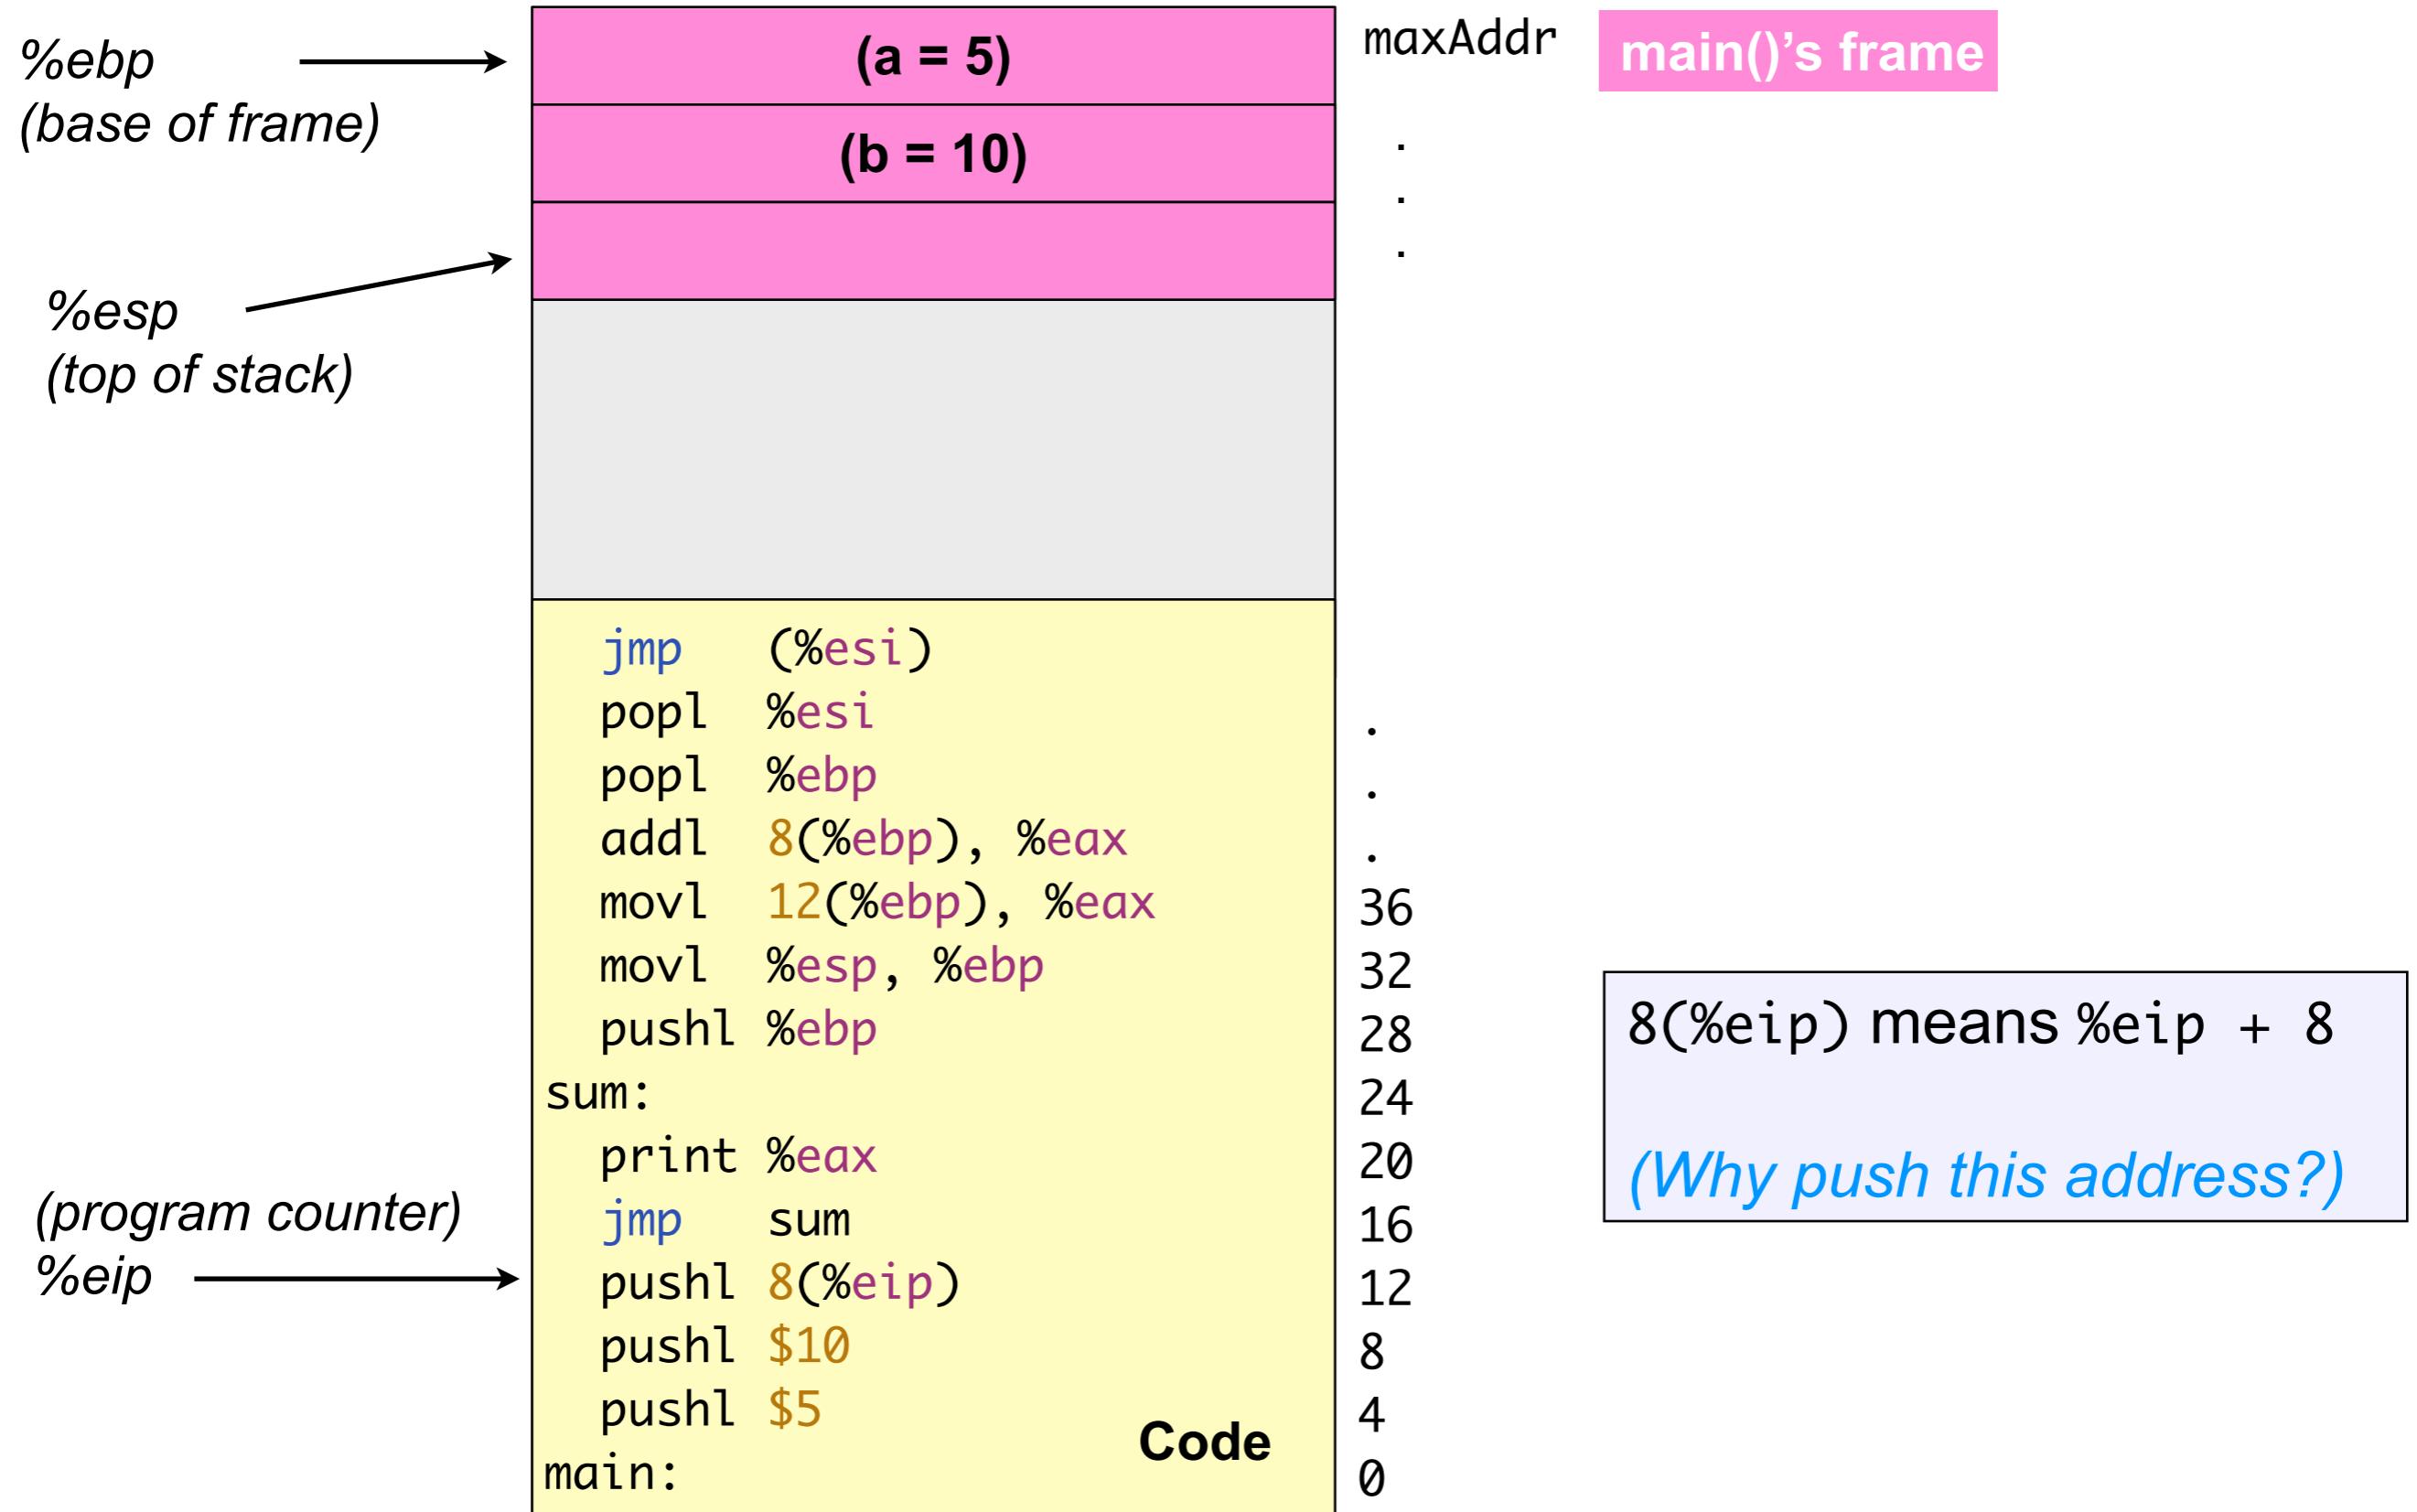
x86 Example (2)



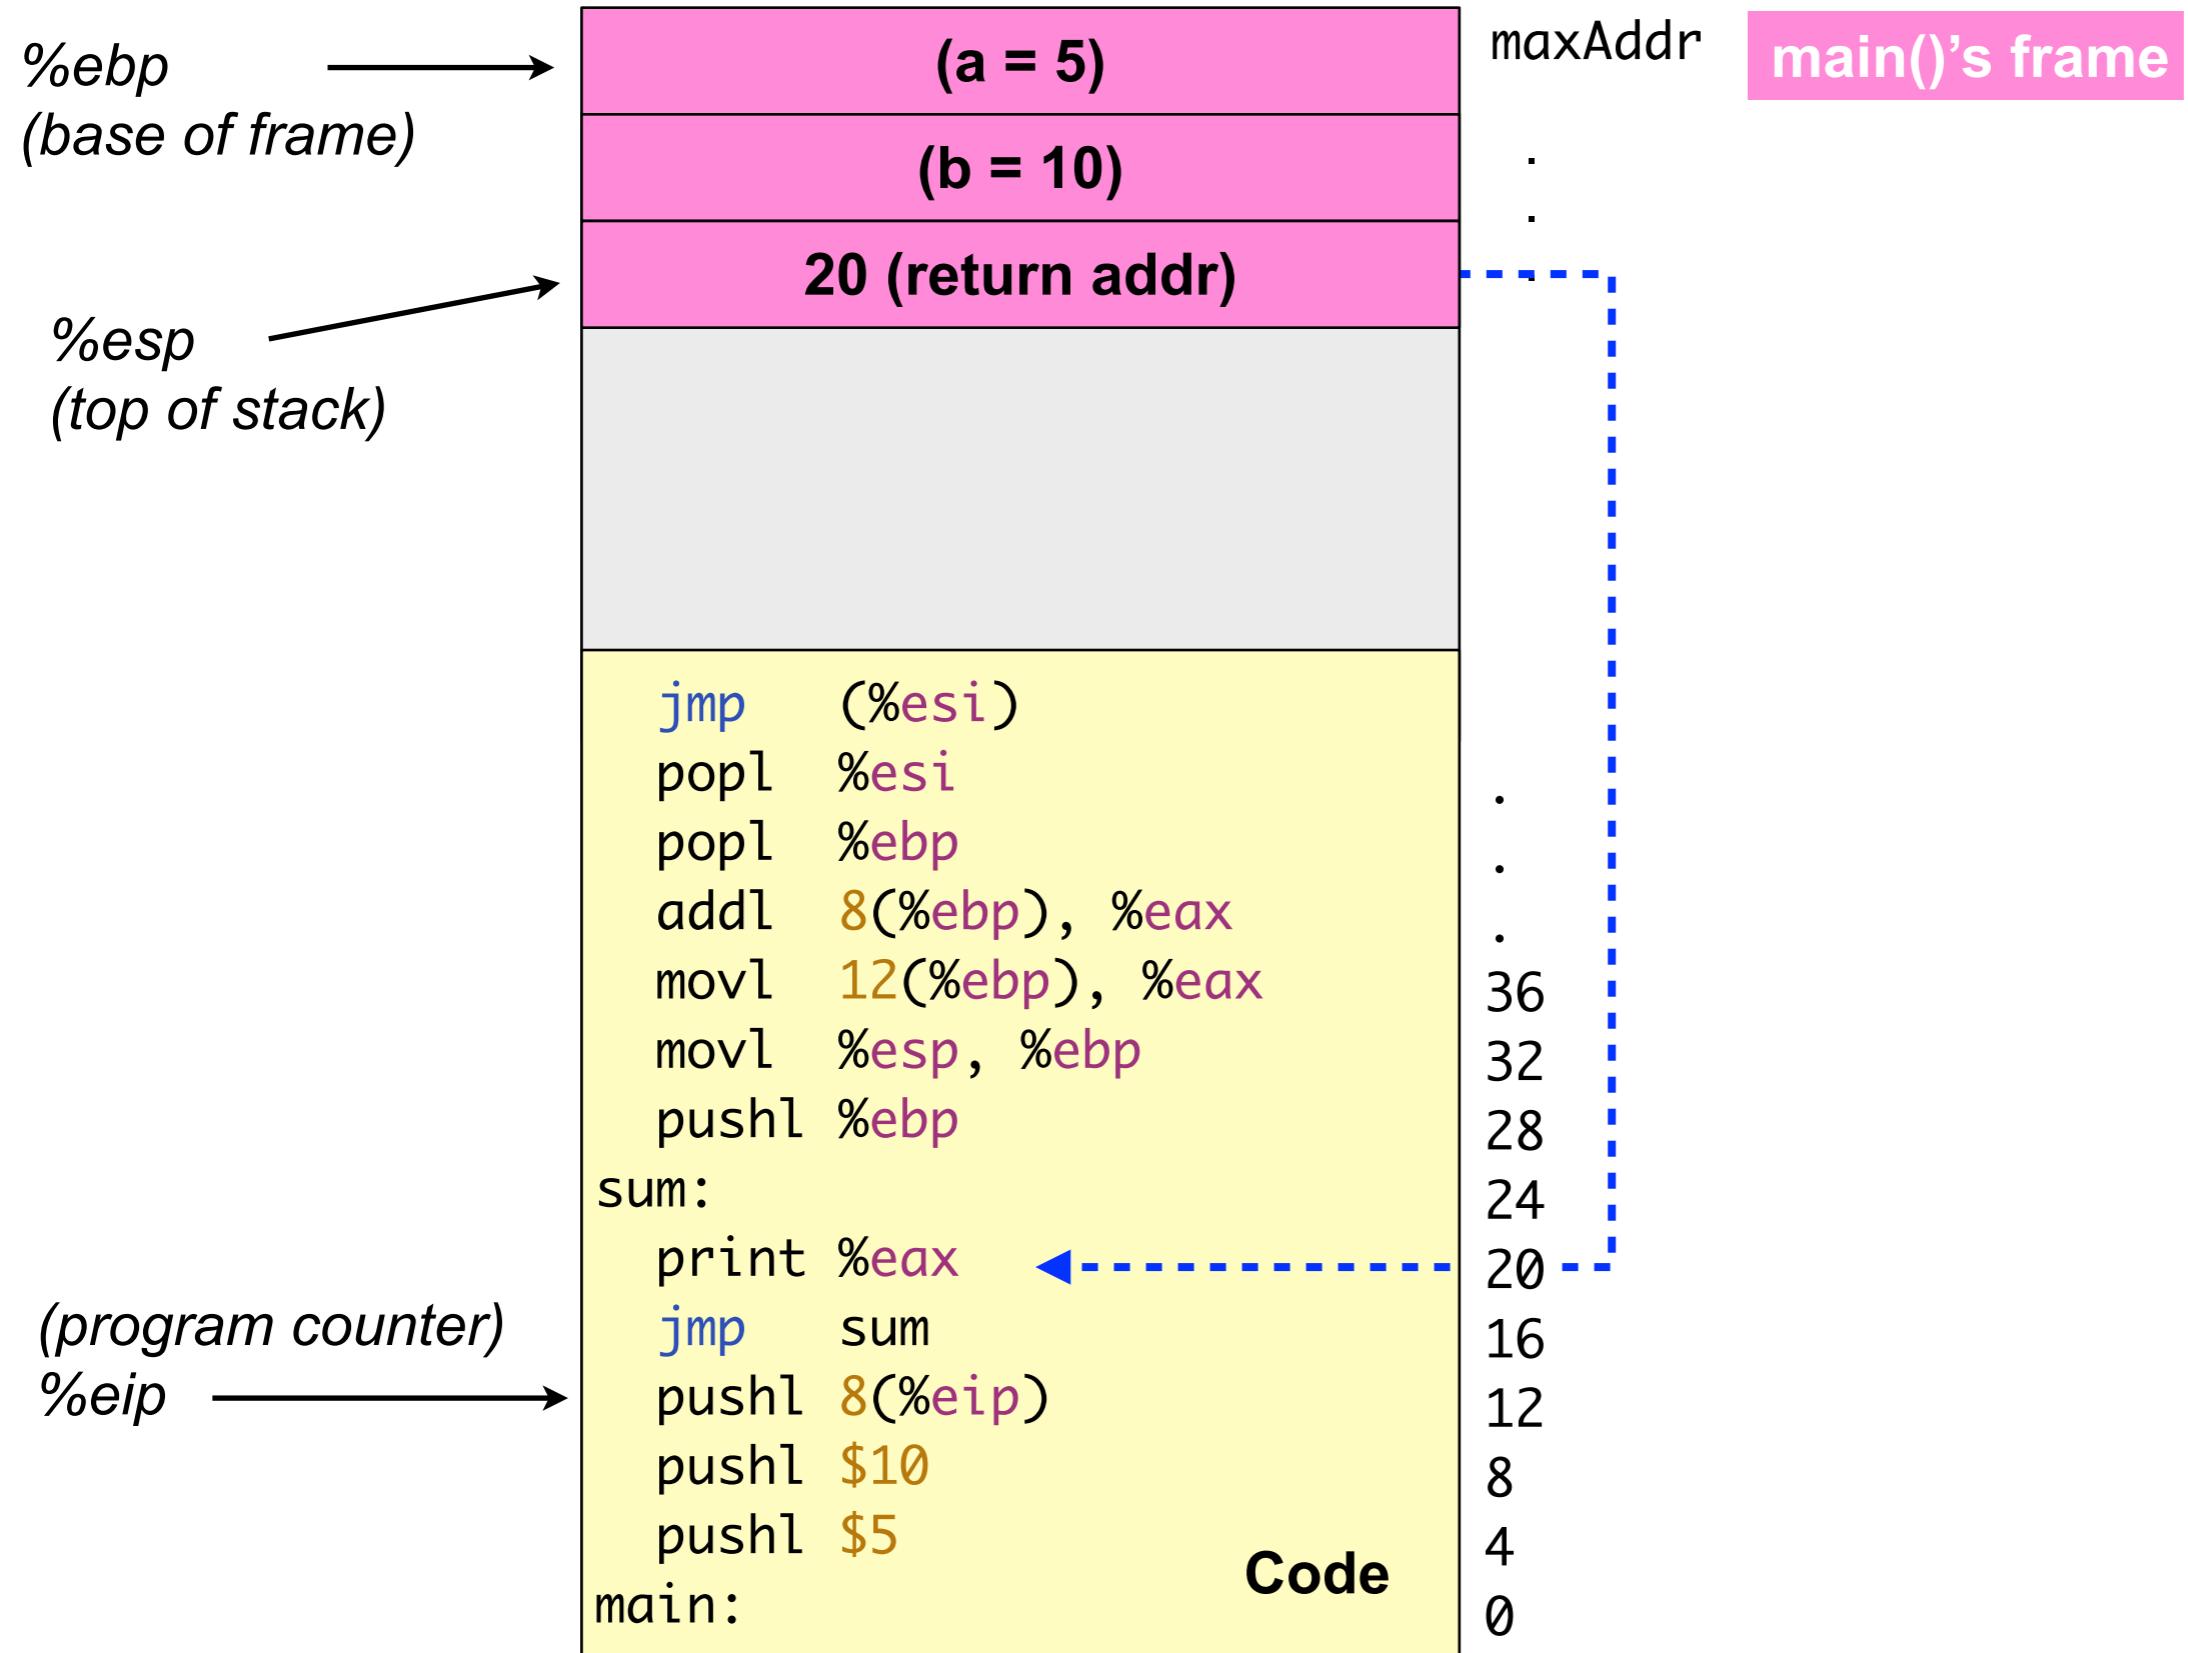
x86 Example (3)



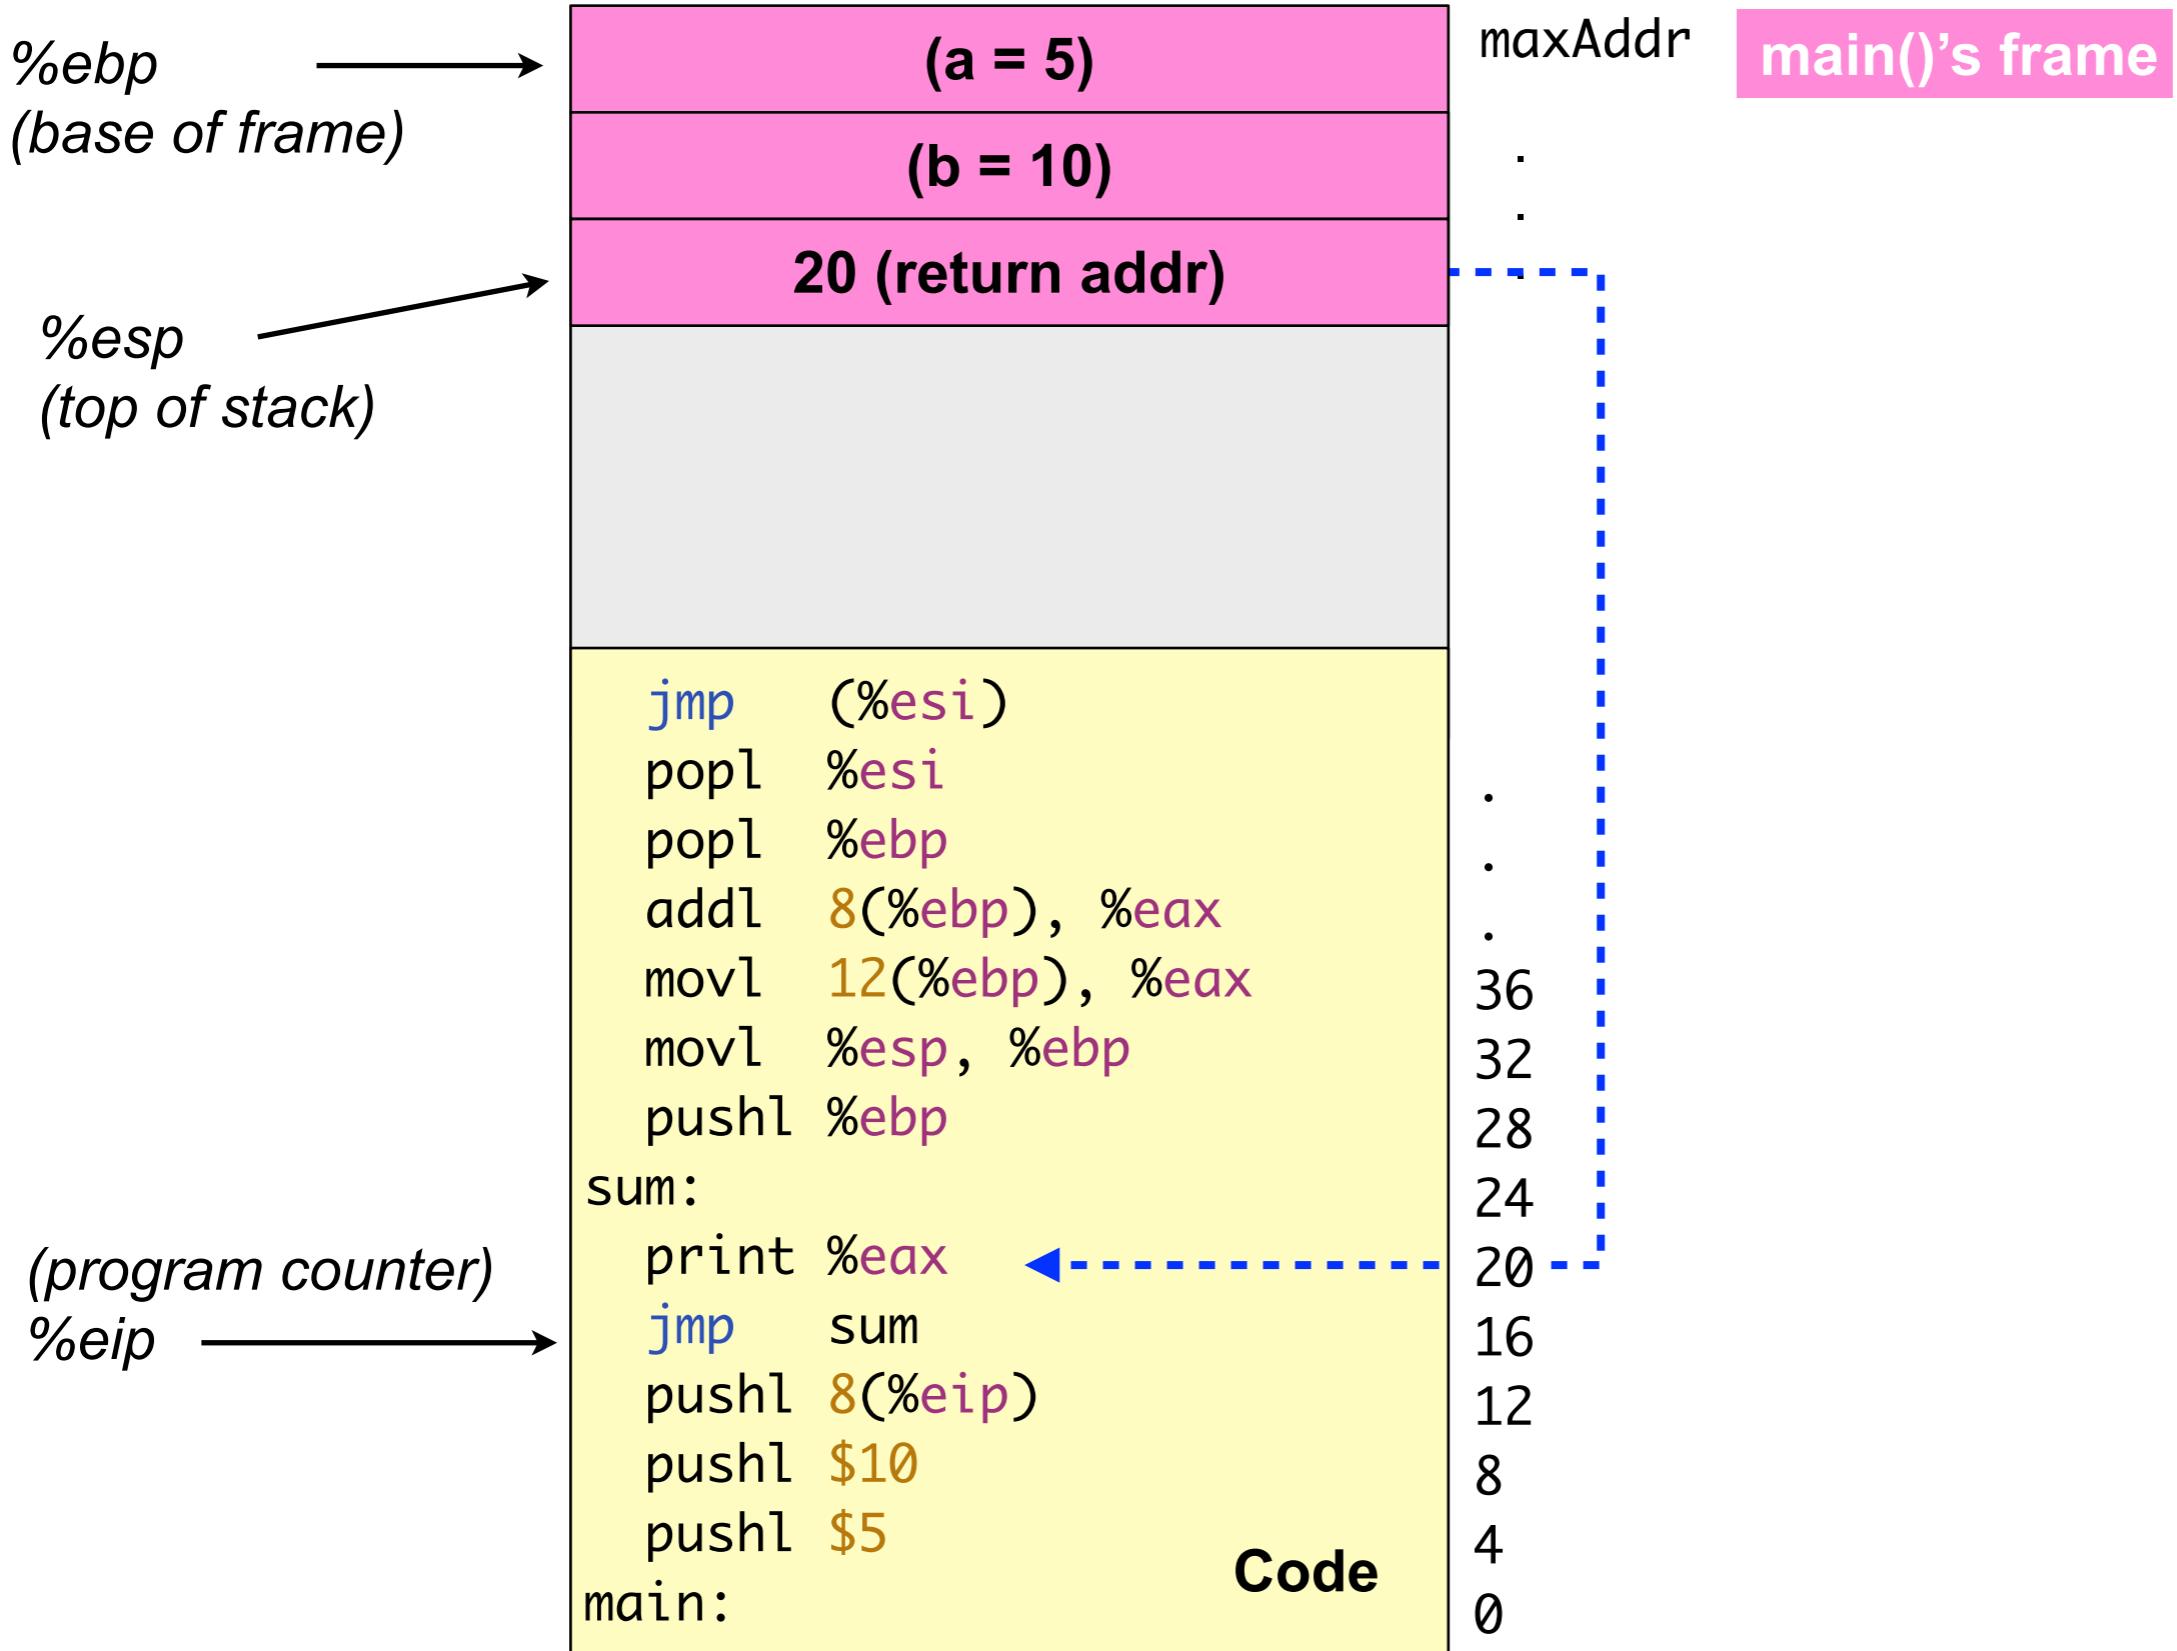
x86 Example (4)



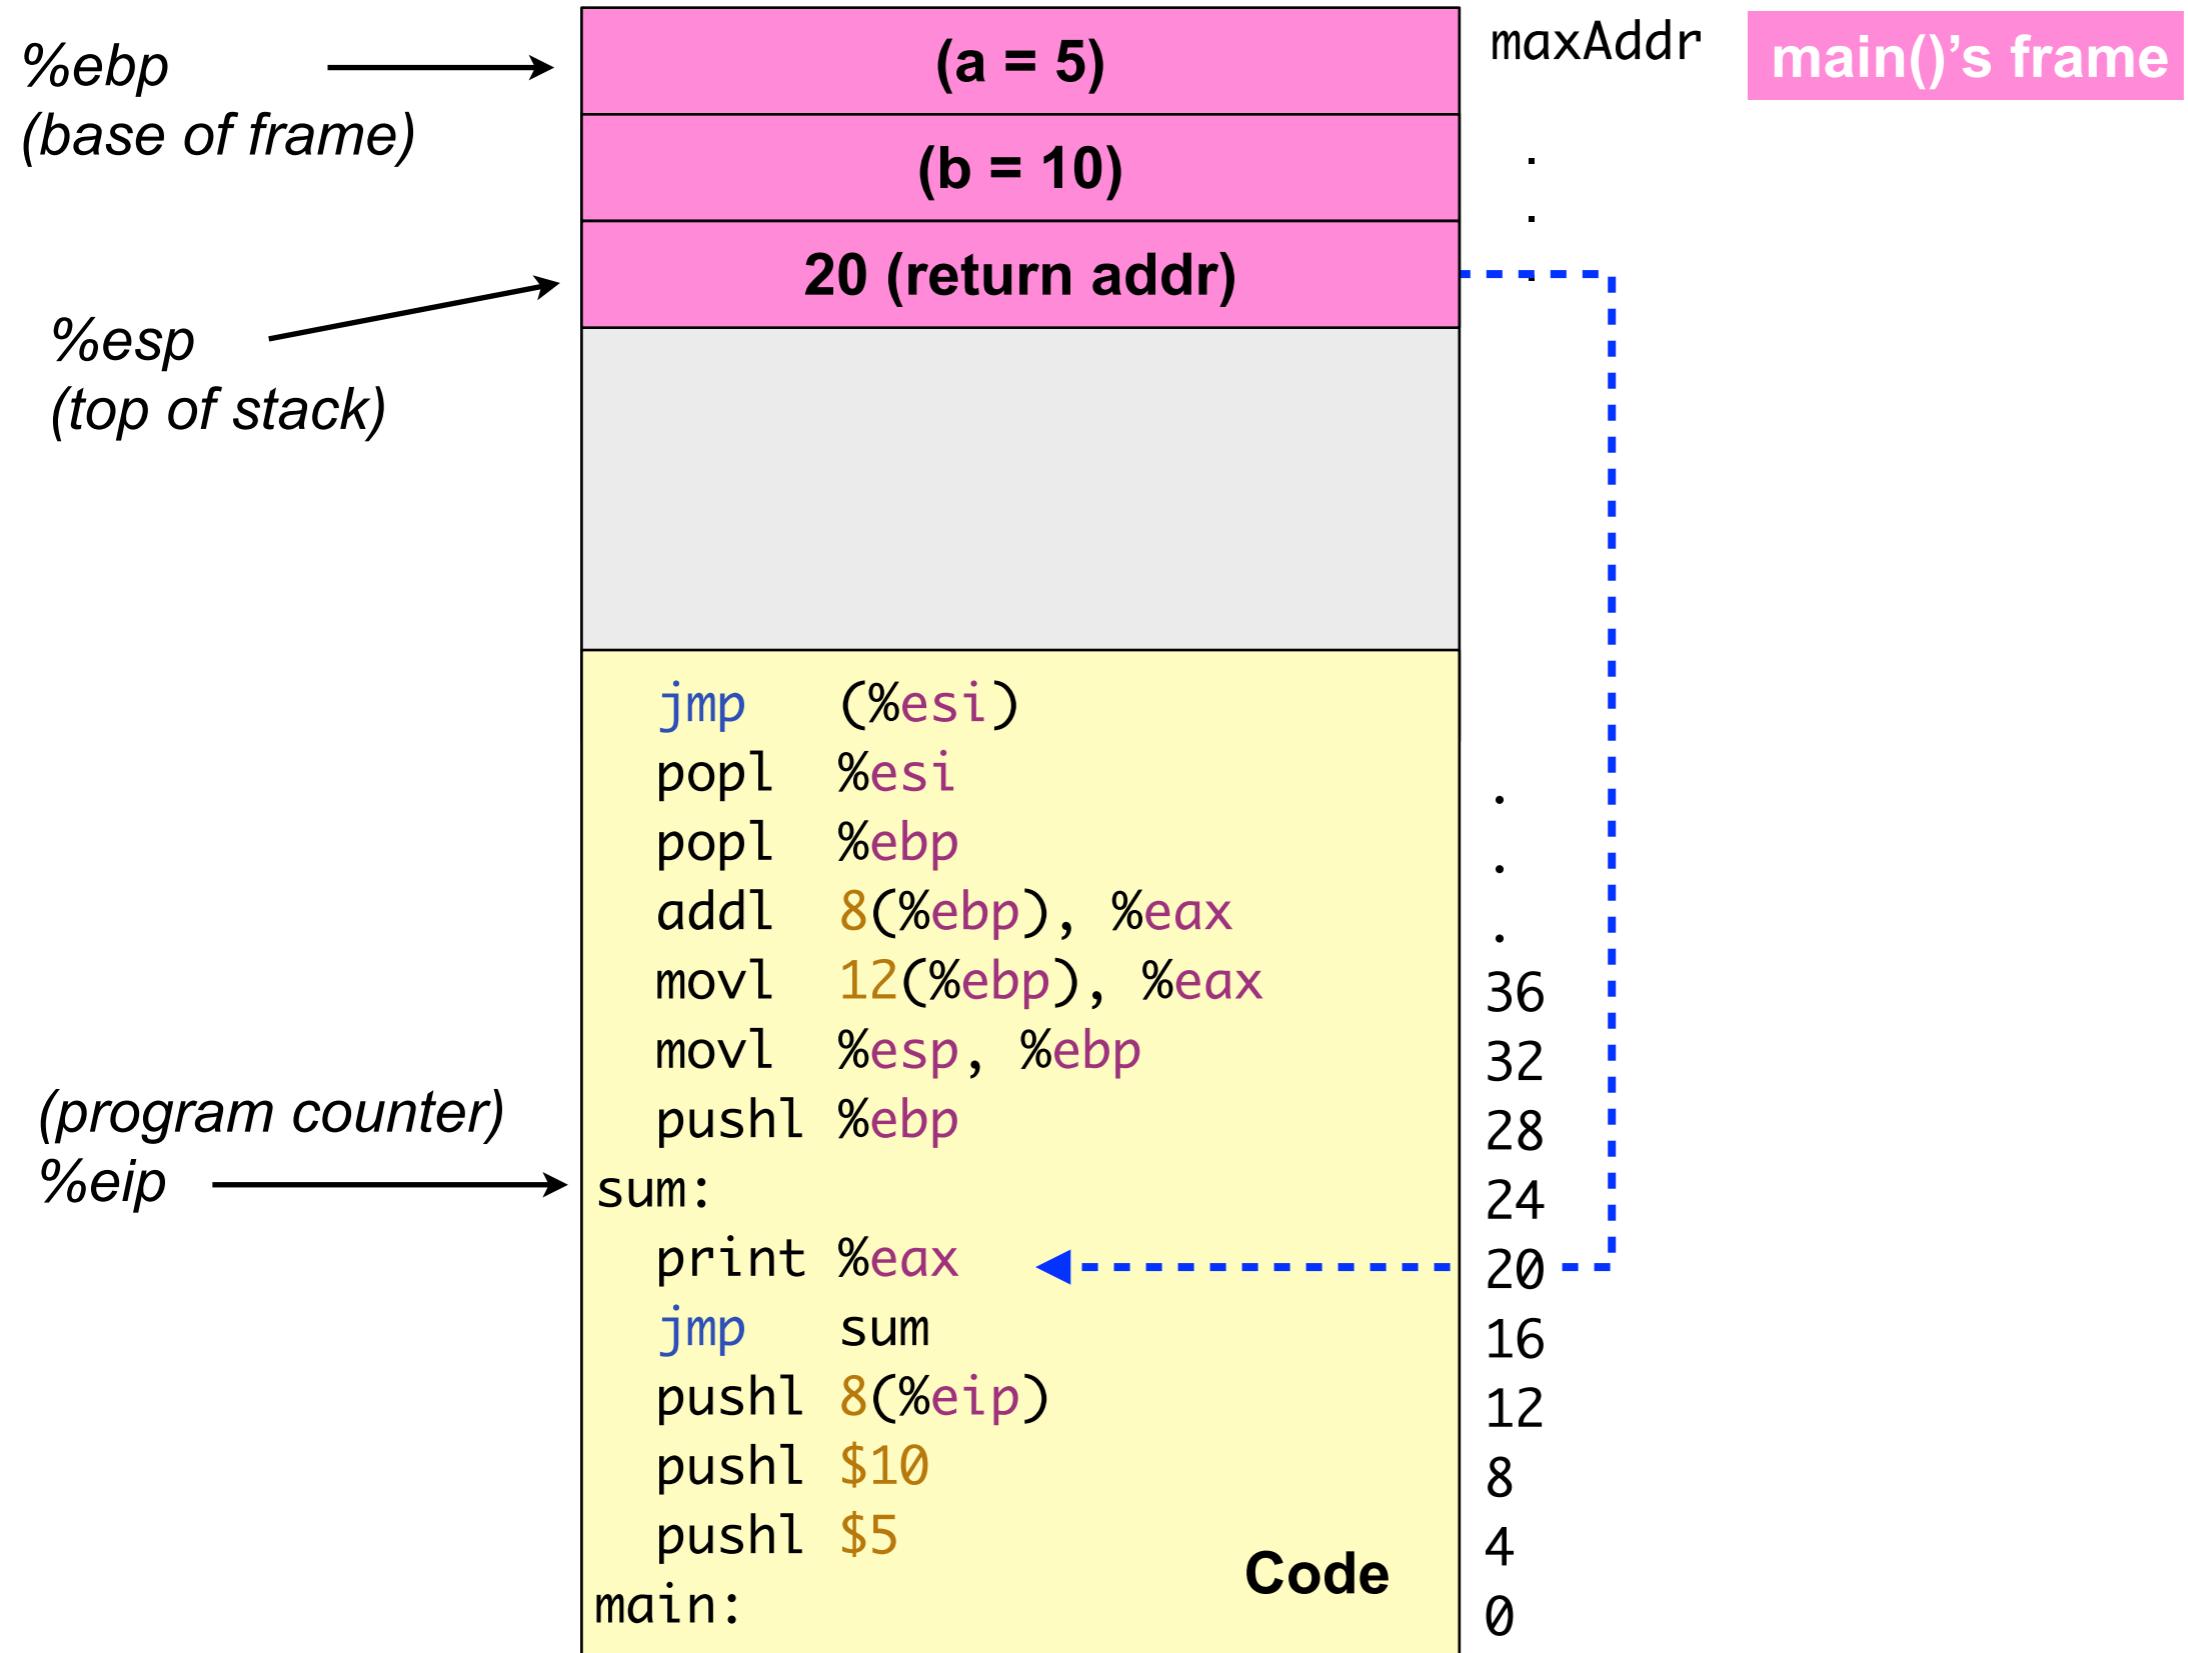
x86 Example (5)



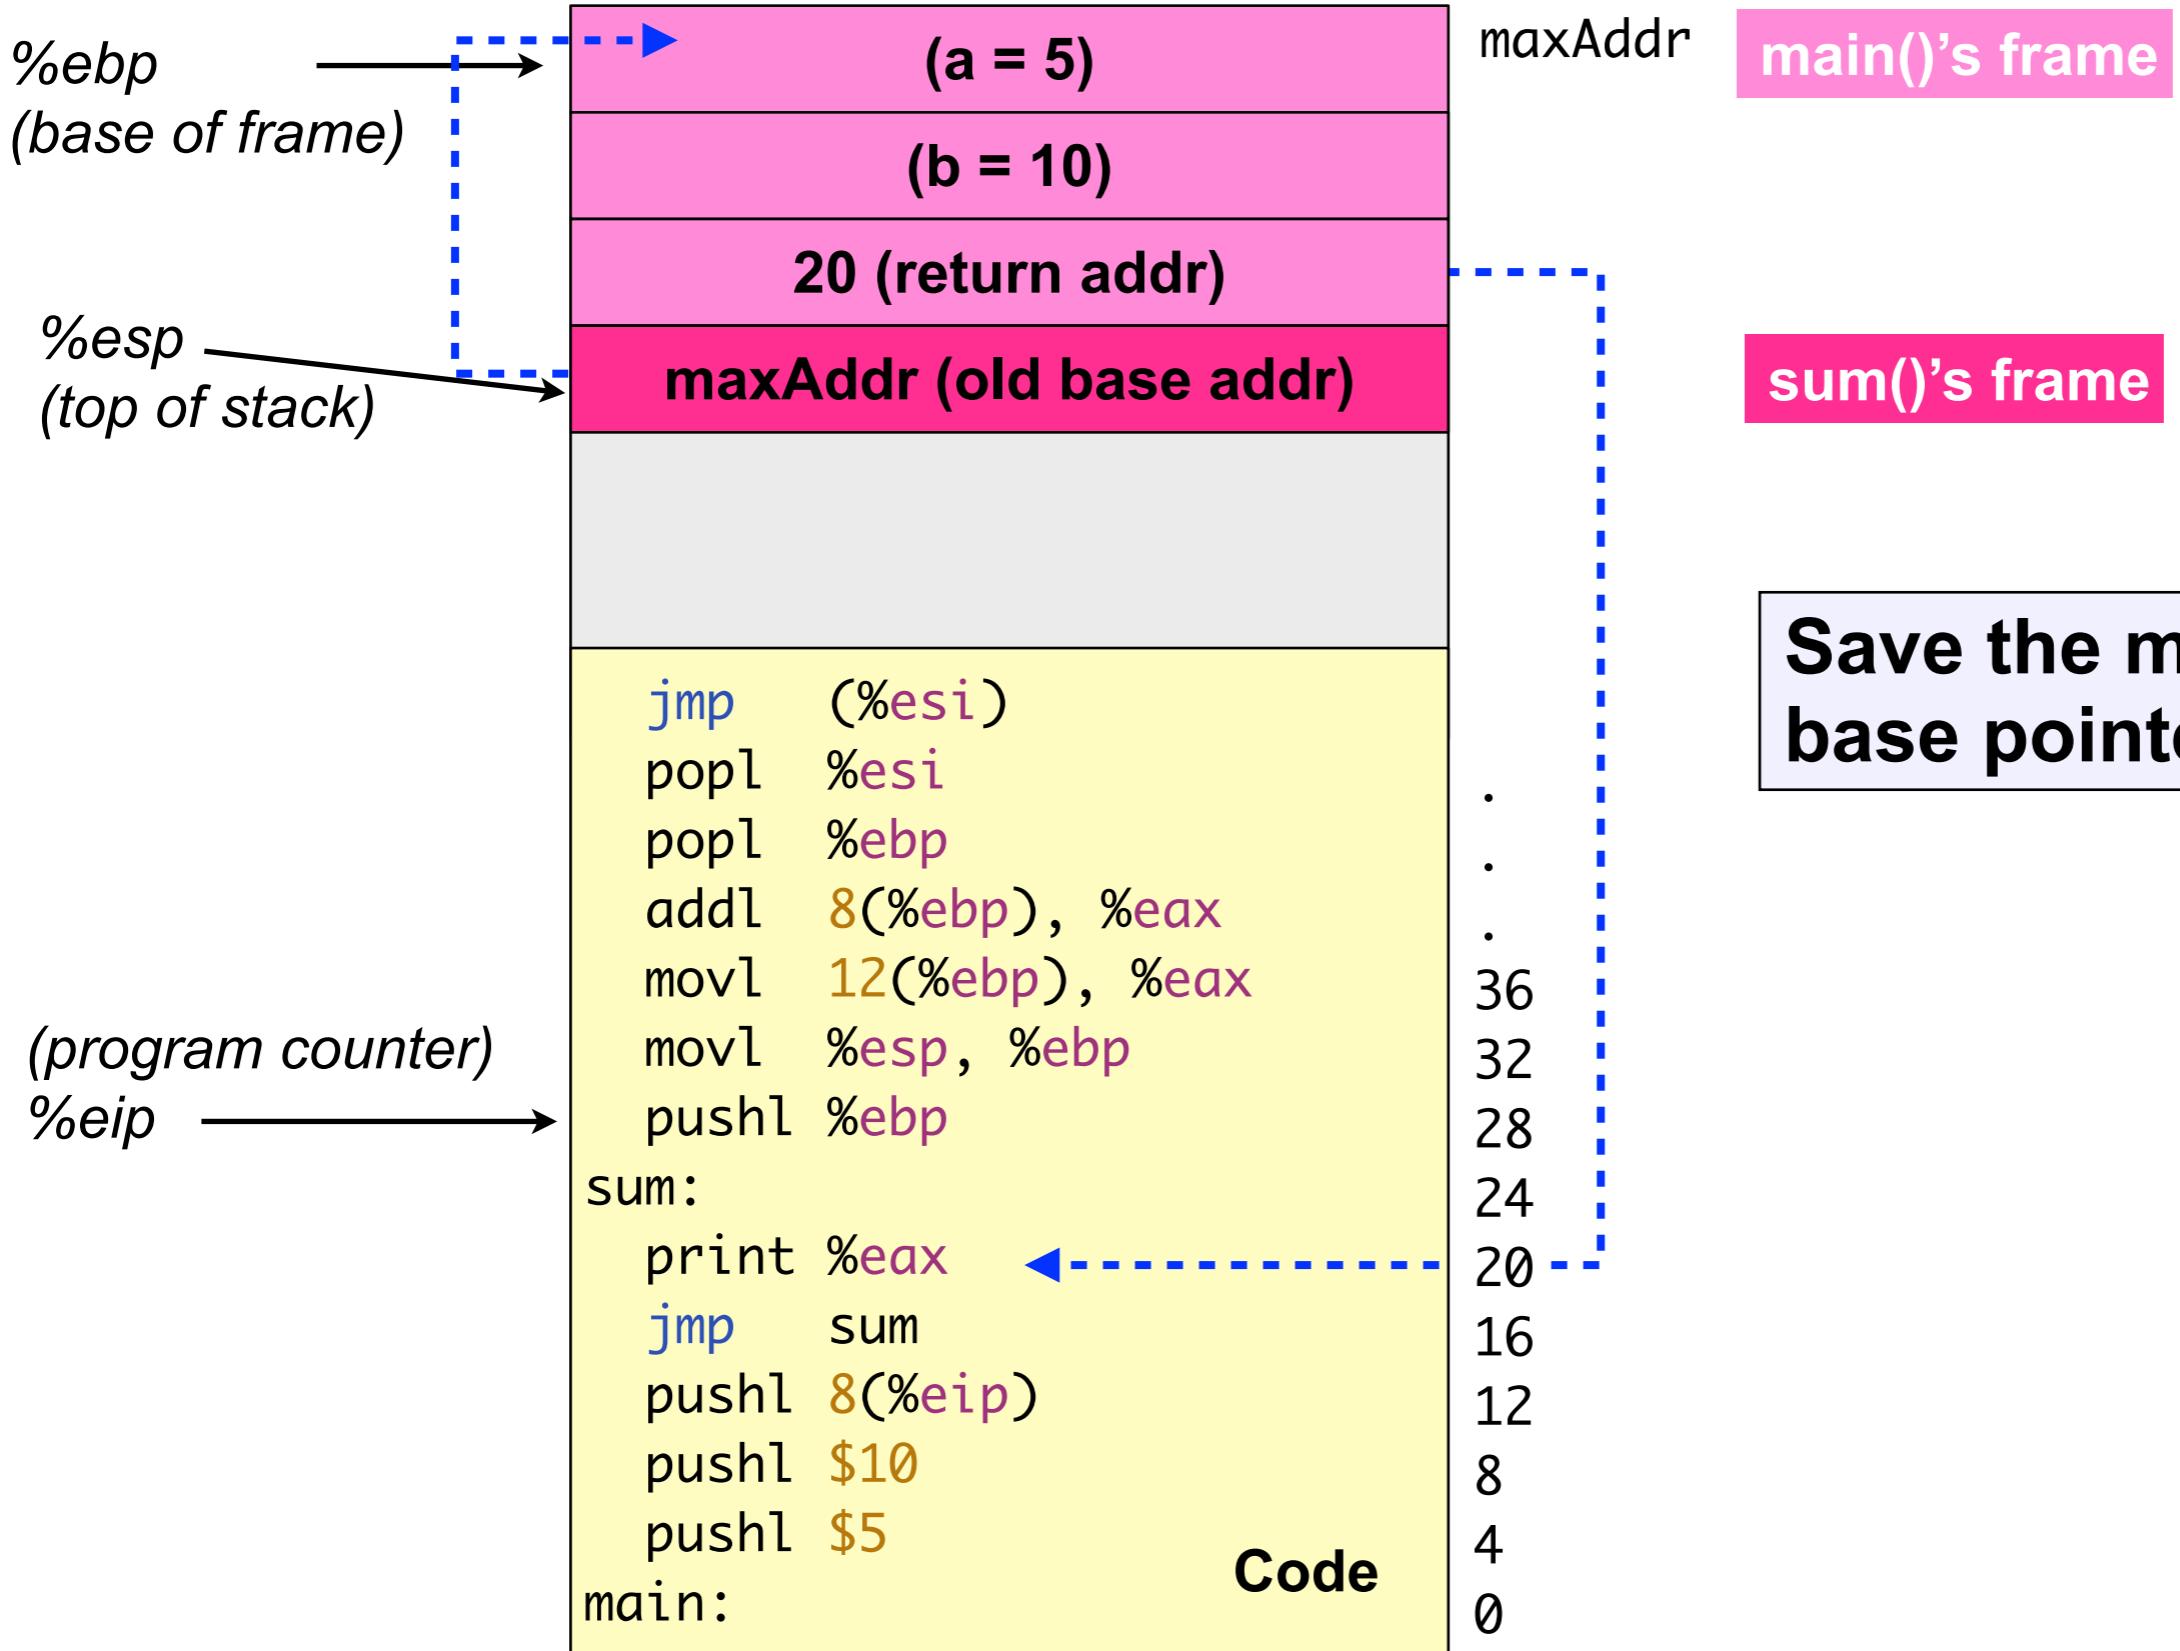
x86 Example (6)



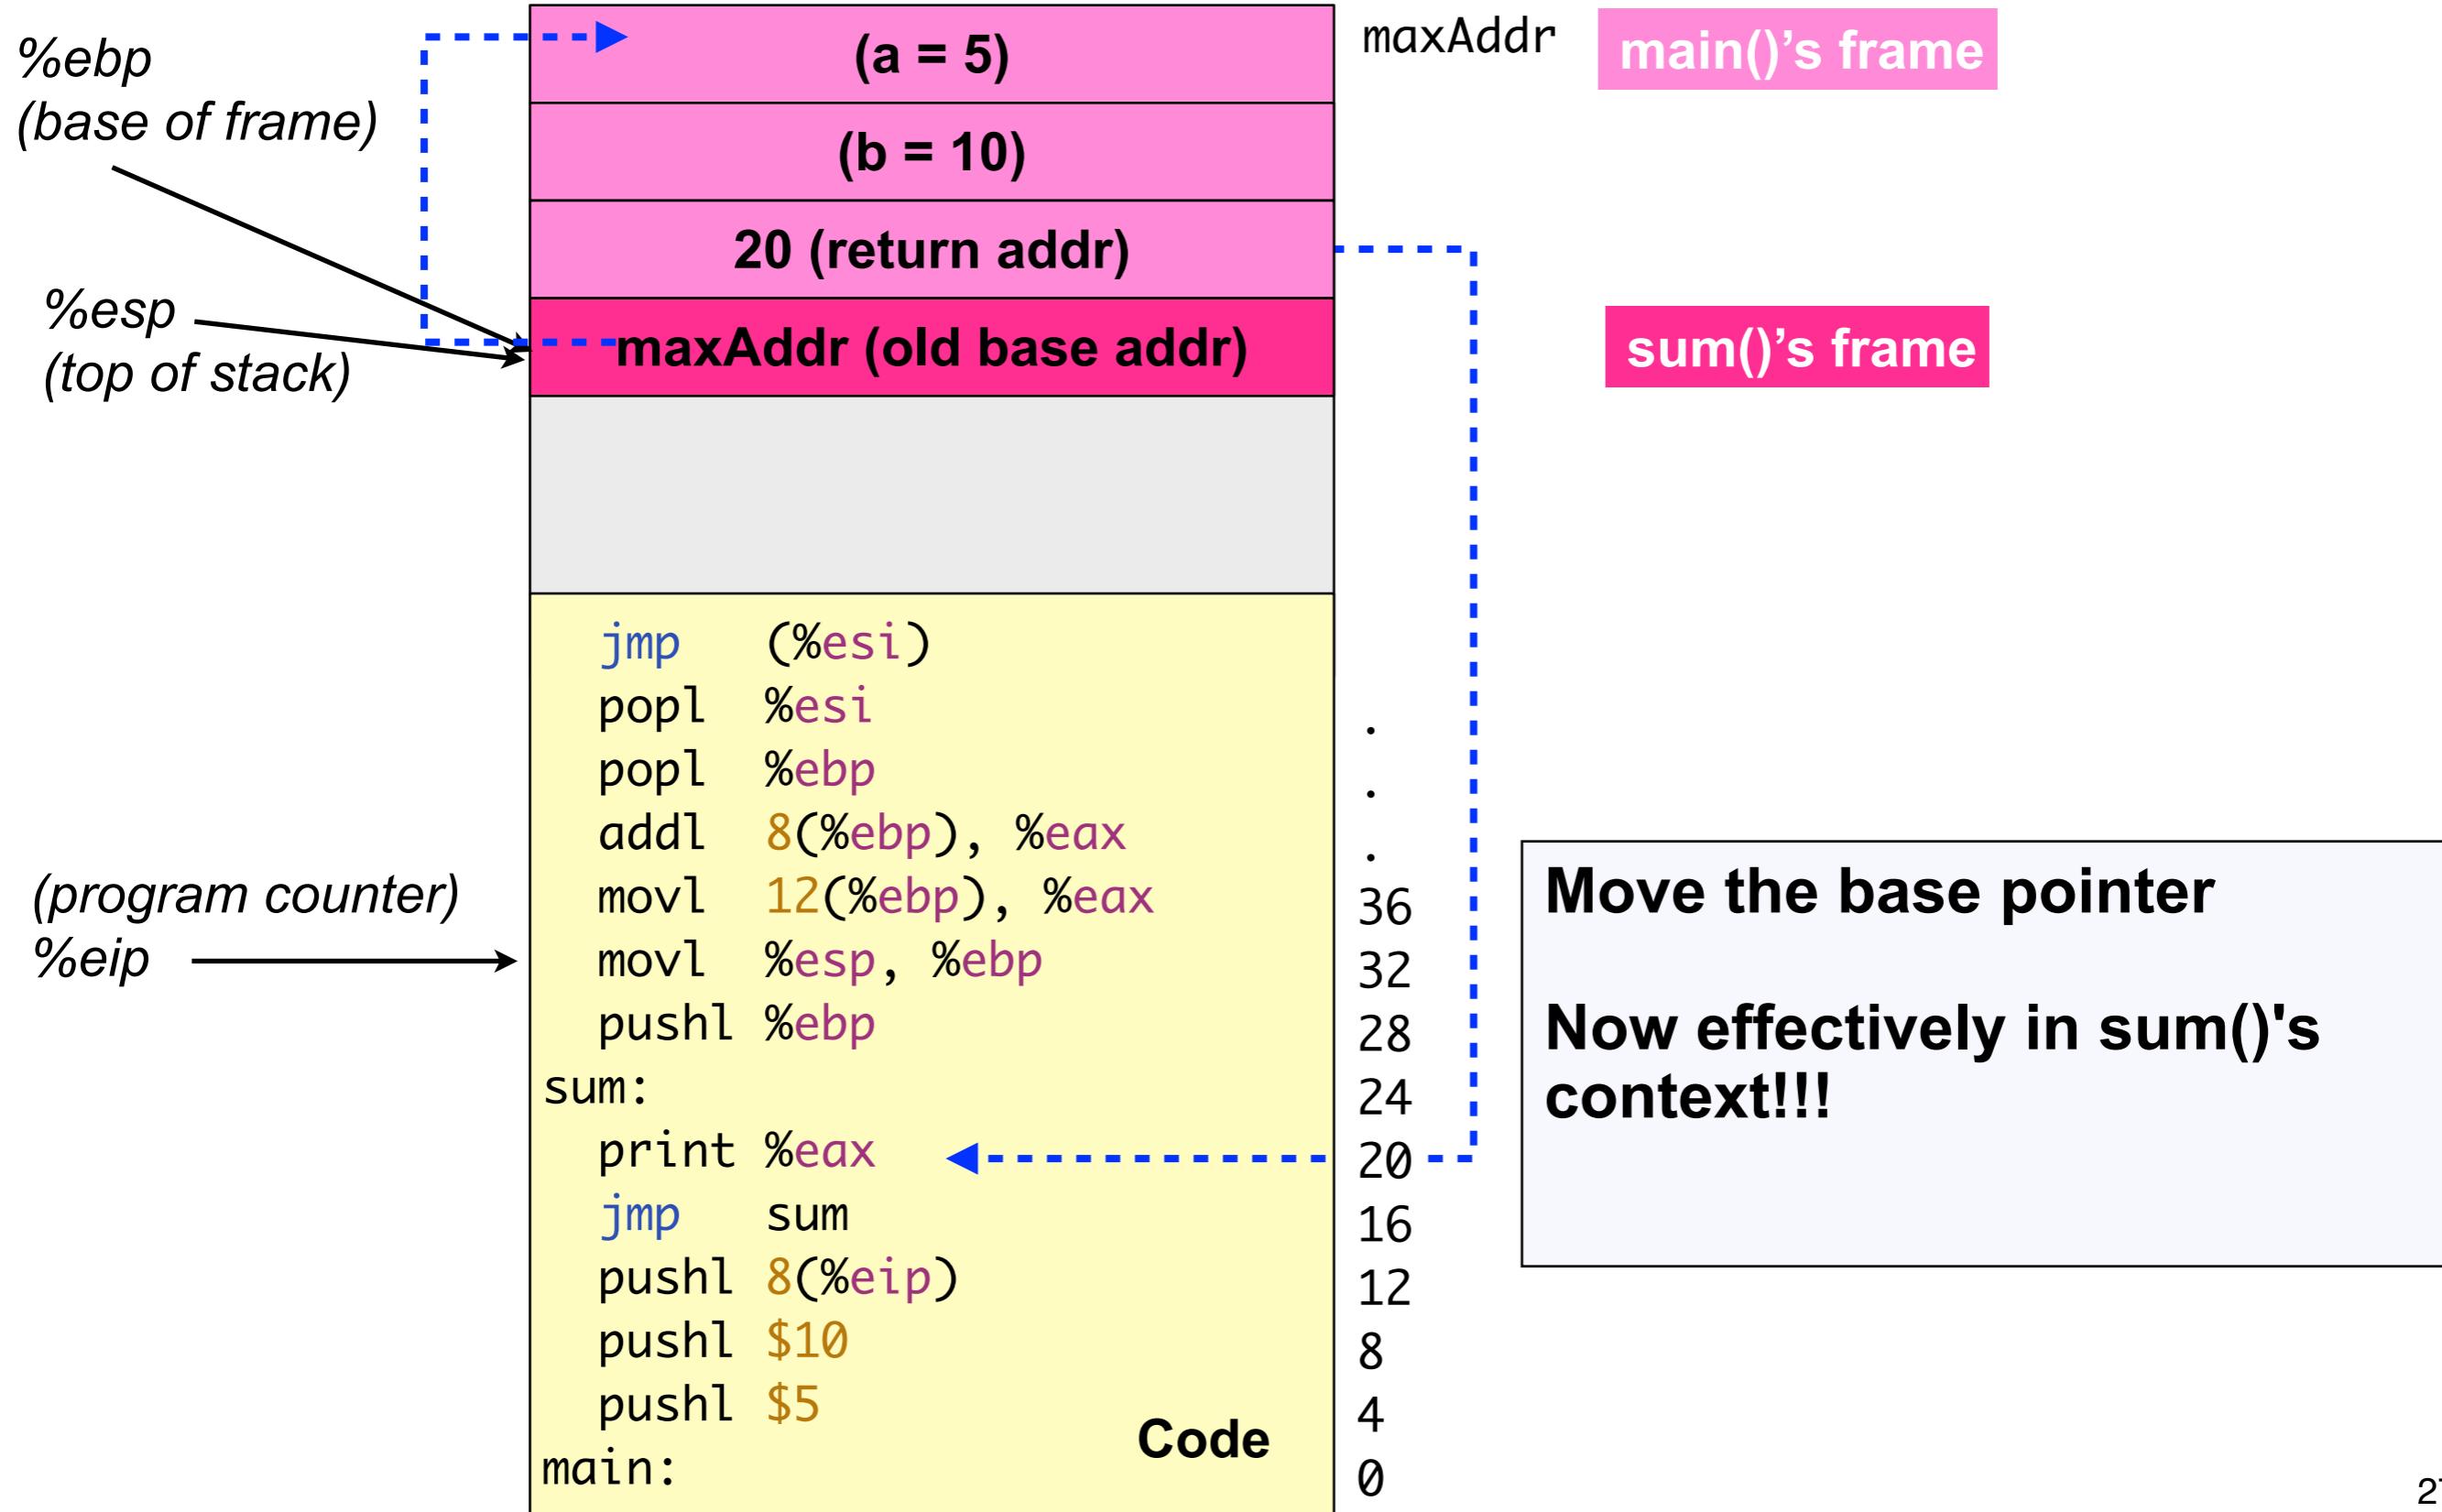
x86 Example (7)



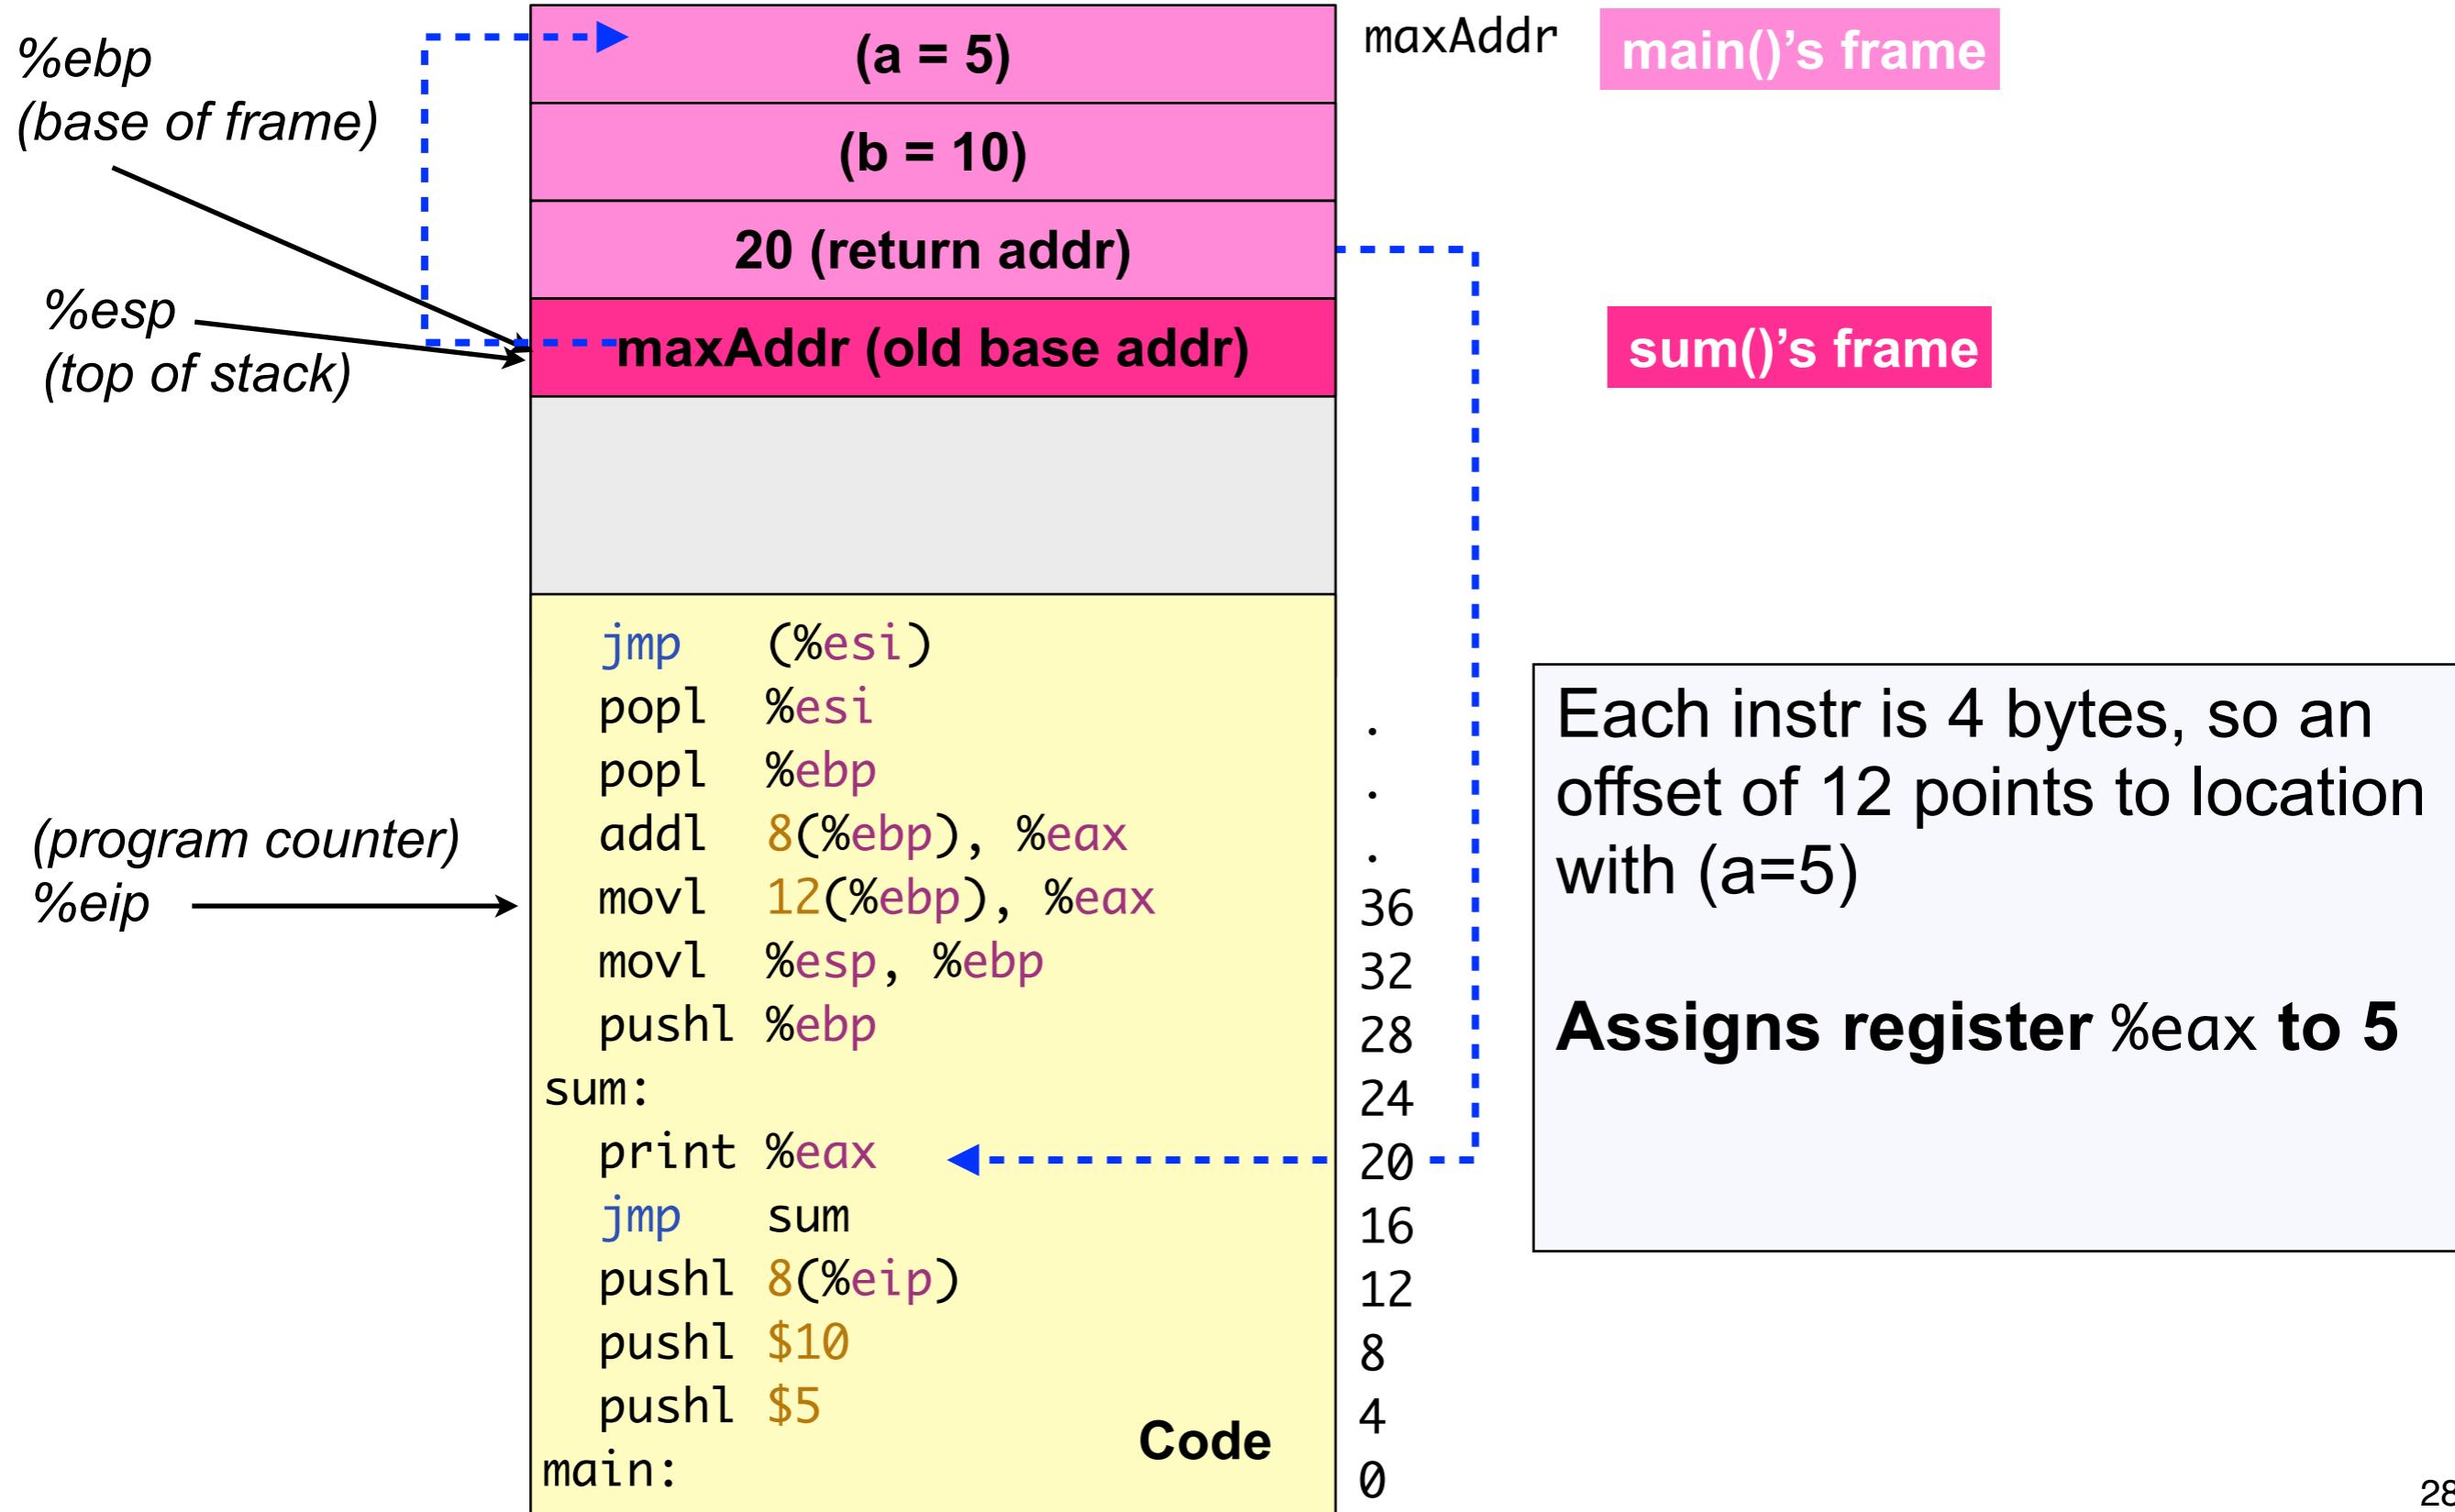
x86 Example (8)



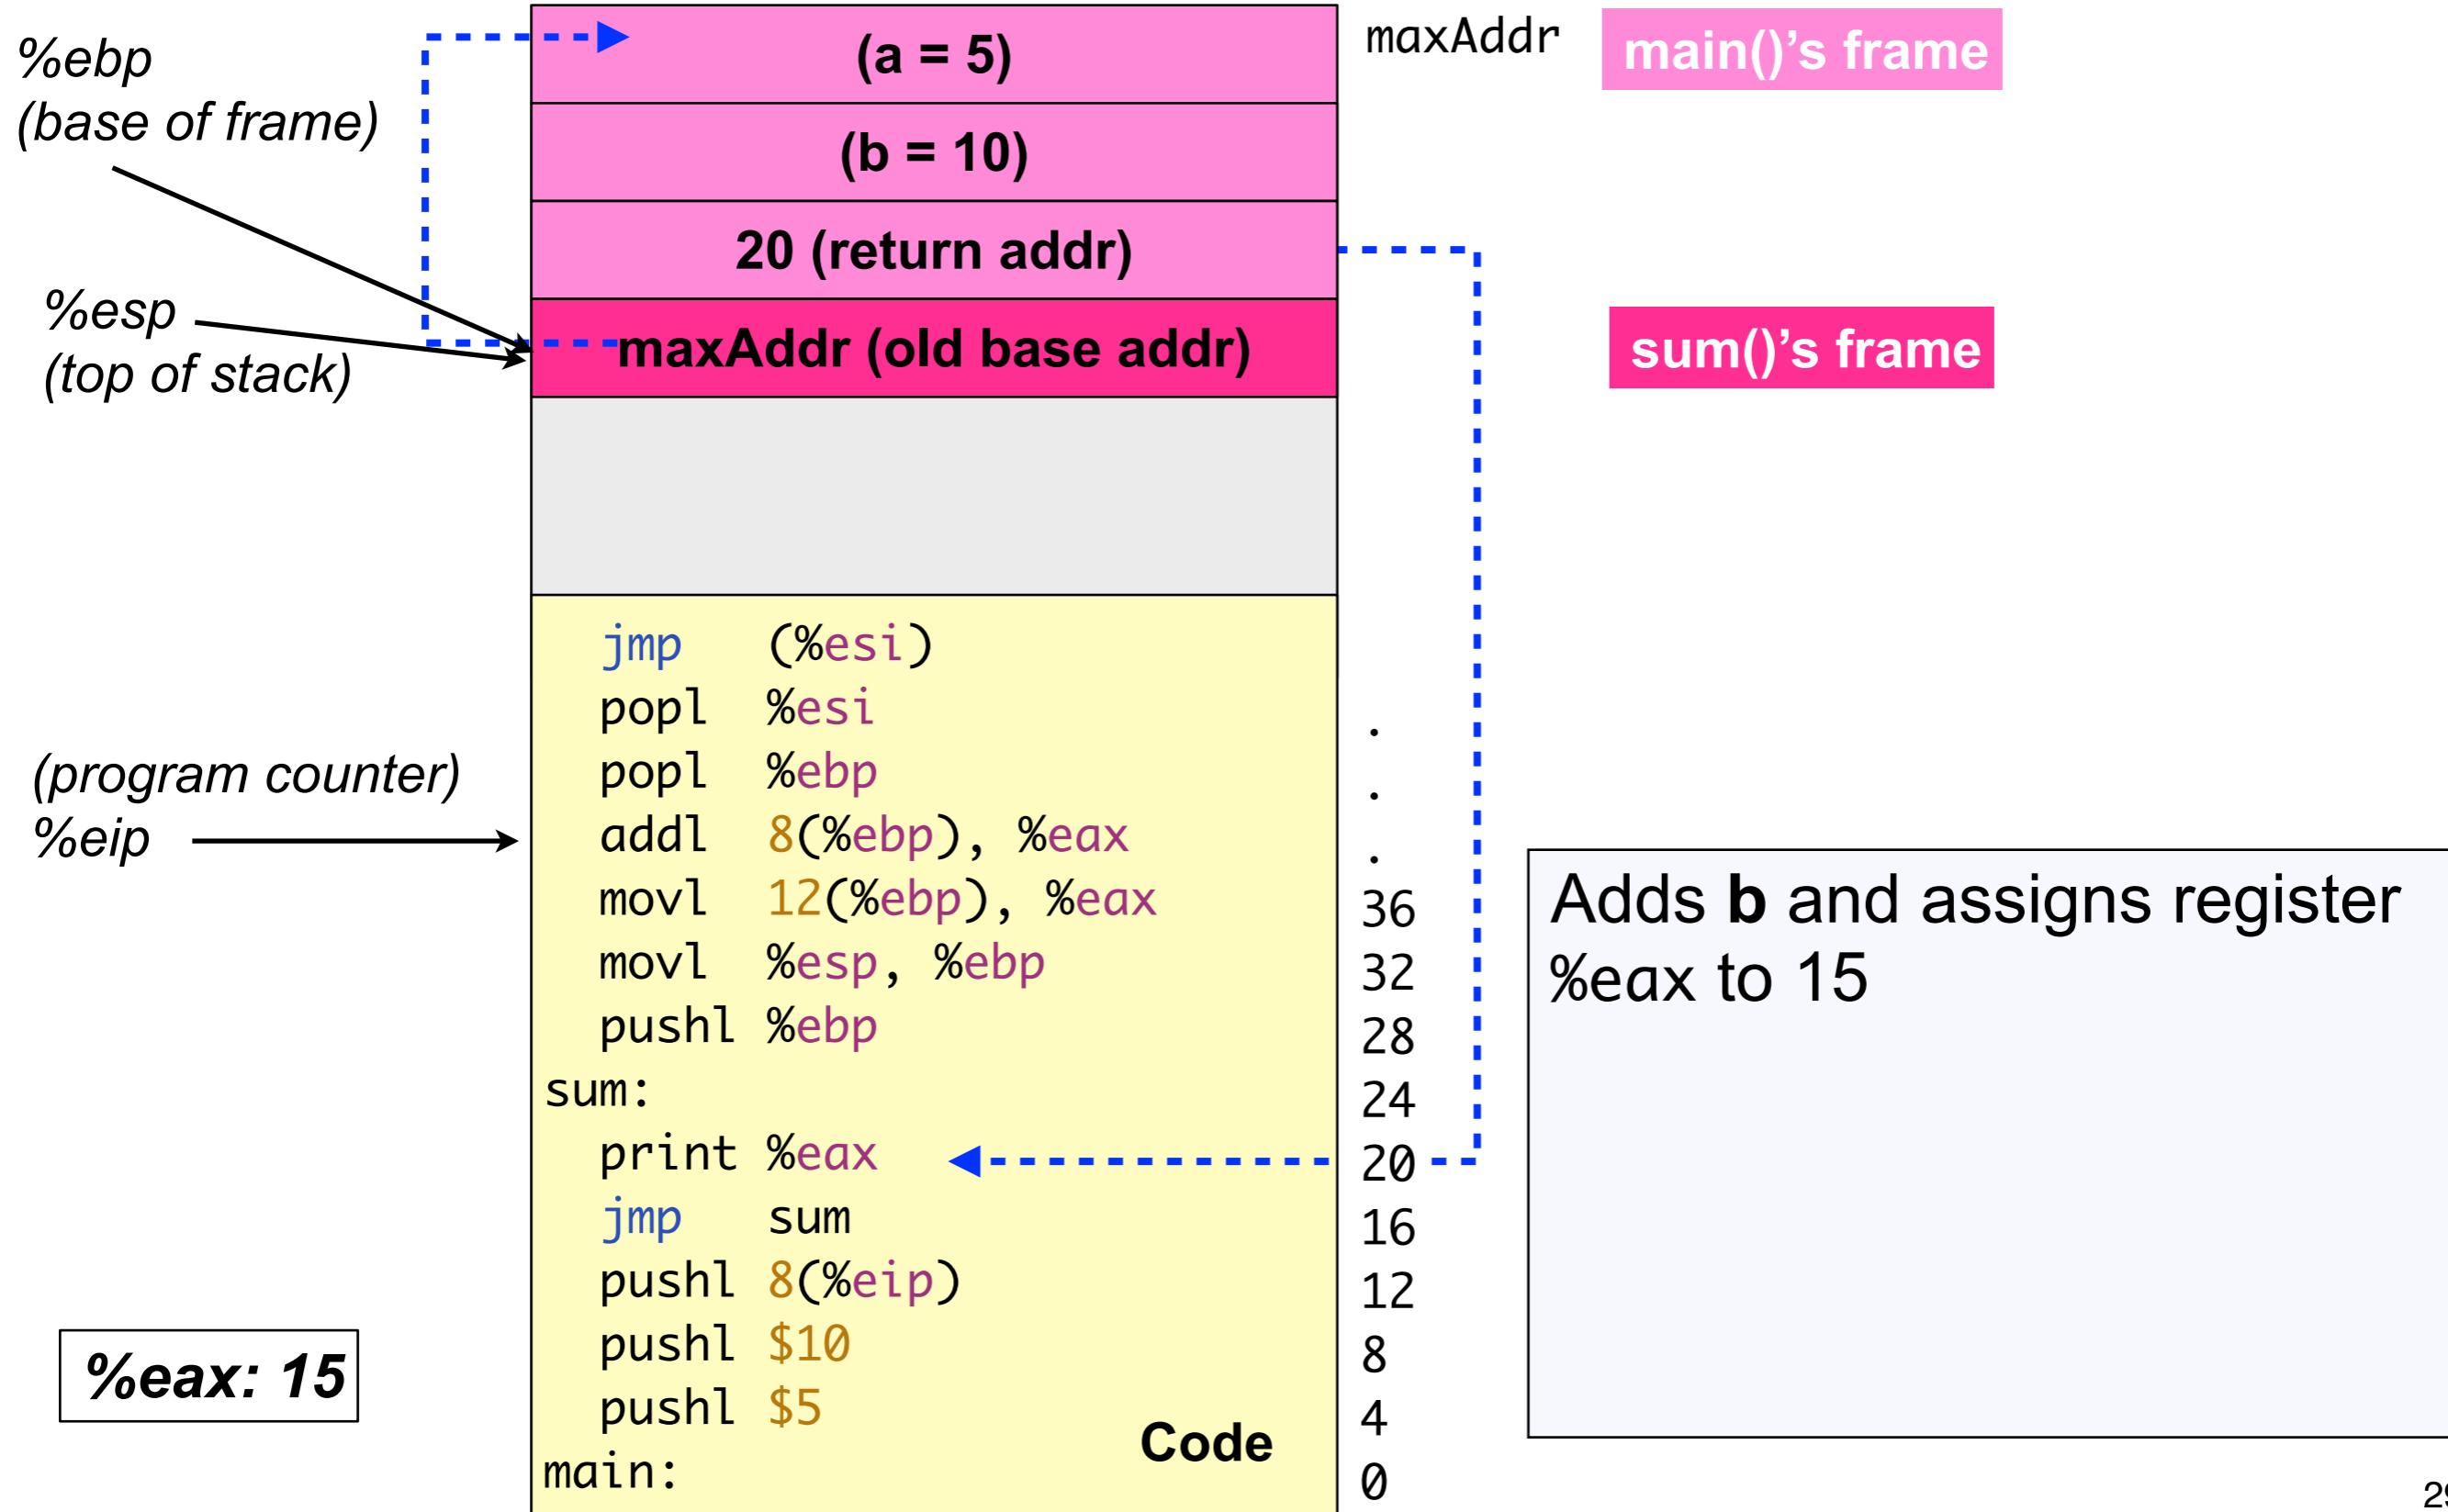
x86 Example (9)



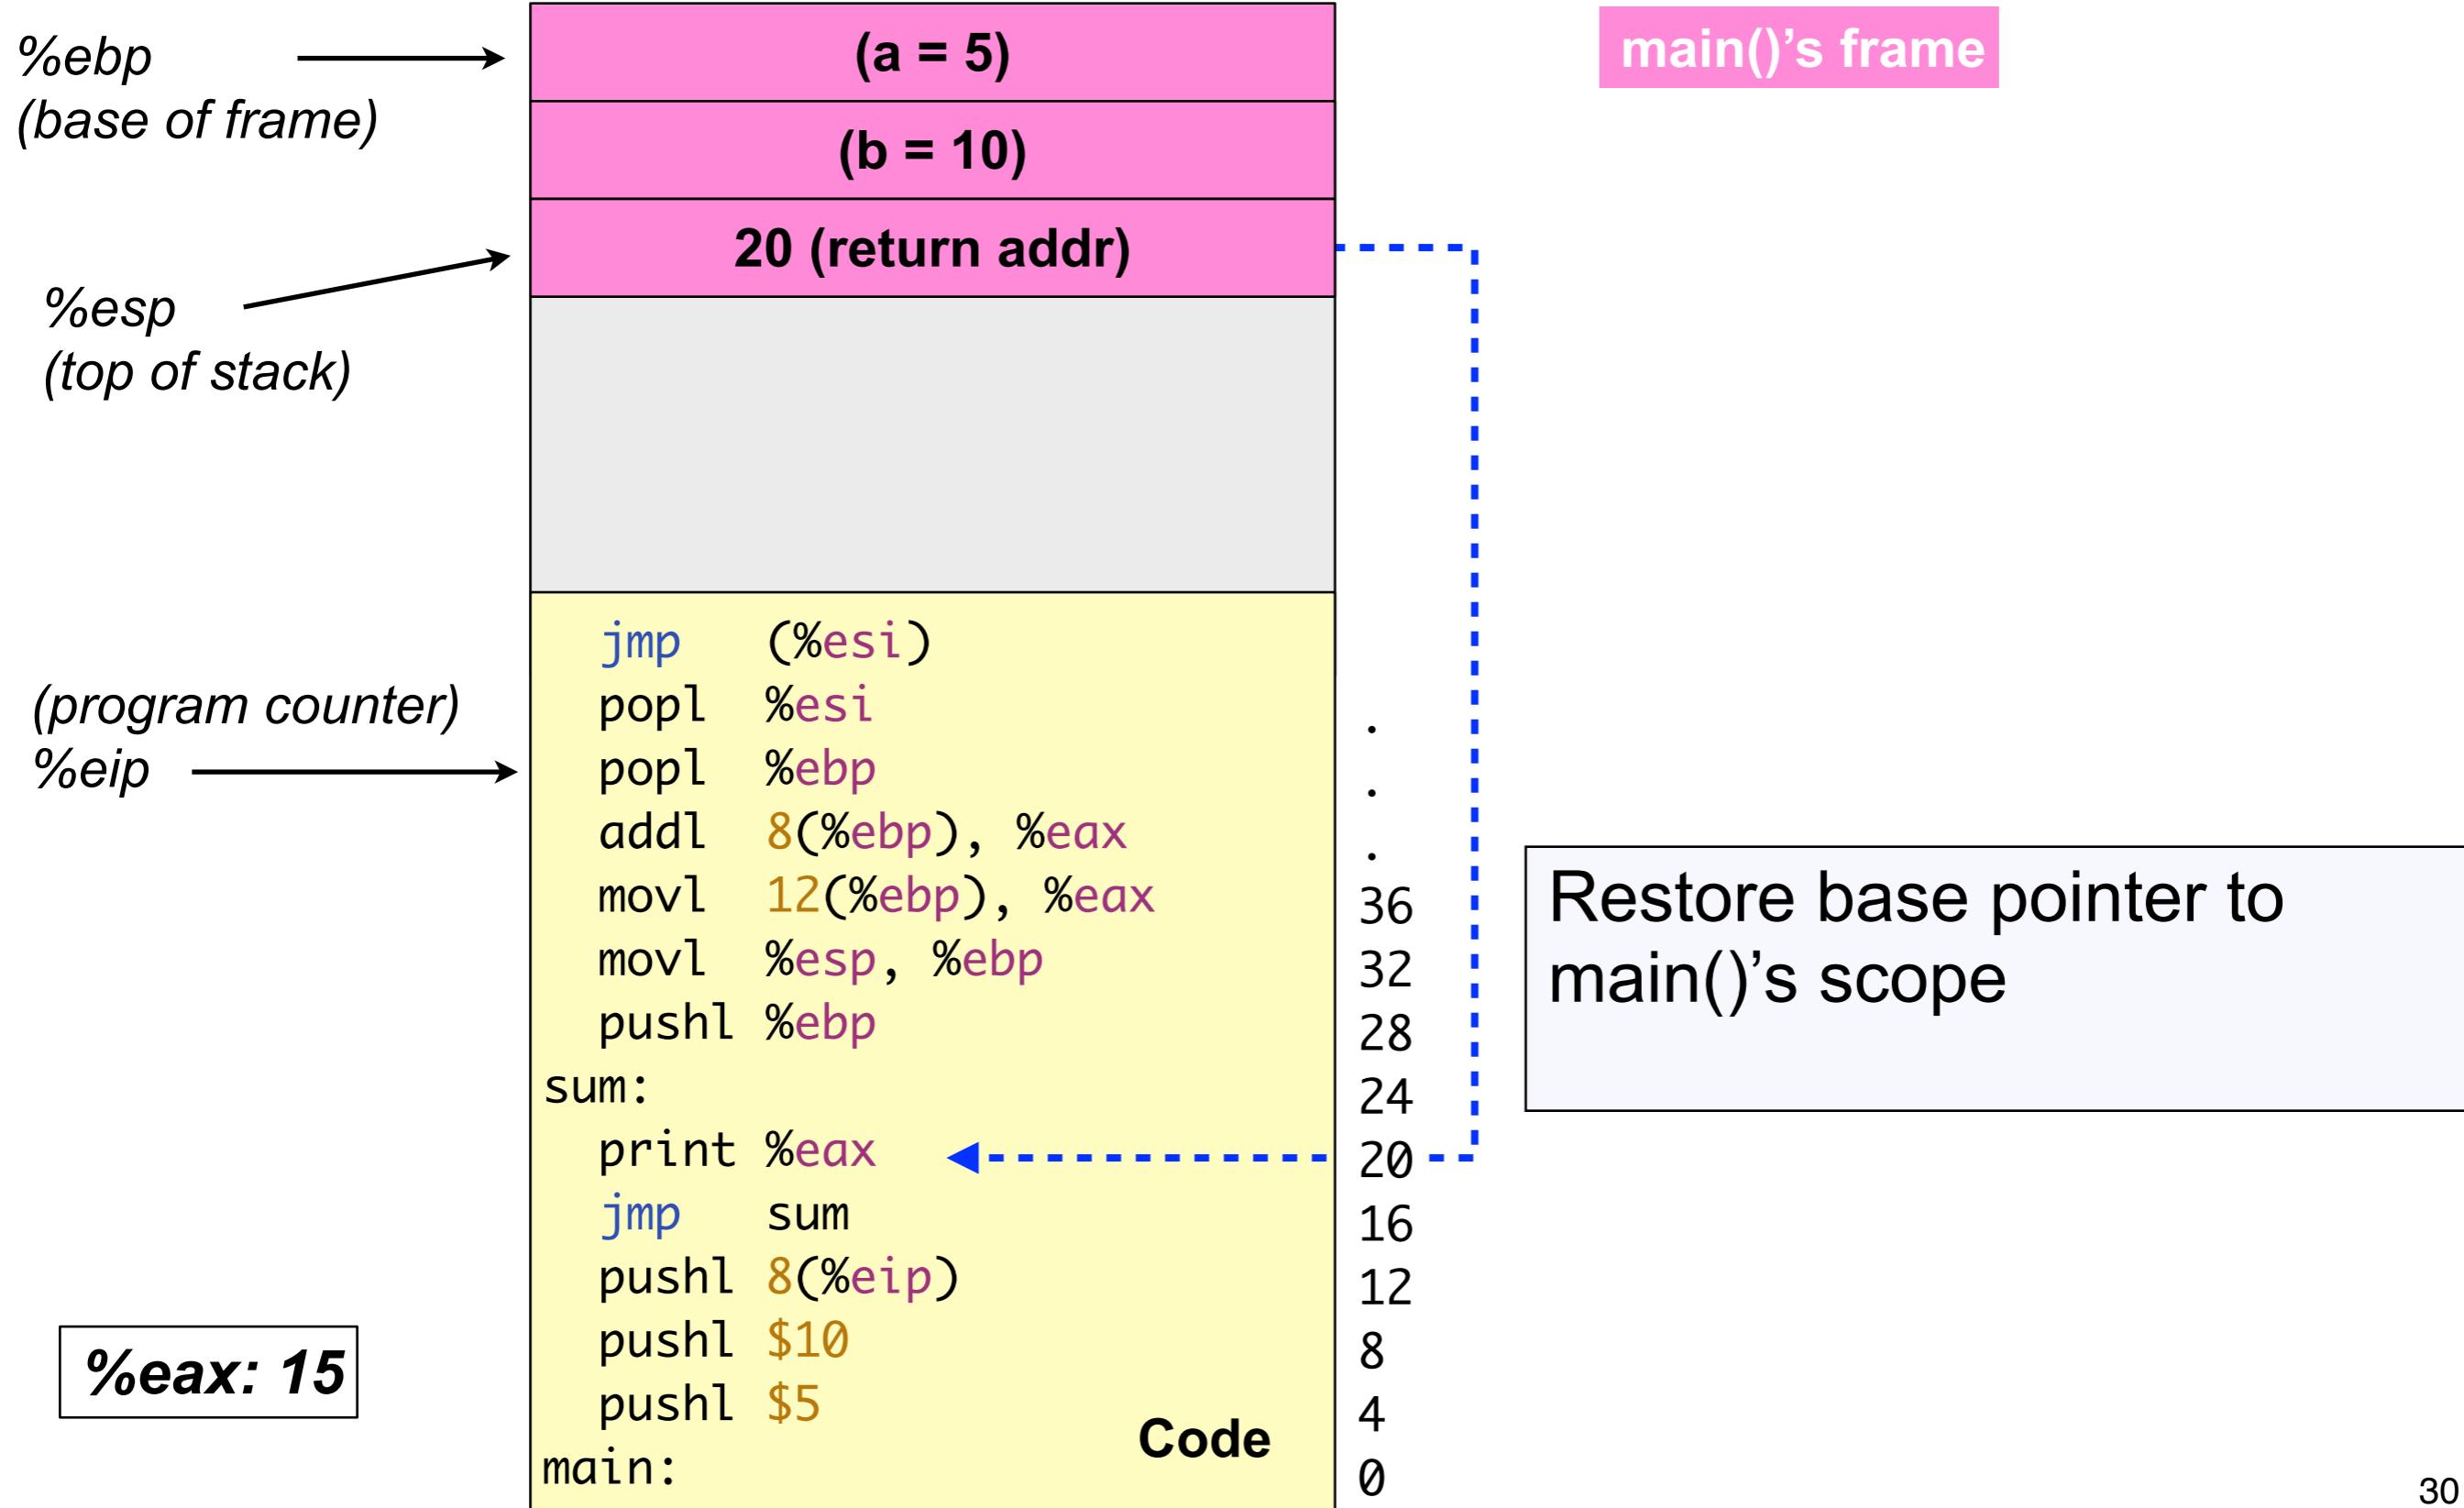
x86 Example (10)



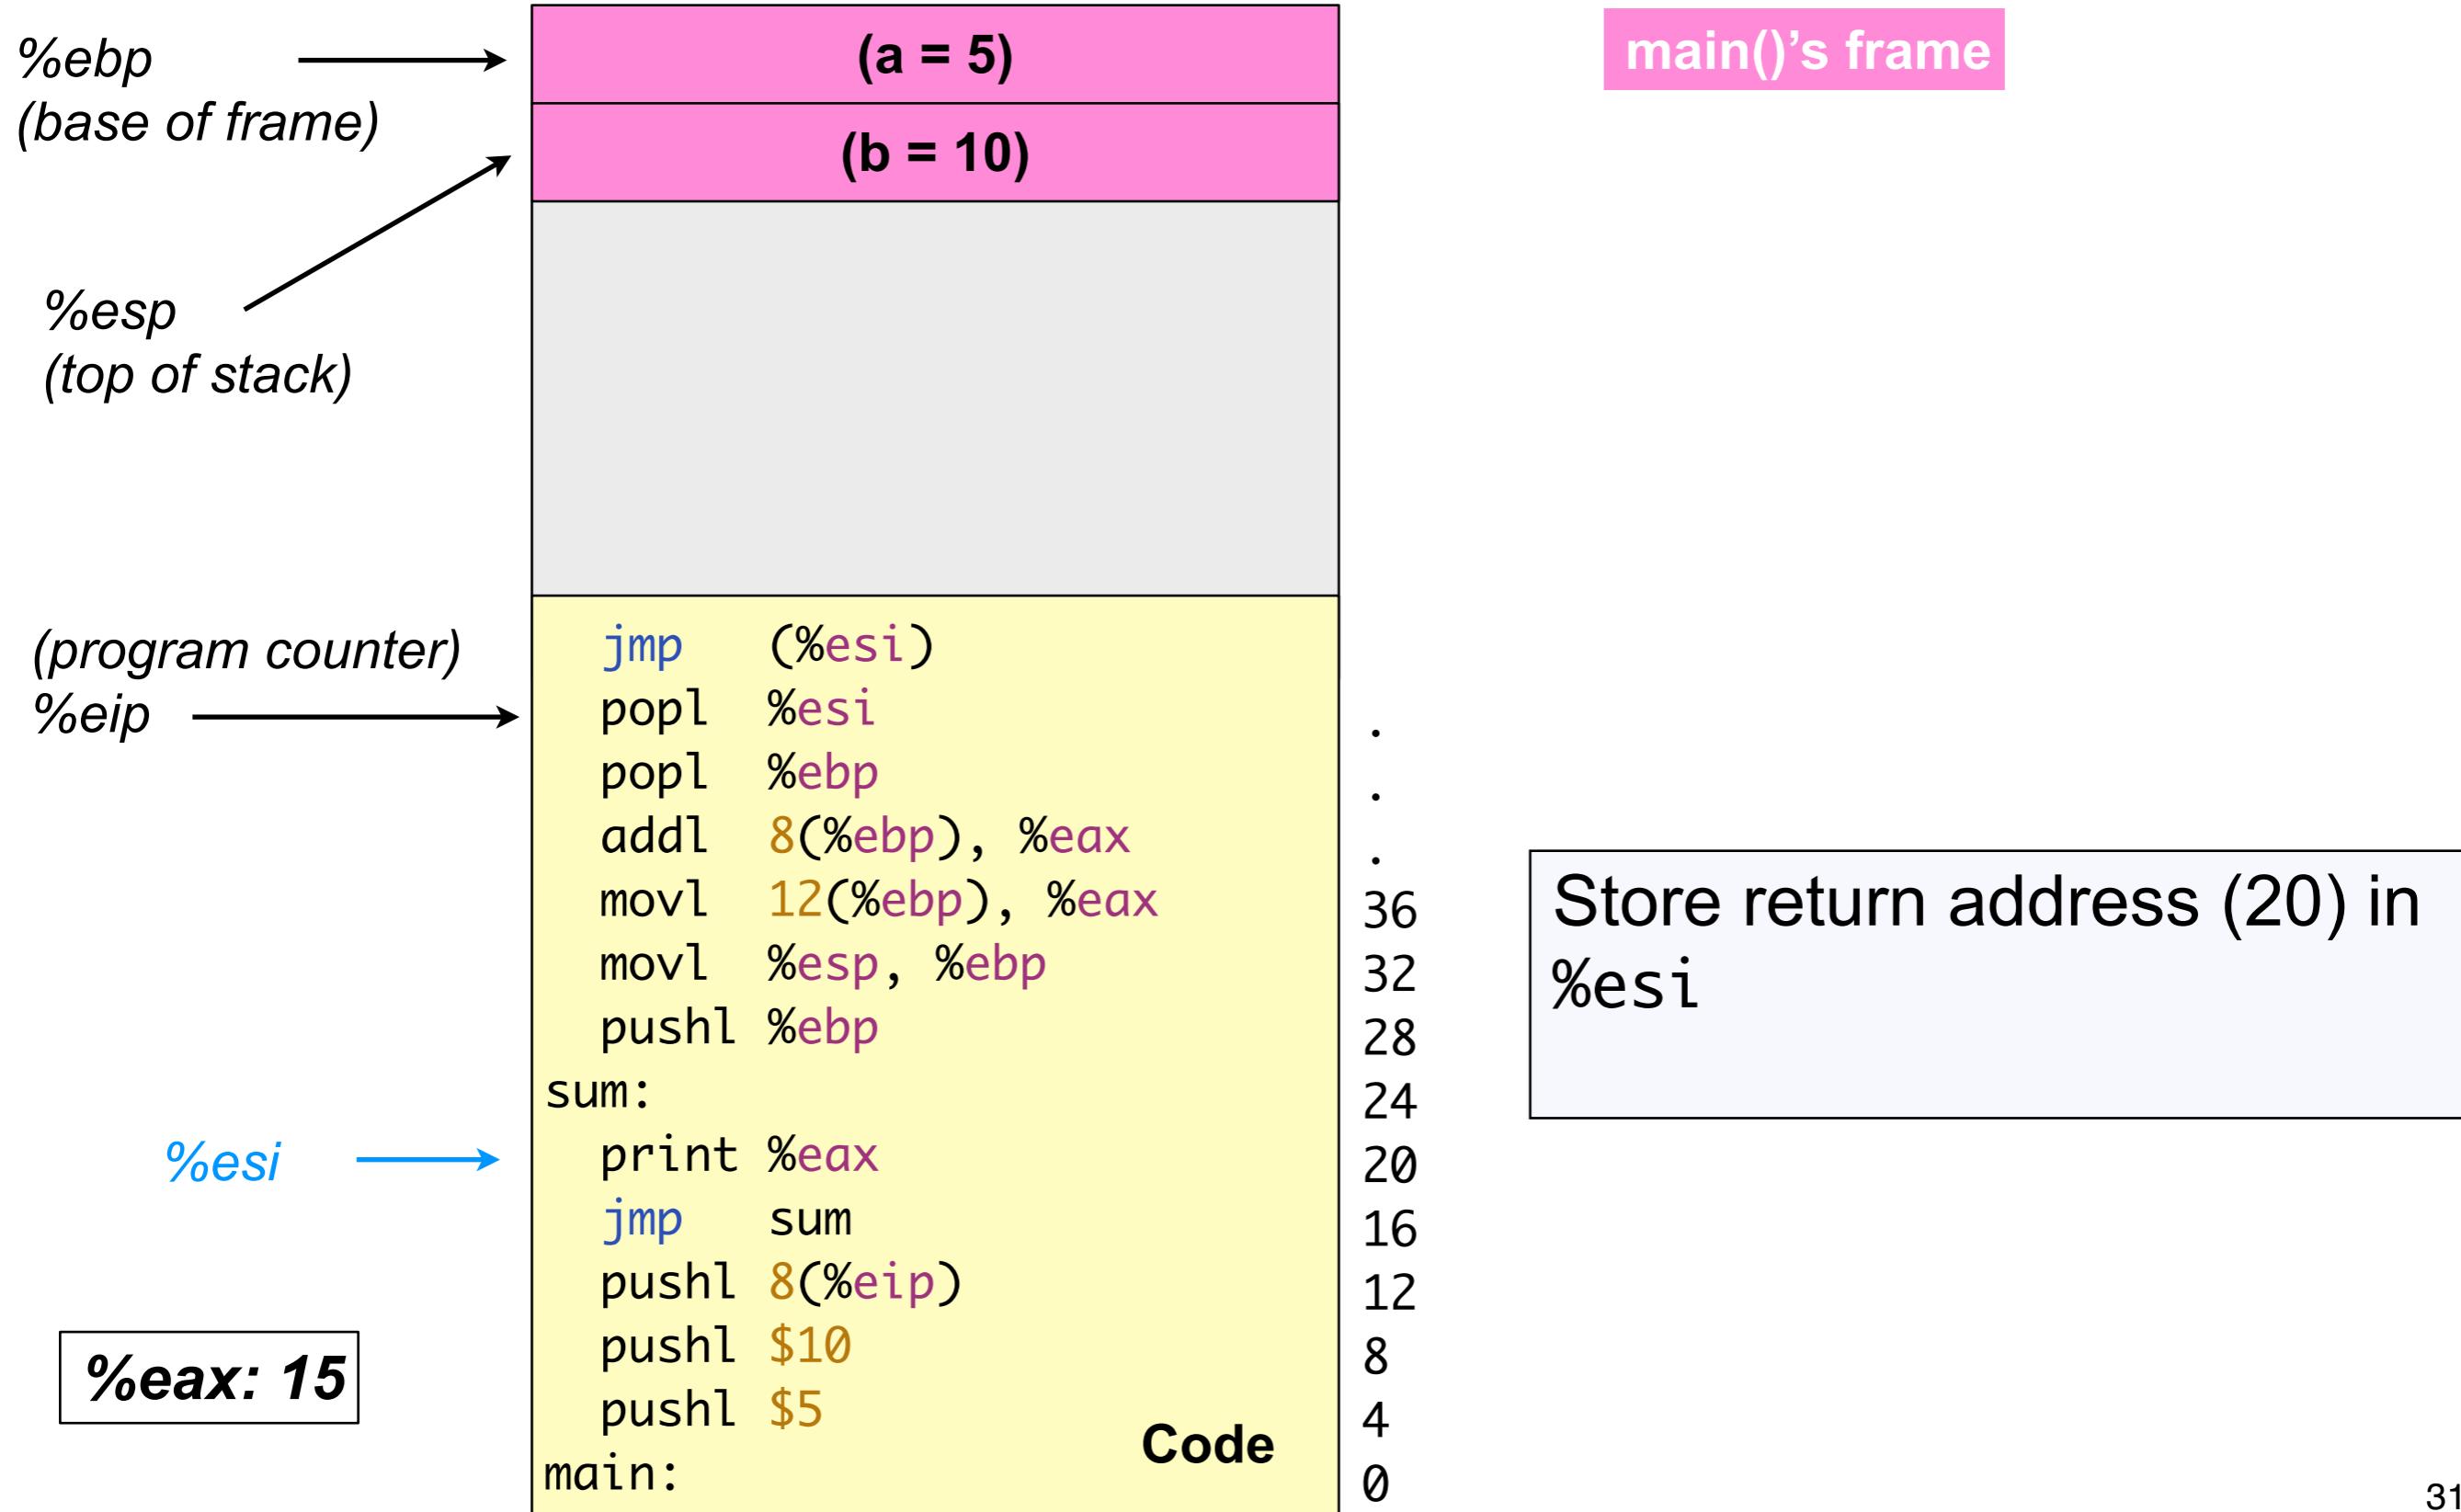
x86 Example (11)



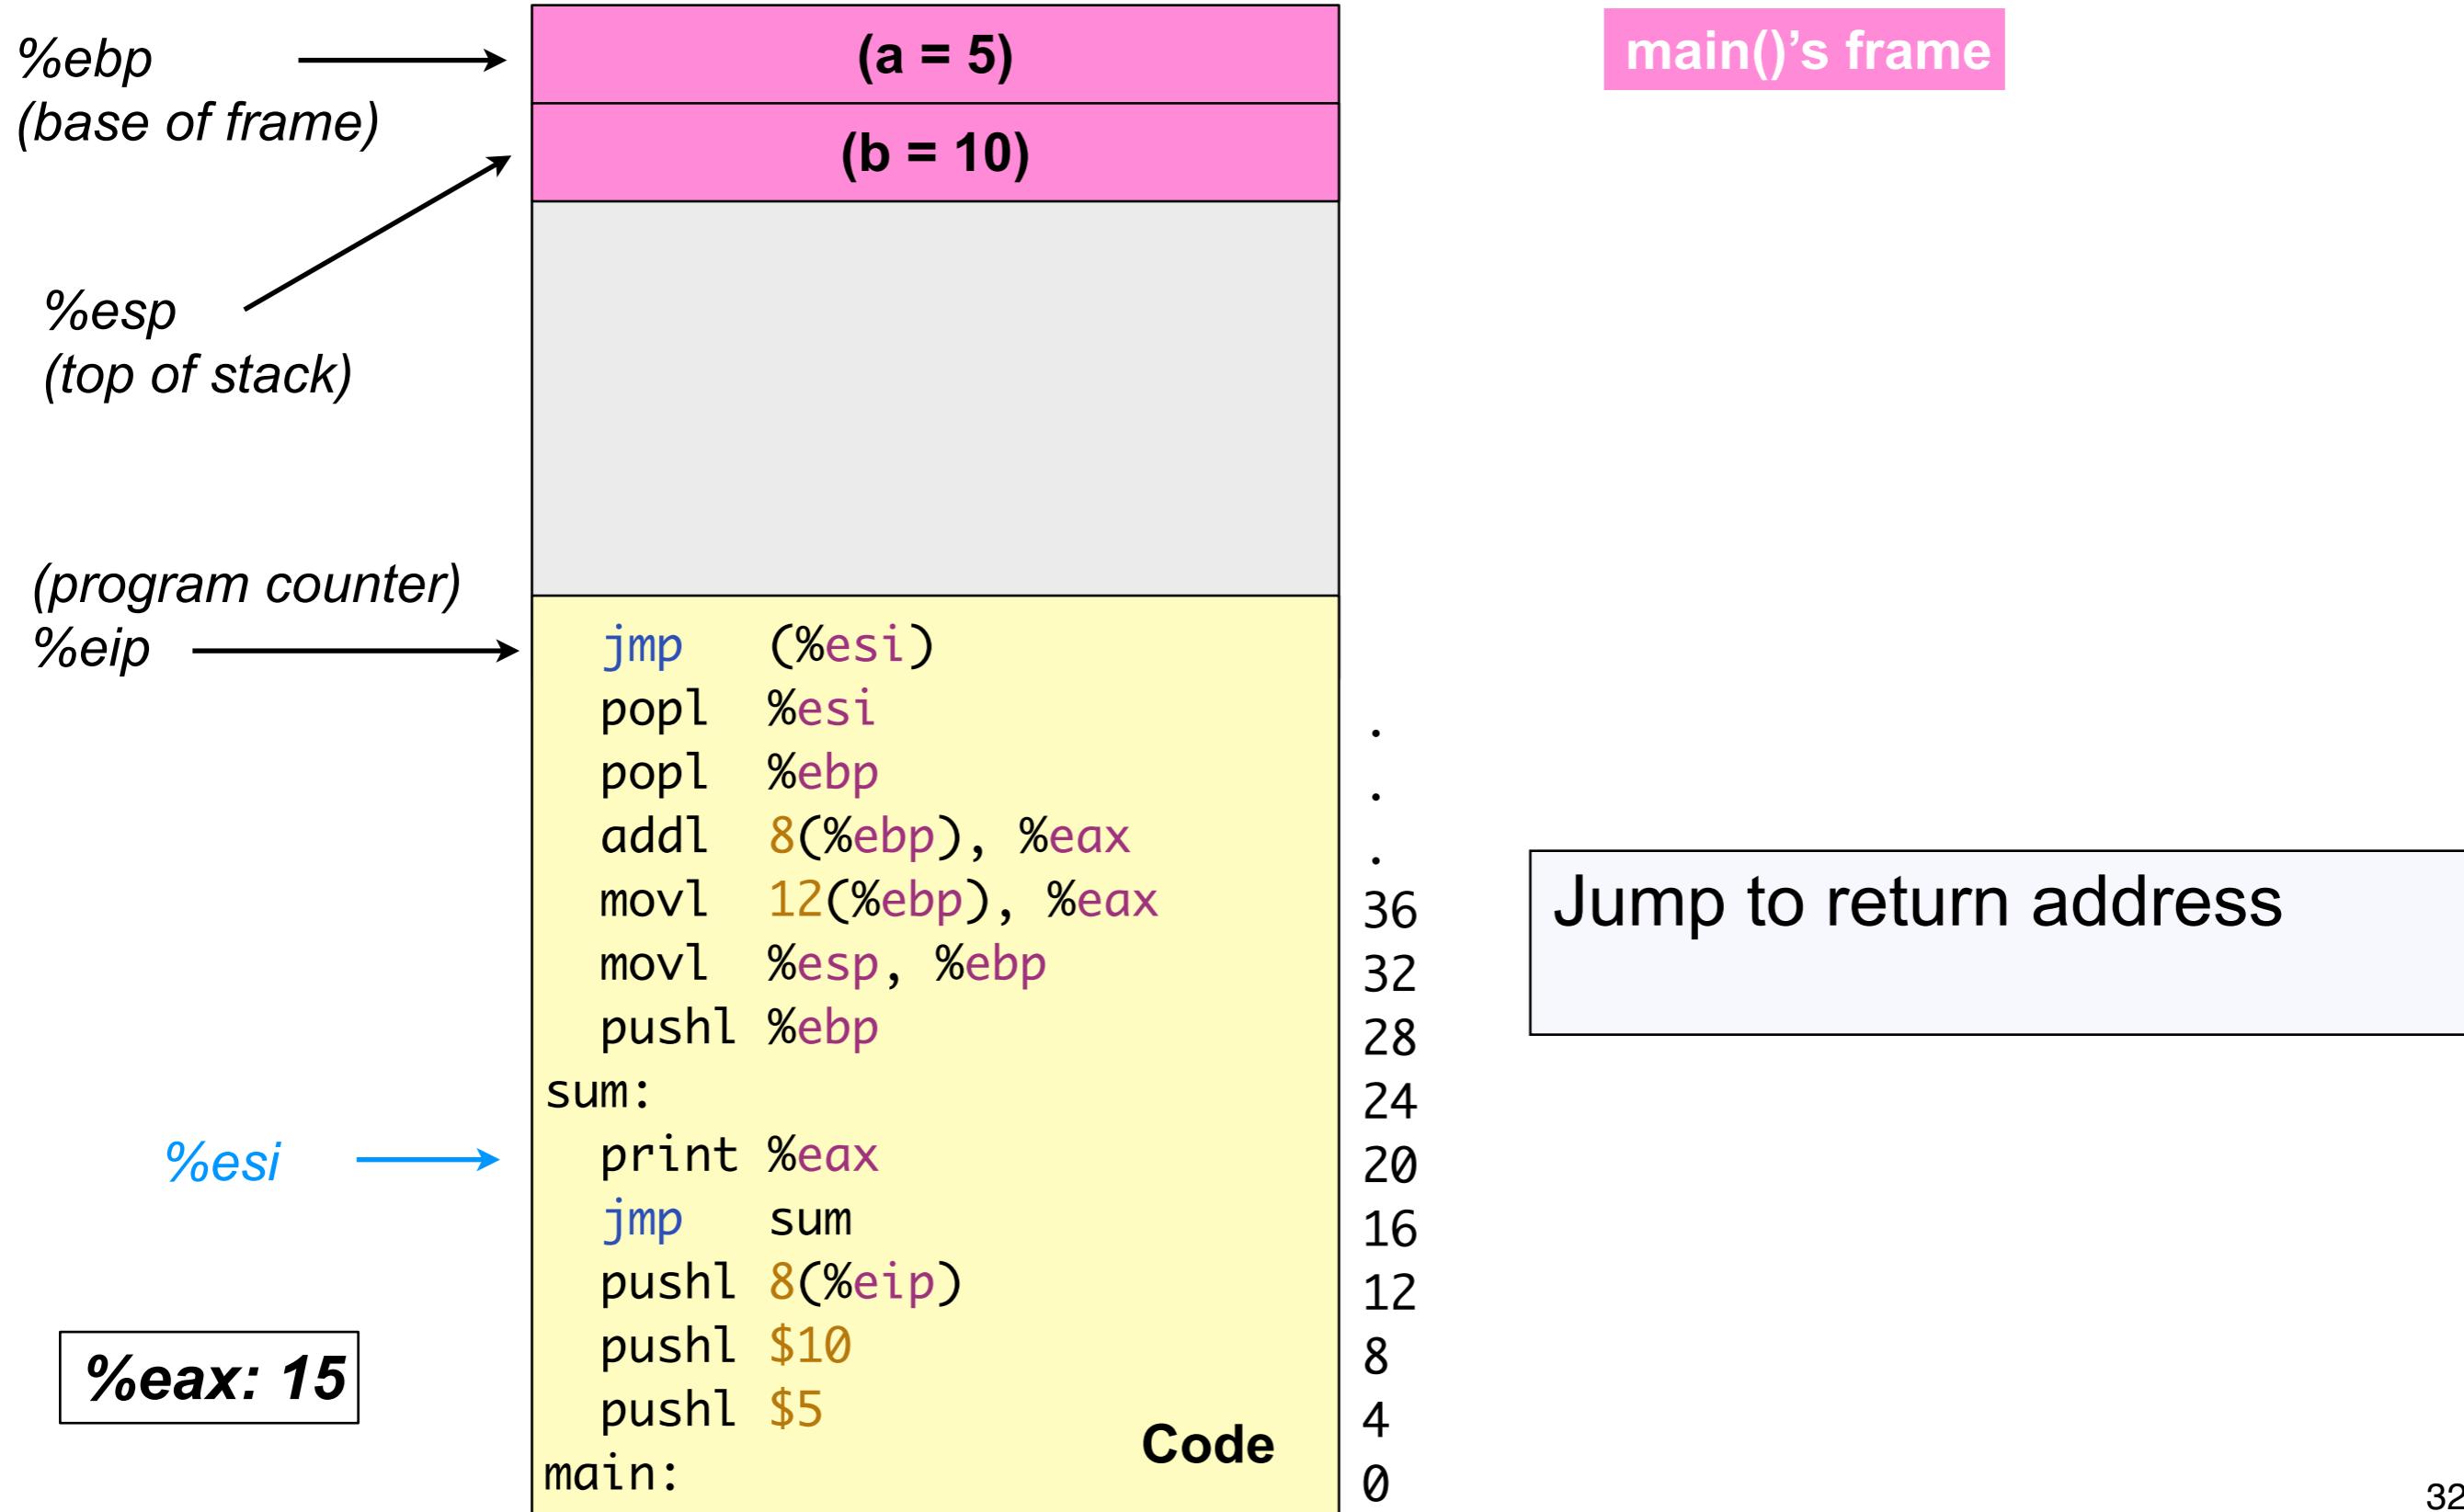
x86 Example (12)



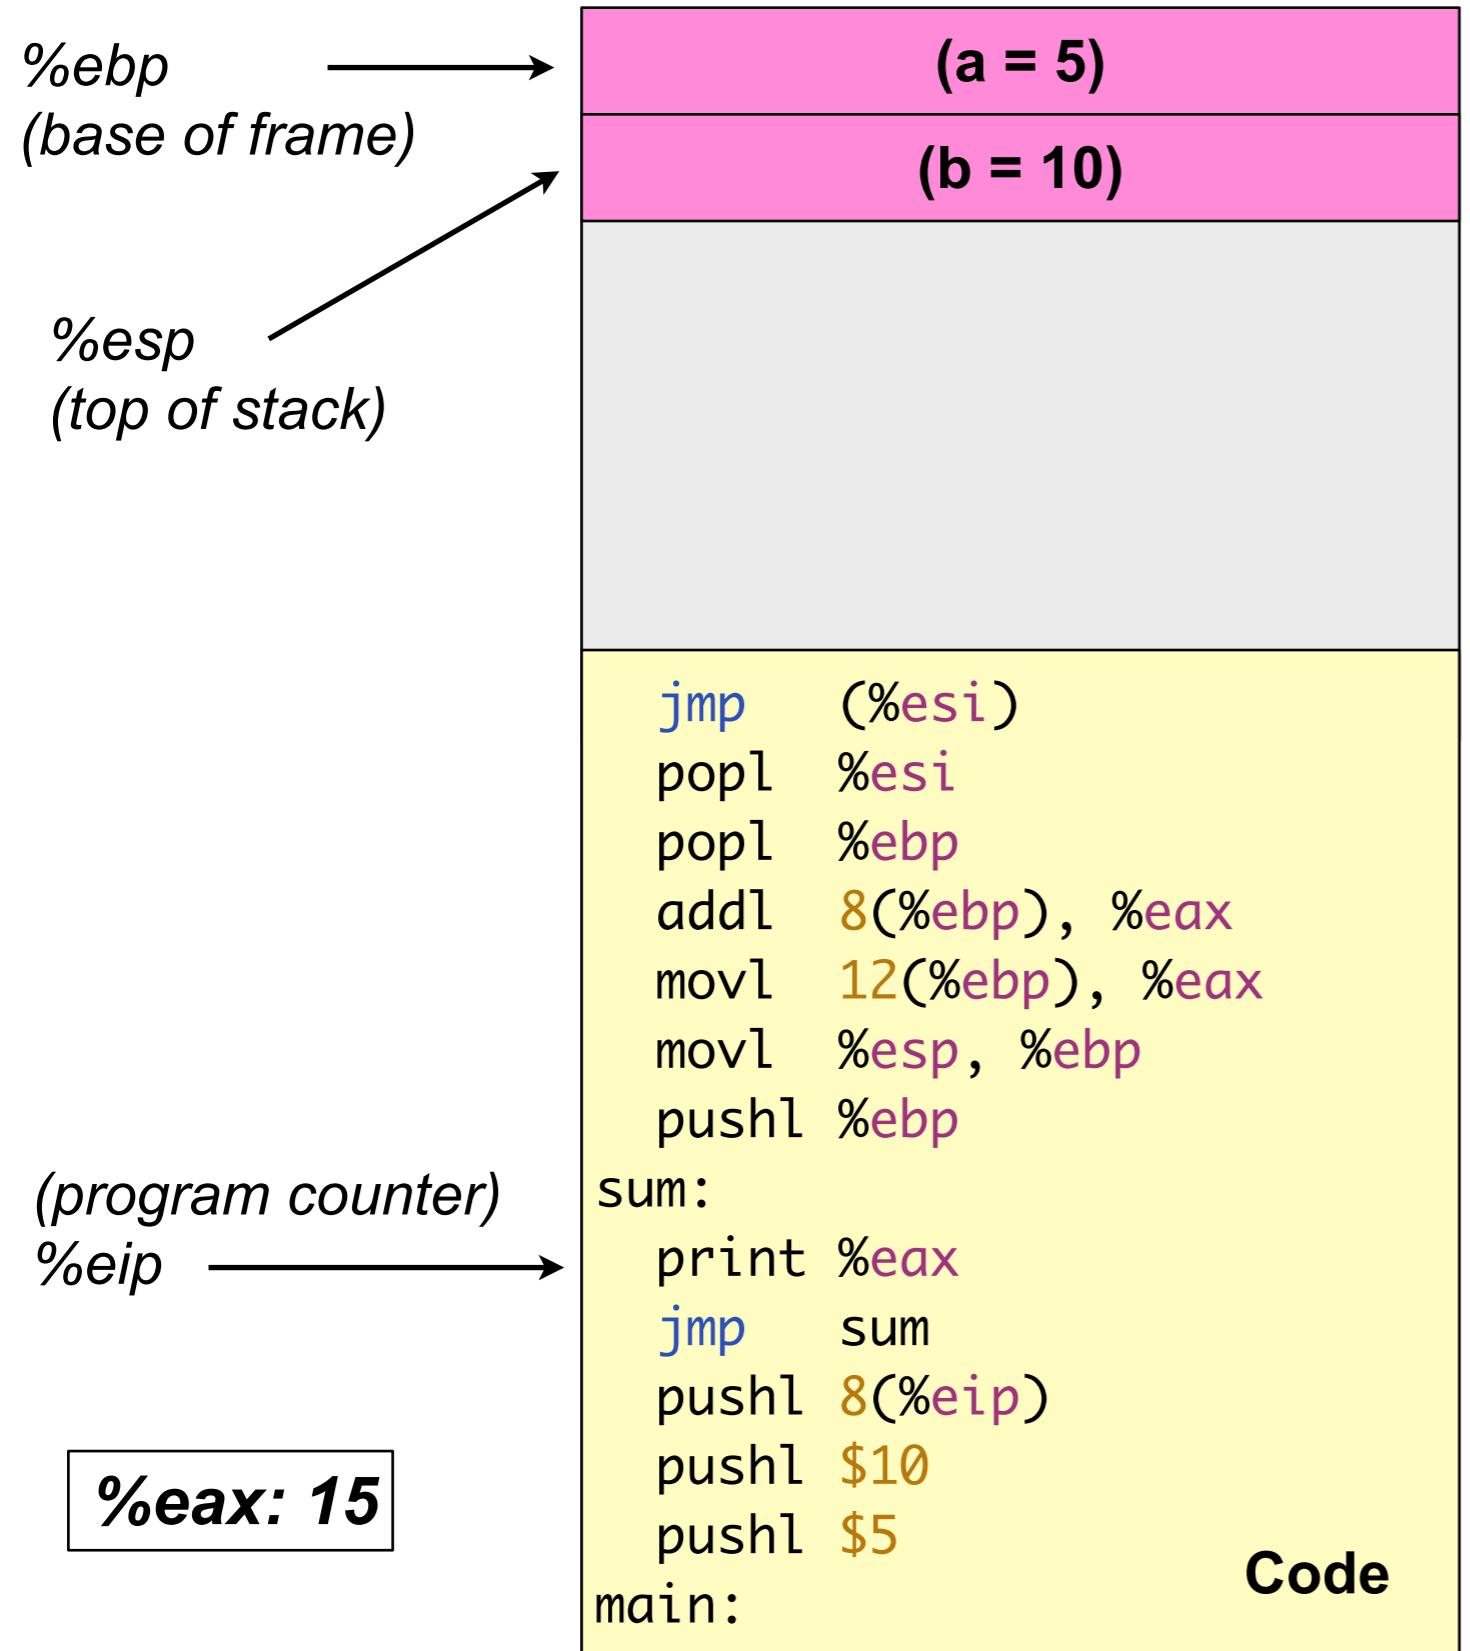
x86 Example (13)



x86 Example (14)



x86 Example (15)



main()'s frame

Continue from where we left off in main()

Print string in by %eax

Aside: Buffer Overflow Exploit

- ▶ One of the most common security vulnerabilities (still!)

```
void foo() {
    char s[8];      // buffer to store input
    getInput(s);   // put stuff in my buffer
}

void getInput(char s[]) {
    char c;
    int i = 0;
    // read keyboard input until EOF is reached
    while ((c = getchar()) != EOF) {
        s[i] = c;
        i++;
    }
    s[i] = '\0'; //terminate string
}
```

- ▶ How to avoid this? Burden is on *the programmer...*
 - Program defensively
 - How would *you* fix `getInput()`?

Goals for This Lecture

- ▶ Motivation: Concurrency
- ▶ Address Space
- ▶ Processes and the PCB
- ▶ Process Management in C
 - Creation: fork()
 - Basic Synchronization: wait()
 - Running another program: exec()

Process Control Block (PCB)

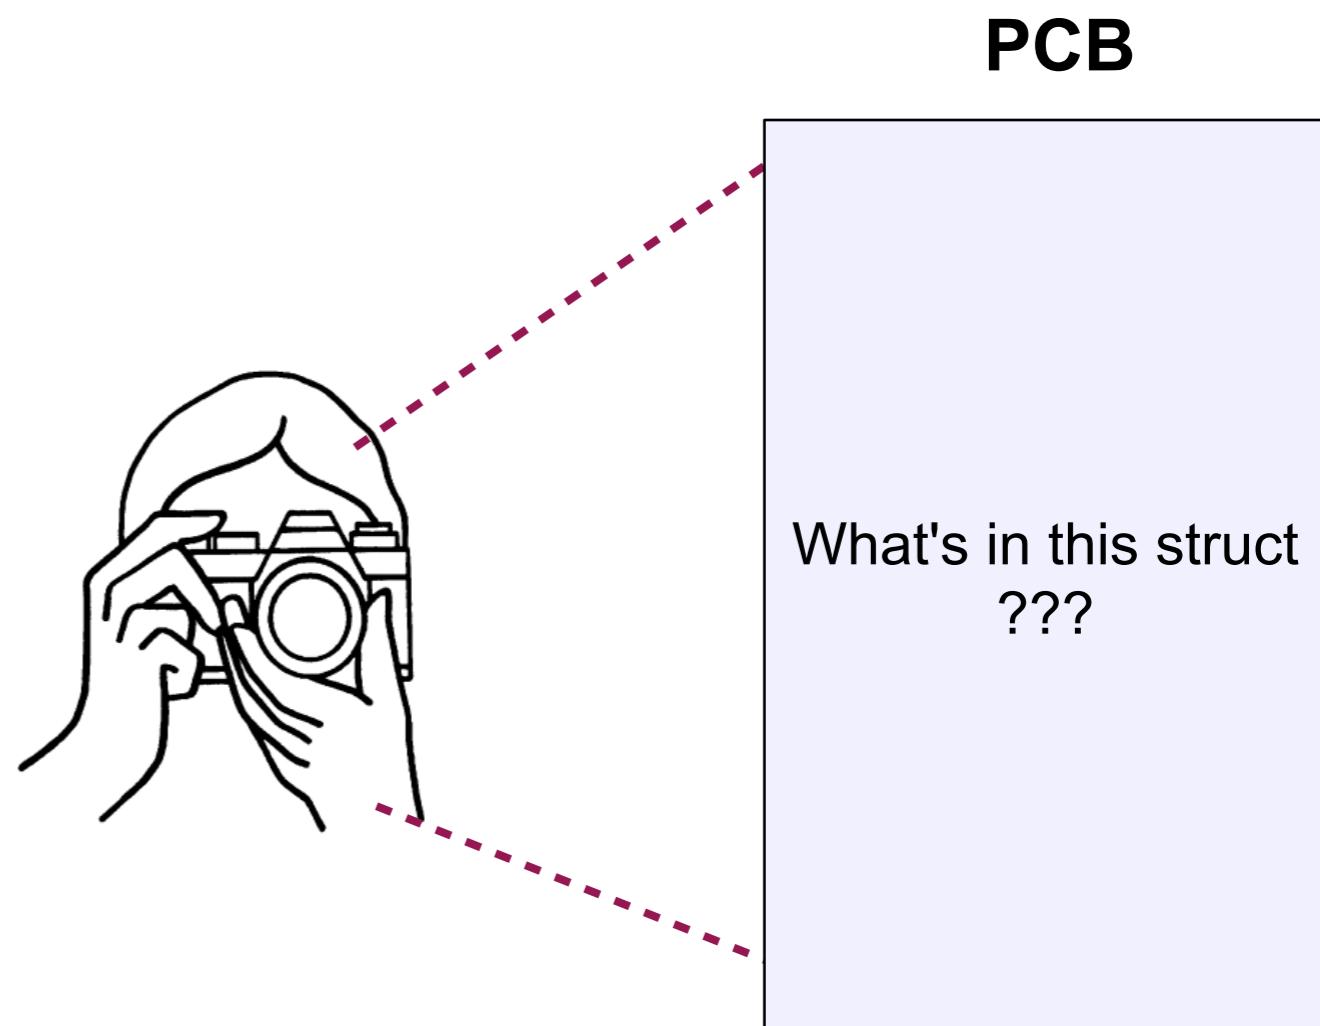
- ▶ We need to capture the state of processes at arbitrary points in time

- *How is process state stored?*

- ▶ OS represents each process as a

Process Control Block (PCB)

- A "snapshot" of process execution
 - What info would a PCB need to store?



Context Switch: Abstractly...

```
void contextSwitch(PCB *currentPCB, PCB *nextPCB) {  
  
    /* save state of current PCB */  
    currentPCB->r0 = CPU[r0];  
    currentPCB->r1 = CPU[r1];  
    currentPCB->r2 = CPU[r2];  
    //...  
    currentPCB->sp = CPU[sp]; // save stack pointer  
    currentPCB->pc = CPU[pc]; // save program counter  
  
    /* save current address space to disk */  
    copy_to_disk(currentPCB->addr_sp);  
  
    /* load new address space from disk */  
    copy_from_disk(nextPCB->addr_sp);  
  
    /* restore state of next PCB */  
    CPU[r0] = nextPCB->r0;  
    CPU[r1] = nextPCB->r1;  
    CPU[r2] = nextPCB->r2;  
    //...  
    CPU[sp] = nextPCB->sp;  
    CPU[pc] = nextPCB->pc;  
}
```



Context Switch: Food for Thought

- ▶ Saving and Loading PCBs aren't too much work on their own...
- ▶ Context switches don't seem like much code but they *are* expensive!
 - First, switches triggered by either a timer interrupt or a trap (a process yields the CPU), so it's slow to handle.
 - Also, picking next process to run (called scheduling) is expensive too.
- ▶ A single context switch in Linux on a modern CPU:
 - Requires ~10,000 CPU cycles
 - In an ideal situation, 1 instruction could be executed per cycle!
 - That's a lot of squandered CPU cycles!

PCB in Linux (No you don't have to know this)

```

struct task_struct {
  long state; ← Each process is in one of several states (next)
  long counter;
  long priority;
  unsigned long signal;
  unsigned long blocked;
  unsigned long flags;
  int errno;
  long debugreg[8];
  struct exec_domain *exec_domain;

  /* various fields */
  struct linux_binfmt *binfmt;
  struct task_struct *next_task, *prev_task;
  struct task_struct *next_run, *prev_run;
  unsigned long saved_kernel_stack;
  unsigned long kernel_stack_page;
  int exit_code, exit_signal;
  unsigned long personality;
  int dumpable:1;
  int did_exec:1;
  int pid; ← Each process has a unique ID
  int pgrp;
  int tty_old_pgrp;
  int session;
  int leader;
  int groups[NGROUPS];
  struct task_struct *p_opptr, *p_pptr, *p_cptr, *p_ysptr, *p_osptr;
  struct wait_queue *wait_chldexit;
  unsigned short uid,euid,suid,fsuid; Start time; time in user mode
  unsigned short gid,egid,sgid,fsgid; time in kernel mode; etc.

  unsigned long timeout, policy, rt_priority;
  unsigned long it_real_value, it_prof_value, it_virt_value;
  unsigned long it_real_incr, it_prof_incr, it_virt_incr;
  struct timer_list real_timer;
  long utime, stime, cutime, cstime, start_time;
}

```

Each process may have a list of open files.

Each process has a pointer to its address space.
Also SP, BP, PC, all stored in *mm struct

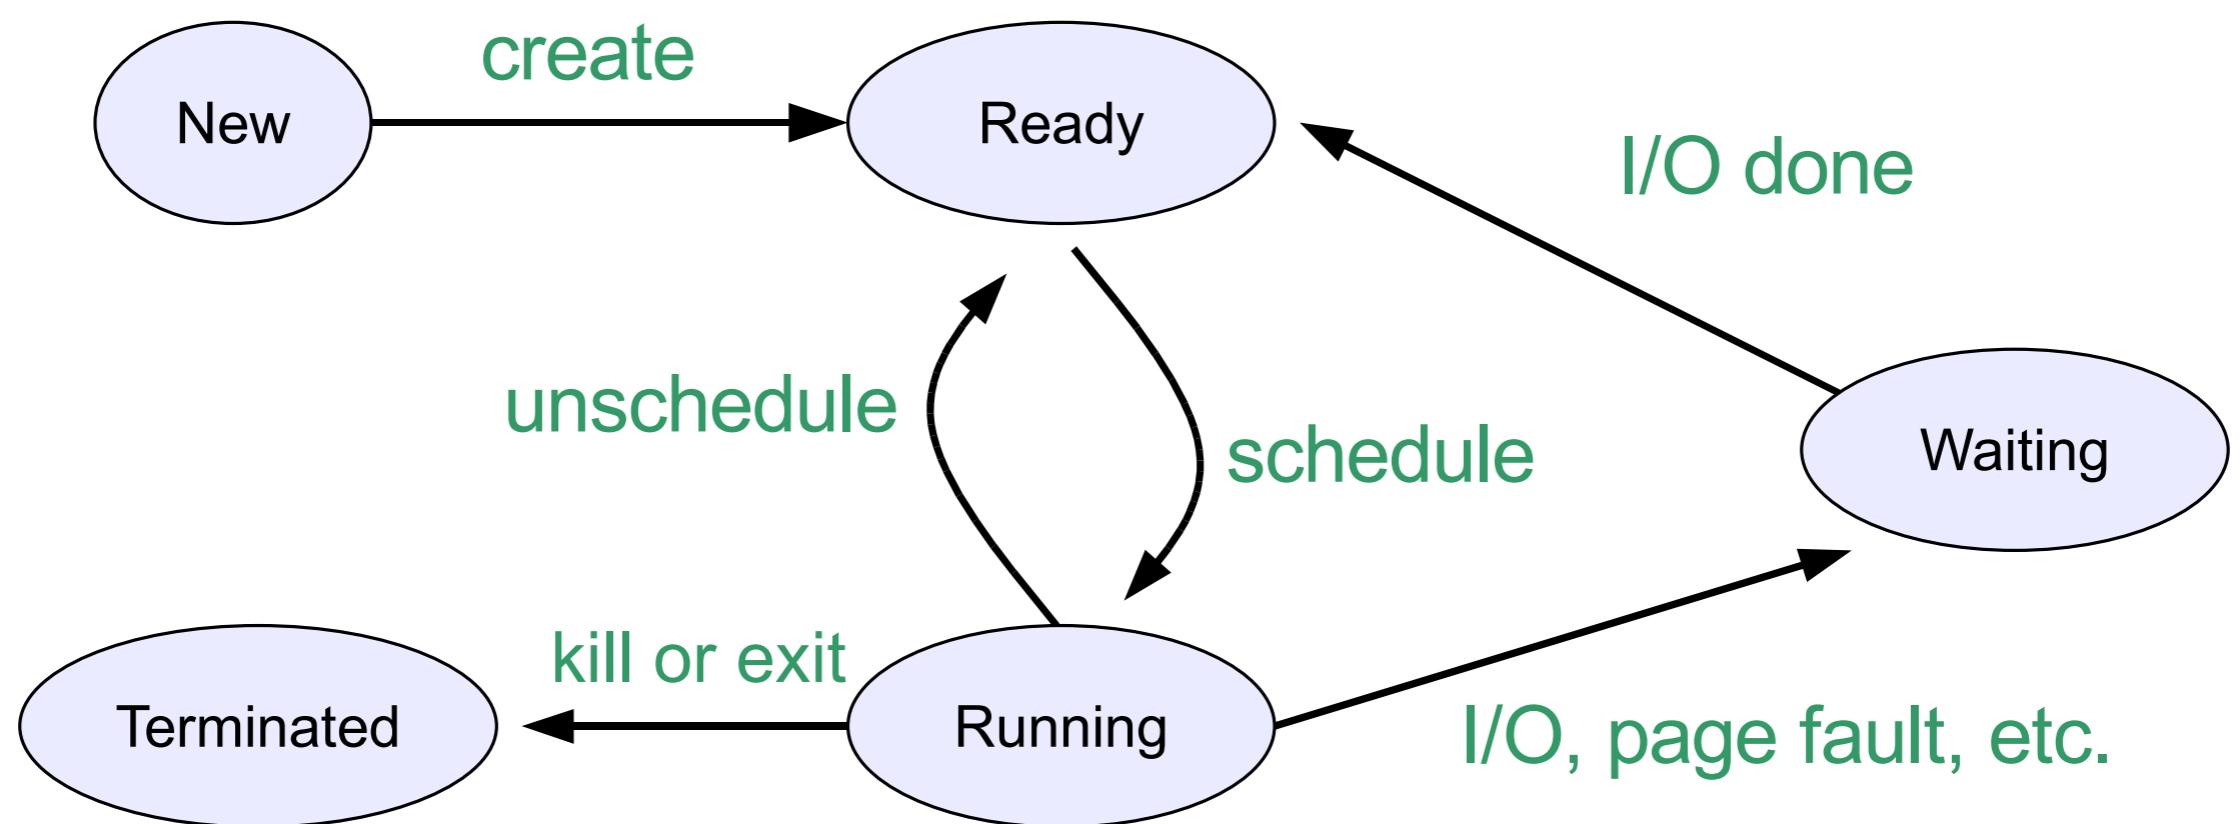
```

/* mm fault and swap info: */
  unsigned long min_flt, maj_flt, nswap,
              cmin_flt, cmaj_flt, cnswap;
  int swappable:1;
  unsigned long swap_address;
  unsigned long old_maj_flt;
  unsigned long dec_flt;
  unsigned long swap_cnt;
/* limits */
  struct rlimit rlim[RLIM_NLIMITS];
  unsigned short used_math;
  char comm[16];
/* file system info */
  int link_count;
  struct tty_struct *tty;
/* ipc stuff */
  struct sem_undo *semundo;
  struct sem_queue *semsleeping;
  struct desc_struct *ldt;
  struct thread_struct tss;
  struct fs_struct *fs;
/* open file information */
  struct files_struct *files;
/* memory management info */
  struct mm_struct *mm;
/* signal handlers */
  struct signal_struct *sig;
};


```

The Life Cycle of Processes: Process States

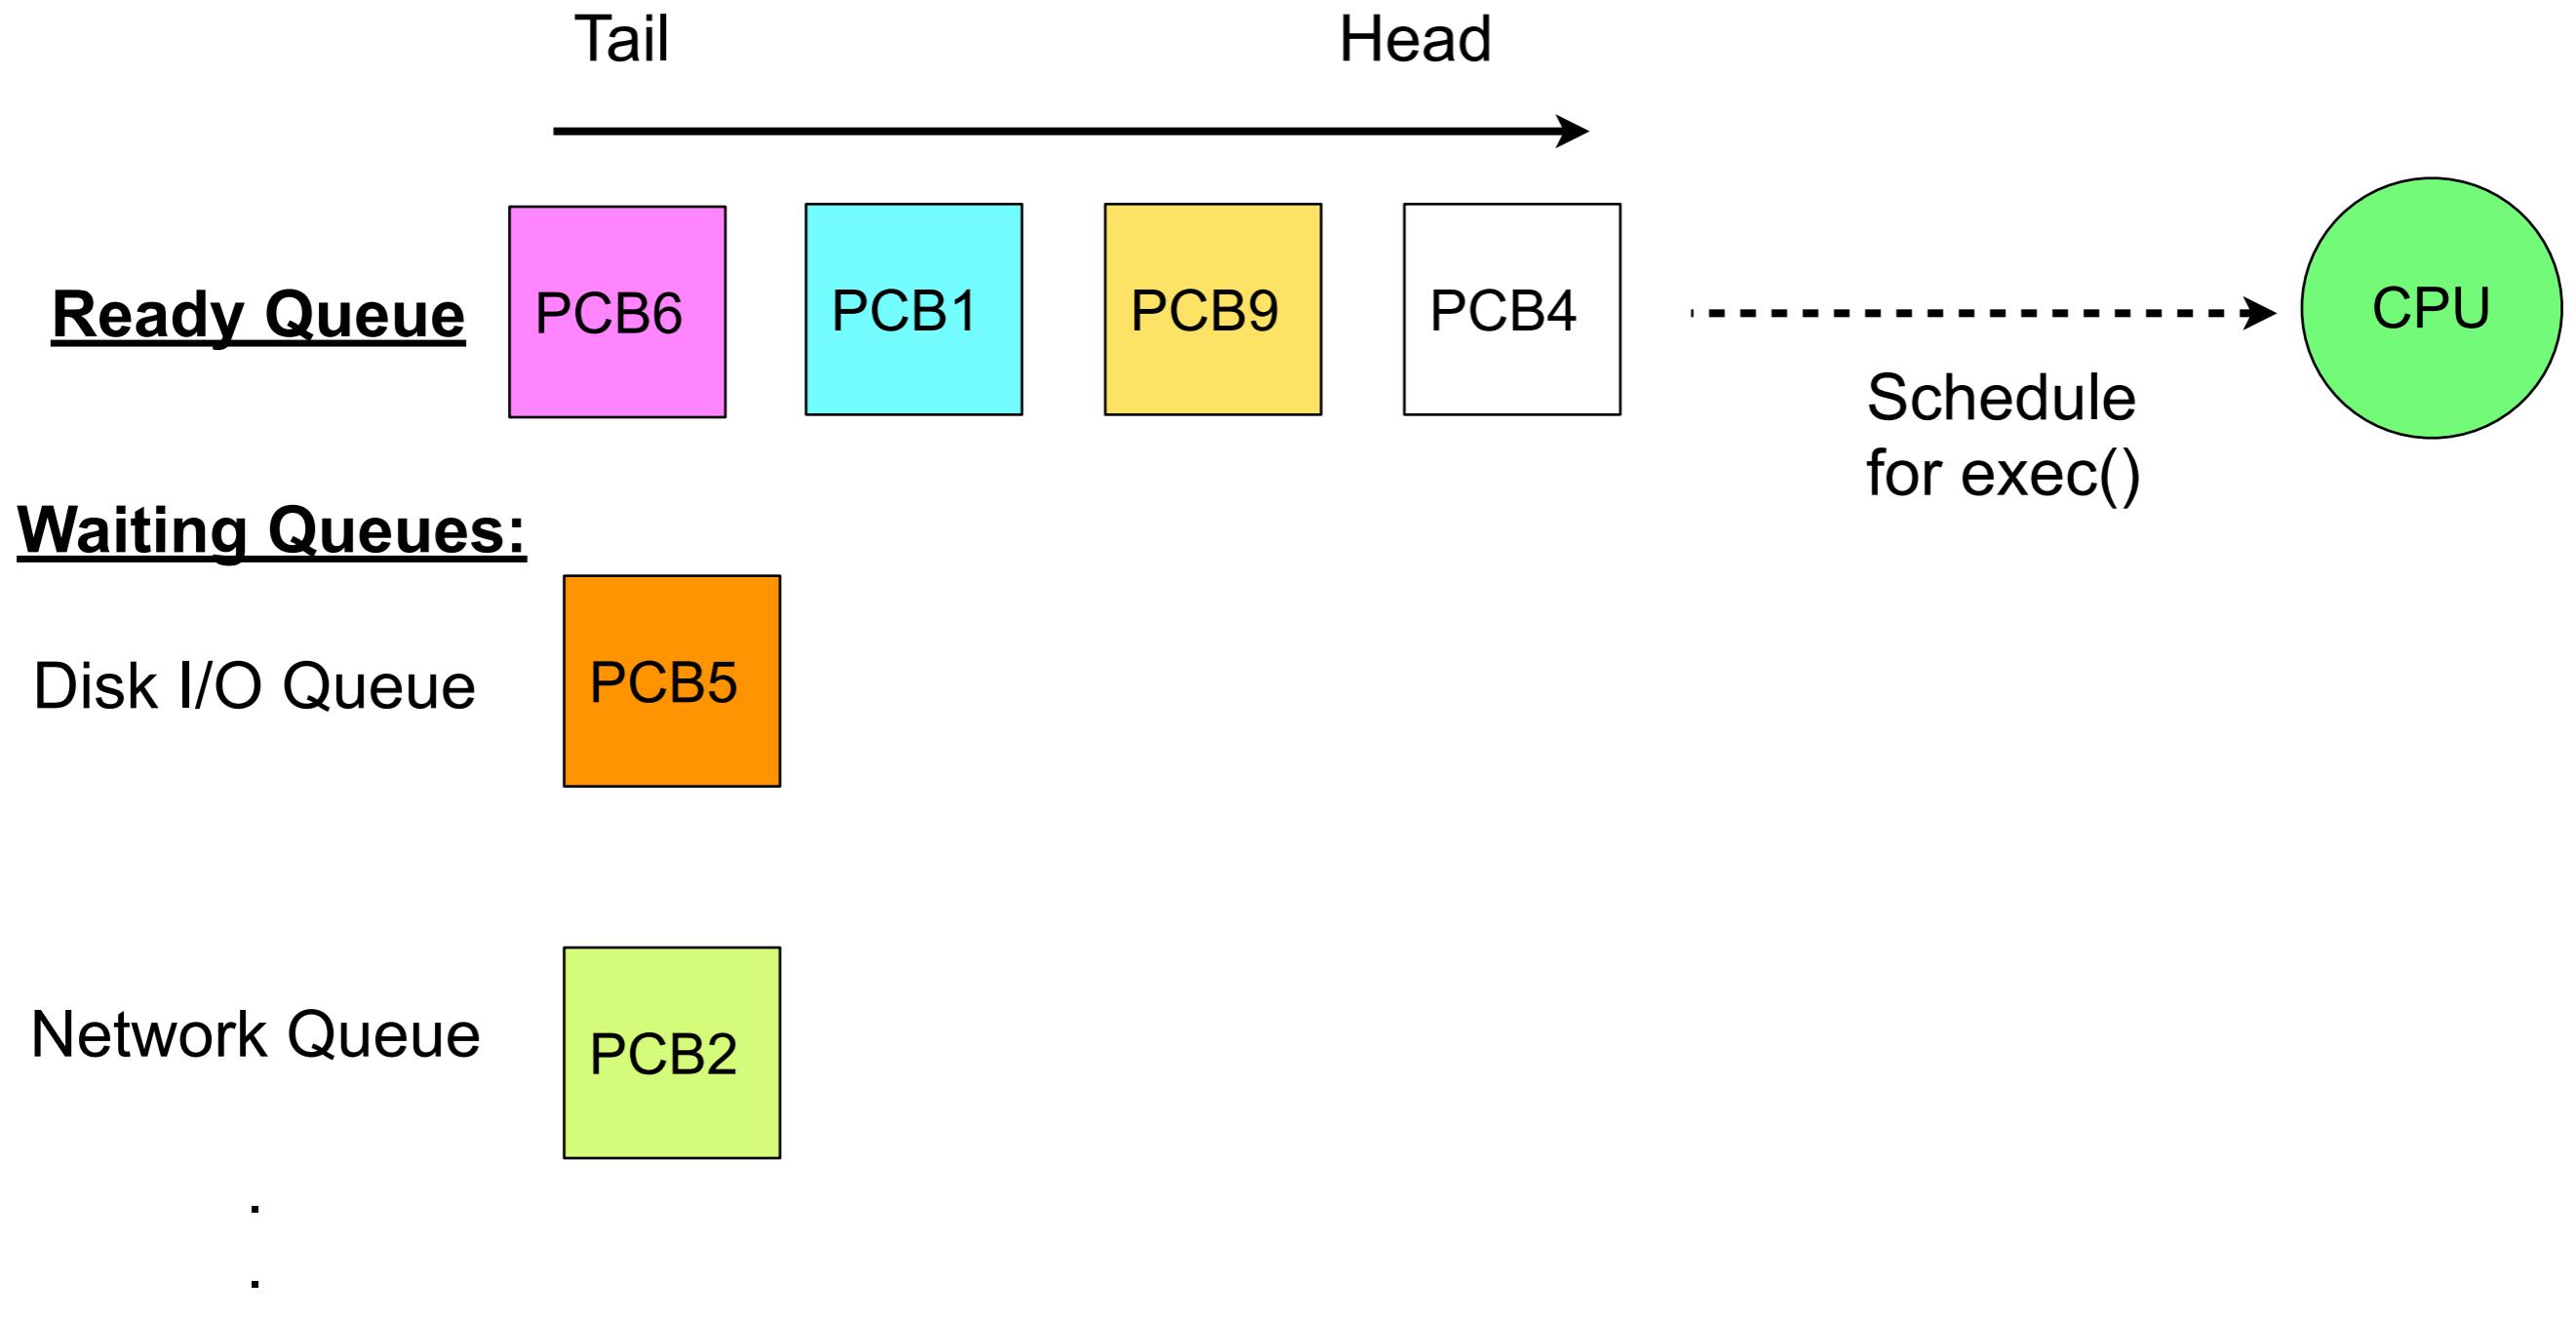
- ▶ As a process run, it transitions among these states
 - A process *cannot* be in two states at once!



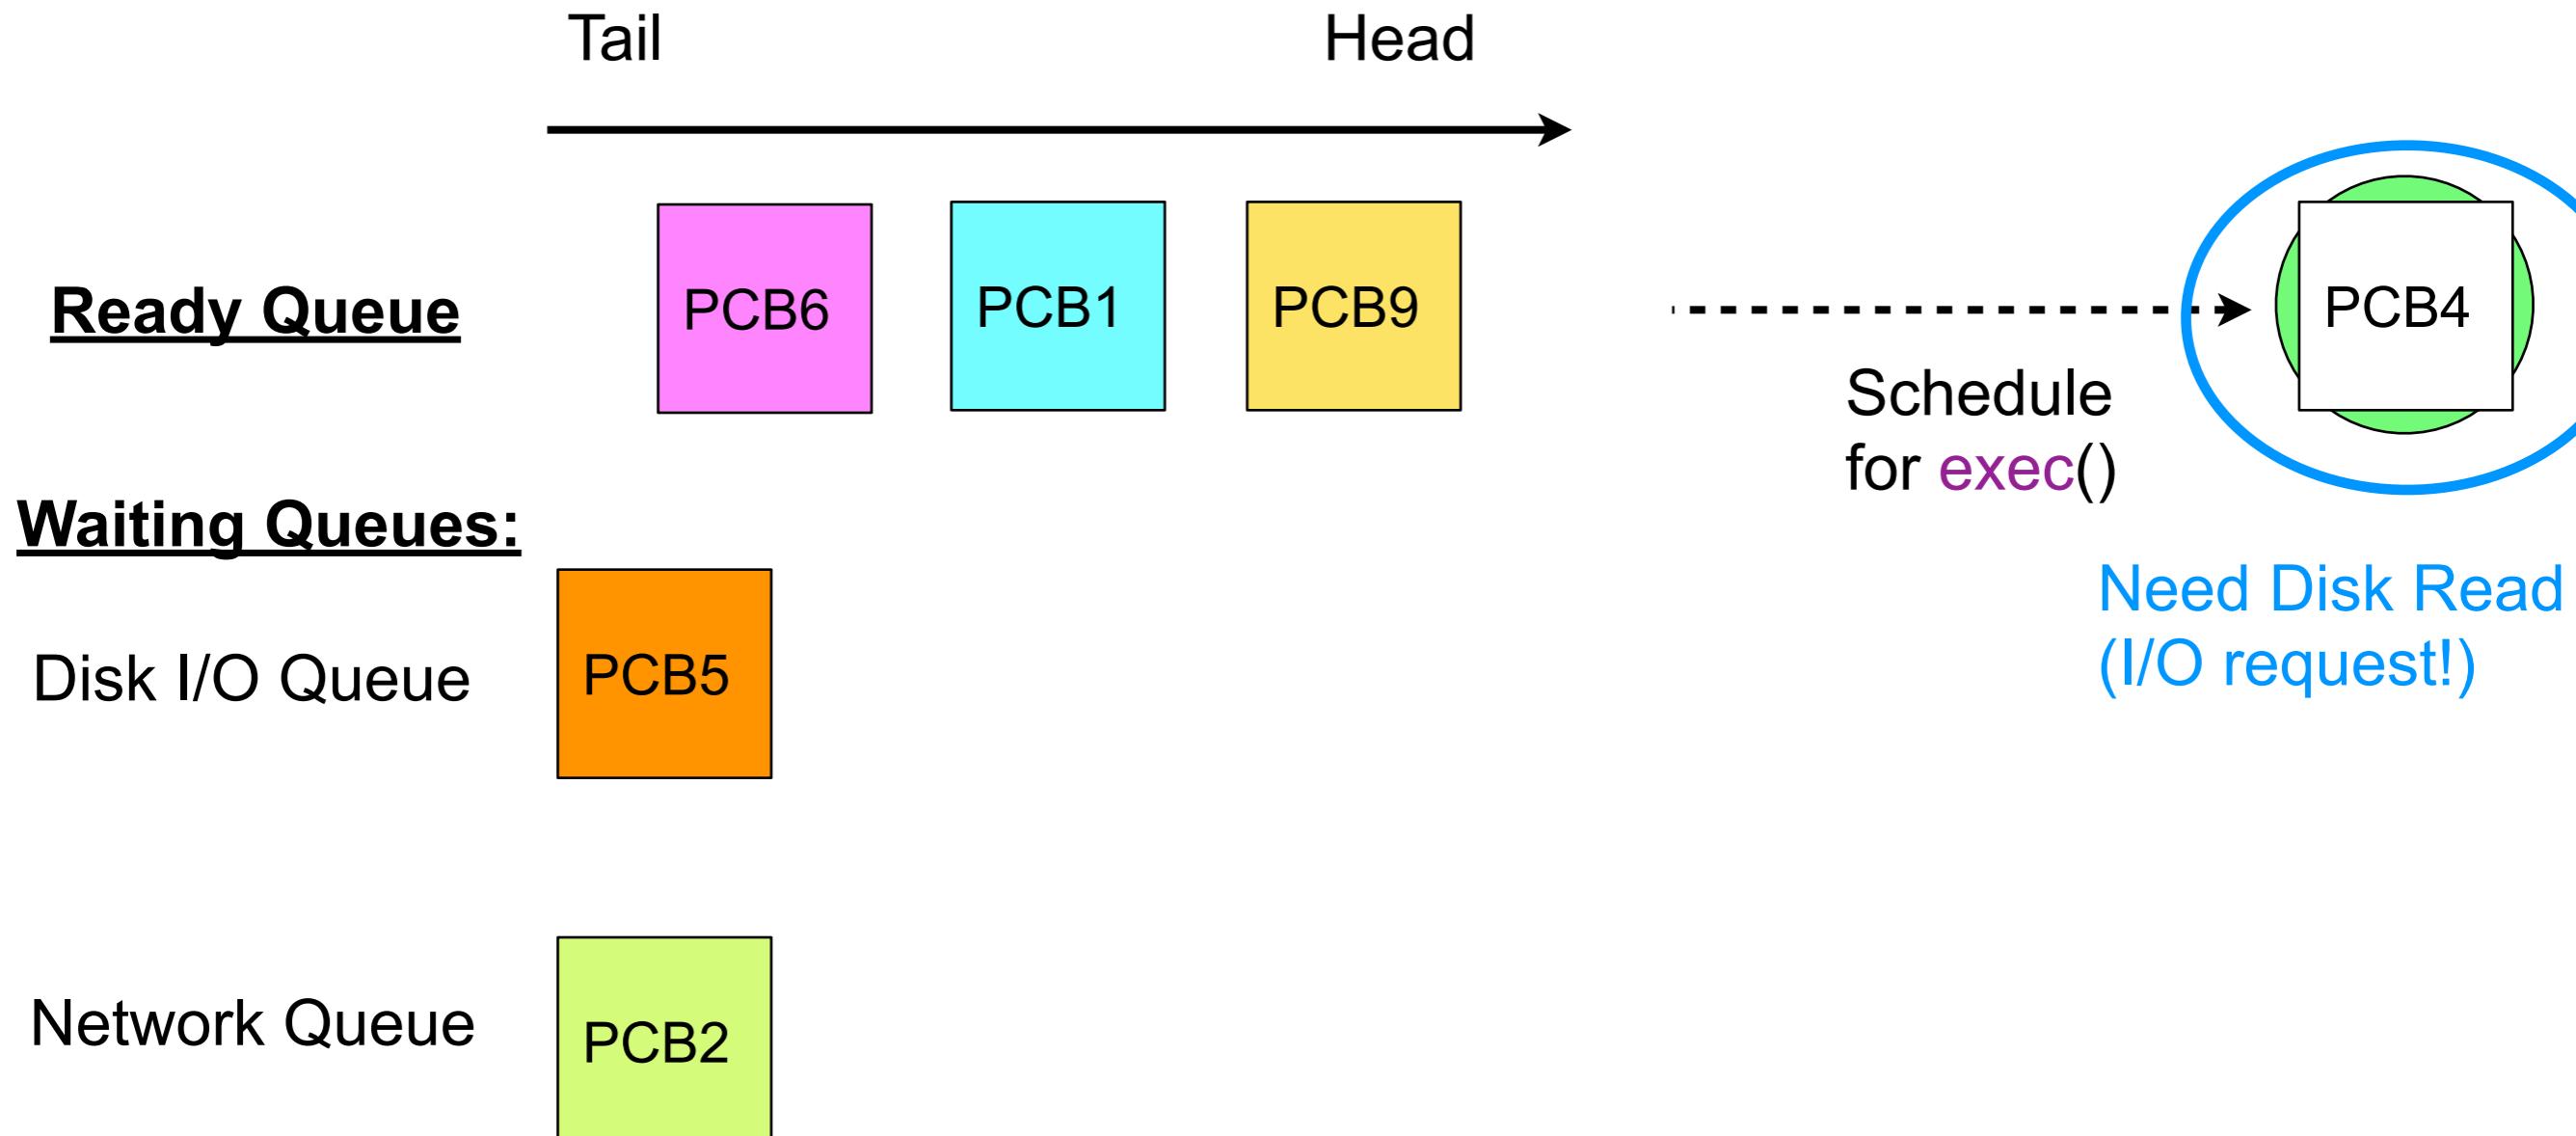
Implementing State Transition with Queues

- ▶ An OS uses *Queues* to implement process states
 - Each queue contains pointers to PCBs
 - PCBs move from queue to queue as they change state
- ▶ Separate queues for *ready* and *waiting* states
 - The ready queue used for processes that are ready to receive CPU time
 - *Why separate ready and waiting states? They seem similar!*
 - Plus a wait queue for each I/O device
 - Disk Queue, Network Queue, Print Queue, etc.

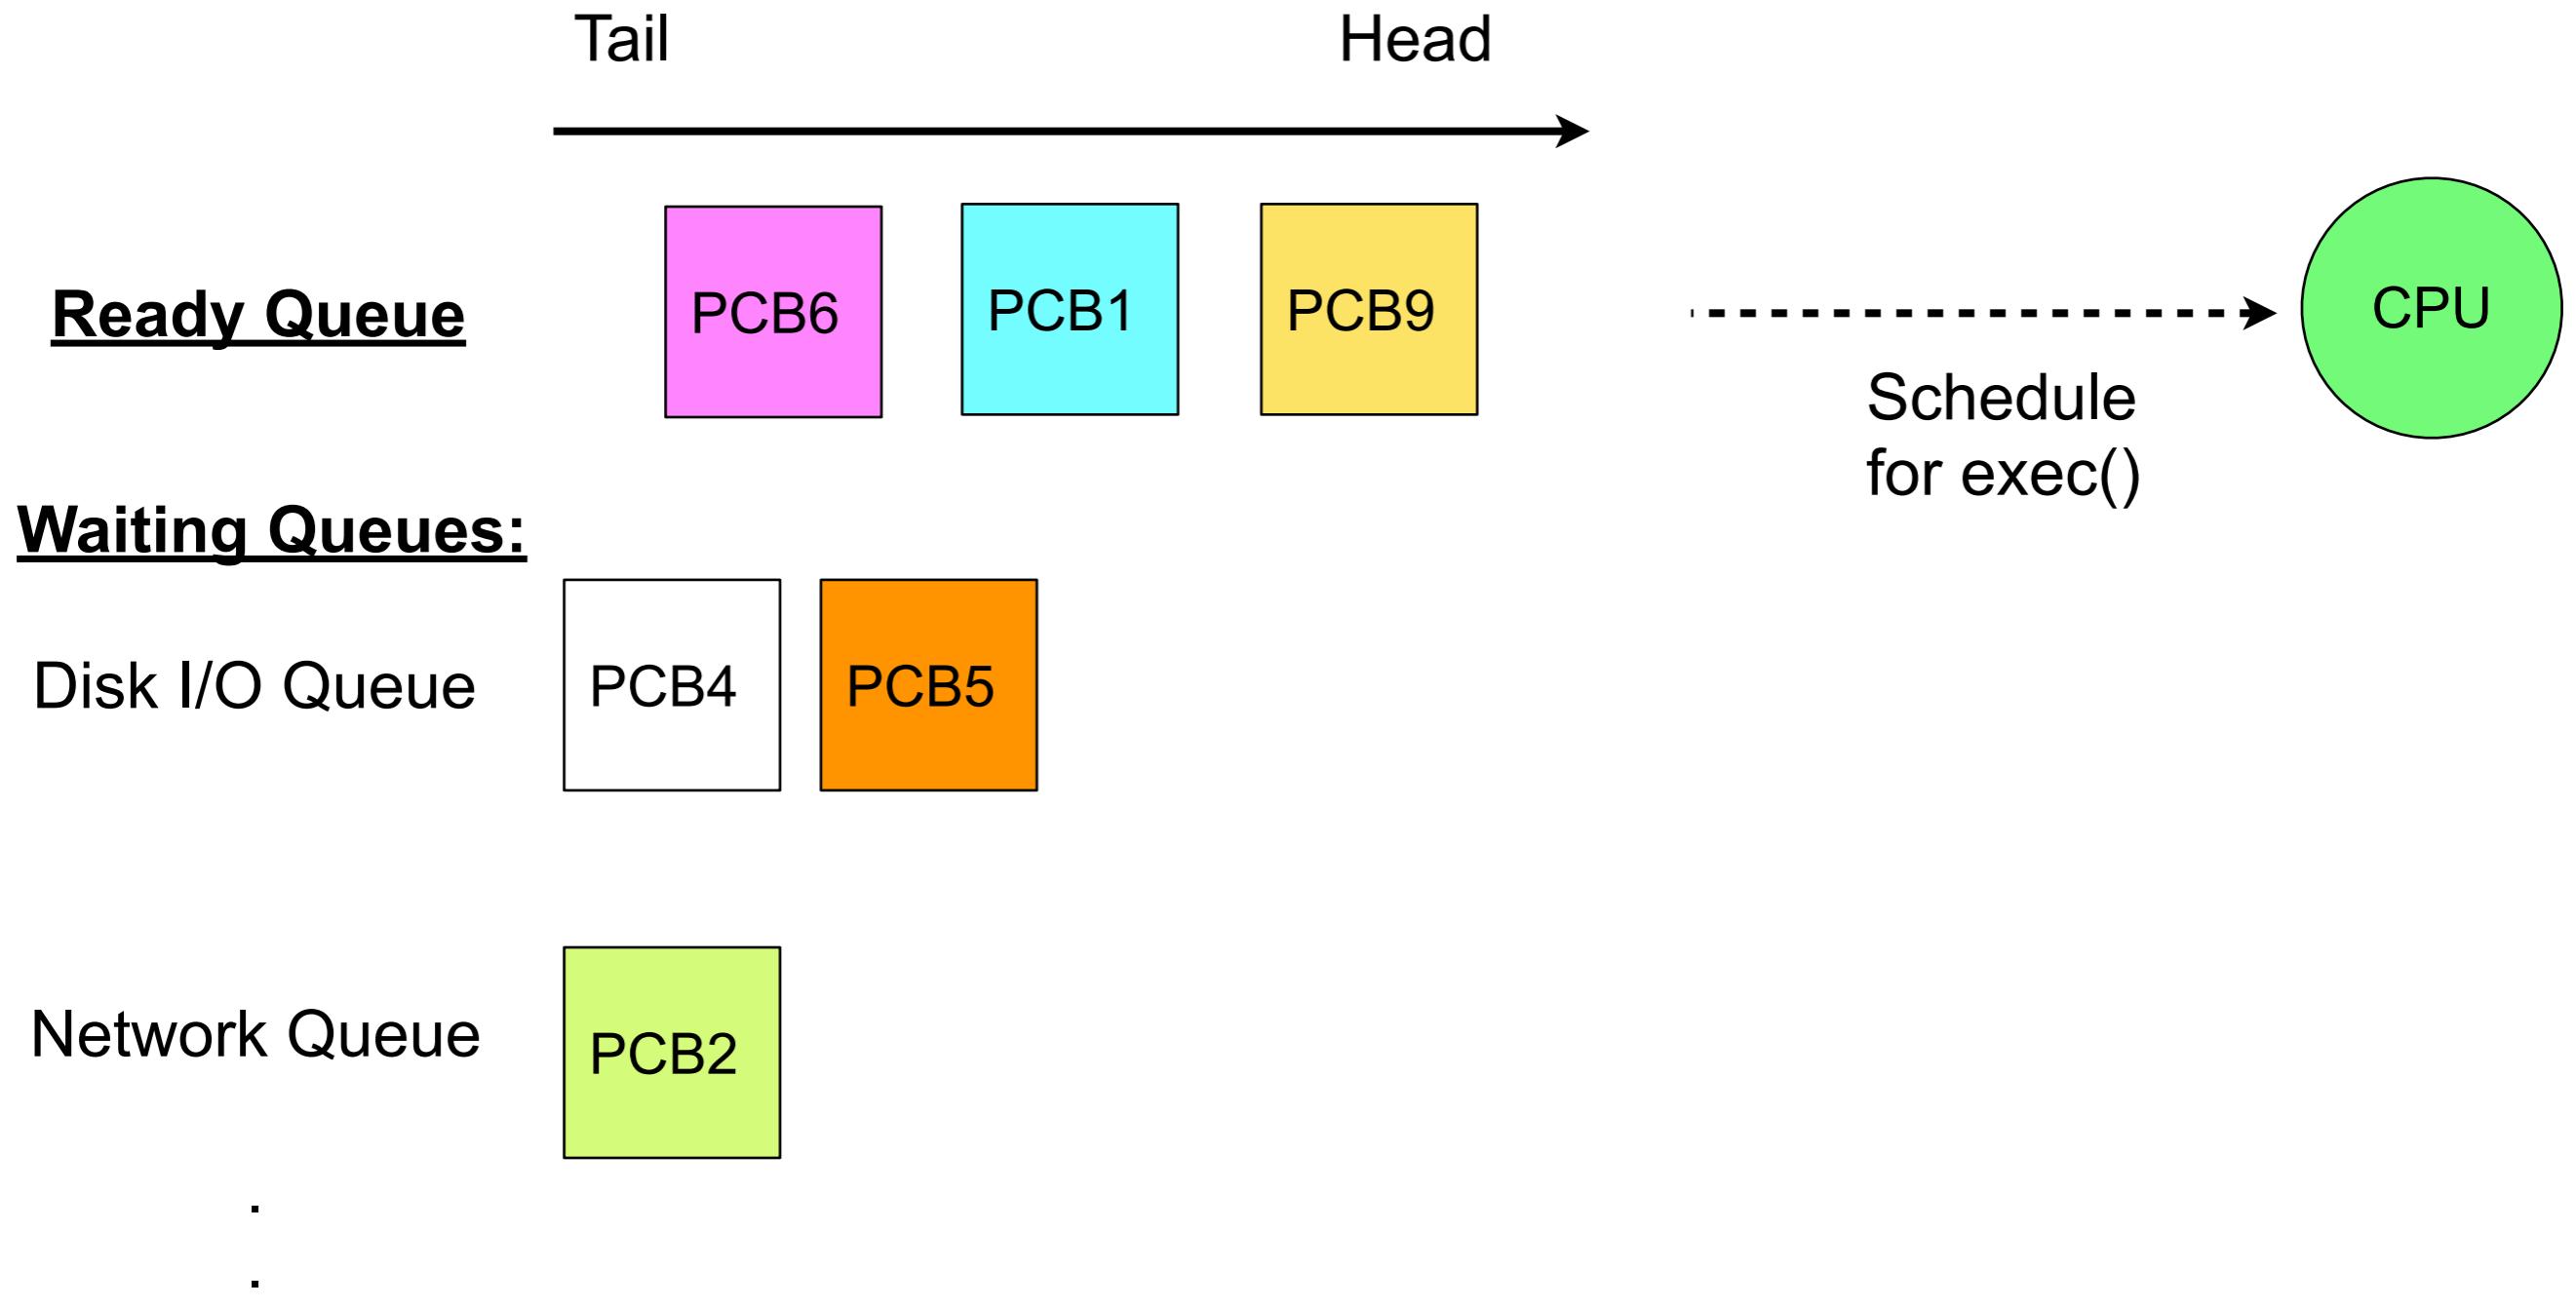
State Transitions: FIFO example



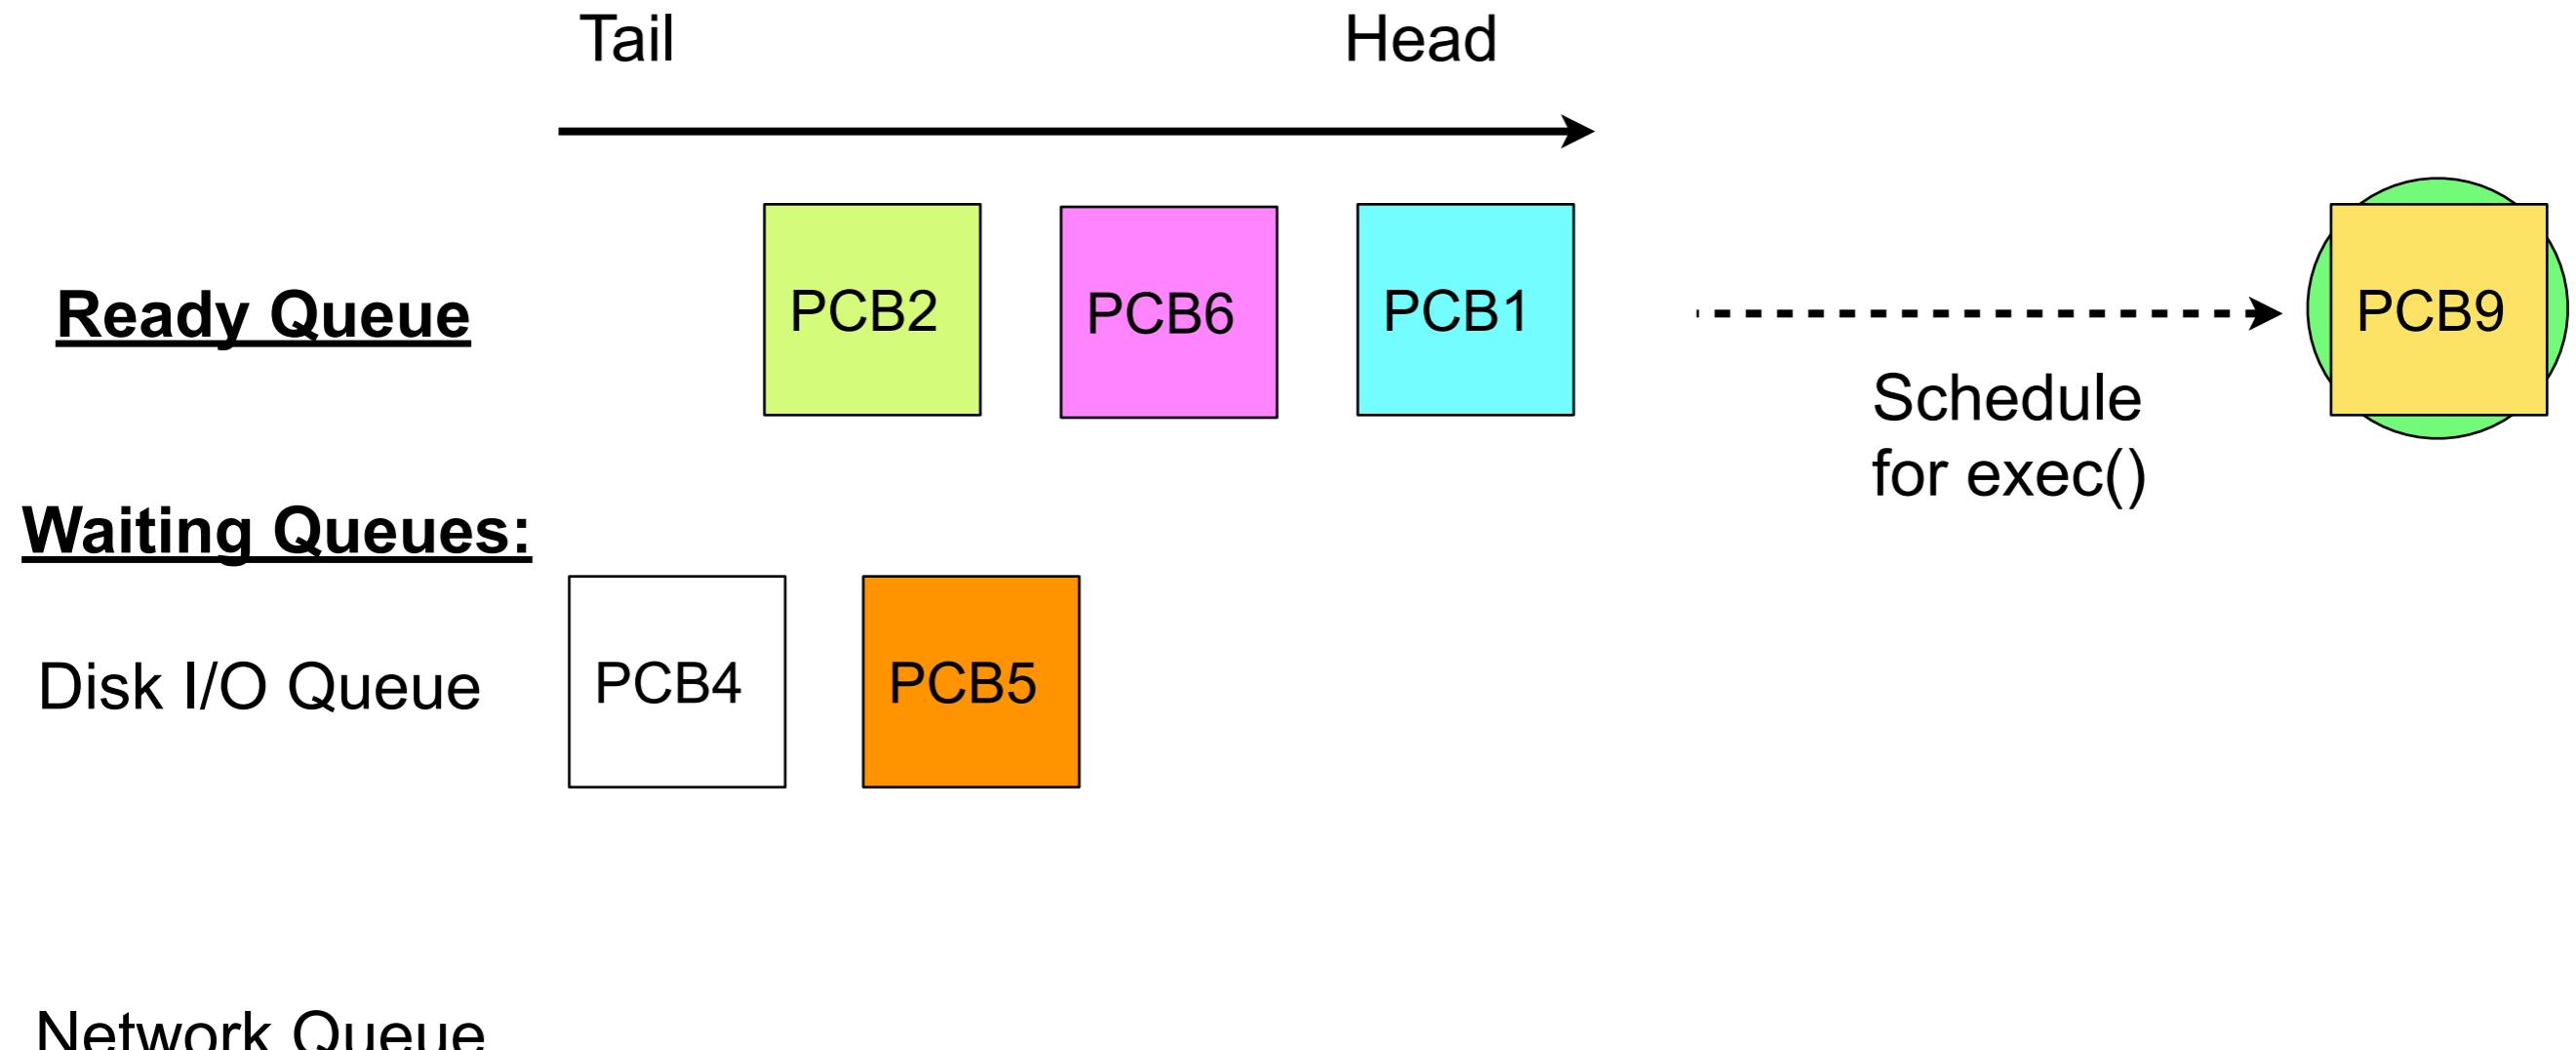
State Transitions: FIFO example



State Transitions: FIFO example



State Transitions: FIFO example



Example: When You Run 'ps au' in linux

Process State (Just focus on the leading term)



USER	PID	%CPU	%MEM	VSZ	RSS	TT	STAT	STARTED	TIME	COMMAND
root	93801	0.0	0.0	34272300	1760	s001	R+	1:49PM	0:00.00	ps au
dchiu	93794	0.0	0.0	34406240	1908	s001	S	1:49PM	0:00.01	-bash
root	93793	0.0	0.0	34447976	5136	s001	Ss	1:49PM	0:00.03	login -fp d
dchiu	99996	0.0	0.0	33794880	1780	s000	Ss+	23Jan23	0:00.32	/bin/bash -

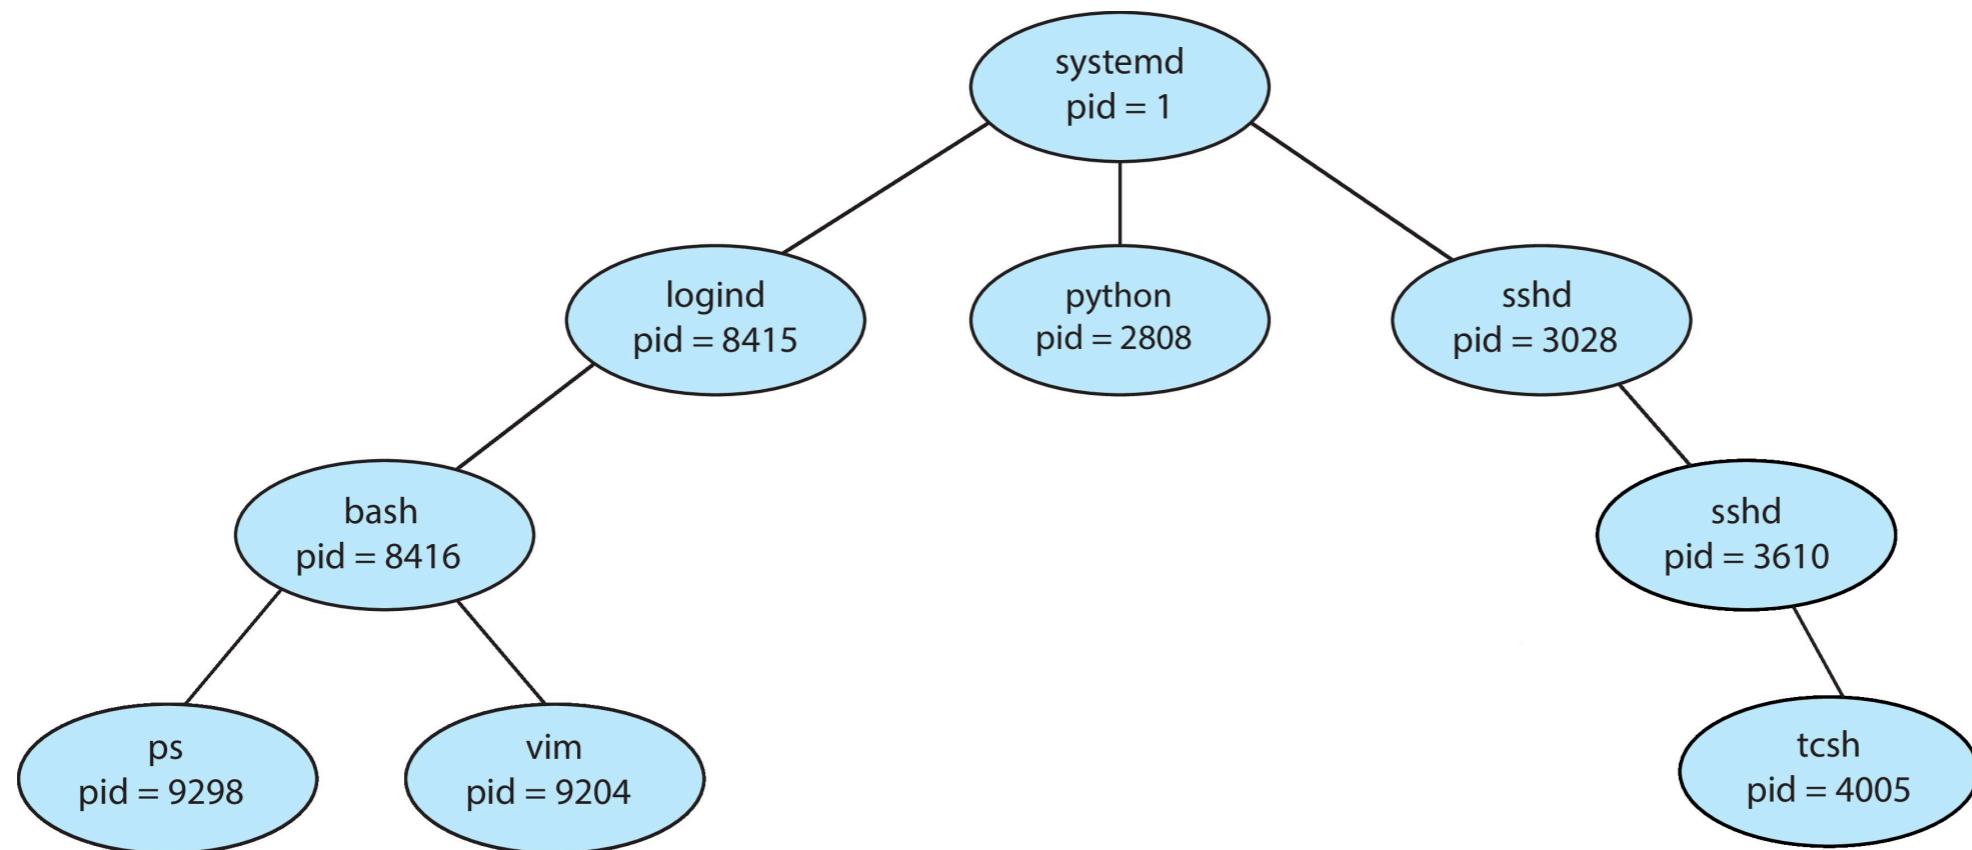
R = Running or Ready
S = Sleeping (Waiting)

Goals for This Lecture

- ▶ Motivation: Concurrency
- ▶ Address Space
- ▶ Processes and the PCB
- ▶ Process Management in C
 - Creation: `fork()`
 - Basic Synchronization: `wait()`
 - Running another program: `exec()`

Process Tree

- ▶ Except for the OS (the first process), all processes are created by another process
 - The lineage of each process is represented with a process tree



Process Creation: fork()

- ▶ To create a new process on the OS, use the system call **fork()**
 - The calling process is called the ***parent process***
 - The new process is called the ***child process***



Process Functions in C

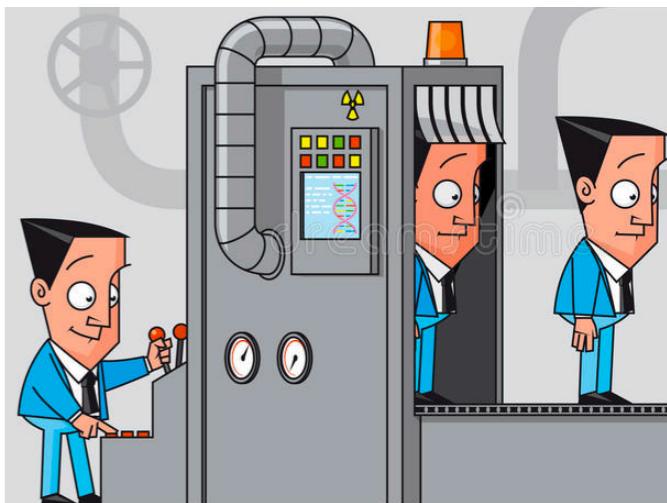
► Syntax in C:

- Need to `#include <unistd.h>`
- `pid_t` is just a `typedef` for `int`

```
/**  
 * @return 0 if it's the child, the child's pid if parentâ€šÀè *  
 * otherwise -1 on error  
 */  
pid_t fork();  
  
/**  
 * @return pid of the calling process  
 */  
pid_t getpid();  
  
/**  
 * @return pid of the calling process' parent  
 */  
pid_t getppid();
```

Process Creation (Cont.)

- ▶ The behavior of `fork()` is a bit odd...
- ▶ `fork()` creates a *near-clone* of the parent process
 - Allocate a new PCB and Address Space for child
 - Copy parent's address space to child's
 - Copy parent's register contents (including PC, SP, BP...)
 - Copy parent's I/O state (e.g., open file and network handles)
 - Child and parent then execute concurrently
 - Then `fork()` returns twice in your code!
 - (*What??*) Once for parent, once for child



Process Creation (Cont.)

► Questions about `fork()`:

- Does child really need a copy of parent's address space?
 - Seems wasteful and expensive to copy!
 - Just share until child makes a change to memory (*copy-on-write policy*)
- "Nearly a clone?" What are some things that are private to the child?
 - What good is a clone? (We'll address that later)
- How do we get the child to run as an independent process?
- When might `fork()` return -1
 - (Indicating error during process creation)?

Process Creation (Cont.)

- ▶ Child's execution starts where the parent's `fork()` returns
 - This is important, because it leads to confusing program traces!

```
#include <unistd.h>
#include <stdio.h>

int main(int argc, char *argv[]) {

    for (int i = 0; i < 10; i++) {
        printf("Process %d: value is %d\n", getpid(), i);

        //fork a child on the 6th iteration
        if (5 == i) {
            printf("Process %d: About to fork...\n", getpid());
            fork();
        }
    }
}
```

Example Output: Non-Deterministic Execution!

```

int main(int argc, char *argv[]) {
    int i;
    for (i = 0; i < 10; i++) {
        printf("Process %d: value is %d\n", getpid(), i);

        if (5 == i) {
            printf("Process %d: About to fork...\n", getpid());
            fork();
        }
    }
}
  
```

Process 4530: value is 0
 Process 4530: value is 1
 Process 4530: value is 2
 Process 4530: value is 3
 Process 4530: value is 4
 Process 4530: value is 5
 Process 4530: About to fork...
Process 4531: value is 6
 Process 4530: value is 6
Process 4531: value is 7
 Process 4530: value is 7
Process 4531: value is 8
 Process 4530: value is 8
Process 4531: value is 9
 Process 4530: value is 9

Process 4530: value is 0
 Process 4530: value is 1
 Process 4530: value is 2
 Process 4530: value is 3
 Process 4530: value is 4
 Process 4530: value is 5
 Process 4530: About to fork...
Process 4531: value is 6
Process 4531: value is 7
Process 4531: value is 8
Process 4531: value is 9
 Process 4530: value is 6
 Process 4530: value is 7
 Process 4530: value is 8
 Process 4530: value is 9

Process 4530: value is 0
 Process 4530: value is 1
 Process 4530: value is 2
 Process 4530: value is 3
 Process 4530: value is 4
 Process 4530: value is 5
 Process 4530: About to fork...
Process 4531: value is 6
 Process 4530: value is 6
 Process 4530: value is 7
 Process 4530: value is 8
 Process 4530: value is 9
Process 4531: value is 7
Process 4531: value is 8
Process 4531: value is 9

Your Turn!

- ▶ How many Hellos get printed??

```
int main() {
    fork();
    fork();
    printf("Hello\n");
    return 0;
}
```

```
int main() {
    fork();
    fork();
    fork();
    printf("Hello\n");
    return 0;
}
```

Telling Parents and Children Apart in Code

- ▶ The return value of `fork()` is the `child's pid` in the parent!
- ▶ The return value of `fork()` is `0` in the child!

```
int child_pid = fork();

if (0 != child_pid) {
    // parent process
    printf("I am the parent proc. My child's proc is %d\n", child_pid);
}
else {
    // child process
    printf("I am the child proc, parent proc is %d\n", getppid());
}
```

Goals for This Lecture

- ▶ Motivation: Concurrency
- ▶ Address Space
- ▶ Processes and the PCB
- ▶ Process Management in C
 - Creation: fork()
 - Basic Process Synchronization: wait()
 - Running another program: exec()

Synchronizing between Parent and Children

- ▶ Need to `#include <sys/wait.h>` and `#include <sys/types.h>`
- ▶ The `wait()` system call makes the parent wait for its child to terminate.

```
pid_t wait(int *status)
```

`@param *status` Output parameter to the status of the terminating child (can be input as `NULL`)
`@return` pid of the terminating child process

- ▶ Sometimes the parent wants to know *why* the child terminated: Normal exit? Was it terminated by OS?
 - More on exit status: man7.org/linux/man-pages/man2/wait.2.html

Synchronizing between Parent and Children

```
#include <stdio.h>
#include <unistd.h>

int main(int argc, char *argv[]) {

    for (int i = 0; i < 10; i++) {

        printf("Process %d: value of i is %d\n", getpid(), i);
        if (5 == i) {
            printf("Process %d about to do a fork!\n", getpid());
            pid_t child_pid = fork();

            //the parent waits for child to finish
            if (0 != child_pid) {
                wait(NULL);
            }
        }
    }
}
```

More C: A More General wait()

- ▶ If a parent created multiple children, can we specify which child to wait for?

```
pid_t waitpid(pid_t pid, int *status, int options);
```

@param pid a child's pid, or -1 for *any* child

@param *status output param to the status of the terminating child (can be **NULL**)

@param options Generally 0 (see man7.org/linux/man-pages/man2/wait.2.html)

@return pid of the child process

Synchronizing between Parent and Children

```

int main(int argc, char *argv[]) {
    for (int i = 0; i < 10; i++) {
        printf("Process %d: value of i is %d\n", getpid(), i);
        if (5 == i) {
            printf("Process %d about to do a fork!\n", getpid());
            pid_t child_pid = fork();

            //the parent waits for child to finish
            if (0 != child_pid) {
                wait(NULL);
            }
        }
    }
}
  
```

Parent waits here while child runs!

Child Terminates

Parent resumes

Process 12312: value of i is 0
 Process 12312: value of i is 1
 Process 12312: value of i is 2
 Process 12312: value of i is 3
 Process 12312: value of i is 4
 Process 12312: value of i is 5
 Process 12312 about to do a fork!
 Process 12313: value of i is 6
 Process 12313: value of i is 7
 Process 12313: value of i is 8
 Process 12313: value of i is 9
 Process 12312: value of i is 6
 Process 12312: value of i is 7
 Process 12312: value of i is 8
 Process 12312: value of i is 9

Linux: Zombie Processes (Z State)

- ▶ What if the child terminates before the parent?
 - Recall all processes supply a **status code** on exit.
 - If **status** has not been read, OS keeps the PCB around until its parent runs `wait()` to read it (called *reaping*)

```
pid_t child_pid = fork(); //fork a child

if (0 != child_pid) {
    sleep(60); // parent sleeps for a minute
}
else {
    // This is the child process. Exit immediately.
    exit(0);
}
```

zombify.c



Example: Zombies

```
$ gcc -o zombify zombify.c
$ ./zombify &
[1] 17788

$ ps au
USER      PID %CPU %MEM      VSZ      RSS      TT STAT STARTED          TIME COMMAND
root    17790  0.0  0.0 2444452   1136   1136 s000 R+  11:21AM 0:00.01 ps au
dchiu  17788  0.0  0.0 2434856    784    784 s000 S   11:21AM 0:00.00 ./zombify
dchiu  10939  0.0  0.0 2467220   952    952 s000 Ss  Wed01PM 0:00.11 -/bin/bash
dchiu  17789  0.0  0.0           0       0 s000 Z   11:21AM 0:00.00 (zombify)
```

60 seconds later....

```
$ ps au
USER      PID %CPU %MEM      VSZ      RSS      TT STAT STARTED          TIME COMMAND
root    17809  0.0  0.0 2445476   1144   1144 s000 R+  11:23AM 0:00.00 ps au
dchiu  10939  0.0  0.0 2467220   952    952 s000 Ss  Wed01PM 0:00.11 -/bin/bash
```

*But in our code, parent never calls **wait()** before it exits, how did the Zombie get reaped?*

(First, we have to understand orphaning and adoption)

Linux: Orphaned Processes

- What if the parent terminates *before* its child terminates?

```

int main(int argc, char *argv[]) {
    pid_t child_pid = fork(); //fork a child
    if (0 != child_pid) {
        sleep(1); //parent sleeps only a second before terminating
        exit(0);
    }
    else {
        //child lives on after the parent terminates
        for (int i = 0; i < 5; i++) {
            printf("mypid=%d, parentpid=%d\n", getpid(), getppid());
            sleep(3);
        }
    }
    return 0;
}
  
```

orphan.c

Parent exits here!
 (Command Prompt returns)

```

$ ./orphan
mypid=10383, parentpid=10382
$
mypid=10383, parentpid=1
mypid=10383, parentpid=1
mypid=10383, parentpid=1
  
```

Child gets a new parent (pid=1)

New parent calls `wait()` on any adopted child, so no zombies!

Recall: Process Tree

```
$ pstree -p
systemd(1)─ NetworkManager(2267)─ dhclient(2417)
|                               └─{NetworkManager}(2418)
|   ├ abrt-dump-oops(2726)
|   ├ abrtd(2718)
|   ├ acpid(2404)
|   ├ atd(2802)
|   ├ automount(2581)─ {automount}(2582)
|   |   ├ {automount}(2583)
|   |   ├ {automount}(2586)
|   |   └ {automount}(2589)
|   ├ avahi-daemon(2279)─ avahi-daemon(2280)
|   ├ bonobo-activati(3169)─ {bonobo-activat}(3170)
|   ├ console-kit-dae(3068)─ {console-kit-da}(3069)
|   |   └ {console-kit-da}(3070)
```

In Linux, all processes are descendants of the `systemd` process, which is the last step in the boot.

`systemd` then reads the system's initialization scripts (in `/etc/init.d`) and executes more programs, eventually completing the boot

Example From Earlier: Who Reaped the Zombie? The OS Did!

```

int main () {
    pid_t child_pid = fork(); //fork a child
    if (0 != child_pid) {
        sleep(60); //parent sleeps for a minute
    }
    else {
        // This is the child process. Exit immediately.
        exit(0);
    }
    return 0;
}
  
```

```

$ ps au
USER      PID %CPU %MEM      VSZ      RSS   TT STAT STARTED      TIME COMMAND
root    17790  0.0  0.0  2444452  1136 s000 R+  11:21AM  0:00.01 ps au
dchiu  17788  0.0  0.0  2434856   784 s000 S   11:21AM  0:00.00 ./zombify
dchiu  17789  0.0  0.0          0       0 s000 Z   11:21AM  0:00.00 (zombify)
  
```

60 seconds later....

```

$ ps au
USER      PID %CPU %MEM      VSZ      RSS   TT STAT STARTED      TIME COMMAND
root    17809  0.0  0.0  2445476  1144 s000 R+  11:23AM  0:00.00 ps au
dchiu 10939  0.0  0.0  2467220   952 s000 Ss  Wed01PM  0:00.11 -/bin/bash
  
```

Rationale: The OS adopted the Zombie (17789) process and reaped it once its parent (17788) exited without reaping

Goals for This Lecture

- ▶ Motivation: Concurrency
- ▶ Address Space
- ▶ Processes and the PCB
- ▶ **Process Management in C**
 - Creation: `fork()`
 - Basic Synchronization: `wait()`
 - Running another program: `exec()`

Why Clone Processes?

- ▶ When might cloning be good?
 - Web Servers like apache `fork()` a child process to deal with every incoming connection
 - Parallel/Distributed Computing - `fork()` to do different parts of the same type of work
 - (e.g., 2 children sorts two halves of a list, then parent merges the halves)
- ▶ Most programs want to spin off *different* programs though (*next*)



The execv() System Call

- ▶ We want to start a process that runs *a different* program
- ▶ Use **exec** family of functions:
 - Need to #include <unistd.h>

```
int execv(const char *filename, char *const argv[]);
```

@param `*filename` Path to the executable

@param `*argv[]` Input arguments for the executable

@return -1 if something went wrong, or doesn't return on normal execution

execv() System Call (Cont.)

- ▶ ***** Important! Important! *****
- ▶ execv() *replaces* the *calling process'* address space and CPU state:
 - OS clears out the calling process' address space.
 - OS loads the given executable file into the code segment.
 - CPU State: Starts over from the beginning of the given executable.

Example: Running ls -l

```
#include <unistd.h>
#include <string.h>

int main(int argc, char *argv[]) {
    /* setting up to run exec() */
    char *command = "/bin/ls";
    char **args = (char**) malloc(10 * sizeof(char*));
    for (int i = 0; i < 10; i++) {
        args[i] = (char*) malloc(64 * sizeof(char));
    }

    strcpy(args[0], command); //store the full path to the command
    strcpy(args[1], "-l");
    args[2] = NULL;          //denote end of args list
    execv(command, args); // exec new program

    /* exec done running */
    printf("Here's more stuff I wanted to do!!!!");
    printf("Here's more stuff I wanted to do!!!!");
    printf("Here's more stuff I wanted to do!!!!");

    return 0;
}
```

testExec.c

Example: Running ls -l (Output)

`execv("/bin/ls")`

called
immediately

**printfs never
run!**

**/bin/ls replaced
execTest code**

```
$ ./execTest
total 21556
-rwxr-xr-x. 1 dchiu domain users 8640 Jun 30 2021 blockinglock
-rw-r--r--. 1 dchiu domain users 1157 Jun 30 2021 blockinglock.c
-rwxr-xr-x. 1 dchiu domain users 1261 Jun 30 2021 BoundedBuffer.java
-rw-r--r--. 1 dchiu domain users 352 Jun 30 2021 collatz.c
-rwxr-xr-x. 1 dchiu domain users 1315 Jun 30 2021 Consumer.java
-rw-r--r--. 1 dchiu domain users 926 Feb 13 14:19 correct_execv.c
-rw-r--r--. 1 dchiu domain users 1428 Jun 30 2021 deadlock.c
-rw-r--r--. 1 dchiu domain users 804 Feb 13 14:26 execv.c
-rw-r--r--. 1 dchiu domain users 441 Jun 30 2021 fork2.c
-rw-r--r--. 1 dchiu domain users 82 Jun 30 2021 forkbomb.c
-rw-r--r--. 1 dchiu domain users 539 Jun 30 2021 fork.c
-rw-r--r--. 1 dchiu domain users 290 Jun 30 2021 fork_wait.c
-rwxr-xr-x. 1 dchiu domain users 698 Jun 30 2021 highlight.css
-rwxr-xr-x. 1 dchiu domain users 220 Jun 30 2021 index.php
drwxr-xr-x. 2 dchiu domain users 113 Jun 30 2021 lec4
-rw-r--r--. 1 dchiu domain users 340 Jun 30 2021 movingAvg.c
-rw-r--r--. 1 dchiu domain users 475 Feb 10 14:22 orphan.c
-rwxr-xr-x. 1 dchiu domain users 18232 Feb 1 15:47 perf
-rw-r--r--. 1 dchiu domain users 742 Feb 1 15:47 perf_syscall.c
-rwxr-xr-x. 1 dchiu domain users 1309 Jun 30 2021 Producer.java
-rw-r--r--. 1 dchiu domain users 313 Jun 30 2021 rtclock.c
-rwxr-xr-x. 1 dchiu domain users 200 Jun 30 2021 rtclock.h
-rw-r--r--. 1 dchiu domain users 201 Jun 30 2021 segfault.c
-rwxr-xr-x. 1 dchiu domain users 8632 Jun 30 2021 spinlock
-rw-r--r--. 1 dchiu domain users 1176 Jun 30 2021 spinlock.c
-rw-r--r--. 1 dchiu domain users 109 Jun 30 2021 test2.c
-rw-r--r--. 1 dchiu domain users 751 Jun 30 2021 test.c
-rwxr-xr-x. 1 dchiu domain users 21889024 Jun 30 2021 test.out
-rw-r--r--. 1 dchiu domain users 253 Feb 8 14:25 zombify.c
```

\$

execv() System Call (Cont.)

- ▶ **Important:** `execv("newProg", ...)` does not `fork()` a new process
 - 0) It replaces the calling process with `newProg`!
 - 1) You have to first explicitly call `fork()`
 - 2) Followed by `execv("newProg", ...)` in the **child** to start `newProg` as a separate process!

Example: Parent Side

```
int main(int argc, char *argv[]) {                                realTestExec.c
    if (0 != fork()) {      //parent
        printf("pid=%d, Parent here... print to 999\n", getpid());
        int i;
        for (i=0; i<1000; i++)
            printf("pid=%d, %d\n", getpid(), i);
    }
}
```

Focus on Parent code first

```
printf("pid=%d, Child proc should never get here\n", getpid());}
```

Example: Child Side

realTestExec.c

```
else {    //child
    printf("pid=%d, I'm the child. I'm going to run echo\n", getpid());
    char *command = "/bin/echo";
    char *newargs[3];
    newargs[0] = command;                      // full path to executable file!
    newargs[1] = "Random String";              // argument to /bin/echo
    newargs[2] = NULL;                         // denote end of args list
    execv(command, newargs); //spawn new program
}
```

Putting It All Together

```
int main(int argc, char *argv[]) {                                realTestExec.c
    if (0 != fork()) {      //parent
        printf("pid=%d, Parent here... print to 999\n", getpid());
        int i;
        for (i=0; i<1000; i++)
            printf("pid=%d, %d\n", getpid(),i);
    }
    else {      //child
        printf("pid=%d, I'm the child. I'm going to run echo\n", getpid());
        char *command = "/bin/echo";
        char *newargs[3];
        newargs[0] = command;           // full path to executable file!
        newargs[1] = "Random String";   // argument to /bin/echo
        newargs[2] = NULL;             // denote end of args list
        execv(command, newargs); //spawn new program
    }
    printf("pid=%d, Child proc should never get here\n", getpid());
}
```

Example Output of `realTestExec`

(1) Parent starts and forks child → pid=12754, Parent here... print to 999
pid=12754, 0
pid=12754, 1
.
. pid=12754, 24
(2) Parent keeps processing
.

(3) Child starts running too → pid=12755, I'm the child. I'm going to run echo
pid=12754, 25
.
. pid=12754, 311
(4) Child fires `execv()` here → Random String
.
. pid=12754, 998
pid=12754, 999
(5) Child finishes → pid=12754, Child proc should never get here
(6) Parent finishes →

Goals for This Lecture

- ▶ Motivation: Concurrency
- ▶ Address Space
- ▶ Processes and the PCB
- ▶ Process Management in C
 - Creation: fork()
 - Basic Synchronization: wait()
 - Running another program: exec()
- ▶ Conclusion

In Conclusion...

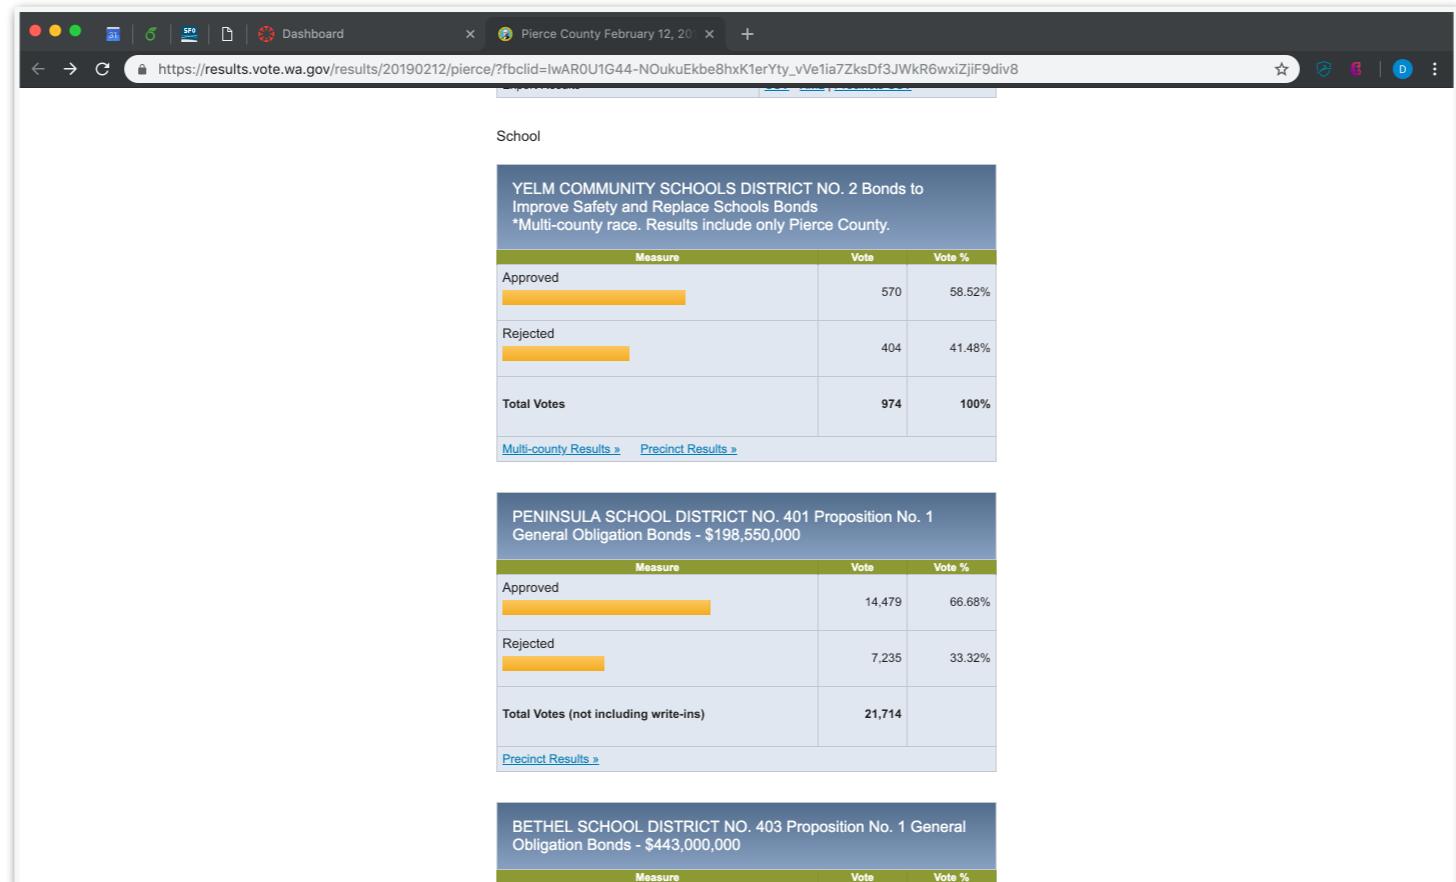
- ▶ To multiprogram and timeshare a computer system, we must allow multiple programs to co-exist and share the CPU
 - Process control blocks (PCB) stores a snapshot of each process state
 - PCBs can be various **states** of execution in its lifecycle
 - Pause and play among the PCBs to multi-task

fork() and exec() in Java?

- ▶ In Java, closest thing to `fork()` is the `ProcessBuilder` class
- ▶ Can invoke an arbitrary program
 - Pass in arguments via the standard input
 - Read results via the standard output
- ▶ Pro - Very easy to use
- ▶ Con - Has a big startup overhead
- ▶ API: <https://docs.oracle.com/javase/8/docs/api/java/lang/ProcessBuilder.html>

Next Time...

- ▶ Problem: Processes are inefficient: fork+exec is expensive
 - Open Chrome (one process), download HTML (fork), download images/css from various locations (fork for each), ...



- ▶ Next Time: Processes are expensive. (*Threads*)

Administrivia 2/6

▶ Announcements

- Hwk 3 due tonight
- Start reading Chapter 3.1-3.5 (dinosaur book)

▶ Last time:

- Finished system calls and traps

▶ Today:

- Process Management
- Address Spaces
 - Program Stack

Administrivia 2/8

- ▶ Reminders
 - Hwk 4 assigned, due Mon 2/20
- ▶ Last time...
 - Policy: CPU virtualization. Mechanism: Context Switch
 - PCBs and context switching
 - Address Spaces; How does the program stack manage scope?
- ▶ Today:
 - Process states
 - Process creation: **fork()** system call

Administrivia 2/10

- ▶ Reminders
 - Hwk 4 assigned, due Mon 2/20
- ▶ Last time...
 - Process states
 - Transitions between states
 - Fork system call
- ▶ Today:
 - The **wait()** system call and reaping children processes
 - New terms: Zombies and Orphans

Administrivia 2/13

- ▶ Reminders
 - Hwk 4 due Mon 2/20
- ▶ Last time...
 - The `wait()` system call
 - Zombie processes and reaping statuses
 - Orphaned processes and adoption
- ▶ Today
 - Finish Processes (Chap 3):
 - The even weirder `exec()` system call
 - Start threads (Chap 4 dinosaur book)