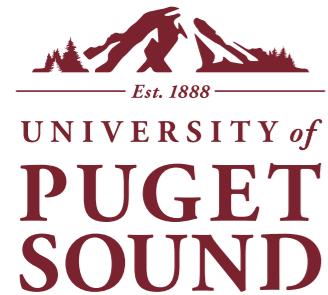


# CS 475

# Operating Systems



Department of Mathematics  
and Computer Science

Lecture 4  
Threads and  
Parallel Computing

# Motivation 1: Process Responsiveness

- ▶ Consider the following C program:
  - First ... I want to compute  $\pi$
  - Then I need to print a funny email I got from America to show students

```
int main(int argc, char *argv[]) {  
    computePi();  
    printAmericasFunnyEmail();  
    return 0;  
}
```

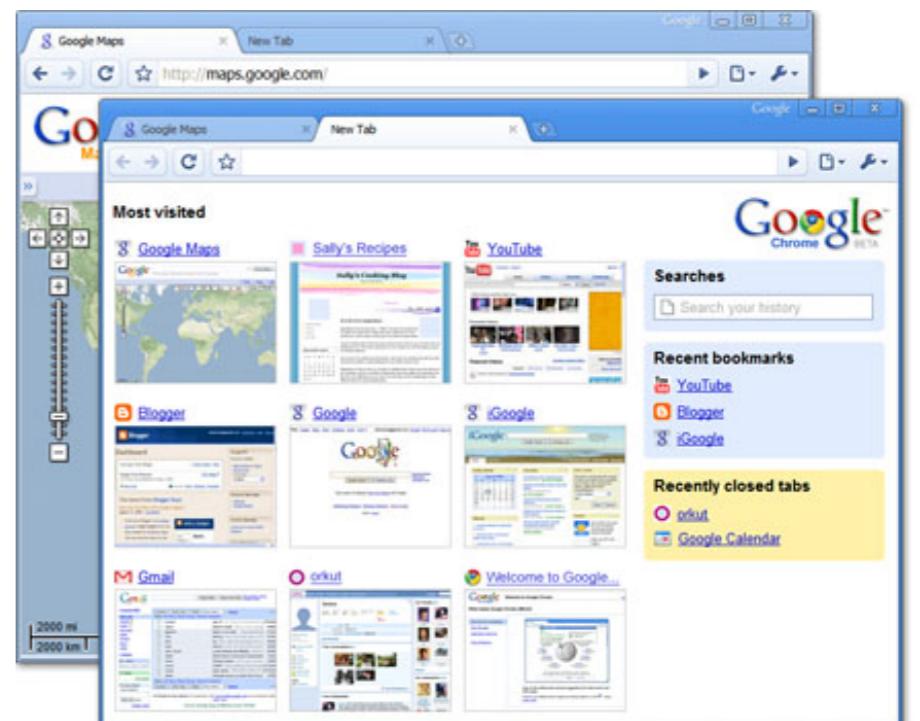
- ▶ Two logically *independent* sub-tasks, but... second task is "stuck" behind the first.
  - `computePI()` will never finish.
  - Even on a timesharing OS, it'll never print America's email.

# Motivation 1: Process Responsiveness (2)

► **Responsiveness:** A single program may need to multitask!

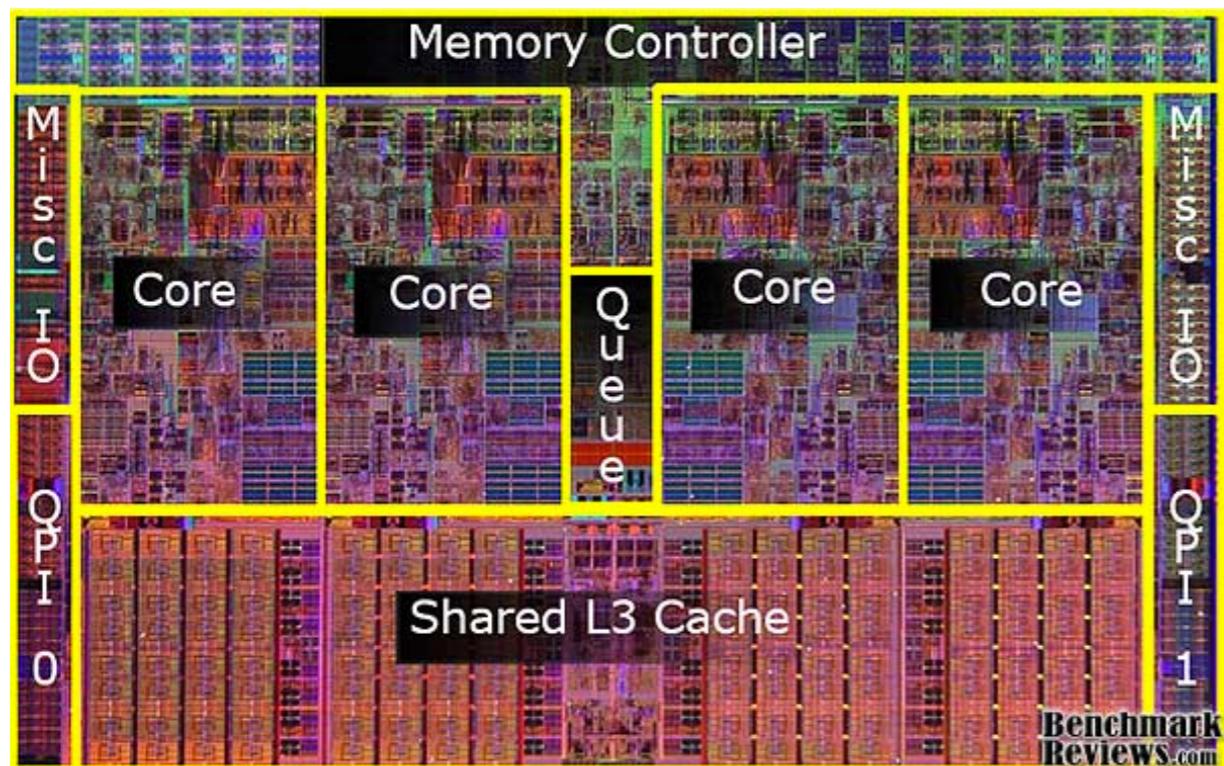
► For instance, web browsers:

- Downloads an HTML page, it accepts user input, fetches from cache, renders page, ... all concurrently!
- In a single-threaded execution, these actions would have to be done sequentially.
  - If browser gets stuck downloading large file, it has to block!



# Motivation 2: Parallel Processing

- ▶ **Problem 1: Increase CPU utilization in the multicore era**
  - Programmers today need to extract parallelism for high performance
    - Can we fetch and load 4 TikTok videos simultaneously?



But can't processes utilize different cores?

Yes, generally if there's a free core, OS will schedule the process on it!

Then why do we need threads?

# Process-Level Parallelism Is Possible, but...

- ▶ We'll just fork a new process for each task. (This *would* work!)

```
int main(int argc, char *argv[]) {
    if (0 != fork()) {
        computePi();    // Let the parent compute PI
    }
    else {
        printAmericasFunnyEmail(); // Let the child print email
    }
}
```

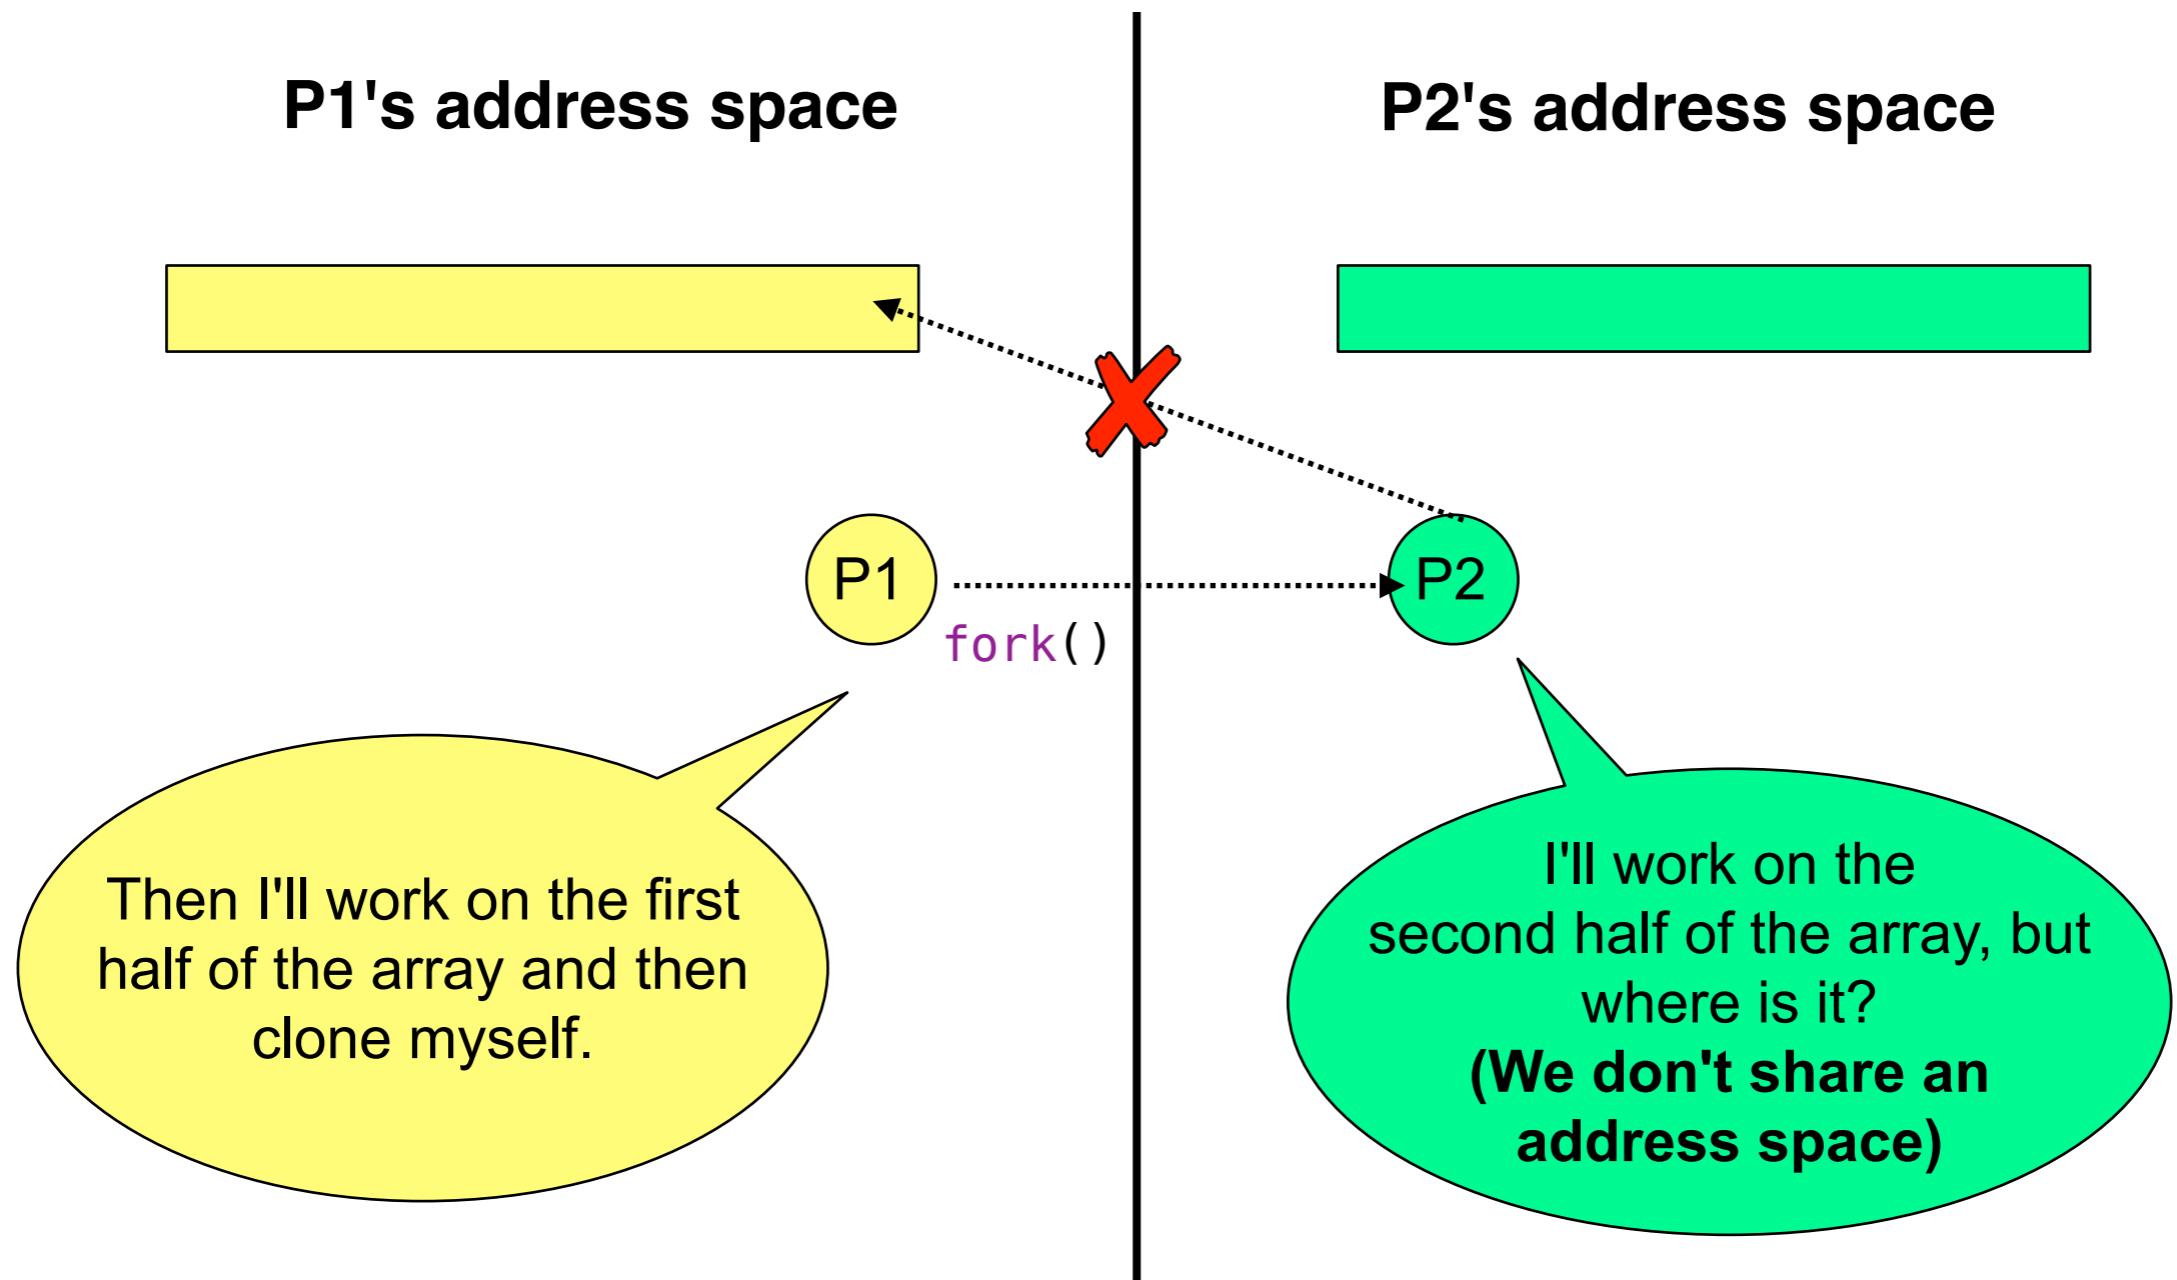
- ▶ Sure, but process creation is expensive 💰 💰 💰!

Recall to create a new process using `fork()`:

- 1) `fork()` is privileged, so you need trap a system call.
- 2) OS allocates a PCB for the new child
- 3) OS allocates child's address space.
- 4) OS copies address space from parent.
- 5) OS copies I/O state from parent.
- 6) Now the new child process can run

# Process-Level Parallelism Is Possible, but...

- ▶ **Problem 2:** Often, processes need to *share* data, but how? (It can be done using the `mmap()` syscall, but it's slow and complicated to setup).



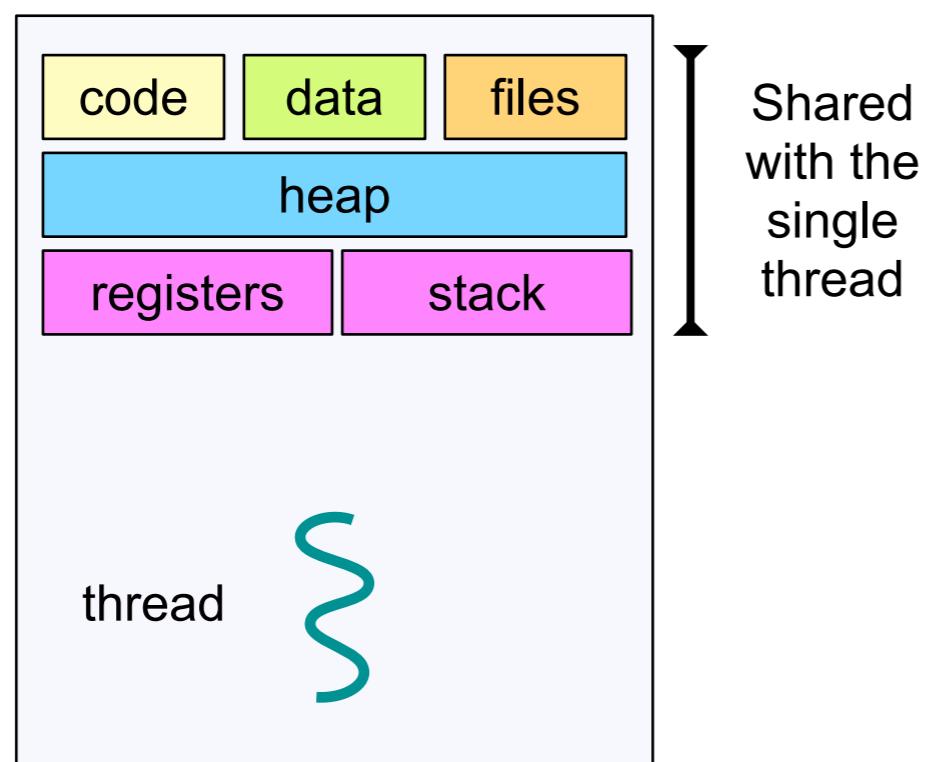
# Goals for Today...

- ▶ Threads
- ▶ **Implementing Threads**
- ▶ Pthread Examples in C
- ▶ Parallel Processing Paradigms
- ▶ User Threads vs. Kernel Threads
- ▶ Conclusion

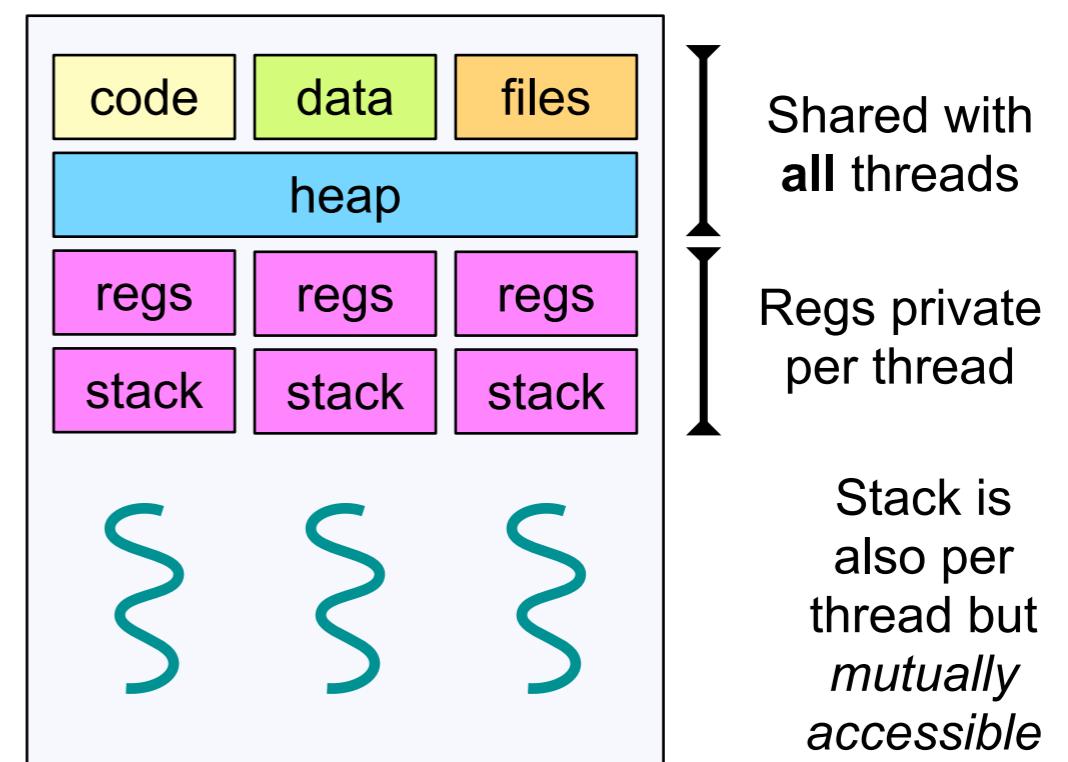
# Definition: Thread of Execution

- ▶ A "*thread*" is an independent instruction stream *inside* a process

- The process *still* contains a single address space
  - Multiple threads all share this address space
- Every thread must be bound to an existing process
- Threads, not processes, are the new units of execution



Single-Threaded Process (What we're used to)



Multi-Threaded Process

# Thread Control Block (TCB)

- ▶ If threads are the new units of execution... PCBs are no longer sufficient!
- ▶ **Idea:** Split the PCB into two parts
  - Piece 1: We still have a PCB, but it's stripped down to hold only the shared process state.
    - Process ID (pid), pointer to address space, open files, sockets
  - Piece 2: *TCBs* hold the data private to each thread (1 TCB per thread)
    - A thread ID (tid)
    - Thread state: CPU regs, SP, BP, PC belonging to that thread
    - Scheduling info: running state, priority of that thread
    - Pointer back to its parent PCB

# Context Switching between Threads

- ▶ Context switches now have two possibilities:
  - Switching two threads in **different** processes. **Slow as before.**
    - Same as switching processes (i.e., it's just like before)
    - Save and restore state of two PCBs (and by extension the two TCBs)
    - **Big deal: Address spaces of 2 processes must be swapped!**
  - Switching two threads in the **same** process. **Faster!**
    - Only have to switch TCBs within same process. No need to save/restore PCB state.
    - The TCB has less state that has to be saved and restored
    - **Big deal: Don't have to switch the address space!**

# Context Switching Threads/Processes (2)

```

void contextSwitch(TCB *currentTCB, TCB *nextTCB) {

  /* save state of current TCB */
  currentTCB->r0 = CPU[r0];
  currentTCB->r1 = CPU[r1];
  currentTCB->r2 = CPU[r2];
  //...
  currentTCB->sp = CPU[sp]; // save stack pointer
  currentTCB->pc = CPU[pc]; // save program counter

  /* we skip exchanging addr space if threads are in the same process */
  if (currentTCB->parentPCB != newTCB->parentPCB) {
    copyToDisk(currentTCB->parentPCB->addrsp);
    copyFromDisk(nextTCB->parentPCB->addrsp);
  }

  /* restore state of next TCB */
  CPU[r0] = nextTCB->r0;
  CPU[r1] = nextTCB->r1;
  CPU[r2] = nextTCB->r2;
  //...
  CPU[sp] = nextTCB->sp;
  CPU[pc] = nextTCB->pc;
}

```



# Advantages of Threads

- ▶ Advantages of threads
  - Much cheaper to use compared to processes!
    - Thread creation time is tiny compared to `fork()`
    - No need to allocate another address space!
    - Don't need an entire PCB either. Just create a new lighter-weight TCB.
    - Reduced time required to context-switch between threads within same process
  - Can improve the responsiveness of programs
  - *Automatic* shared memory (Heap, Data, and Code segments)!
  - *Parallelism*: Scales to multicore CPUs
    - Each thread of execution can be dispatched to an available CPU core

# Disadvantages of Threads

## ► Disadvantages:

- Synchronizing threads is hard.
  - We have a whole book chapter on it!
- You can monopolize the CPU by just creating more threads.
  - More threads means that more "portions" of a process are in the OS queues!
- Overuse of threads isn't always a good thing.
  - Diminishing returns of parallel programs. Would 1000 threads sort a list significantly faster than 4?

# Goals for Today...

- ▶ Threads
- ▶ Implementing Threads
- ▶ **Pthread Examples in C**
- ▶ Parallel Processing Paradigms
- ▶ User Threads vs. Kernel Threads
- ▶ Conclusion

# First, Define the Thread's Work Function

- ▶ What does each thread do? Define its routine using this template:
  - Note that **void\*** is a pointer to a generic
    - (You have to *cast* it later to make sense of it)

```
/**  
 * Defines a single thread's routine.  
 *  
 * @param *arg Generic pointer to some data or structure (can be NULL)  
 * @return a generic pointer to some data or structure (can be NULL)  
 */  
void* routine(void* arg) {  
  
    // 1. if arg != NULL, cast *arg in order to understand its input  
    // 2. Do some work.  
    // 3. return a pointer to the data (or NULL)  
}
```

# pthread Library: Creating a Thread

- ▶ To use threads in C, **#include <pthread.h>** in your program.

```
int pthread_create(  
    pthread_t      *tid,  
    pthread_attr_t *attr,  
    void*          *routine(void *), // what the? It's a function ptr  
    void           *arg);
```

**@param** `*tid`: A pointer to the new thread's unique ID

**@param** `*attr`: Set of attributes for the new thread, use `NULL` for default attributes

**@param** `*routine`: The new thread starts by invoking `routine()`

**@param** `*arg`: Pointer to argument passed to `routine()`

**@return** 0 if successful, otherwise an error number

# pthread Library: Terminating a Thread

- Sometimes we may want to force a thread to terminate.

```
int pthread_cancel(pthread_t tid)
```

`@param tid`: The target thread ID

`@return` 0 if successful, otherwise an error number

- Two cancellation protocols:

- ***Asynchronous***: target thread is immediately terminated
  - What if it was in the middle of updating a shared structure?
- ***Deferred***: target thread periodically checks whether there is a termination request
  - If so, wrap up any critical work and self-terminate (**pthread** does this)

# pthread Library: Terminating a Thread (2)

## ► Deferred Cancellation

- Threads check for cancellation request at various cancellation points.

## ► *Cancellation Points* are predefined locations where it is deemed safe to terminate a thread, such as:

- Blocking on an I/O
- During sleep
  - A list of cancellation points are found here: <https://man7.org/linux/man-pages/man7/pthreads.7.html>
- or at an explicitly defined cancellation points in a thread's code
  - [https://man7.org/linux/man-pages/man3/pthread\\_testcancel.3.html](https://man7.org/linux/man-pages/man3/pthread_testcancel.3.html)

# pthread Library: Thread ID

- ▶ To get the calling thread's ID
  - Usually some random non-negative value
  - Don't just assume it's 0, 1, 2, ...

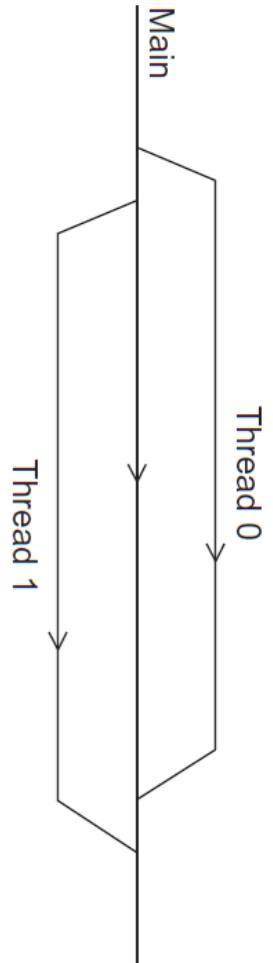
```
pthread_t pthread_self()
```

**@return** thread's identifier. This function always succeeds.

# pthread Library: Reaping a Thread

- ▶ The calling thread waits for thread **tid** to terminate
  - When a dead thread is reaped, its resources are reclaimed
    - *e.g.*, TCB, thread's stack, and register contents

```
int pthread_join(  
    pthread_t tid,  
    void **thread_return  
)
```



**@param** **tid** The thread ID

**@param** **\*\*thread\_return** Pointer to the thread's return pointer.

**@return** 0 if successful, otherwise an error number

# Passing a Single Argument to a Thread

```
void* print_msg(void *arg) {
    char *msg = (char*) arg; // Important: cast the input argument!
    printf("%s\n", msg);
    return NULL;
}

int main(int argc, char* argv[]) {
    // Important: put it on the heap so it's shared!
    char *message = (char*) malloc(10 * sizeof(char));
    strcpy(message, "Hello!!!");

    pthread_t thread1, thread2;
    pthread_create(&thread1, NULL, print_msg, message);
    pthread_create(&thread2, NULL, print_msg, message);
    pthread_join(thread1, NULL);
    pthread_join(thread2, NULL);
    return 0;
}
```

**Output: (Note that we're not sure which thread printed first)**

Hello!!!  
Hello!!!

# Passing Multiple Arguments to a Thread

```
/** Multiple args? Declare a struct to encapsulate input arguments a thread */
typedef struct arg_t {
    int num;
    char *msg;
} arg_t;

/** thread routine */
void* print_msg(void *arg) {
    // cast the input argument!
    arg_t* input = (arg_t*) arg;
    for (int i = 0; i < input->num; i++) {
        printf("%s\n", input->msg);
    }
    return NULL;
}

int main(int argc, char* argv[]) {
    // allocate and populate the input struct on the heap!
    arg_t *params = (arg_t*) malloc(sizeof(arg_t));
    params->num = 500;
    params->msg = (char*) malloc(sizeof(char) * 10)
    strcpy(params->msg, "Hello!!!");

    pthread_t thread;
    pthread_create(&thread, NULL, print_msg, params);
    pthread_join(thread, NULL);
}
```

# Reading Thread's return Value

```
/** thread routine */
void* work(void *arg) {
    int *val = (int*) malloc(sizeof(int));
    *val = 100;
    return val;
}

int main(int argc, char* argv[]) {
    pthread_t thread;
    pthread_create(&thread, NULL, work, NULL);

    // thread return value held in a void* pointer
    void *returned;
    pthread_join(thread, &returned);

    //cast void* to an int*, then de-reference
    int actual = *((int*) returned);
    printf("Here was what the thread returned: %d\n", actual);
}
```

## Output:

```
Here was what the thread returned: 100
```

# Why Is This `work()` Code Broken?

```
/** thread routine */
void* work(void *arg) {
    int val = 100;
    return &val;
}
```

**val** is on the thread's stack, and stacks are private per thread.  
**val**'s address is inaccessible from another thread (**main**).

```
int main(int argc, char* argv[]) {
    pthread_t thread;
    pthread_create(&thread, NULL, work, NULL);

    // thread return value held in a void* pointer
    void *returned;
    pthread_join(thread, &returned);

    //cast void* to an int*, then dereference
    int actual = *((int*) returned);
    printf("Here was what the thread returned %d\n", actual);
}
```

*Dereferencing **val** leads to segfault!*

## Output:

*Segmentation fault*

# What's Going On?

```

int x = 1; // global (shared)

void* print_msg(void *arg) {
    x++; // increment x!
    if (x == 2) {
        printf("Thread %lu: I win!\n", pthread_self());
    }
    else {
        printf("Thread %lu: I lost!\n", pthread_self());
    }
    return NULL;
}

```

*Code that threads will run*

Begin threads  
 Thread 1: I Win  
 Thread 2: I lost  
  
*(Blue won the race!)*

```

int main(int argc, char* argv[]) {
    /* fork phase */
    printf("Begin threads\n");
    pthread_t thread1, thread2;
    pthread_create(&thread1, NULL, print_msg, NULL);
    pthread_create(&thread2, NULL, print_msg, NULL);

    /* join phase */
    pthread_join(thread1, NULL);
    pthread_join(thread2, NULL);
    return 0;
}

```

Begin threads  
 Thread 2: I Win  
 Thread 1: I lost  
  
*(Pink won the race!)*

Begin threads  
 Thread 1: I lost  
 Thread 2: I lost  
  
*(What the? Yes it can happen)*

# Goals for Today...

- ▶ Threads
- ▶ Implementing Threads
- ▶ Pthread Examples in C
- ▶ **Parallel Processing Paradigms**
- ▶ User Threads vs. Kernel Threads
- ▶ Conclusion

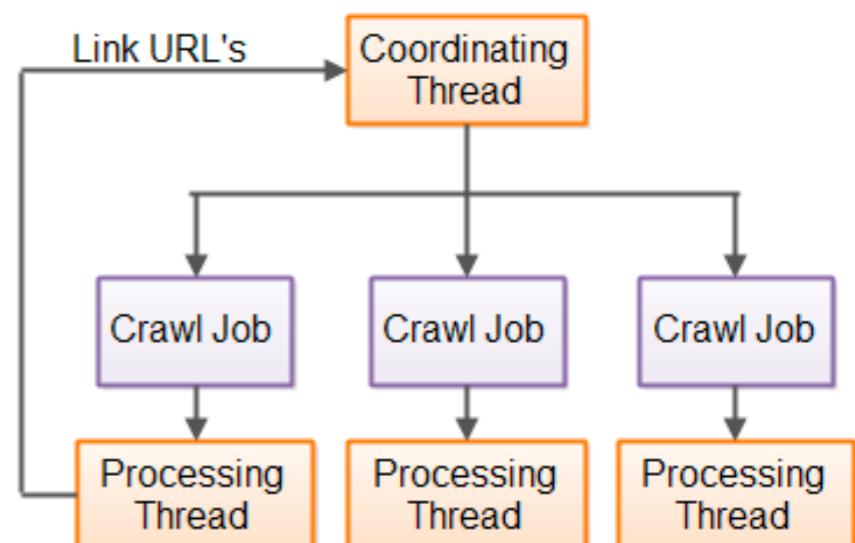
# Two Approaches to Parallel Computing

- ▶ *Task Parallelism* when the problem can be decomposed into subtasks. Each thread works on an independent subtask.
  - Git
    - To push all the files in your git repository onto Github
    - Create a thread to upload each file.

# Two Approaches to Parallel Computing

- ▶ ***Task Parallelism*** when the problem can be decomposed into subtasks. Each thread works on an independent subtask.

- Git
  - To push all the files in your git repository onto Github
  - Create a thread to upload each file.
- Web Crawler
  - To discover all the webpages out there, start with a random page
  - Create a new thread to retrieve each link found on the page. Repeat.



# Two Approaches to Parallel Computing

- ▶ *Data Parallelism*: when the data set can be decomposed into smaller chunks. Each thread works on a subset of data.
  - Sort a massive list:
    - View the list as multiple sublists.
    - Threads sort each sublist.
    - Main thread merges all sublists afterwards.
  - Sudoku game validation:
    - View the game board as multiple columns, rows, and 3x3 boxes
    - Threads validate each column, row, box
    - If none of the threads return false, the game is valid!

7	3	5	6	1	4	8	9	2
8	4	2	9	7	3	5	6	1
9	6	1	2	8	5	3	7	4
2	8	6	3	4	9	1	5	7
4	1	3	8	5	7	9	2	6
5	7	9	1	2	6	4	3	8
1	5	7	4	9	2	6	8	3
6	9	4	7	3	8	2	1	5
3	2	8	5	6	1	7	4	9

# Two Approaches to Parallel Computing

## ► *Another Data Parallel example:*

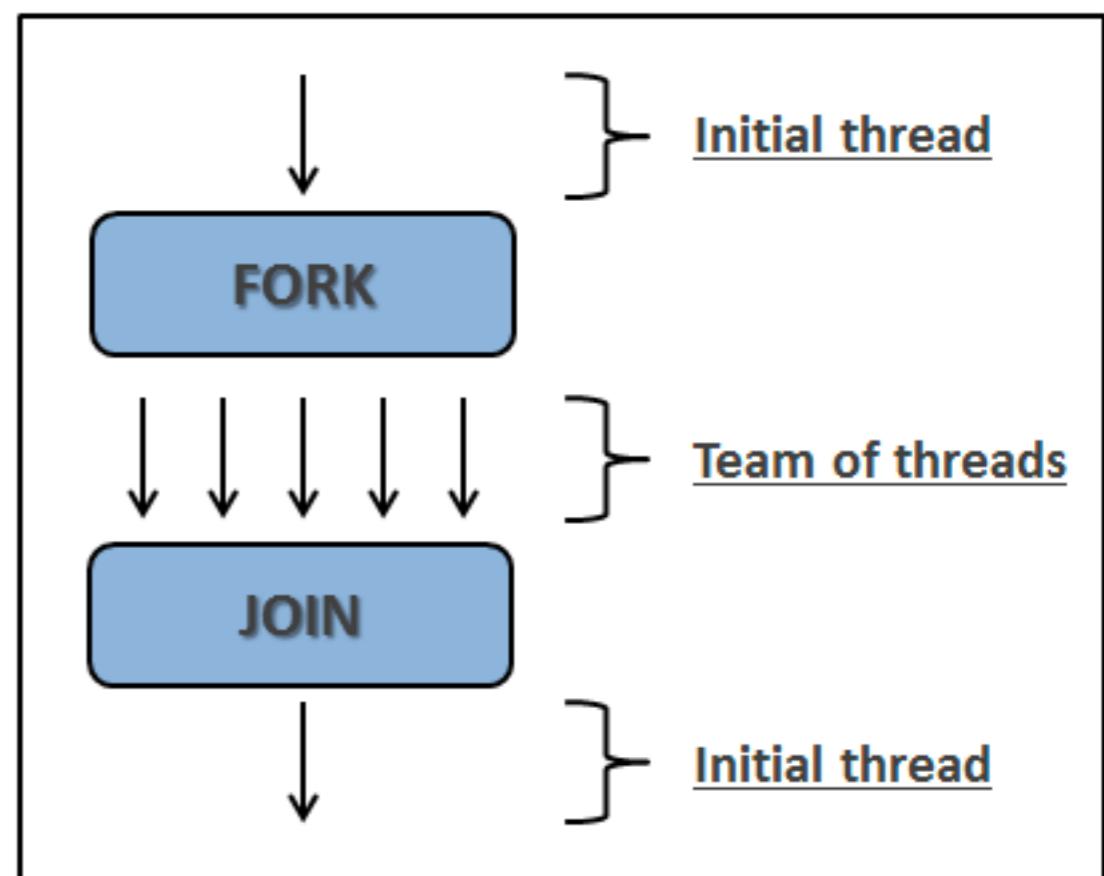
- Matrix multiplication
  - Two input matrices A, B, and a result matrix C are in shared space.
  - Threads calculate the product(s) defined for their "region."
    - (Maybe each thread calculates a chunk of rows in C? Up to programmer)
  - Main thread simply waits for all worker threads to finish, and exits.

$$\begin{bmatrix} a_1 & a_2 & a_3 \\ a_4 & a_5 & a_6 \\ a_7 & a_8 & a_9 \end{bmatrix}
 \begin{bmatrix} b_1 & b_2 & b_3 \\ b_4 & b_5 & b_6 \\ b_7 & b_8 & b_9 \end{bmatrix}
 =
 \begin{bmatrix} c_1 & c_2 & c_3 \\ c_4 & c_5 & c_6 \\ c_7 & c_8 & c_9 \end{bmatrix}$$

# Fork-Join Paradigm of Parallel Programming

► *In both approaches, the Fork-Join approach can be applied.*

- Given a problem of size  $n$
- Main thread creates  $m$  worker threads
  - Each worker does some fraction of work of the **same nature**
  - Each worker might derive some partial result
- Main thread waits for all threads to finish and reduces partial results to a final one.



# Example: Parallel Count Odds

- ▶ Let's start with a small ***Data Parallel*** example:
  - Count up all the odd numbers in a given array of 100,000 numbers
- ▶ Fork-Join setup:
  - Malloc the array (so it's shared).
  - Assign the array with random numbers.
  - **Fork:** Create 2 threads to count the first and second halves of the array.
  - **Join:** Add up both counts and print.

# Parallel Count Odds (Preamble)

## ► Preamble code area

```
#include <pthread.h>
#include <stdio.h>
#include <unistd.h>
#include <stdlib.h>

const int N = 1000000;
int *array; // points to the array of ints on the heap
int *results;
```

# Parallel Count Odds (Two Threads)

```

void* worker(void *param) {
    int *whoAmI = (int*) param; // either thread 0 or 1
    Cast the parameter to make sense of it! (It's just a thread ID)
    int begin = (*whoAmI % 2 == 0) ? 0 : N/2;
    int end = (*whoAmI % 2 == 0) ? N/2 : N;
    printf("thread %d start at [%d] and end at [%d]\n", *whoAmI, begin, end);

    int count = 0;
    for (int i = begin; i < end; i++) {
        if (array[i] % 2 == 1) {
            count++;
        }
    }
    results[*whoAmI] = count;
    free(param);
    return NULL;
}

```

# Parallel Count Odds (Two Threads)

```

void* worker(void *param) {
    int *whoAmI = (int*) param; // either thread 0 or 1

    int begin = (*whoAmI % 2 == 0) ? 0 : N/2;
    int end = (*whoAmI % 2 == 0) ? N/2 : N;
    printf("thread %d start at [%d] and end at [%d]\n", *whoAmI, begin, end);

    int count = 0;
    for (int i = begin; i < end; i++) {
        if (array[i] % 2 == 1) {
            count++;
        }
    }
    results[*whoAmI] = count;
    free(param);
    return NULL;
}

```

*Depending on which thread I am (0 or 1), count the 1st or 2nd half of the array.*

# Parallel Count Odds (Two Threads)

```

void* worker(void *param) {
  int *whoAmI = (int*) param; // either thread 0 or 1

  int begin = (*whoAmI % 2 == 0) ? 0 : N/2;
  int end =   (*whoAmI % 2 == 0) ? N/2 : N;
  printf("thread %d start at [%d] and end at [%d]\n", *whoAmI, begin, end);

  int count = 0;
  for (int i = begin; i < end; i++) {
    if (array[i] % 2 == 1) {
      count++;
    }
  }
  results[*whoAmI] = count;
  free(param);
  return NULL;
}

```

*Work on my just half!*

# Parallel Count Odds (Two Threads)

```

void* worker(void *param) {
  int *whoAmI = (int*) param; // either thread 0 or 1

  int begin = (*whoAmI % 2 == 0) ? 0 : N/2;
  int end =   (*whoAmI % 2 == 0) ? N/2 : N;
  printf("thread %d start at [%d] and end at [%d]\n", *whoAmI, begin, end);

  int count = 0;
  for (int i = begin; i < end; i++) {
    if (array[i] % 2 == 1) {
      count++;
    }
  }
  results[*whoAmI] = count;
  free(param);
  return NULL;
}

```

*Each thread deposits their counts in the corresponding spot.*

# Parallel Count Odds (main function)

```

int main(int argc, char *argv[]) {
    srand(0);      // "seed" the random number generator
    array = (int*) malloc(sizeof(int)*N);
    for (int i = 0; i < N; i++) {
        array[i] = rand();
    }

    // malloc an array for threads to deposit their results
    results = (int*) malloc(sizeof(int) * 2);

    // spin up the two worker threads!
    pthread_t tid[2];
    for (int i = 0; i < 2; i++) {
        int *threadID = (int*) malloc(sizeof(int));
        *threadID = i;
        pthread_create(&tid[i], NULL, worker, threadID);
    }

    // wait for all the threads to finish
    int sum = 0;
    for (int i = 0; i < 2; i++) {
        pthread_join(tid[i], NULL);
        sum += results[i];
    }
    printf("Total count: %d\n", sum);
    free(results);
    free(array);
    return 0;
}

```

*Allocate the array on the heap.  
Then populate it with random nums*

# Parallel Count Odds (main function)

```

int main(int argc, char *argv[]) {
    srand(0); // "seed" the random number generator
    array = (int*) malloc(sizeof(int)*N);
    for (int i = 0; i < N; i++) {
        array[i] = rand();
    }

    // malloc an array for threads to deposit their results
    results = (int*) malloc(sizeof(int) * 2);

    // spin up the two worker threads!
    pthread_t tid[2];
    for (int i = 0; i < 2; i++) {
        int *threadID = (int*) malloc(sizeof(int));
        *threadID = i;
        pthread_create(&tid[i], NULL, worker, threadID);
    }

    // wait for all the threads to finish
    int sum = 0;
    for (int i = 0; i < 2; i++) {
        pthread_join(tid[i], NULL);
        sum += results[i];
    }
    printf("Total count: %d\n", sum);
    free(results);
    free(array);
    return 0;
}

```

*Allocate the array on the heap.  
Then populate it.*

*results[...] will store each  
threads' partial counts.*

*Prepare thread inputs: Just their IDs.  
Create both threads.*

# Parallel Count Odds (main function)

```

int main(int argc, char *argv[]) {
    srand(0); // "seed" the random number generator
    array = (int*) malloc(sizeof(int)*N);
    for (int i = 0; i < N; i++) {
        array[i] = rand();
    }

    // malloc an array for threads to deposit their results
    results = (int*) malloc(sizeof(int) * 2);

    // spin up the two worker threads!
    pthread_t tid[2];
    for (int i = 0; i < 2; i++) {
        int *threadID = (int*) malloc(sizeof(int));
        *threadID = i;
        pthread_create(&tid[i], NULL, worker, threadID);
    }

    // wait for all the threads to finish
    int sum = 0;
    for (int i = 0; i < 2; i++) {
        pthread_join(tid[i], NULL);
        sum += results[i];
    }
    printf("Total count: %d\n", sum);
    free(results);
    free(array);
    return 0;
}

```

*Allocate the array on the heap.  
Then populate it.*

*results[...] will store each  
threads' partial counts.*

*Prepare thread inputs: Just their IDs.  
Create both threads.*

*Join the threads, add up their partial  
counts, which are stored in  
**results**[...]*

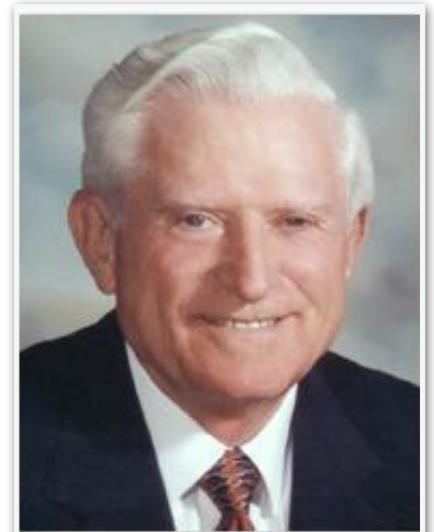
# Amdahl's Law

- ▶ 1967 landmark keynote speech
  - *"Validity of the Single Processor Approach to Achieving Large Scale Computing Capabilities."* 1967 AFIPS Spring Joint Computer Conference (SJCC'67). [\[link here\]](#)
- ▶ How to express "parallel speedup?"
  - Diminishing returns of parallel architectures

$$\text{Speedup} = \frac{1}{(1 - f) + \frac{f}{n}}$$

$f$  = fraction of code that can be parallelized

$n$  = degree of parallelization (threads/CPU Cores)



Gene Amdahl

# Example: Diminishing Returns.

- ▶ Suppose 40% of your code can be parallelized by a factor of 8. What's the overall speedup of the parallel version?

$$Speedup = \frac{1}{(1-f) + \frac{f}{n}} = \frac{1}{(1-0.4) + \frac{0.4}{8}} = 1.54 = 54\%$$

- (Yikes that's no where near an 8x (= 800%) speedup!)
- ▶ What if we had a 16 core machine? That's twice the number of cores!

$$Speedup = \frac{1}{(1-0.4) + \frac{0.4}{16}} = 1.6 = 60\%$$

# Reality: Opposing Laws

## ► **Question:** Spend lots of time + money investing in parallel computing?

- *Amdahl's Law says:* "No, extracting parallelism is hard!"
  - Diminishing returns! Error prone!
  - Amount of parallelism must overcome thread-management overhead.
- *Moore's Law:*
  - Pre-2007: Amdahl is right... just wait for the faster processor to come out
  - After-2007: Processors getting slower, but we now have more cores.
    - Parallel computing becomes a necessity for current generation of coders
    - (ACM ugrad curriculum now calls for parallel computing as part of "Core CS")

# Checklist for Making Parallel Programs

- ▶ Decide how to split up the work, so that each thread:
  - Performs the same nature of work, but works on a different "portion"
  - Write that function using
    - What info does each thread need to know?
    - (For instance, a pointer to an array? The start and end index of a portion?)
- ▶ Back to **main()**:
  - Where do threads put their results?
    - Do they return it? Do they write it to place in the heap?
  - Prepares the input data for each thread, creates threads, joins them up
  - What to do with threads' results?

# Your Turn!

- ▶ Some exercises you can try
  - Parallel sort algorithm
    - Split list into N segments
    - Each segment is sorted (insertion sort) by a different thread
    - Merge all sorted sublists into final sorted list
    - **Solution:** <https://github.com/davidtchiu/cs475-parInsertionSort>
  - Parallel search algorithm
    - Split list into N parts, assign each thread to search a different chunk of the array
  - Parallel Riemann Sum for computing integrals

# Goals for Today...

- ▶ Threads
- ▶ Implementing Threads
- ▶ Pthread Examples in C
- ▶ Parallel Processing Paradigms
- ▶ **User Threads vs. Kernel Threads**
- ▶ Conclusion

# User Threads

- ▶ In early systems, the OS only knew about processes.
  - Processes were a single execution stream
  - But users needed intra-process responsiveness:
    - Like, `computePI(); printAmericasFunnyEmail();`
- ▶ "*User Threads*" are completely managed by the programmer in unprivileged user mode
  - The OS wouldn't even know that multiple threads exist!
  - Users had to implement "thread support functions" in a **user library**
    - `CreateThread()`, `SwitchThread()`, `JoinThread()`, ...

# User Threads (Pros)

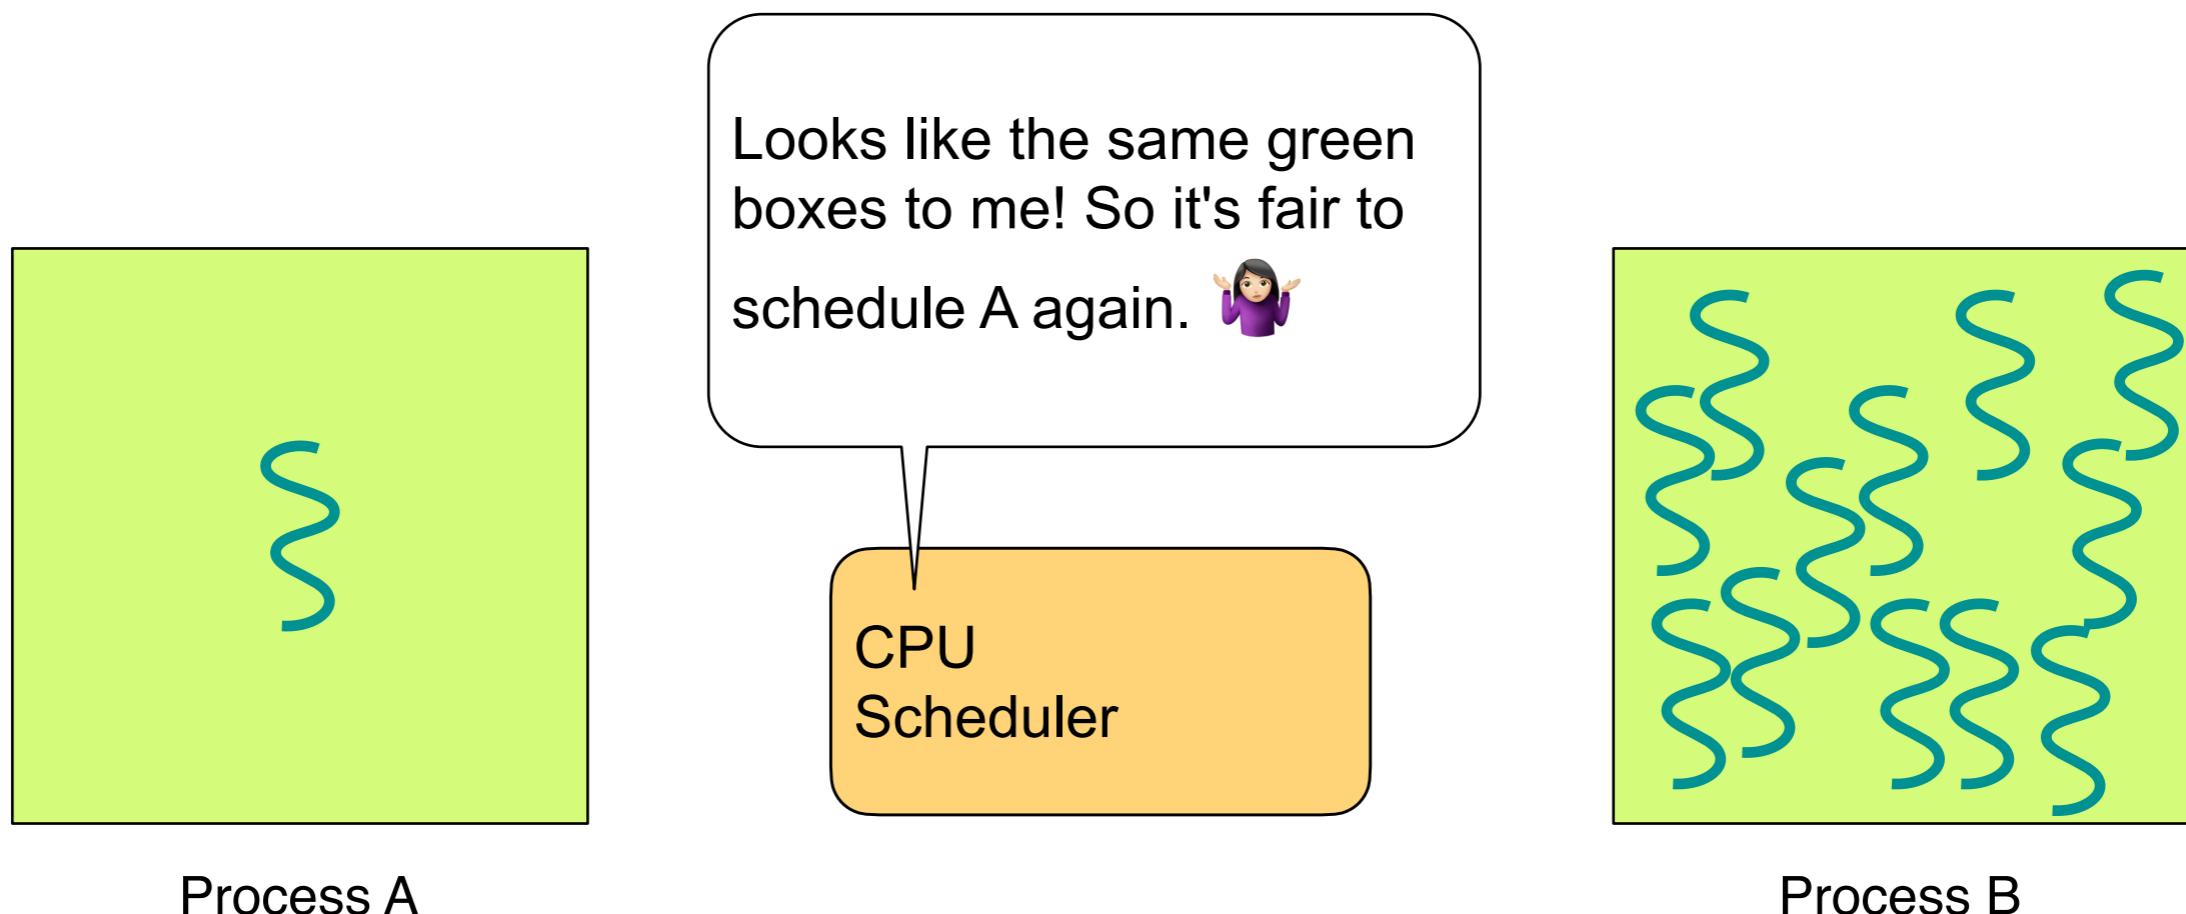
- ▶ Does not require any modification to the OS
  - Even if an OS doesn't support threads directly, a user can still write programs with threads.
- ▶ Thread operations (creation, termination, switching) are very fast
  - Orders of magnitude faster than switching processes
  - (Because the user code is doing the "switching," not the OS!)

# User Threads (Cons)

- ▶ If one user thread blocks on I/O, *it stalls the entire process*
  - A thread makes a system call, all other threads wait with it! (*Why?*)
- ▶ Can't leverage parallel architectures
  - Without kernel support, the OS still dispatches jobs on the granularity of a process.
- ▶ No OS protection of private data
  - A user thread's stack & registers not private from another user thread
- ▶ Fairness in CPU Scheduling is now a concern
  - (Next slide)

# User Threads (Cons)

- ▶ OS may not make *fair, insightful* decisions when scheduling
  - Say each process gets a fixed slice of CPU time
  - OS assigns same time slice, whether process has 1 thread or 1000!



# Kernel Threads

- ▶ *Kernel Threads* (In use today)
  - OS *knows* about all threads in execution.
  - Thread management functions (creation, join, termination) done via system calls.
  - OS views kernel threads as schedule-able units and knows to context switch between them.
- ▶ All modern OS provides a **Thread API** to users to create and manage threads
  - In C, it's called the **pthread.h** library.

# Kernel Threads (Cons)

- ▶ Creating, destroying threads would be system calls (slow)
- ▶ A context switch between two threads in same process is over an order of magnitude slower than in user threads!

### **"Null Fork" Times on VAX uniprocessor**

User Thread	Kernel Thread	Processes
34 us	948 us	11,300 us

"NULL fork" means: Create, schedule, execute, complete a no-op

# OS Examples

- ▶ The first OS's did not have native kernel thread support.
  - For responsiveness, you'd have to implement a user-thread library!
  - Fun fact: C's pthread library was originally a user library for thread support in early Unix systems
- ▶ Switch to kernel support in 1980-1990s
- ▶ All modern OS now support kernel threads.
- ▶ But some OS kept user threads around additionally!
  - (Example: Windows 10 "fibers")



# OS Examples: Windows 10

- ▶ Windows 10 added support for fibers.
- ▶ What are "*fibers*"?
  - Fibers are user threads!
  - Users make a call to `yield()` CPU to another fiber.

# Abstraction: Thread Assignment Models

## ► *N-to-1 Model*

- Multiple user threads associated with one kernel thread
- User-Level threads used this model
- Examples: Xinu, early Unix, DOS

## ► *1-to-1 Model*

- One user-level thread is associated with one kernel thread
- Most operating systems today: Linux, OS X, BSD, iOS

# Goals for Today...

- ▶ Threads
- ▶ Implementing Threads
- ▶ Pthread Examples in C
- ▶ Parallel Processing Paradigms
- ▶ User Threads vs. Kernel Threads
- ▶ Conclusion

# In Conclusion

- ▶ Why threads?
  - Processes are inefficient (in both space and time)
  - But programs need to multitask and share data
  - Threads are fast, and they are now the OS's unit of execution
- ▶ So... how does `runScheduler()` work? (*Next Topic: CPU Scheduling*)

```
while (1) {  
    curTCB  = getRunThread();  
    nextTCB = runScheduler();  
    ctxsw(curTCB, nextTCB);  
}
```

# Administrivia 2/24

- ▶ Midterm exam Friday.
- ▶ Hwk 4 Monday 3/3.
- ▶ Last time... Finished process management
  - New topic: Limitations of processes (and why we need threads)
  - Implementation of threads
- ▶ Today: Read Chap 4.1-4.4 (Dinosaur book)
  - The POSIX thread (**pthread**) library for C
  - User threads vs. Kernel threads

# Administrivia 3/3

- ▶ Hwk 4 due tonight!
- ▶ Hwk 5 posted, due Friday 3/14
- ▶ Last time...
  - Read Chap 4.1-4.4 (Dinosaur book)
  - The POSIX thread (**pthread**) library for C
- ▶ Today
  - Pthread example: Parallel Counter
  - Amdahl's Law on parallel speedup limitations
  - Where to implement threads?
    - User space vs. Kernel space