

CS 475

Operating Systems



Department of Mathematics
and Computer Science

Lecture 6
Synchronization (Part I)


Problem: Race Conditions

- ▶ *Race Condition*: When 2 or more threads read and modify a shared resource without *synchronization*.
- ▶ The **correctness of the end result** depends on:
 - Which context switches occurred, if any
 - Which thread ran after the context switch
 - Which thread got switched out
 - What the threads were doing
 - Reading the shared variable? Writing to the shared variable? Neither?
 - *Difficulty: Nondeterministic runtime sequencing; may not be repeatable (race conditions are super hard to debug)*



Too Much Milk Problem



- ▶ To understand *race conditions*... let's consider this scenario
 - Ren and Stimpy live together, but don't like to communicate 😞

Time	Ren	Stimpy
3:00	Look in fridge	
3:05	"Out of milk!"	
3:10	Leave for store	Look in fridge
3:15	Arrive at store	"Out of milk!"
3:20	Buy Milk	Leave for store
3:25	Arrive at home, put milk away	Arrive at store
3:30		Buy Milk
3:35		Arrive at home, put milk away

Too Much Milk Problem (2)



Actual picture of fridge
from a 2016 OS student



Code Example of a Race Condition

- ▶ Assume that `y` is shared between two threads:

Initially, `y = 12`;

Thread 1

```
x = 1;  
x = y + 1;
```

Thread 2

```
y = 5;  
y = y * 2;
```

- ▶ Let the threads run concurrently. *What is the value of `x` after both threads are done executing?*

Code Example of a Race Condition

- ▶ Assume that `y` is shared between two threads:

Initially, `y = 12`;

Thread 1

```
x = 1;  
x = y + 1;
```

Thread 2

```
y = 5;  
y = y * 2;
```

- ▶ Let the threads run concurrently. *What is the value of `x` after both threads are done executing?*
- ▶ *How many different schedules could there be? (In this case, 6 schedules!)*

$$\frac{\left(\sum_i O_i\right)!}{\prod_i O_i!}$$
 where O_i is the number of operations in the i th thread

More Fundamentally...

- ▶ Assume variable **x** is shared between two threads.

Thread 1

```
x = 1;
```

Thread 2

```
x = 2;
```

- ▶ Can **x** be 3 after the two threads run?


How could it be 3? What if register access is interleave-able?

Can x Be 3? (No - phew!)

► **Definition:** *Atomic Operations*

- An operating is *atomic* if it cannot be broken up into multiple smaller operations
- In modern CPU architectures, register references (e.g., loads and stores) are guaranteed to be atomic (*Phew!*)

(But this still doesn't fix nondeterminism caused by race conditions between atomic operations!)



Another Example: What Gets Printed?

```
#include <pthread.h>
#include <stdio.h>

int x = 0; // shared global

void* doStuff(void* args) {
    for (int i = 0; i < 1000000; i++) {
        x++;
    }
    return NULL;
}

int main() {
    pthread_t t1, t2;
    pthread_create(&t1, NULL, doStuff, NULL);
    pthread_create(&t2, NULL, doStuff, NULL);
    pthread_join(t1, NULL);
    pthread_join(t2, NULL);
    printf("Value of x is %d\n", x);
    return 0;
}
```

add2m.c

I want to increment **x** to 2,000,000 but split the work across 2 threads to do it in half the time!

Output:

```
$ gcc add2m.c -lpthread -o add2m
$ ./add2m
Value of x is 1029184

$ ./add2m
Value of x is 1224682

$ ./add2m
Value of x is 1004207

$ ./add2m
Value of x is 1003645
```

Reason: The **++** operator is not atomic!

Goals for This Lecture...

- ▶ **Basic Problem Definition**
 - Race Conditions and the Critical Section Problem
- ▶ **Mechanisms to Control Access to Shared Resources**
 - Low Level Mechanisms:
 - Busy-Waiting (Spin) Locks
 - Self-Blocking Locks
 - High Level Mechanisms:
 - Semaphores
 - Condition Variables and Monitors


Critical Section Problem

► Definitions:

- ***"Critical Section"***
 - A fragment of code where multiple threads read+write shared data.
 - Its execution must be controlled to prevent race conditions.
- ***"Critical Section Problem"***
 - To ensure that multiple threads access shared resources in a way that prevents race conditions while maintaining
 - Mutual exclusion
 - Progress
 - Bounded wait

Critical Section Problem (2)


- ▶ An analogy: Roads and intersections
 - Roads = Threads
 - Intersections = Shared Resource
 - Its *access* represents the critical section
- ▶ How do we go about solving the **intersection** problem in real life?
 - Signals, stop signs, yield signs, etc.
 - All are used to **synchronize** access to the intersection
- ▶ 3 desirable features of a CS solution:
 - (Next)



Critical Section Feature: Mutual Exclusion

► Feature 1: Mutual Exclusion (this property is non-negotiable)

- Only **one** thread is allowed to access the CS at any point in time
- Implies all other threads wanting access to the critical section must wait




Does the **traffic light** provide mutual exclusion of the intersection?

Yes, it allows traffic on one road to be active in the critical section. All other traffic must wait.

Critical Section Feature: Progress

► Feature 2: Progress

- If the CS is available (no other thread is executing it), a requesting thread must be permitted entry without delay.
 - Otherwise, the thread is waiting for no reason, which hurts performance!



Does the **traffic light** provide progress?


Not really. Sometimes you have to wait at a red light even when there's no traffic on the other road.

Real-world solution? A signal tripper?

Critical Section Feature: Bounded Wait

▶ Feature 3: Bounded Wait

- A requesting thread cannot be delayed indefinitely.



Does the **traffic light** provide bounded wait?

Yes. Cars converging at the intersection will eventually get its turn


Back to the Too Much Milk Problem

- ▶ Identify the **critical section** of code:

- Actions that we don't want two roommates to be doing at the same time

Time	Ren	Stimpy
3:00	Open fridge	
3:05	<i>Look for milk -- Out of milk!</i>	
3:10	Leave for store	Open fridge
3:15	Arrive at store	<i>Look for milk -- Out of milk!</i>
3:20	Purchase Milk	Leave for store
3:25	Arrive home, put milk away	Arrive at store
3:30		Purchase Milk
3:35		Arrive home, put milk away

These "code blocks" are not atomic, and thus should not be allowed to interleave.




Solution #0: Lock the Fridge!

- ▶ What can be used as our "traffic light"?

```
while (1) {
    Lock the fridge & take the key
    if (noMilk) {
        buyMilk();
    }
    Unlock fridge
    Replace key in common area
}
```

- ▶ Is the pad-lock a correct CS solution?

- Mutual exclusion? Yes.
- Progress? Yes, grab the key from the common area.
- Bounded Wait? Assuming Ren or Stimpy don't lose key and returns home, yes.
- **It works!** But it's a poor solution. Poor performance to shared resource.



Solution #1: Leave a Note

- ▶ Leave note instead! Still allows fridge to be accessed for other items.

```
while (1) {  
    if (noMilk) {  
        if (noNote) {  
            leaveNote();  
            buyMilk();  
            removeNote();  
        }  
    }  
}
```



- ▶ *Is the code correct?*
 - Mutual exclusion? *No! Could still have too much milk!*
 - Progress? *Doesn't matter*
 - Bounded Wait? *Doesn't matter*

Solution #2: Leave Note First!

- ▶ We left the note too late! Instead, let's:
 - Leave a note first
 - (Don't let my roommate inspect the fridge while I'm inspecting)
 - Check if no milk, then go buy
 - Remove note after buying
 - Don't buy milk if there is a note.

```
while (1) {  
    leaveNote();  
    if (noMilk) {  
        if (noNote) {  
            buyMilk();  
        }  
    }  
    removeNote();  
}
```

*Mutual Exclusion? Well, yes, because
no thread ever enters.*

Solution #3? Use Colored Notes!

- ▶ Distinguish the notes (lesson from Solution #2)
 - Ren uses blue notes
 - Stimpy uses red notes
- ▶ Leave a note first (lesson from Solution #1)
 - Check if no note from *other* roommate
 - Check if no milk, then buy
 - Remove note after buying
- ▶ Don't buy milk if there is a note



Solution #3 (Cont.)

- ▶ This solution guarantees **mutual exclusion**.
 - Is there still a problem with **progress** and **bounded wait**?

Stimpy's critical section

```
while (1) {  
    leaveNote(stimpy);  
    if (noNote(ren)) {  
        if (noMilk)  
            buyMilk();  
    }  
    removeNote(stimpy);  
}
```

Ren's critical section

```
while (1) {  
    leaveNote(ren);  
    if (noNote(stimpy)) {  
        if (noMilk)  
            buyMilk();  
    }  
    removeNote(ren);  
}
```

Solution #3 (Cont.)

- ▶ This solution guarantees **mutual exclusion**.
 - How about **progress** and **bounded wait**?
 - Progress is not guaranteed.
 - Both threads leave notes right before either thread checks for note
 - Bounded wait is not guaranteed.
 - Say Stimpy wants in, but always lags behind Ren, so he always sees Ren's note

Stimpy's critical section

```
while (1) {
    leaveNote(stimpy);
    if (noNote(ren)) {
        if (noMilk)
            buyMilk();
    }
    removeNote(stimpy);
}
```

Ren's critical section

```
while (1) {
    leaveNote(ren);
    if (noNote(stimpy)) {
        if (noMilk)
            buyMilk();
    }
    removeNote(ren);
}
```

Locks

- ▶ **Bottom Line:** We need a simple mechanism that can scale to any number of threads!
 - Desired usage:

```
acquire(lock);
//critical section starts
if (noMilk) {
    buyMilk();
}
//critical section ends
release(lock);
```



- Preferably the two operations are in user space so we don't need to bother the OS each time we want to (un)lock.

Goals for This Lecture...

- ▶ Basic Problem Definition
- ▶ Mechanisms to Control Access to Shared Resources
 - Low Level Mechanisms:
 - Disabling (masking) interrupts
 - Busy-Waiting (Spin) Locks
 - Self-Blocking Locks
 - High Level Mechanisms:
 - Semaphores
 - Condition Variables and Monitors

Implement Locks by Disabling Interrupts?

- ▶ First try at implementing locks in user space.
 - Attack the source: interrupts!
 - We need to stop context switches from occurring, and they occur on interrupt!
 - Idea: Let users disable and enable interrupts!
 - **cli** is the assembly instruction to disable interrupts on Intel x86.
 - **sti** is the assembly instruction to enable them.
 - Should we let **cli** and **sti** be unprivileged so anyone can call them??

Implement Locks by Disabling Interrupts (3)

► Problems Abound!

- What if the thread needs to do I/O or make syscall in the CS? **Can't!**
- What if programmer forgets to unlock the CS before leaving? **Stuck.**
- Or, what if a user makes a silly mistake like the following... **Stuck.**

```
void davidsThread() {
    cli();
    while(1); // Oops I put a ; here
    {
        /* important critical
           section code here */
    }
    sti();
}
```

Summary of Using `cli` and `sti` as Locks

- ▶ System responsiveness?
 - A programmer can freeze up the whole system
- ▶ Effects on multicore CPUs?
 - Disabling interrupts is done on granularity of a single core!
 - Other cores unaffected; Can still run threads that access critical section
 - Must disable interrupts on all cores at once (There's no such thing)
- ▶ In summary, this is too aggressive (like locking out the whole fridge)
 - Any user thread could bring the system to a crawl or break it!
 - `cli` and `sti` must be privileged, and we need to find a better way

Goals for This Lecture...

- ▶ Basic Problem Definition
- ▶ Mechanisms to Control Access to Shared Resources
 - Low Level Mechanisms:
 - Disabling interrupts
 - Busy-Waiting (Spin) Locks
 - Self-Blocking Locks
 - High Level Mechanisms:
 - Semaphores
 - Condition Variables and Monitors

Another Attempt: Spin Locks (User Mode)

- ▶ A 2nd attempt at a **user-space** locking solution: *Spin locks*
 - "Busy waiting"
- ▶ Spin lock definition:
- ▶ Example Usage:

```

typedef struct lock_t {
    int held = 0;
} lock_t;

void acquire(lock_t *L) {
    while (L->held)
        ; // Busy wait (spin)

    // We're in! Lock up!
    L->held = 1;
}

void release(lock_t *L) {
    L->held = 0;
}
  
```

```

// myLock shared among threads
lock_t myLock;

// pthreads run this
void *threadWork(void *args) {
    acquire(&myLock);

    // <do critical stuff>

    release(&myLock);
}
  
```

Spin Locks (Cont.)

- ▶ But does the busy-wait loop work?
- ▶ Let's say a thread **T1** is in the critical section
 - OS switches to **T2**
 - Then **acquire()** is called by thread **T2**
 - **T2** waits... eventually switches back to **T1**
 - **T1** calls **release()** on the lock:
 - **l->held** now 0
 - **T2** exits spin loop
 - But what if a switch back to **T1** occurs *right before* **T2** sets **l->held** to 1?

```
void acquire(lock_t *L) {  
    while (L->held)  
        ;  
    Context Switch here (Noooooooo)  
    L->held = 1;  
}
```

acquire() and release() Need to Be Atomic!

- ▶ **Problem:** the code for `acquire()` and `release()` are also *critical sections!*
- *Can't be broken up. There is a race condition!*

```
typedef struct lock_t {
    int held = 0; //shared
} lock_t;
```

```
void acquire(lock_t *L) {
    while (L->held) {
        ;
    }
    L->held = 1;
}
```

```
void release(lock_t *L) {
    L->held = 0;
}
```

There's a critical section in `acquire()` !!

This block of "*test it, then set it*" code needs to be atomic!

`release()` is already atomic.

(Recall assignment statements are guaranteed atomic by the CPU!)

Need User-Mode Hardware Support!!!

- Most CPUs provide a **test_and_set()** op to be *an atomic user operation!*

```
int test_and_set(int *flag) {
    int res = *flag;
    *flag = 1;
    return res;
}
```

CPU hardware guarantees this code block cannot be broken up (i.e., it is atomic)

- Need to combine the **while-test** and the **assignment statement** into **one** indivisible (atomic) CPU instruction!

```
void acquire(lock_t *L) {
    while (L->held) {
        ;
    }
    L->held = 1;
}
```



*re-written
atomically*

```
void acquire(lock_t *L) {
    while (test_and_set(&L->held))
        ;
}
```

A working user-mode spin lock! (Yes!!)

Using Spin-Locks in C: API

- ▶ The **pthread** Library supports spin locks
 - (OS X doesn't support it due to energy efficiency)

```
#include <pthread.h>

/** 
 * Initializes a spin lock.
 * @param pshared PTHREAD_PROCESS_SHARED if shared in multiple processes, or
 *                 PTHREAD_PROCESS_PRIVATE if lock is shared between threads
 *                 in same process.
 */
pthread_spin_init(pthread_spinlock_t *lock, int pshared);

/** Destroys (frees resources for) a lock */
pthread_spin_destroy(pthread_spinlock_t *lock);

/** Acquire */
pthread_spin_lock(pthread_spinlock_t *lock);

/** Release */
pthread_spin_unlock(pthread_spinlock_t *lock);
```

Summary of Spin Locks

▶ Pros:

- Easy to implement with help from CPU hardware
 - Exploit atomic user operations like `test_and_set()`, `cpxchg()`, ...
- **All done in user mode!** Fast!!
- Don't have to manually disable interrupts!
 - Works on multi-core processors
 - Can be used to lock out individual data elements
 - Use a different lock for each array element (instead of locking out entire array)

▶ Cons:

- Busy waiting is wasteful!
- Susceptible to the *Priority Inversion Problem*



Big Problem with Spin-Locks: Priority Inversion

▶ *Priority-Inversion Problem*

- High priority thread **starves** because a low priority thread has a spin-lock that the high priority thread needs.

▶ Mars Pathfinder/Sojourner Mission (1997)

- Repeated system reset while on Mars.



Priority Inversion (2)

- Shared resource: **InfoBus** (just a shared malloc'd array)
- Threads:
 - T_c : **Communication thread (Med priority)**
 - Slow, long running. Sends data and images to earth
 - » 1 packet @ speed of light, 3 min (closest) to 22 min (farthest)
 - Work is entirely unrelated to the **InfoBus**
 - T_I : **Info thread (High priority)**
 - Moves data in and out of **InfoBus** (need synchronization)
 - **If this thread doesn't get to run periodically, system resets!**
 - T_w : **Weather gathering thread (Low priority)**
 - Puts data into **InfoBus** (need synchronization)
 - Runs infrequently

Priority Inversion (3)

Tw (Low Priority)	Ti (High Priority)	Tc (Med Priority)
<code>acquire();</code>		
<i>// Got in! Write weather data to InfoBus</i>		
	(I need to access the InfoBus, pronto!)	
	<code>acquire();</code>	
	<i>; (spin wait)</i>	
	<i>; (spin wait)</i>	<i>send cool image to earth</i>
<i>T_w starves because of priority scheduling.</i>		<i>still sending (it's big)</i>
<i>T_i also starves because of T_w's hold on spin lock.</i>	<i>; (T_i gets scheduling priority over T_w)</i>	
<i>OS detects T_i isn't running. Reboot!</i>	<i>; (more spinning)</i>	
	<i>; (It can't make progress!)</i>	<i>Still sending...</i>
	<i>; (It can't make progress!)</i>	
	<i>SYSTEM RESET</i>	


Big Problem with Spin-Locks: Priority Inversion

▶ *Priority Inversion Problem*

- High priority thread **starves** because a low priority thread has lock on item that the high priority thread needs
 - (And Low priority thread can't get scheduled to release the lock!)

▶ Solution: use a *priority inheritance* policy

- Low priority threads gets highest priority of all threads waiting on that same lock.




Goals for This Lecture...

- ▶ Basic Problem Definition
- ▶ Mechanisms to Control Access to Shared Resources
 - Low Level Mechanisms:
 - Disabling interrupts
 - Busy-Waiting (Spin) Locks
 - Self-Blocking Locks
 - High Level Mechanisms:
 - Semaphores
 - Condition Variables and Monitors

Alternative to Spinning? Blocking

- ▶ Spin locks are wasteful and susceptible to priority inversion
 - Instead of spinning (thread in *Running & Ready* state)
 - Make threads waiting on locks go to sleep (or "self-block") to put thread in *Waiting* state.



Implementing Blocking Locks in User Mode

- ▶ Need to have a queue of process/threads waiting on the lock

```
typedef struct blockinglock_t {  
    int held = 0;  
    struct queue *waitq; //need a wait-queue of thread ids  
} blockinglock_t;
```

```
void acquire(blockinglock_t *L) {  
    if (L->held) {  
        enqueue(getpid(), L->waitq); //put myself on the lock's wait queue  
        sleep(); //put myself to sleep  
    }  
    else  
        L->held = 1;  
}
```

Implementing Blocking Locks in User Mode

- ▶ Need to have a queue of process/threads waiting on the lock

```
typedef struct blockinglock_t {
    int held = 0;
    struct queue *waitq; //need a wait-queue of thread ids
} blockinglock_t;
```

```
void acquire(blockinglock_t *L) {
    if (L->held) {
        enqueue(getpid(), L->waitq); //put myself on the lock's wait queue
        sleep(); //put myself to sleep
    }
    else
        L->held = 1;
}

void release(blockinglock_t *L) {
    if (isEmpty(L->waitq))
        L->held = 0;
    else
        wakeup(dequeue(l->waitq)); //dequeue next process, ask OS to make it ready
                                    //don't set l->held to 0 in else clause
}
```

But acquire and release
have critical sections! Lock them up!

Implementing Blocking Locks in User Mode (Try 0)

- ▶ Use a **spinlock** to guard the critical sections of the blocking lock!

```
typedef struct blockinglock_t {  
    int held = 0;  
    spinlock_t *mutex; // Use a spinlock to guard acquire() and release()  
    struct queue *waitq; //need a wait-queue of thread ids  
} blockinglock_t;
```

```
void acquire(blockinglock_t *L) {  
    spinlock_acquire(L->mutex);  
    if (L->held) {  
        enqueue(getpid(), L->waitq);  
        sleep();  
    }  
    else  
        L->held = 1;  
    spinlock_release(L->mutex);  
}
```

```
void release(blockinglock_t *L) {  
    spinlock_acquire(L->mutex);  
    if (isEmpty(L->waitq))  
        L->held = 0;  
    else  
        wakeup(dequeue(L->waitq));  
    spinlock_release(L->mutex);  
}
```

The Problem with "Try 0": Deadlock

T1	T2
<code>acquire() { ...</code>	
<code> spinlock_acq()</code>	
<code> L->held = 1</code>	
<code> spinlock_rel()</code>	
<code>}</code>	
<code><<Crit. Sec.>></code>	
	<code>acquire() { ...</code>
	<code> spinlock_acq()</code>
	<code> enqueue(...)</code>
	<code> sleep();</code>
<code>release() {</code>	
	<code> spinlock_acq()</code>
<code>(spin forever)</code>	<code>(sleep forever)</code>

```

void acquire(blockinglock_t *L) {
  spinlock_acquire(L->mutex);
  if (L->held) {
    enqueue(getpid(), L->waitq);
    sleep();
  }
  else
    L->held = 1;
  spinlock_release(L->mutex);
}

void release(blockinglock_t *L) {
  spinlock_acquire(L->mutex);
  if (isEmpty(L->waitq))
    L->held = 0;
  else
    wakeup(dequeue(L->waitq));
  spinlock_release(L->mutex);
}

```

Implementing Blocking Locks in User Mode (Try 1)

- ▶ All right, then swap sleep and spinlock_release!

```
typedef struct blockinglock_t {  
    int held = 0;  
    spinlock_t *mutex; // Use a spinlock to guard acquire() and release()  
    struct queue *waitq; //need a wait-queue of thread ids  
} blockinglock_t;
```

```
void acquire(blockinglock_t *L) {  
    spinlock_acquire(L->mutex);  
    if (L->held) {  
        enqueue(getpid(), L->waitq);  
        spinlock_release(L->mutex);  
        sleep();  
    } else {  
        L->held = 1;  
        spinlock_release(L->mutex);  
    }  
}
```

```
void release(blockinglock_t *L) {  
    spinlock_acquire(L->mutex)  
    if (isEmpty(L->waitq))  
        L->held = 0;  
    else  
        wakeup(dequeue(L->waitq));  
    spinlock_release(L->mutex);  
}
```

Ah HA! Release spinlock first, *then* sleep!

The Problem with "Try 1": Starvation

T1	T2
<code>acquire() { ...</code>	
<code>spinlock_acq()</code>	
<code>L->held = 1</code>	
<code>spinlock_rel()</code>	
<code>}</code>	
<code><<Crit. Sec.>></code>	
	<code>acquire() { ...</code>
	<code>spinlock_acq()</code>
	<code>enqueue(...)</code>
	<code>spinlock_rel()</code>
<code>release() {</code>	
<code>spinlock_acq()</code>	
<code>dequeue(...)</code>	
<code>wakeup(...)</code>	<i><-- is lost! T2 not asleep yet!</i>
<code>spinlock_rel()</code>	<code>sleep();</code>
<code>}</code>	(sleep forever)

```

void acquire(blockinglock_t *L) {
    spinlock_acquire(L->mutex);
    if (L->held) {
        enqueue(getpid(), L->waitq);
        spinlock_release(L->mutex);
        sleep();
    } else {
        L->held = 1;
        spinlock_release(L->mutex);
    }
}

void release(blockinglock_t *L) {
    spinlock_acquire(L->mutex)
    if (isEmpty(L->waitq))
        L->held = 0;
    else
        wakeup(dequeue(L->waitq));
    spinlock_release(L->mutex);
}

```

Implementing Blocking Locks in User Mode (Correct)

```
typedef struct blockinglock_t {
    int held = 0;
    spinlock_t *mutex; // Use a spinlock to guard acquire() and release()
    struct queue *waitq; // Wait queue of pids
} blockinglock_t;
```

```
void acquire(blockinglock_t *L) {
    spinlock_acquire(L->mutex);
    if (L->held) {
        enqueue(getpid(), L->waitq);
        spinlock_release_and_sleep(L->mutex); // sleep() needs to release() all locks
                                                // being held by the thread atomically
    }
    else {
        L->held = 1;
        spinlock_release(L->mutex);
    }
}

void release(blockinglock_t *L) {
    spinlock_acquire(L->mutex);
    if (isEmpty(L->waitq))
        L->held = 0;
    else
        wakeup(dequeue(L->waitq));
    spinlock_release(L->mutex);
}
```

Need `sleep()` to also `release()` atomically!
(Ask CPU maker for atomic support)

Summary of Blocking Locks

▶ Pros:

- No more busy-waiting. CPU is less taxed.
- Avoids priority inversion problem! (*why?*)
- Some of the implementation is in user-mode. Transitions to kernel-mode only when `sleep()` and `wakeup()` are called.

▶ Cons:

- Need to manage a queue per lock
 - Queue management must be atomic, but it's a small critical section
 - Most OS will just use a spin lock for lock-queue management
- `sleep()` and `wakeup()` need system calls, so there's an overhead
 - *Think: When would you prefer a spin lock over a blocking lock? When would you prefer the opposite?*

No Priority Inversion?

Tw (Low Priority)	Ti (High Priority)	Tc (Med Priority)
<code>acquire();</code>		
write to InfoBus		
...		
	<code>acquire() { ...</code> <code>sleep(); // zzzzzz</code>	
	Ti is no longer in contention with Tw in the ready queue!	
		<code>send cool image to earth</code>
...		
...		
<code>release();</code>	<code>(You woke me!)</code> <code>Yay I get to run. Rover happy.</code>	

Using Blocking Locks in C

- ▶ The Blocking Lock type is **`pthread_mutex_t`**

```
#include <pthread.h>

pthread_mutex_init(pthread_mutex_t *mutex, pthread_mutex_attr *attr)

pthread_mutex_destroy(pthread_mutex_t *mutex)

pthread_mutex_lock(pthread_mutex_t *mutex)

pthread_mutex_unlock(pthread_mutex_t *mutex)
```

Usage Example

```
//shared lock
pthread_mutex_t *lock;

int main() {
    lock = (pthread_mutex_t*) malloc(sizeof(pthread_mutex_t));
    pthread_mutex_init(lock, NULL); //sets it to "unlocked" state

    // ... more code ...
    // (creating threads...)

    // ... more code ...
    // (joining threads)
    pthread_destroy(lock);
    return 0;
}

void* threadWork(void* args) {
    pthread_mutex_lock(lock); //try to acquire the lock

    /////////////////////////////////
    // << critical section >>
    ///////////////////////////////

    pthread_mutex_unlock(lock); //release
}
```

Full Circle (Remember This Race Condition?)

```
#include <pthread.h>
#include <stdio.h>

int x = 0;

void* doStuff(void*);

int main() {
    pthread_t t1, t2;
    pthread_create(&t1, NULL, doStuff, NULL);
    pthread_create(&t2, NULL, doStuff, NULL);
    pthread_join(t1, NULL);
    pthread_join(t2, NULL);
    printf("Value of x is %d\n", x);

    return 0;
}

void* doStuff(void* args) {
    int i;
    for (i = 0; i < 1000000; i++)
        x++;
    return NULL;
}
```

I want to increment x to 2,000,000 but split the work across 2 cores!

Output:

```
$ gcc add2m.c -lpthread -o add2m
$ ./add2m
Value of x is 1029184

$ ./add2m
Value of x is 1224682

$ ./add2m
Value of x is 1004207

$ ./add2m
Value of x is 1003645
```

Full Circle - Using Self-Blocking Locks

```

int x = 0;
pthread_mutex_t *lock;

void* doStuff(void*);

int main() {
    //initialize the lock
    lock = (pthread_mutex_t*) malloc(sizeof(pthread_mutex_t));
    pthread_mutex_init(lock, NULL);

    pthread_t t1, t2;
    pthread_create(&t1, NULL, doStuff, NULL);
    pthread_create(&t2, NULL, doStuff, NULL);
    pthread_join(t1, NULL);
    pthread_join(t2, NULL);
    printf("Value of x is %d\n", x);
    pthread_mutex_destroy(lock);
    return 0;
}

void* doStuff(void* args) {
    int i;
    for (i = 0; i < 1000000; i++) {
        pthread_mutex_lock(lock);
        x++;
        pthread_mutex_unlock(lock);
    }
    return NULL;
}
  
```

*Fixes race condition, but now much slower than sequential version!
 (Ugh... Amdahl's Law)*

Full Circle - Using Spin Locks

```

int x = 0;
pthread_spinlock_t *lock;

void *doStuff(void *);

int main() {
    // initialize the lock
    lock = (pthread_spinlock_t *)malloc(sizeof(pthread_spinlock_t));
    pthread_spin_init(lock, PTHREAD_PROCESS_PRIVATE);

    pthread_t t1, t2;
    pthread_create(&t1, NULL, doStuff, NULL);
    pthread_create(&t2, NULL, doStuff, NULL);
    pthread_join(t1, NULL);
    pthread_join(t2, NULL);
    printf("Value of x is %d\n", x);
    pthread_spin_destroy(lock);

    return 0;
}

void *doStuff(void *args) {
    int i;
    for (i = 0; i < 1000000; i++)
    {
        pthread_spin_lock(lock);
        x++;
        pthread_spin_unlock(lock);
    }
    return NULL;
}

```

This version will run faster than blocking locks.... (why?)

Administrivia 3/14

- ▶ Today: Start Chapter 6
 - Race conditions
 - Concurrency leads to nondeterministic outcomes
 - Interleaving schedules (i.e., context switching is the culprit)
 - But context switching is a good thing
 - We can't just go back to batch processing!
 - Critical section problem

Administrivia 3/24

- ▶ Announcements:
 - Hwk 6 posted! Due Wednesday, 4/9
- ▶ Last time...
 - Enforcing atomicity
 - The critical section problem
- ▶ Today:
 - Review the "Too Much Milk" Problem
 - Building locks: Spin locks (busy waiting)
 - Building self-blocking locks (not covered in book)

Administrivia 3/26

- ▶ Announcements:
 - Hwk 6 posted! Due Wednesday, 4/9
- ▶ Last time...
 - Peterson's Solution to the TMM problem (only 2 threads)
 - Building user-mode locks: Spin locks (busy waiting)
 - Importance of test-and-set() as an unprivileged atomic instruction
- ▶ Today
 - Spinlocks continued: Priority inversion problem
 - In C: pthread_spinlock_t
 - Building self-blocking locks (not covered in book)
 - In C, pthread_mutex_t