

CS 475

Operating Systems



Department of Mathematics
and Computer Science

Lecture 6
Synchronization (Part II)


Review of Locks

- ▶ Locks are the simplest synchronization primitives, allowing 1 thread in a critical section at a time.
 - (The view is that locks are a "barrier" for threads to get past)
- ▶ Types
 - Spinning (busy waiting -- requires TaS(), doesn't make syscalls)
 - Blocking (put self to sleep -- requires TaS(), and making syscalls)
- ▶ *What if we needed more complicated coordination of threads?*
 - *What if I want one thread to tell another thread to go?*

Motivation

► Application Example: Thread Pools

- To save overhead of creating and re-creating threads, some applications create a "pool" of threads at startup.
- Threads are initially idle and waiting for work
 - How do you get them all to **wait**?
- If a task becomes available, dispatch a thread to handle it.
 - How do you **signal** a waiting thread to wake up?



Motivation (2)

▶ Example: Sleeping Barber Problem

- 1 barber thread
- N chairs for customer threads

▶ If there are no customers, the barber:

- Goes to sleep.

▶ If a customer enters and:

- The barber is asleep, the customer wakes them up.
- The barber is busy, and there's a free chair, the customer waits.
- The waiting room is full, the customer leaves.

The Sleeping Barber Problem




Goals for This Lecture...

- ▶ Motivation for Higher Level Synchronization
- ▶ **Semaphores**
 - Definition
 - Implementation of Blocking Semaphores
- ▶ Classical Synchronization Problems (Chap 7)
- ▶ Semaphore support in C

Semaphores

- ▶ We want a more general synchronization primitive for threads to coordinate and signal to each other when to go, when to stop.
- ▶ Semaphores can provide robust *coordination* between threads
 - Method of visual signaling, usually by means of flags or lights.
 - Before the invention of the telegraph, semaphore signaling from high towers was used to transmit messages between distant points.



Fun fact:
Semaphores were invented by Edsger Dijkstra for the "T.H.E. Operating System"

Semaphores (Struct)

Semaphore Structure

```
typedef struct sem_t {  
    int val;  
} sem_t;
```

Semaphores (Creation/Initialization)

Semaphore Structure

```
typedef struct sem_t {  
    int val;  
} sem_t;
```

```
/** Create */  
sem_t* sem_open(int init_val) {  
    sem_t *S = (sem_t*) malloc(sizeof(sem_t));  
    S->val = init_val;  
    return S;  
}
```

Semaphores (Wait/Decrement)

Semaphore Structure

```
typedef struct sem_t {  
    int val;  
} sem_t;
```

```
/** Create */  
sem_t* sem_open(int init_val) {  
    sem_t *S = (sem_t*) malloc(sizeof(sem_t));  
    S->val = init_val;  
    return S;  
}  
  
/** Wait (spinning semaphore code below is assumed atomic) */  
void wait(sem_t* S) {  
    while (S->val == 0)  
        ;  
    S->val--;  
}
```

Semaphores (Signal/Increment)


Semaphore Structure

```
typedef struct sem_t {  
    int val;  
} sem_t;
```

```
/** Create */  
sem_t* sem_open(int init_val) {  
    sem_t *S = (sem_t*) malloc(sizeof(sem_t));  
    S->val = init_val;  
    return S;  
}  
  
/** Wait (spinning semaphore code below is assumed atomic) */  
void wait(sem_t* S) {  
    while (S->val == 0)  
        ;  
    S->val--;  
}  
  
/** Signal (code below is assumed atomic) */  
void signal(sem_t* S) {  
    S->val++;  
}
```

A Coordination Mechanism?

```
// suppose we want to let two trains in critical section
sem_t *sem = sem_open(2);
```



Binary Semaphores (When initialized to 1...)

- Recall the semaphore operations (assume they're atomic):

```
/** Wait */
void wait(sem_t* S) {
    while (S->val == 0)
        ;
    S->val--;
}
```


```
/** Signal */
void signal(sem_t* S) {
    S->val++;
}
```

- A *binary semaphore* oscillates between 0 and 1
 - How does a semaphore that's been initialized to 1 behave?
 - How many threads can get into the critical section at once..?

```
sem_t *S = sem_open(1);

wait(S);
<< critical section >>

signal(S);
```



```
lock_t *lock = newLock();

acquire(lock);
<< critical section >>

release(lock);
```

Too Much Milk Problem Revisited

- ▶ Semaphore-based solution
 - (Semaphore as a lock -- initialized to 1)

```
sem_t *ok_to_buy_milk = sem_open(1);

void roommate() {
    while (1) {
        wait(ok_to_buy_milk); // Try to buy milk
        if (noMilk) {
            buyMilk();
        }
        signal(ok_to_buy_milk); // Let someone else check fridge
    }
}
```



Binary Semaphores (When Initialized to 0)

- Recall the semaphore operations (assume they're atomic):

```
/** Wait */
void wait(sem_t* S) {
    while (S->val <= 0)
        ;
    S->val--;
}
```

```
/** Signal */
void signal(sem_t* S) {
    S->val++;
}
```

- We know: Semaphores initialized to 1 is equivalent to a lock.
 - What if we initialize a semaphore to 0?
 - What would that allow us to do?

Binary Semaphores (Initialized to 0)

- ▶ Binary semaphores initialized to *0*?
 - Think of it as a *Gate Keeper*
 - Everyone waits here. I tell you when to pass.
- ▶ For instance,
 - In `pthread_join()`, the calling thread just wait until another thread exits.



```
sem_t* sem = sem_open(0);

void pthread_join() {
    wait(sem);
    reapThread();
}

void pthread_exit() {
    signal(sem);
}
```

We've Relied on Semaphores Before!

- ▶ Anytime we joined threads, the main (calling) thread had to wait until all threads exited!

```
// spawn threads
for (int i = 0; i < NUM_THREADS; i++) {
    pthread_create(&threads[i], NULL, parallelFunc, NULL);
}

// wait for threads to finish
for (int i = 0; i < NUM_THREADS; i++) {
    pthread_join(threads[i], NULL);
}
```


Goals for This Lecture...

- ▶ Motivation for Higher Level Synchronization
- ▶ Semaphores
 - Definition
 - Implementation of Blocking Semaphores
- ▶ Classical Synchronization Problems (Chap 7)
- ▶ Semaphore support in C

Motivation (2)

▶ Example: Sleeping Barber Problem

- 1 barber thread
 - N chairs for customer threads
- ▶ If there are no customers, the barber:
- Goes to sleep.
- ▶ If a customer enters and:
- The barber is asleep, the customer wakes them up.
 - The barber is busy, and there's a free chair, the customer waits.
 - The waiting room is full, the customer leaves.



Synchronize This!

- ▶ Hint: You'll need 3 semaphores
 - One for coordinating the barber's action
 - One for coordinating the customer's action
 - One for locking up the shared variable: waitingCustomers

```
const int CHAIRS = 5;
int waitingCustomers = 0;
```

```
void barber() {
    while (1) {
        sleep() when no customers
        waitingCustomers -= 1;
        cutHair();
    }
}
```

```
void customer() {
    if (waitingCustomers < CHAIRS) {
        waitingCustomers++;
        wakeup(barber);
        getHaircut();
    }
    leave();
}
```


Helpful Professor Problem

► Correctness Properties: When do you wait?

- Professor (One thread)

- Waits for students to arrive
- When there's a student waiting at the door, prof panics
 - Lets one student in at a time
- Waits for student to ask question
- After question is asked, prof yells at the student
- Waits for student to leave
- After student leaves, prof waves bye.

```
void prof() {
    while (1) {
        panic();
        yellAtStudent();
        waveGoodbye();
    }
}
```



► Full working code here:


- <https://github.com/davidtchiu/cs475-lec-helpfulprof>

Helpful Professor Problem

► Correctness Properties: When do you wait?

- **Students (N threads)**

- If Professor isn't ready to take another student in his office, they must **wait**
- After student enters office, they ask questions
- **Waits** for Professor's answer
- Leaves Professor's office



What Are All the Problems?

- ▶ What could go wrong without synchronization? (Hint: Everything)

```
void prof() {  
    while (1) {  
        panic();  
        yellAtStudent();  
        waveGoodbye();  
    }  
}
```

```
void student() {  
    askQuestion();  
    leave();  
}
```

First Try (prof's Play: Acts I and II)

- ▶ Use a semaphore to coordinate 1-student access to my office:

```
sem_t* prof_available = sem_open(0); // track whether prof is ready
```

```
void prof() {
    while (1) {
        panic();
        signal(prof_available);
        yellAtStudent();
        waveGoodbye();
    }
}
```

```
void student() {
    wait(prof_available);
    askQuestion();
    leave();
}
```

Problems? prof shouldn't continuously signal that they're available! (That lets more than one student in!)

--- prof needs to wait if a student is not outside the office!

Second Attempt

- ▶ Suppose I define the following semaphores:

```
sem_t* student_outside = sem_open(0); // track if student is waiting
sem_t* prof_available = sem_open(0); // track whether prof is ready
```

```
void prof() {
    while (1) {
        wait(student_outside);
        panic();
        signal(prof_available);
        yellAtStudent();
        waveGoodbye();
    }
}
```

```
void student() {
    signal(student_outside);
    wait(prof_available);
    askQuestion();
    leave();
}
```

Problems? Sometimes prof yells at a student before they ask the question!

Sometimes prof waves goodbye before student leaves (or before they ask the question)

Sometimes student leaves before getting yelled at!

Helpful Professor (Correct)

```
sem_t* student_outside = sem_open(0);
sem_t* prof_available = sem_open(0);
sem_t* question_asked = sem_open(0); //wait for question
sem_t* question_answered = sem_open(0); //wait for prof's answer
sem_t* student_leaving = sem_open(0); //wait for student to leave
```

```
void prof() {
    while (1) {
        wait(student_outside);
        panic();
        signal(prof_available);
        wait(question_asked);
        yellAtStudent();
        signal(question_answered);
        wait(student_leaving);
        waveGoodbye();
    }
}
```

```
void student() {
    signal(student_outside);
    wait(prof_available);
    askQuestion();
    signal(question_asked);
    wait(question_answered);
    leave();
    signal(student_leaving);
}
```

Evaluation of Semaphores

▶ Pros:

- More general and robust than locks
- Can be used to control event handling



▶ Cons:

- Asymmetric and extremely error prone
 - Flipping the order of `wait()` and `signal()` in the bounded-buffer problem leads to a deadlock
(Was it obvious at the time?)
- Would be nice if we had even higher-level language support to deal with synchronization.

Administrivia 3/28

- ▶ Reminders
 - Hwk 6 due 4/9! (No semaphores needed, but they can be used)
- ▶ Last time...
 - Blocking locks (vs. Spinlocks)
 - Implementation of locks
- ▶ Today:
 - Semaphores
 - Helpful Professor Problem