

CS 475

Operating Systems



Department of Mathematics
and Computer Science

Lecture 7
Deadlocks

Review of Semaphores (Last year's exam)

- ▶ This function pops an item off each stack and pushes the items onto the opposite stack.

- If either stack was empty, both stacks should be left unchanged.

- ▶ Each stack object associates with a semaphore initialized to 1.

- ▶ Thread-Safe Version

- There are 3 bugs in the code.
Name them.

```
// 2+ threads run this function
void stackSwap(*stack1, *stack2) {

    wait(stack1->sem);
    thing1 = pop(stack1);
    if (thing1 != NULL) {

        wait(stack2->sem);
        thing2 = pop(stack2);
        if (thing2 != NULL) {
            push(stack2, thing1);
            push(stack1, thing2);
            signal(stack2->sem);
            signal(stack1->sem);
        }
    }
}
```

One Big Bug ... Mutual Starvation (Deadlock)

- ▶ What *could* happen if ... we have two stacks: **S_a** and **S_b**.
 - Thread 1 calls **stackSwap(S_a,S_b)** and Thread 2 calls **stackSwap(S_b,S_a)** concurrently?

```
void stackSwap(*stack1, *stack2) {
    wait(stack1->sem);
    thing1 = pop(stack1);
    if (thing1 != NULL) {
        wait(stack2->sem);
        thing2 = pop(stack2);
        if (thing2 != NULL) {
            push(stack2, thing1);
            push(stack1, thing2);
            signal(stack2->sem);
            signal(stack1->sem);
        }
    }
}
```

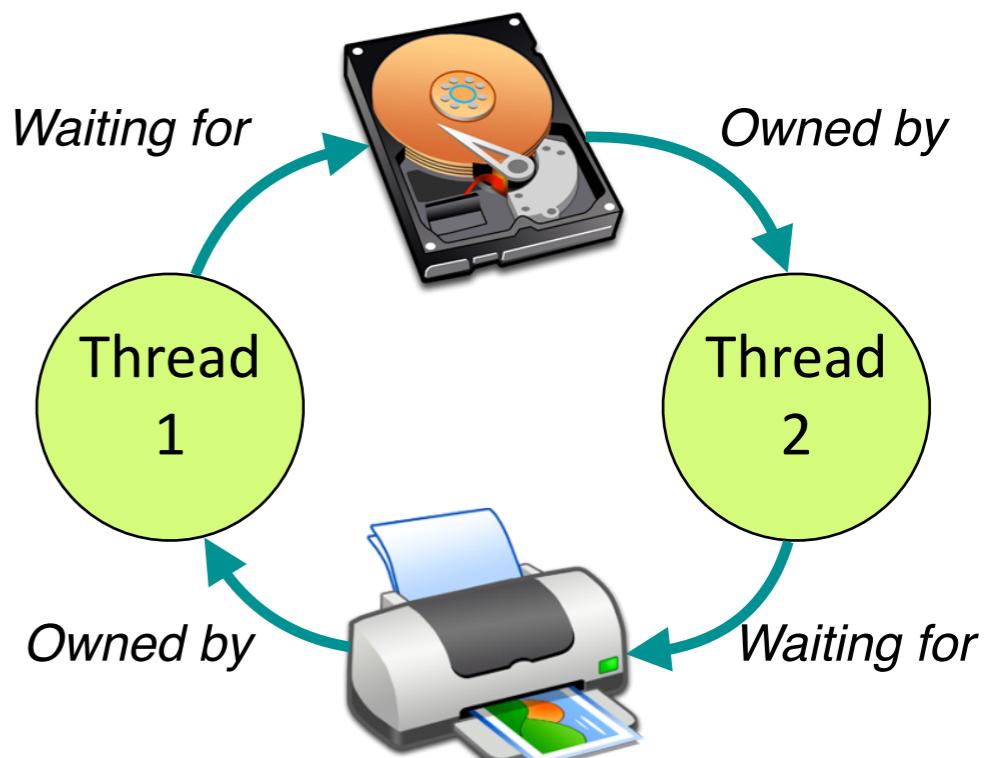
T1	T2
stackSwap(S _a ,S _b)	
	wait(S _a ->sem);
	thing1 = pop(S _a);
	stackSwap(S _b ,S _a)
	wait(S _b ->sem);
	thing1 = pop(S _b)
	if (thing1 != ...)
	if (thing1 != ...)
	wait(S _b ->sem);
	wait(S _a ->sem);
(wait forever)	(wait forever)

Deadlock and Starvation

- ▶ When synchronizing access across multiple threads...
 - **Starvation:** When a thread waits indefinitely for CPU execution.
 - **Deadlock:** When a set of threads are stuck waiting on an event that can only be caused by the another thread in the same set.

deadlock \implies *starvation* (why?)

starvation $\not\implies$ *deadlock* (why?)



Another Example Deadlock

- ▶ Deadlocks are incredibly easy to produce.
 - Here, **lock1** and **lock2** are acquired in the opposite ordering.
 - [David's Code Example: deadlock.c]

```
void* doStuff(void* args) {  
    pthread_spin_lock(lock1);  
    sleep(1);  
    pthread_spin_lock(lock2);  
    printf("Hi!\n");  
    pthread_spin_unlock(lock2);  
    pthread_spin_unlock(lock1);  
}
```

Thread A

```
void* doStuff2(void* args) {  
    pthread_spin_lock(lock2);  
    sleep(1);  
    pthread_spin_lock(lock1);  
    printf("Hi 2!\n");  
    pthread_spin_unlock(lock1);  
    pthread_spin_unlock(lock2);  
}
```

Thread B

Topics for Today...

- ▶ **Deadlock Conditions**
- ▶ Ways to deal with deadlocks:
 - Deadlock Detection
 - Single Resource Allocation Graph
 - Multi-Resource Allocation Graph
 - Deadlock Avoidance
 - Banker's Algorithm

Four Deadlock Conditions

- ▶ **Necessary and Sufficient** conditions for deadlocks to occur:
 - "Necessary:" All 4 must exist for there to be a deadlock
 - If we take any of these conditions away, the system will be deadlock-free
 - "Sufficient:" If the 4 are shown to exist currently, then there *must* be a deadlock
- ▶ The four deadlock conditions include:
 - Mutual Exclusion
 - No Preemption
 - Hold and Wait
 - Circular Wait

Deadlock Conditions 1 & 2

▶ Condition 1: Mutual Exclusion

- Only one thread can access the resource at any time

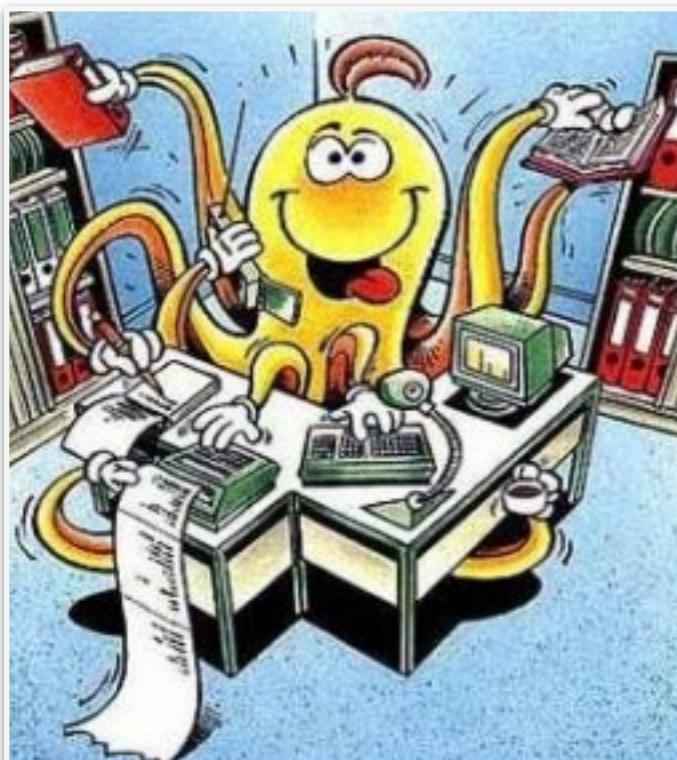
▶ Condition 2: No Preemption

- A thread can't be coerced to release its hold on a resource.
- Why would preemption be considered unsafe?
 - *e.g., try forcibly taking a lock away from a thread*

Deadlock Condition 3: Hold and Wait

▶ Condition 3: Hold and Wait

- When threads are allowed to:
 - *Hold* on to one or more resources, and be...
 - *Waiting* to acquire additional resources



To do my taxes, I need to acquire the following resources.

- A calculator
- Tax form 1040EZ
- My W2
- My spouse's W2
- My 1098

For each one I obtain, I don't put it back down for someone else to take.

Deadlock Condition 4: Circular Wait

▶ Condition 4: Circular Wait

T_1 is waiting for a resource that is held by T_2

T_2 is waiting for a resource that is held by T_3

...

T_{n-1} is waiting for a resource that is held by T_n

T_n is waiting for a resource that is held by T_1

My spouse needs calculator and her own W2 for something else.

I obtained the calculator.

She has her W2, which I need.

We're waiting for each other to give one up.

Example: Existence of 4 DL Conditions

- ▶ Mutual Exclusion
- ▶ No-Preemption
- ▶ Hold-and-Wait
- ▶ Circular Wait

Thread A

```
void* doStuff(void* args) {  
    pthread_spin_lock(lock1);  
    sleep(1);  
    pthread_spin_lock(lock2);  
    printf("Hi!\n");  
    pthread_spin_unlock(lock2);  
    pthread_spin_unlock(lock1);  
}
```

Thread B

```
void* doStuff2(void* args) {  
    pthread_spin_lock(lock2);  
    sleep(1);  
    pthread_spin_lock(lock1);  
    printf("Hi 2!\n");  
    pthread_spin_unlock(lock1);  
    pthread_spin_unlock(lock2);  
}
```

Another Example: Dining Philosophers

- ▶ 5 forks, 5 philosophers. Need both *left* and *right* forks to eat

- ▶ Each philosopher has only two states:

- **Eating**

- Need to acquire both forks when eating

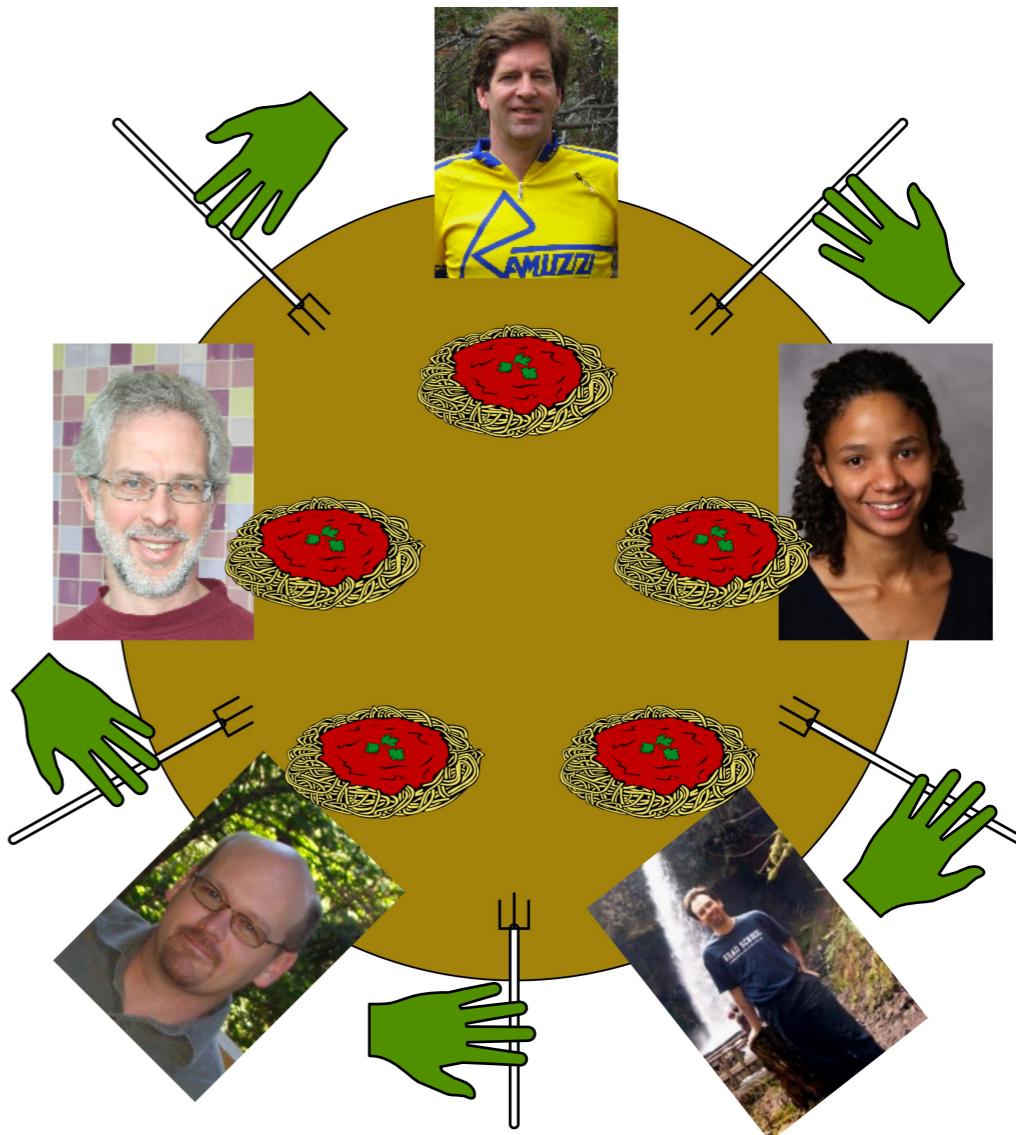
- **Thinking**

- Don't need forks to think
 - Put them both down

Deadlock:

*Everyone requests **right fork**, then...*

*... tries to grab their **left fork** (deadlock)*



What Can the OS Do?

- ▶ The OS should intervene when there's a deadlock.
- ▶ Several approaches toward handling deadlocks
 - **Prevent** deadlocks from ever happening by removing DL condition(s)
 - OS only allocates resources in such a way that it avoids unsafe situations (**Avoidance**)
 - **Detection:** Allow the OS to enter a deadlocked state(!)
 - Detect it
 - Recover from it

Topics for Today...

- ▶ Deadlock Conditions
- ▶ Ways to deal with deadlocks:
 - Deadlock Detection
 - Single Resource Allocation Graph
 - Multi-Resource Allocation Graph
 - Deadlock Avoidance
 - Banker's Algorithm

Modeling Resource Allocation (Graphs)

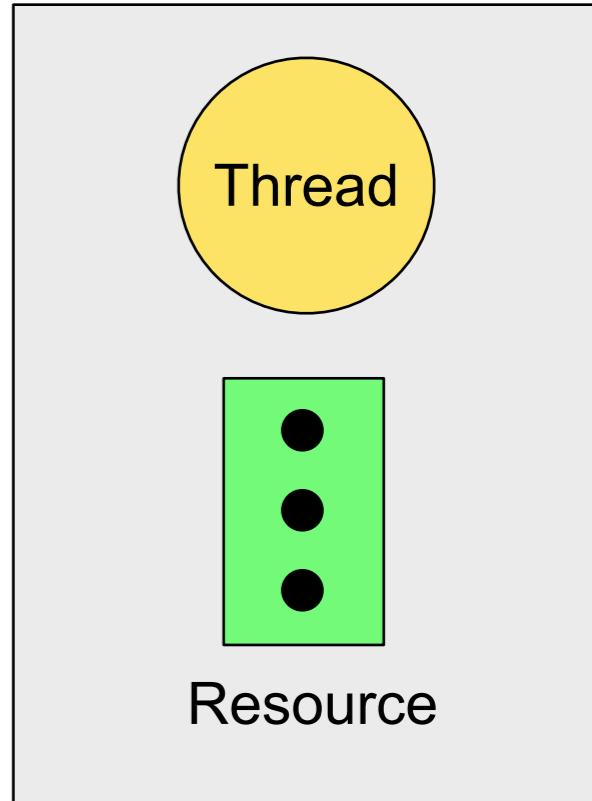
- ▶ Modeling deadlocks as graphs. Nodes can be:

- Threads, or
- Resource types
 - CPU, memory, disk, printer, semaphores, ...

- ▶ Each resource type may have multiple instances

- (Each instance is represented by ●)
- There's one instance of a lock
- A quad-core processor means we have 4 CPU instances
- We may have 2 or more printers

Types of Nodes/Vertices



Resource Allocation Graph (RAG)

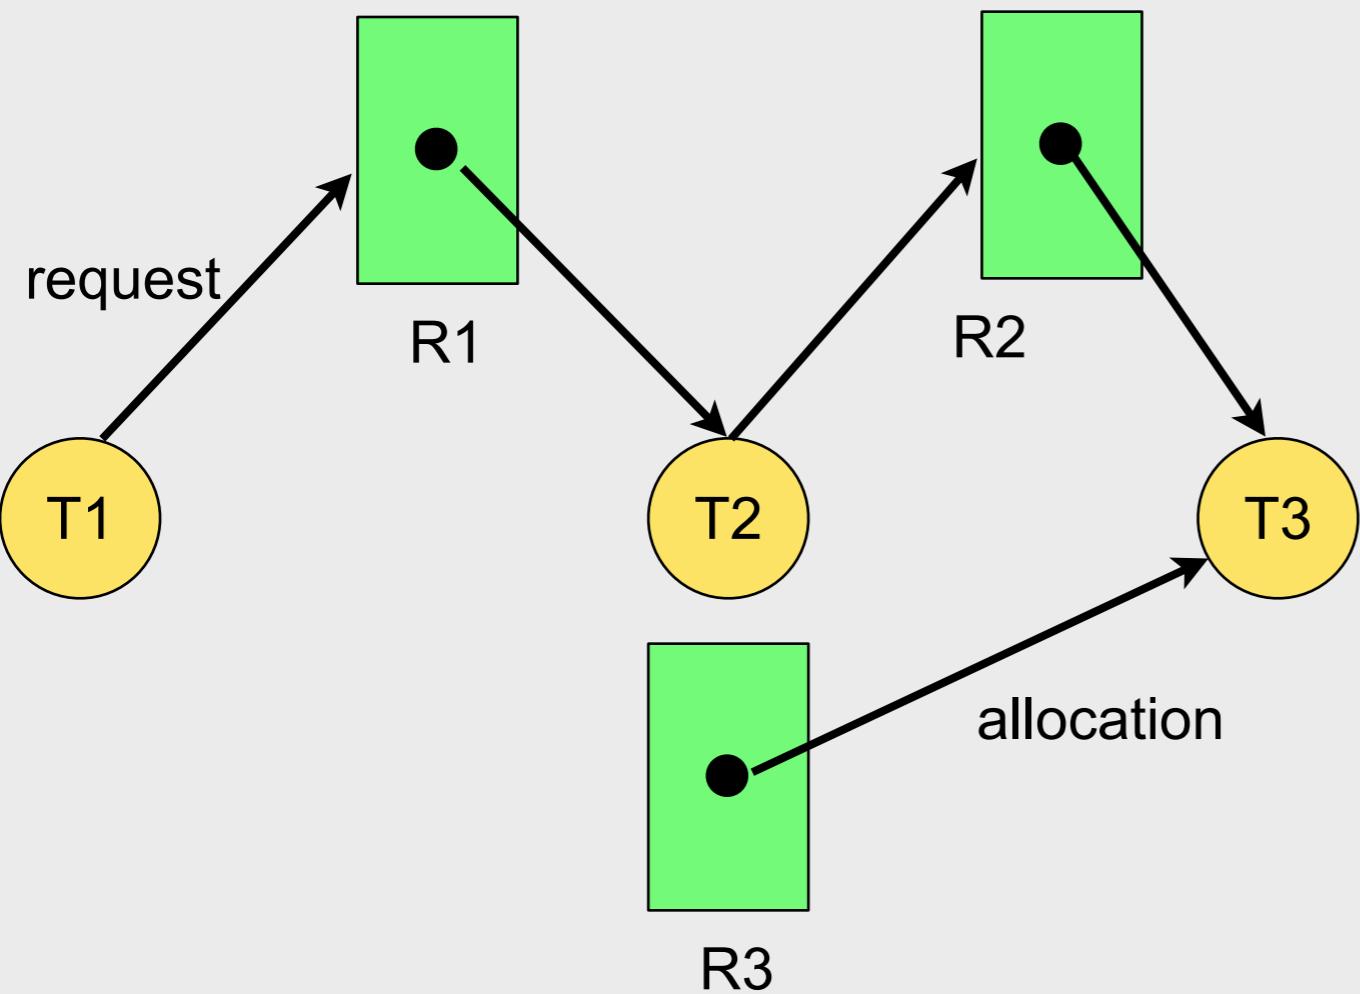
- ▶ We can model resources' state using a *Directed Graph* $G = (V, E)$
- ▶ $V = T \cup R$ is partitioned into two types of vertices (nodes):
 - $T = \{t_1, \dots, t_n\}$ is the current set of threads running in the system
 - $R = \{r_1, \dots, r_m\}$ is the set of resource types in the system
- ▶ E (the set of edges) is partitioned into two types of directed edges:
 - Request edge: (t_i, r_j) exists when t_i requests an instance of resource r_j
 - Allocation edge: (r_j, t_i) exists when r_j is assigned to t_i

Resource Allocation Graph (Cont.)

- ▶ How to build a RAG?
- ▶ Event-driven
 - When a request for a resource is made, **add** request edge
 - When a request is fulfilled, **change** request edge to allocation edge
 - When a thread releases a resource, **remove** allocation edge
- ▶ Therefore, the OS needs to update allocation and deallocation requests
 - **acquire()** and **release()** must be system calls

Example with Single-Instance Resources

- For simplicity, suppose our system does not have multiple resource instances: One CPU, One Printer,



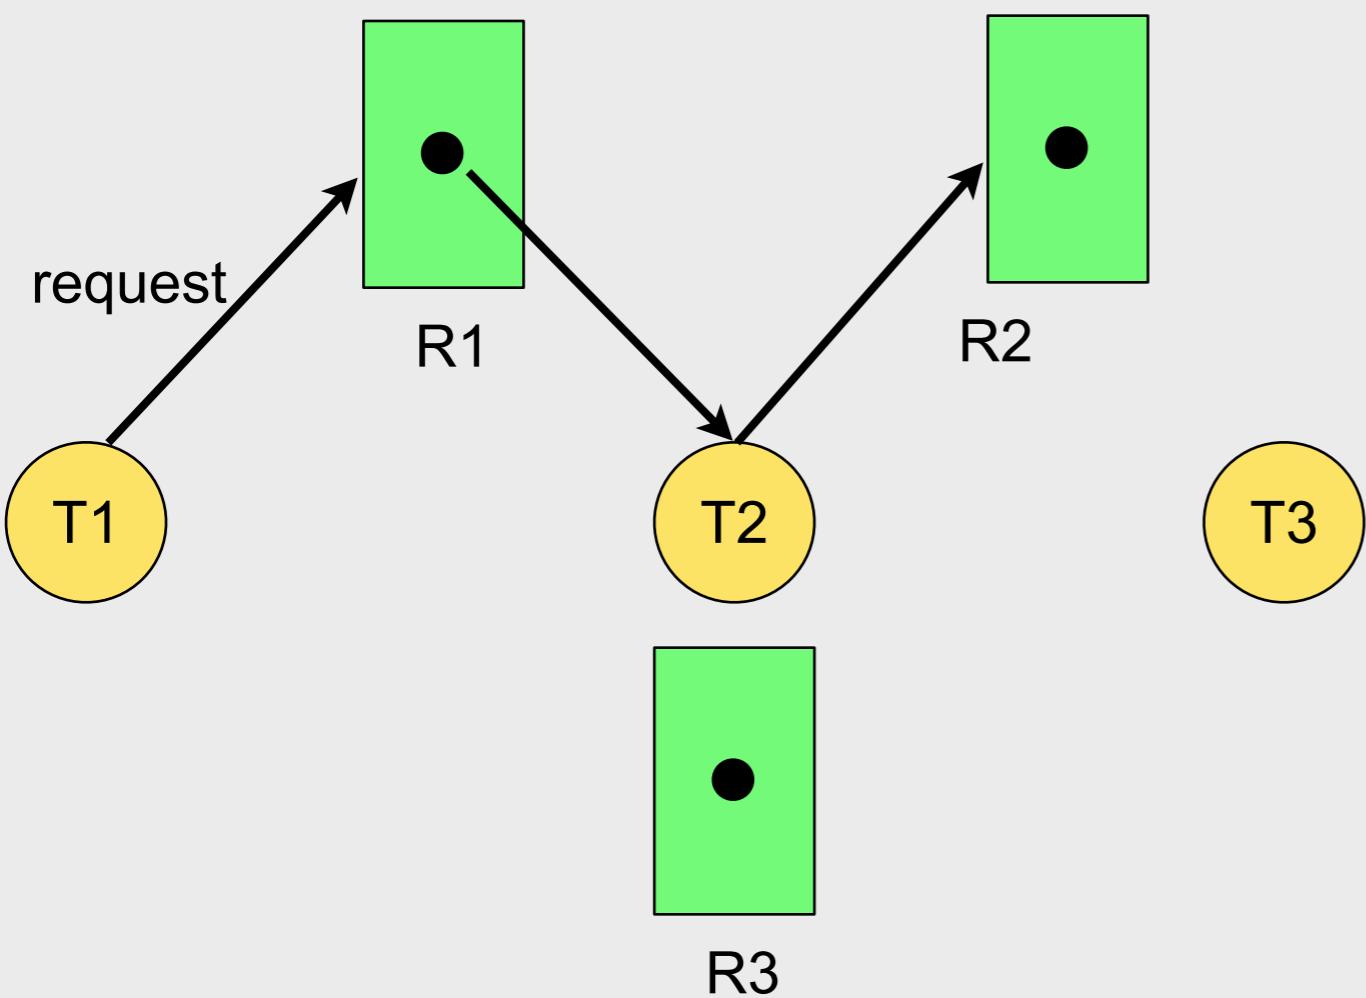
Is there a Deadlock in this RAG?

How to check?

- Can all threads finish execution?
- Simulate execution and try to "reduce the edges"

Example with Single-Instance Resources

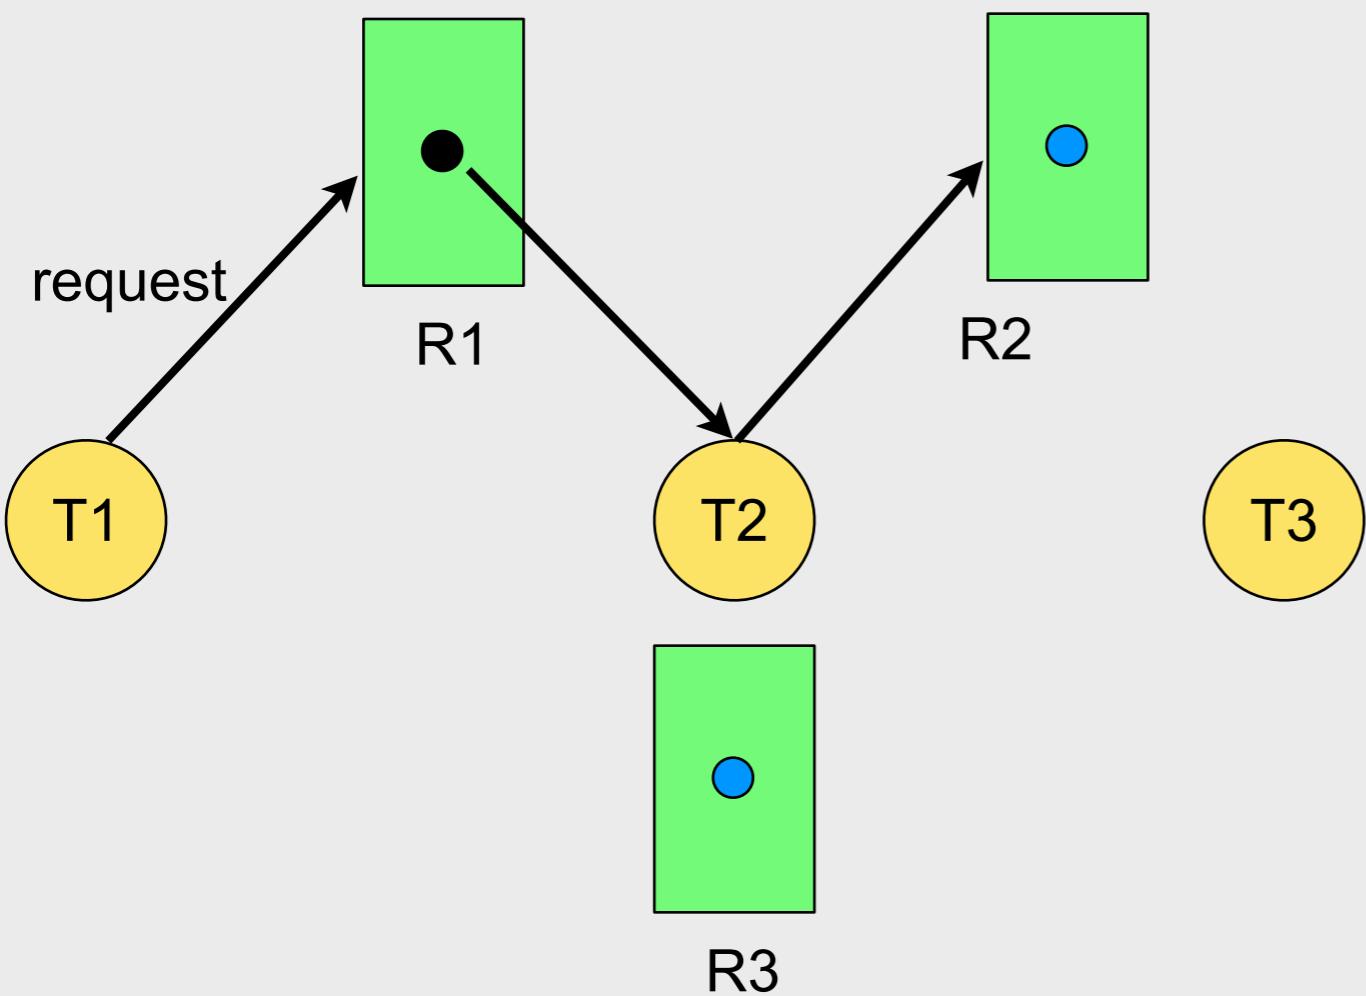
- ▶ For simplicity, suppose our system does not have multiple resource instances: One CPU, One Printer,



T3 has all the resources it needs and can finish execution!

Example with Single-Instance Resources

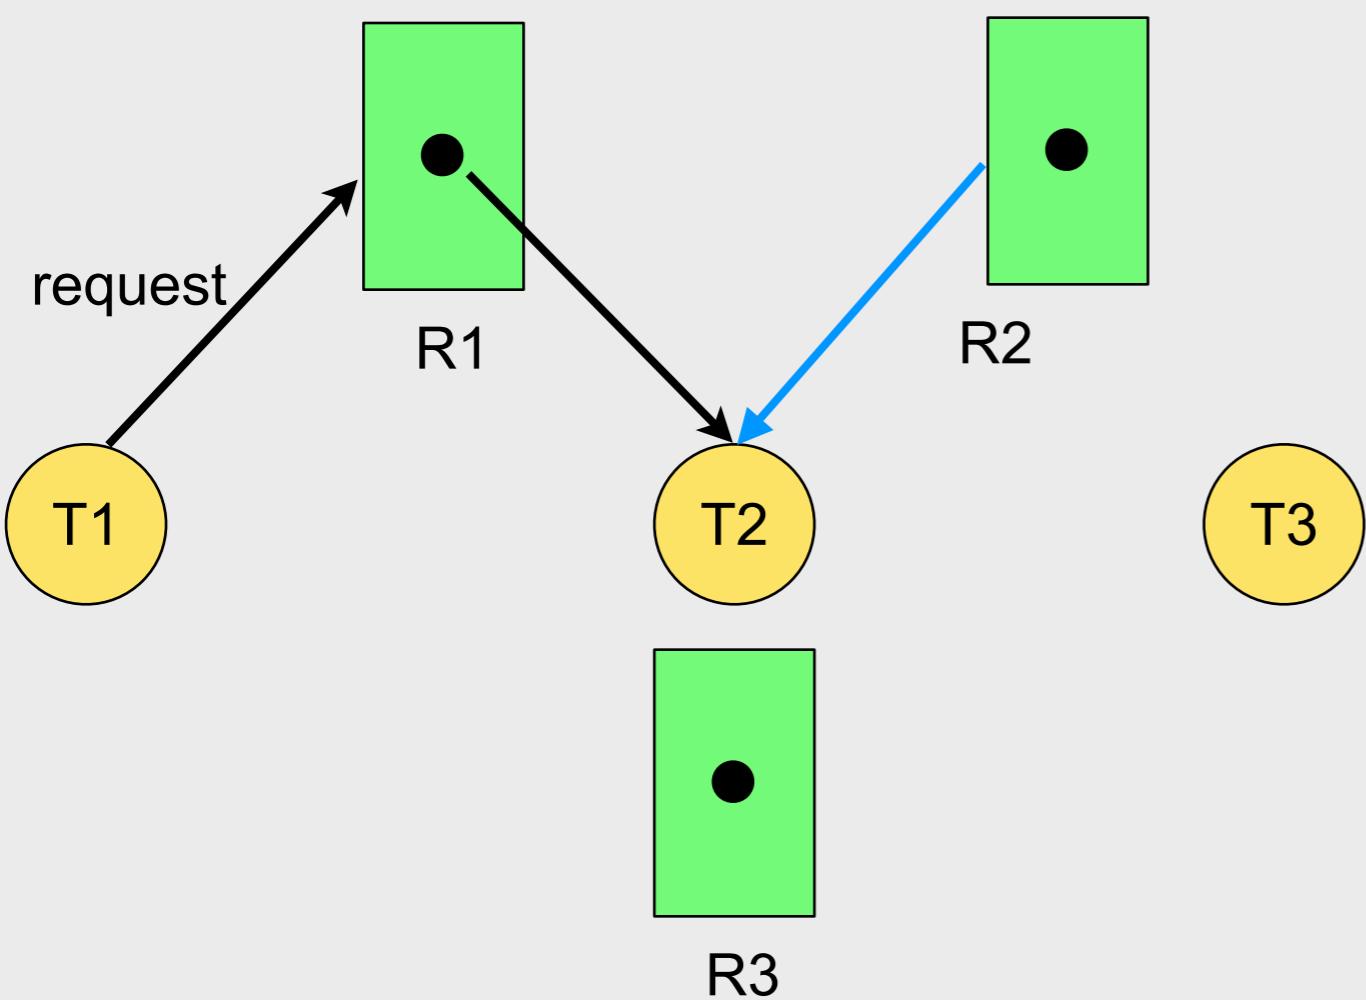
- ▶ For simplicity, suppose our system does not have multiple resource instances: One CPU, One Printer,



An instance of R2 and R3 have been freed!

Example with Single-Instance Resources

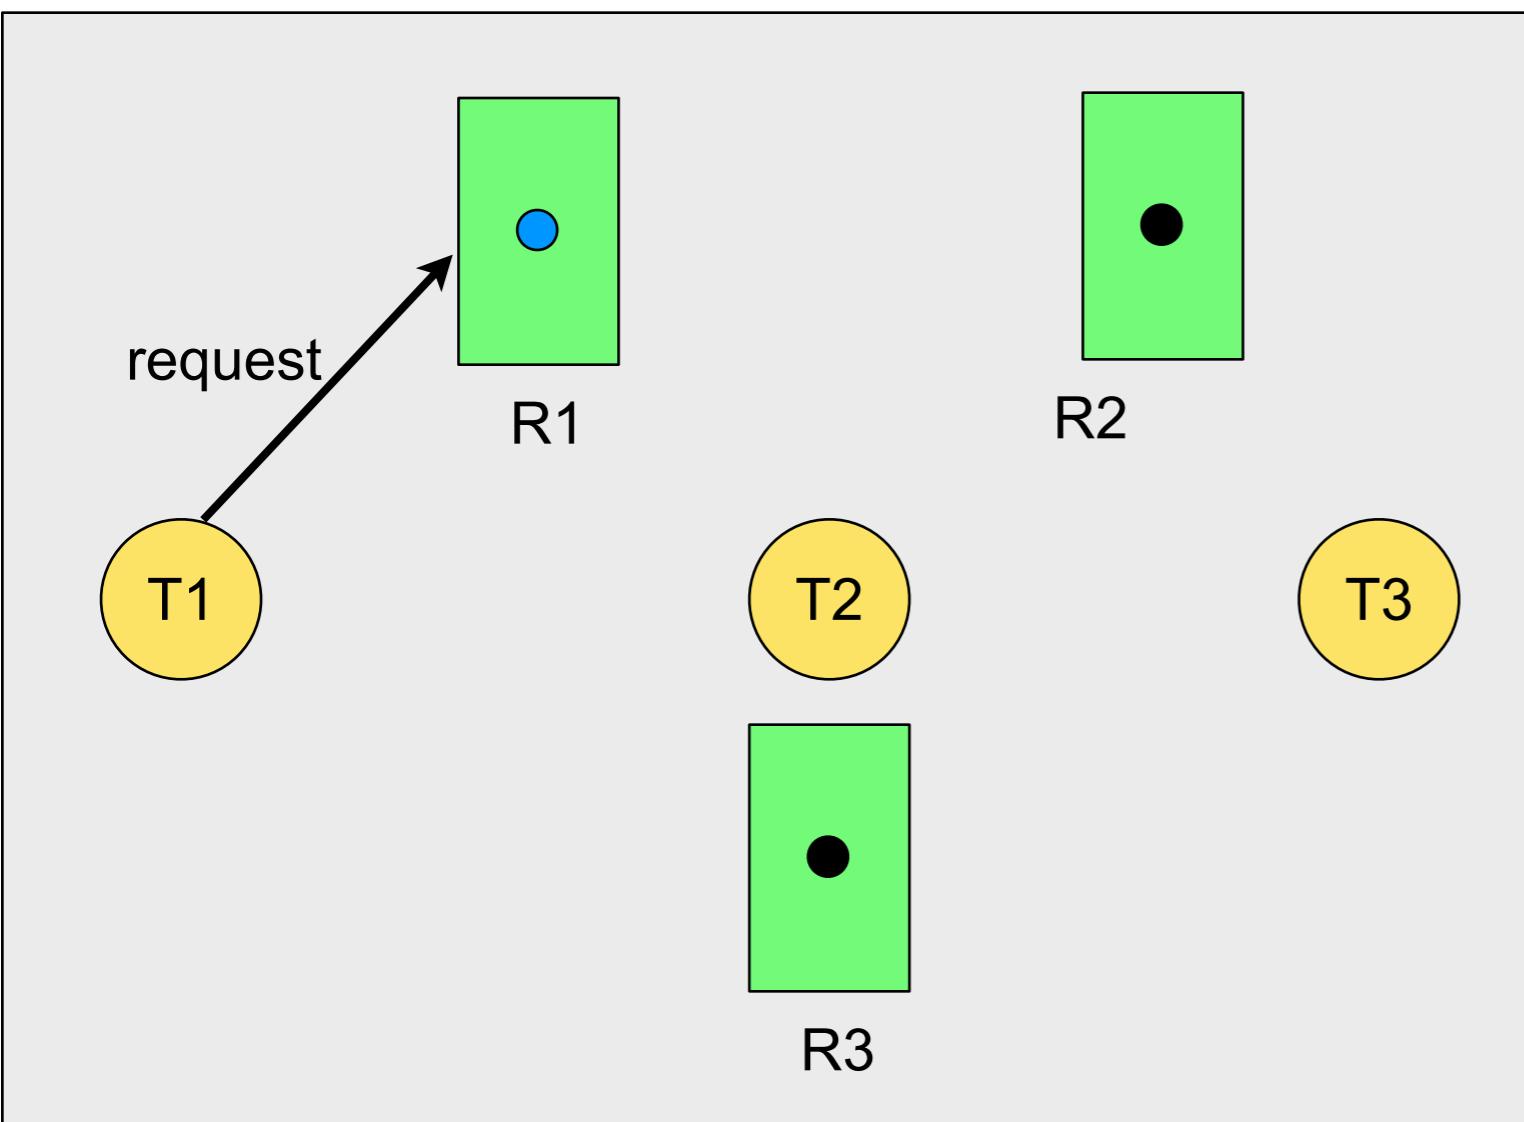
- For simplicity, suppose our system does not have multiple resource instances: One CPU, One Printer,



An instance of R2 can now be assigned to T2.

Example with Single-Instance Resources

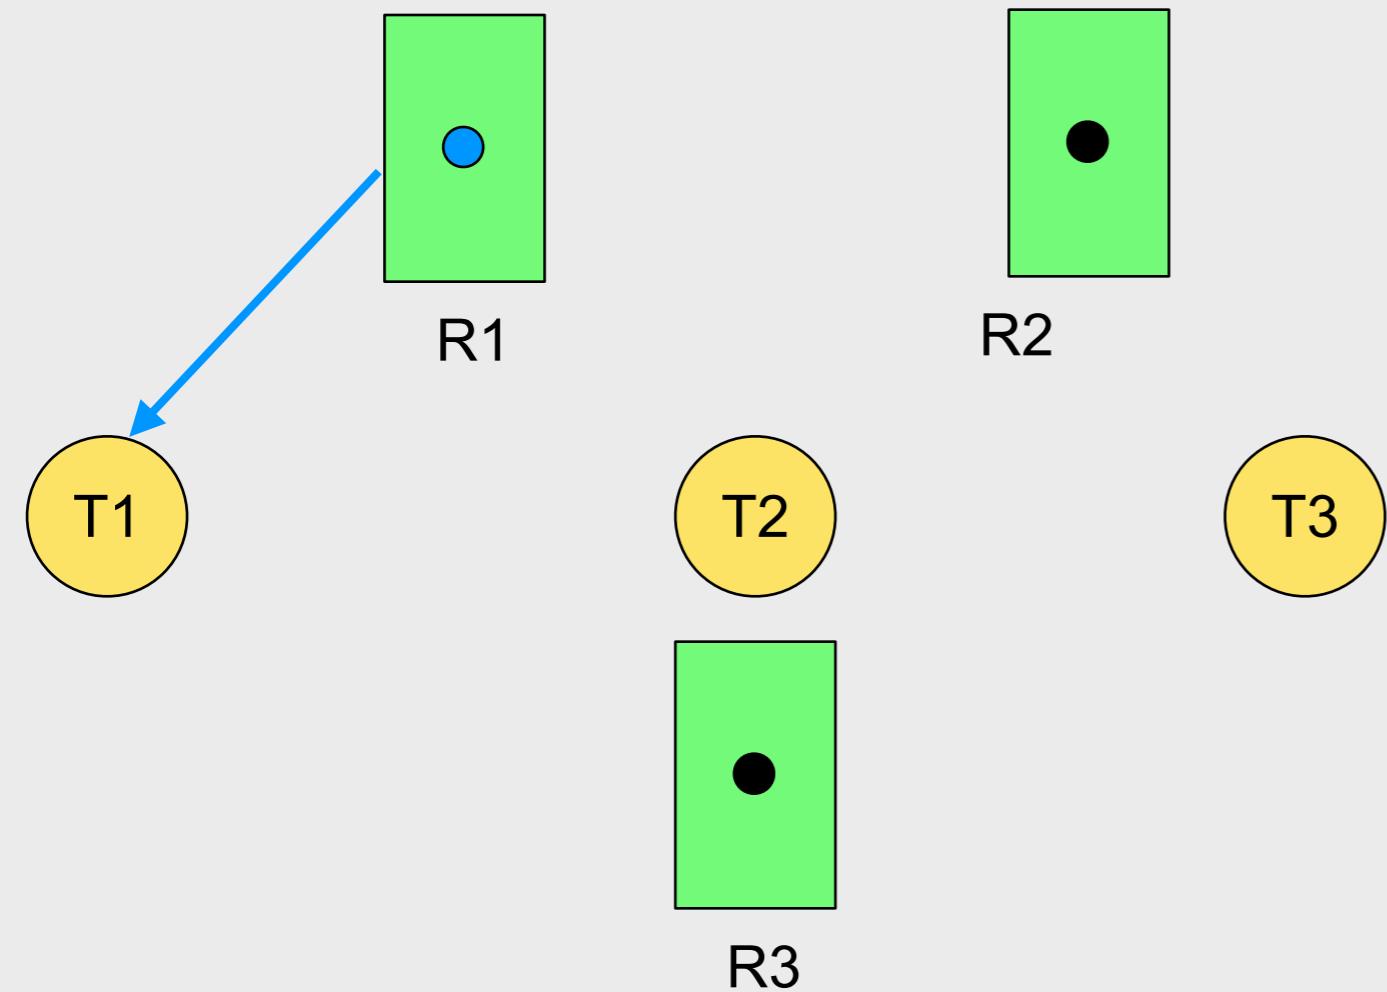
- ▶ For simplicity, suppose our system does not have multiple resource instances: One CPU, One Printer,



T2 has all the resources it needs and can finish execution!

Example with Single-Instance Resources

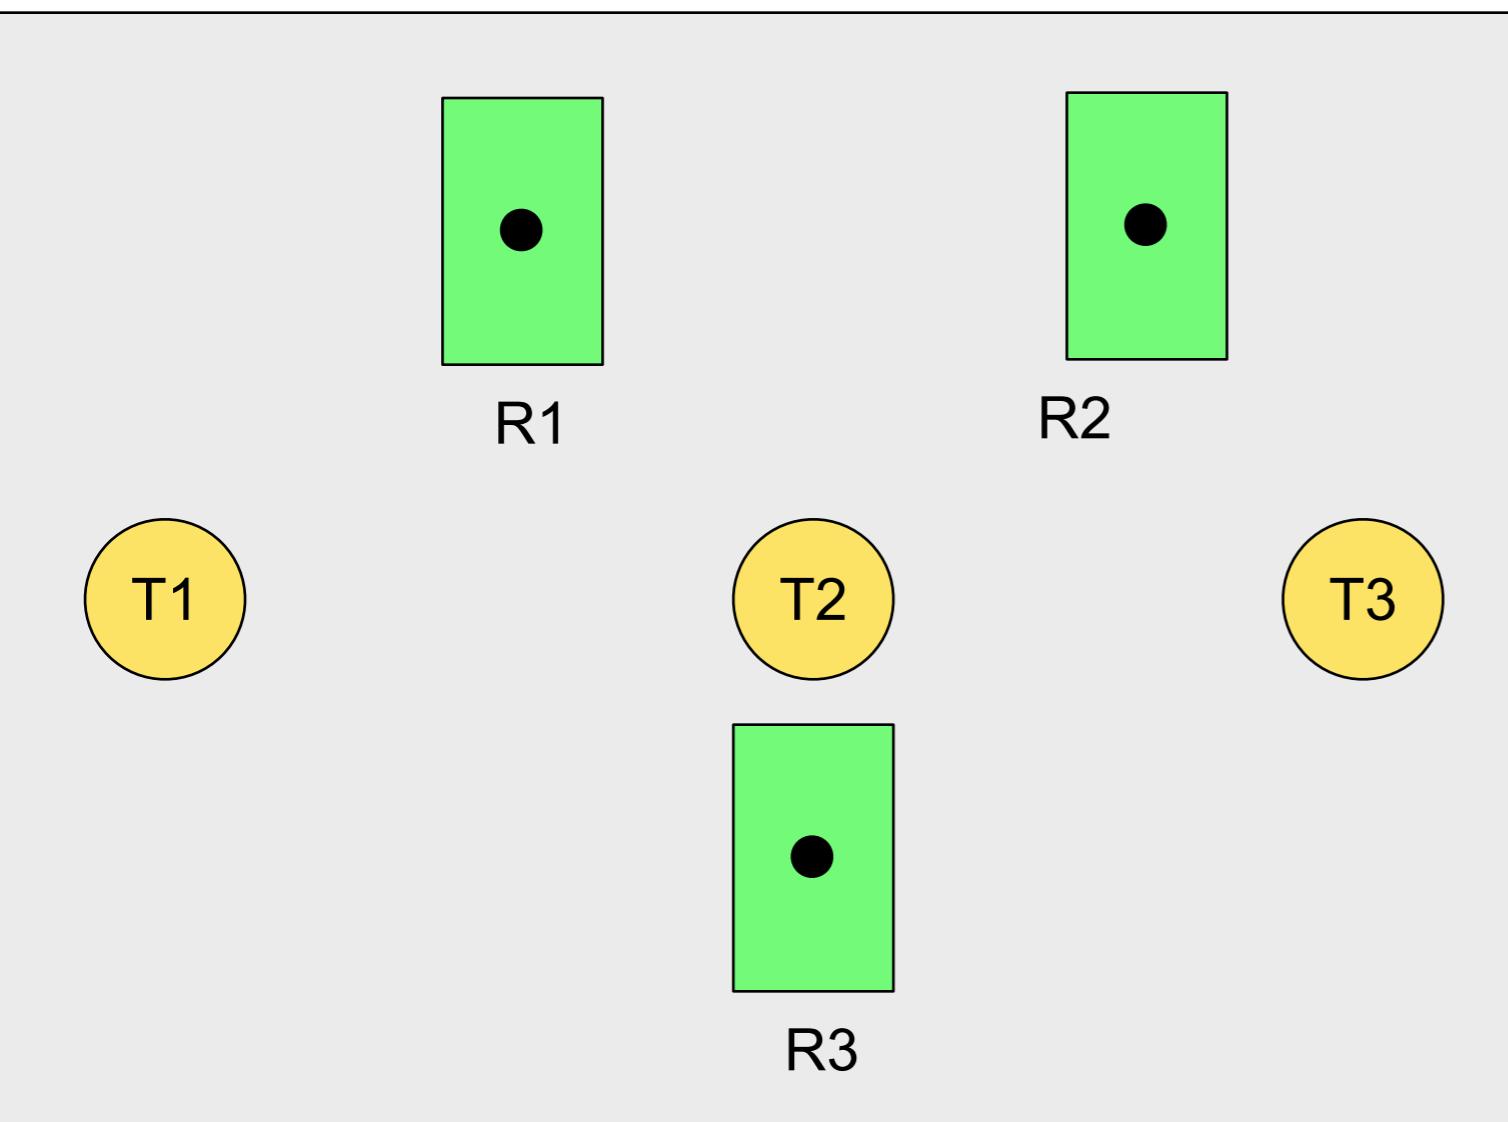
- For simplicity, suppose our system does not have multiple resource instances: One CPU, One Printer,



An instance of R1 can be allocated to T1.

Example with Single-Instance Resources

- ▶ For simplicity, suppose our system does not have multiple resource instances: One CPU, One Printer,

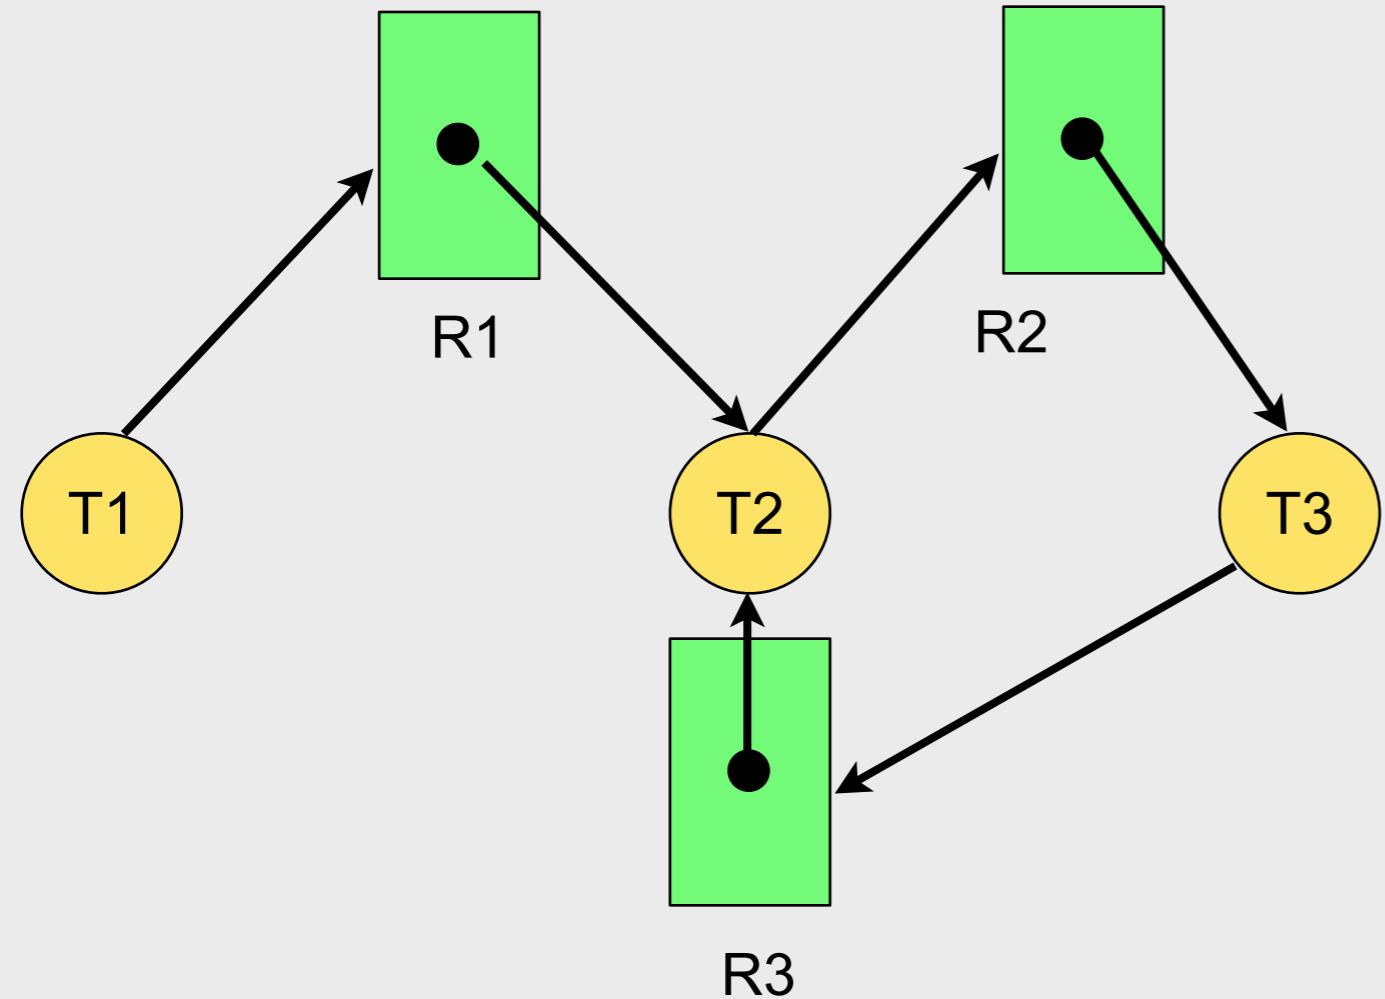


T1 can finish.

T1, T2, T3 all finished: deadlock free!

Another Example

- ▶ How about now? Try reducing edges like earlier.



Is there a deadlock now?

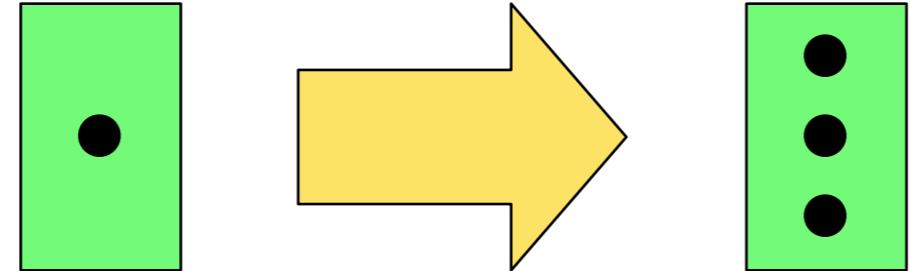
What is different about this and the previous RAG?

Can you identify the four deadlock conditions in this RAG?

RAG for Single-Instance Resources

- ▶ Deadlock detection is trivial for single-instance resources:
 - No Cycle => No Deadlock 
 - Because *Circular Wait* is a necessary condition for a deadlock, and you show it doesn't exist!
 - Cycle => Deadlock 
 - All four conditions would hold in the presence of a cycle
 - ... which is **sufficient** for the presence of a deadlock
- ▶ Cycle detection is just a variant of the Depth First Search (DFS) algorithm.
 - Complexity: $O(|V| + |E|)$

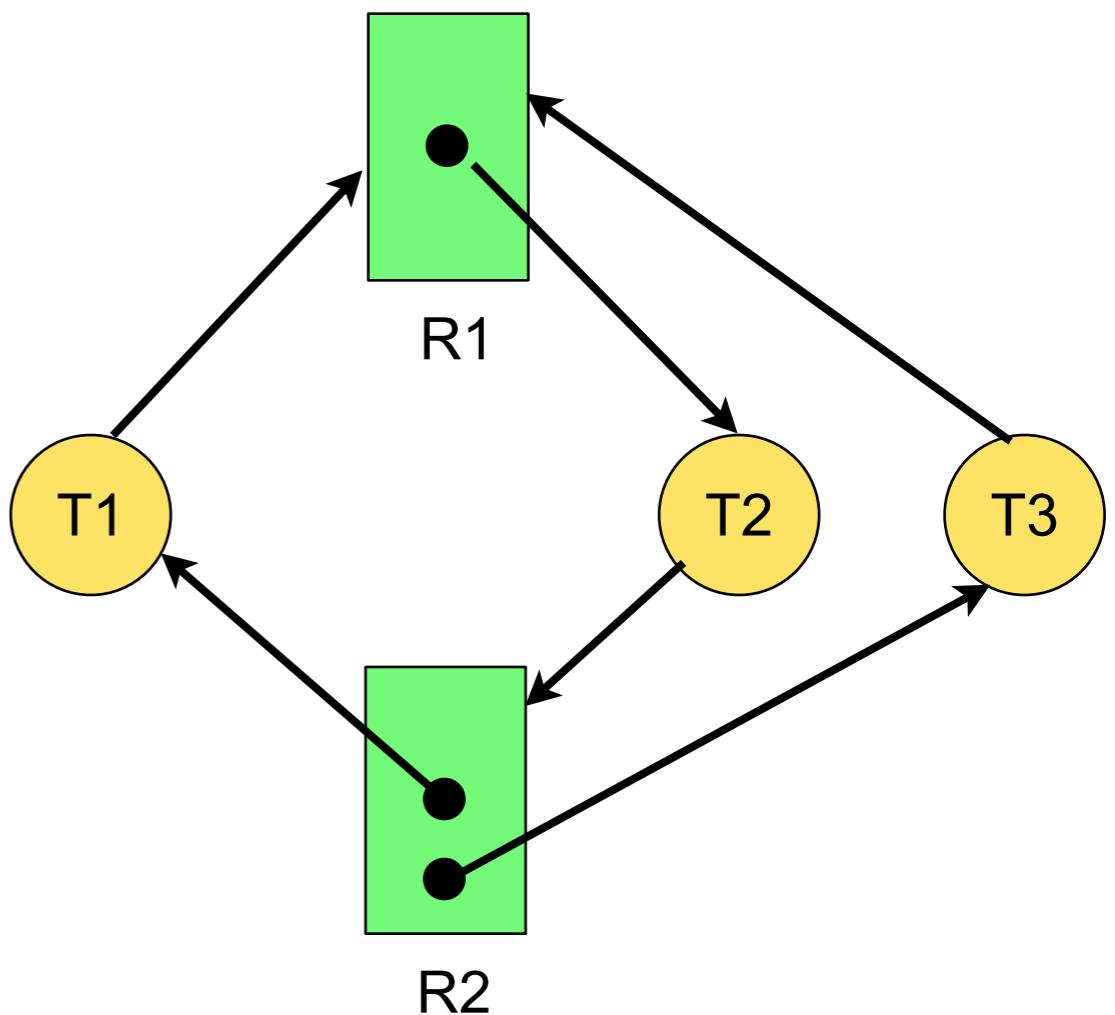
RAGs with Multi-Instance Resources (mRAG)



- ▶ In general, we should consider multiple resource instances
 - In practice, we would have multiple CPU cores
 - We might have multiple disks, printers, graphics cards
- ▶ *New question: Would the same cycle-detection algorithm scale to multi-instance resources?*

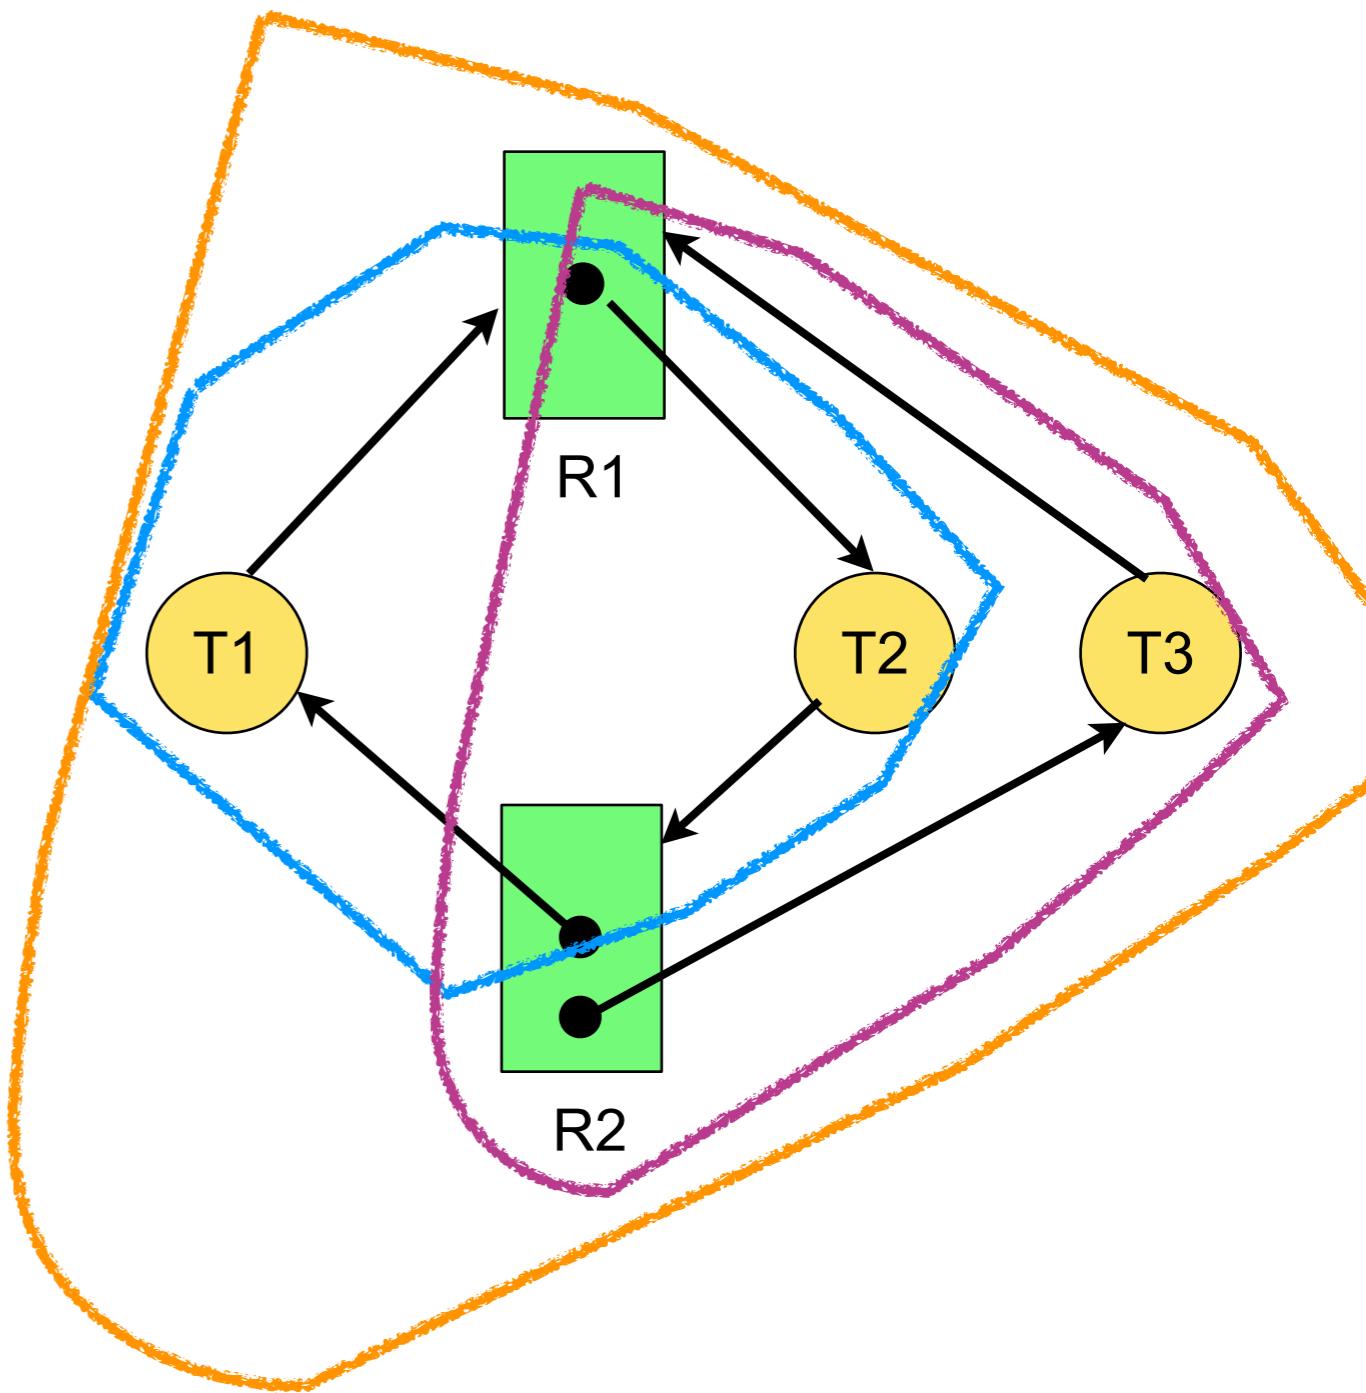
mRAG Example

- ▶ Check whether cycle exists.
- ▶ Is there a deadlock?

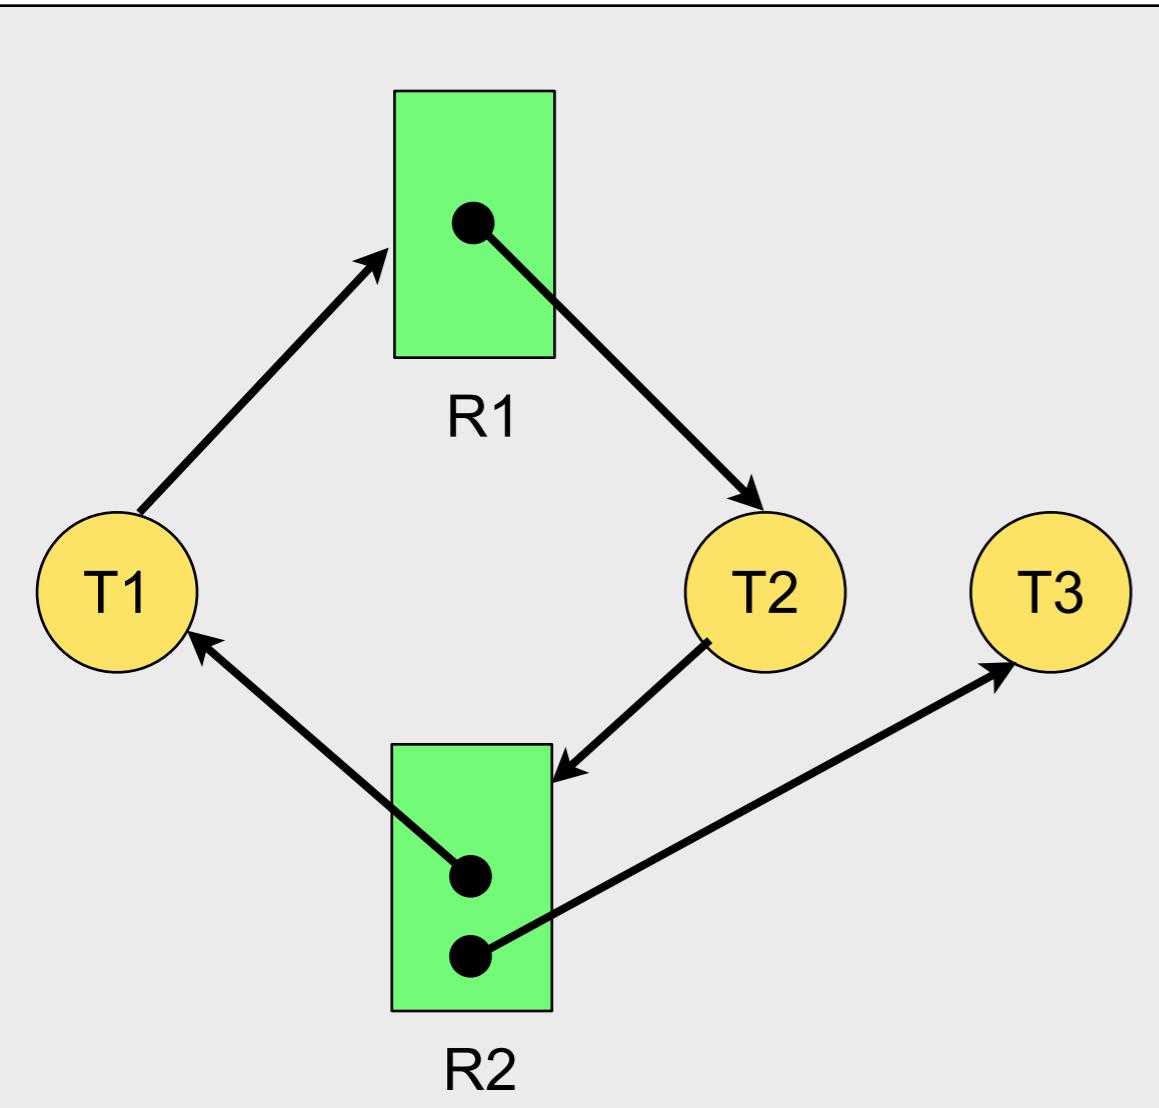


mRAG Example (2)

- ▶ Cycle exists... (actually several)
...and there is a deadlock!
- ▶ The cycle-detection algorithm
seems to work for mRAGs also!



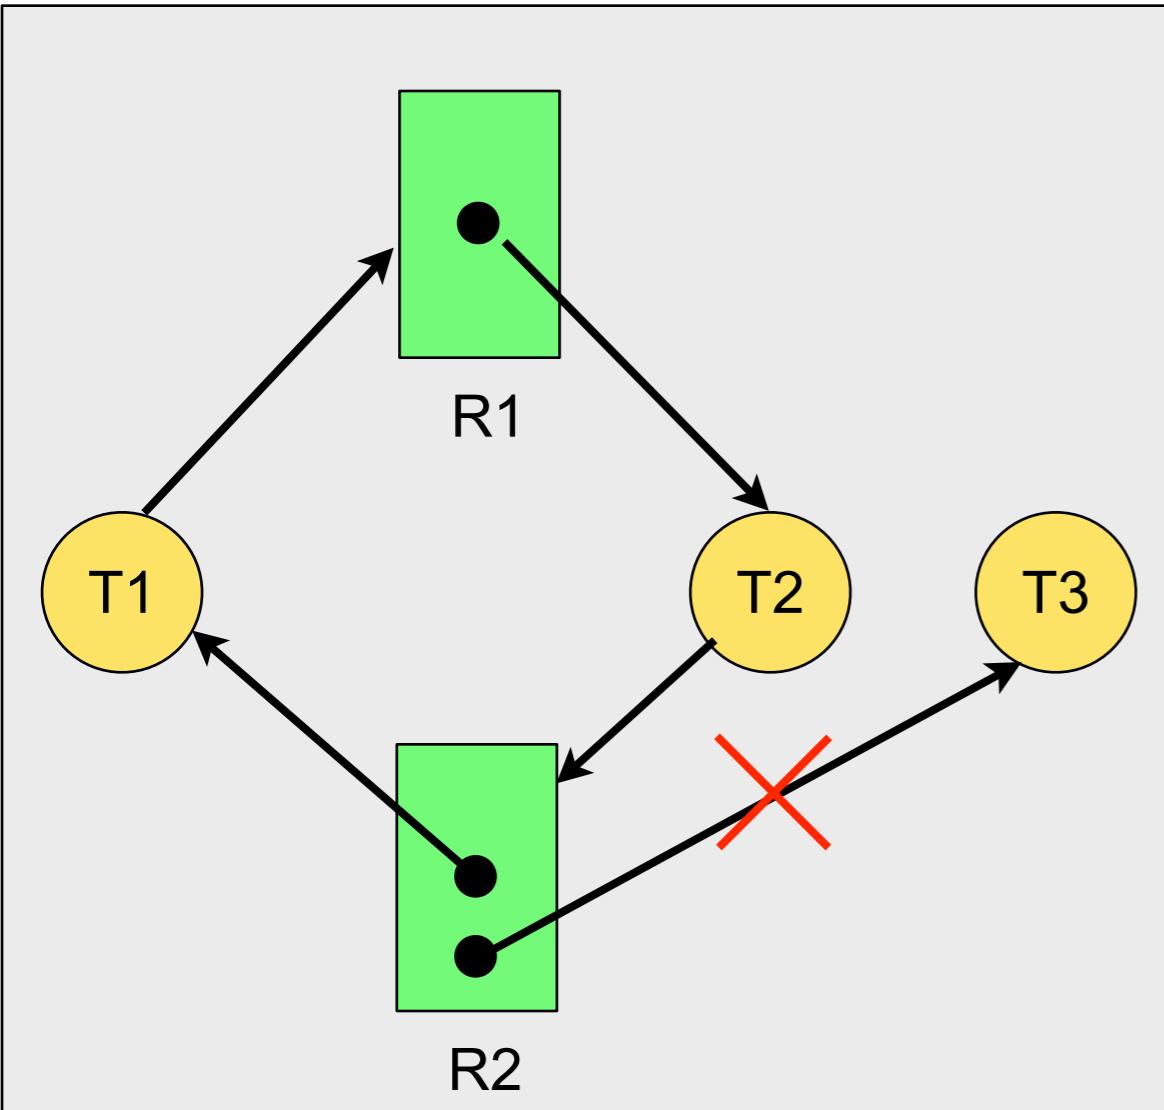
Another mRAG Example (1)



- ▶ Let's make a small change
 - Remove request edge (t_3, r_1)
- ▶ A cycle *still* exists.
- ▶ Deadlock?

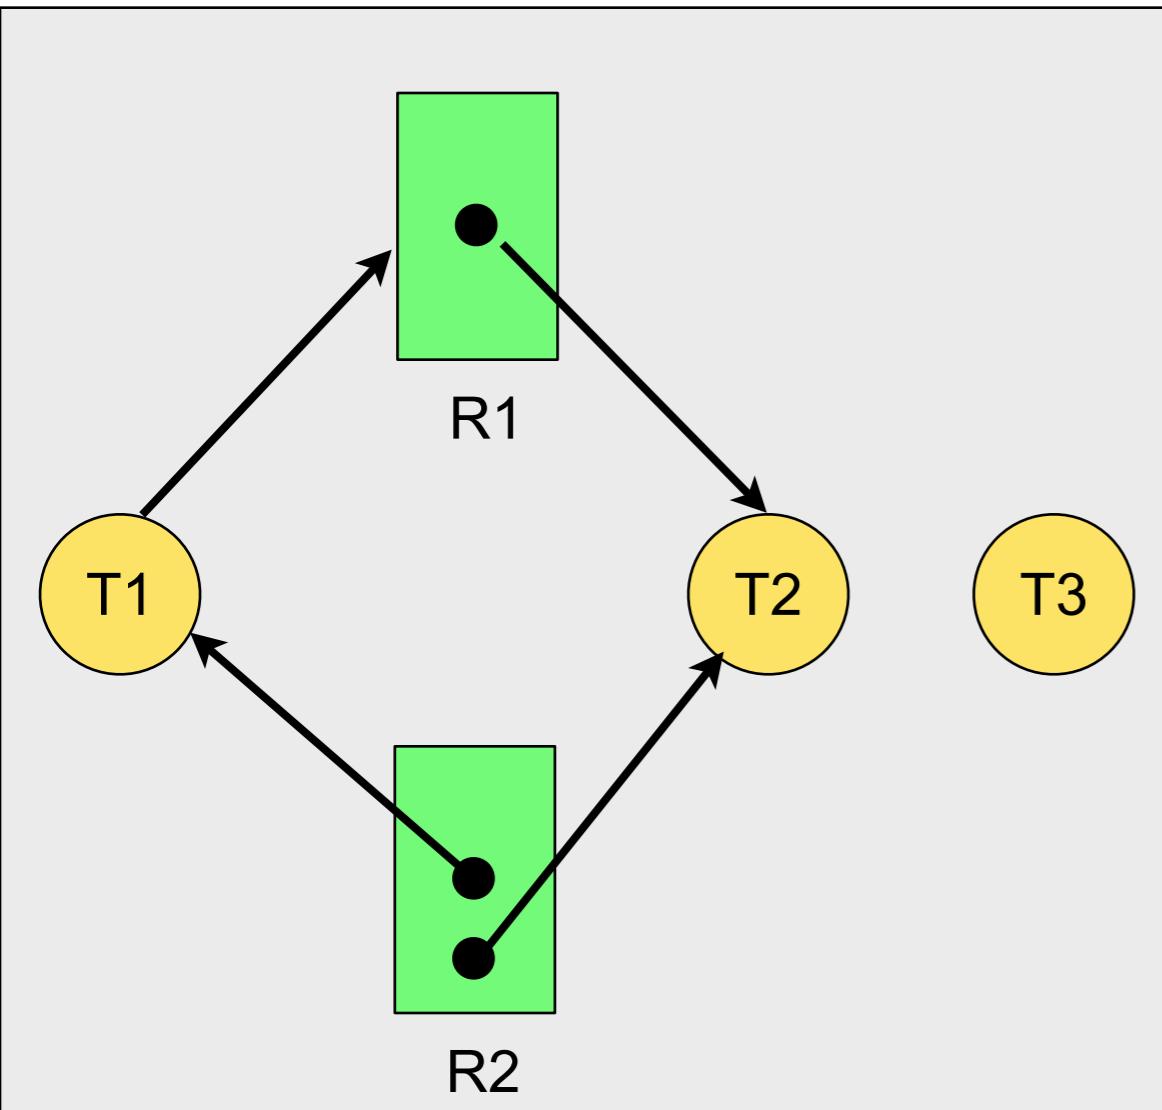
Another mRAG Example (2)

- ▶ T3 is not waiting, so it finishes
 - Allocation edge removed,
 - Instance of R2 reclaimed by OS



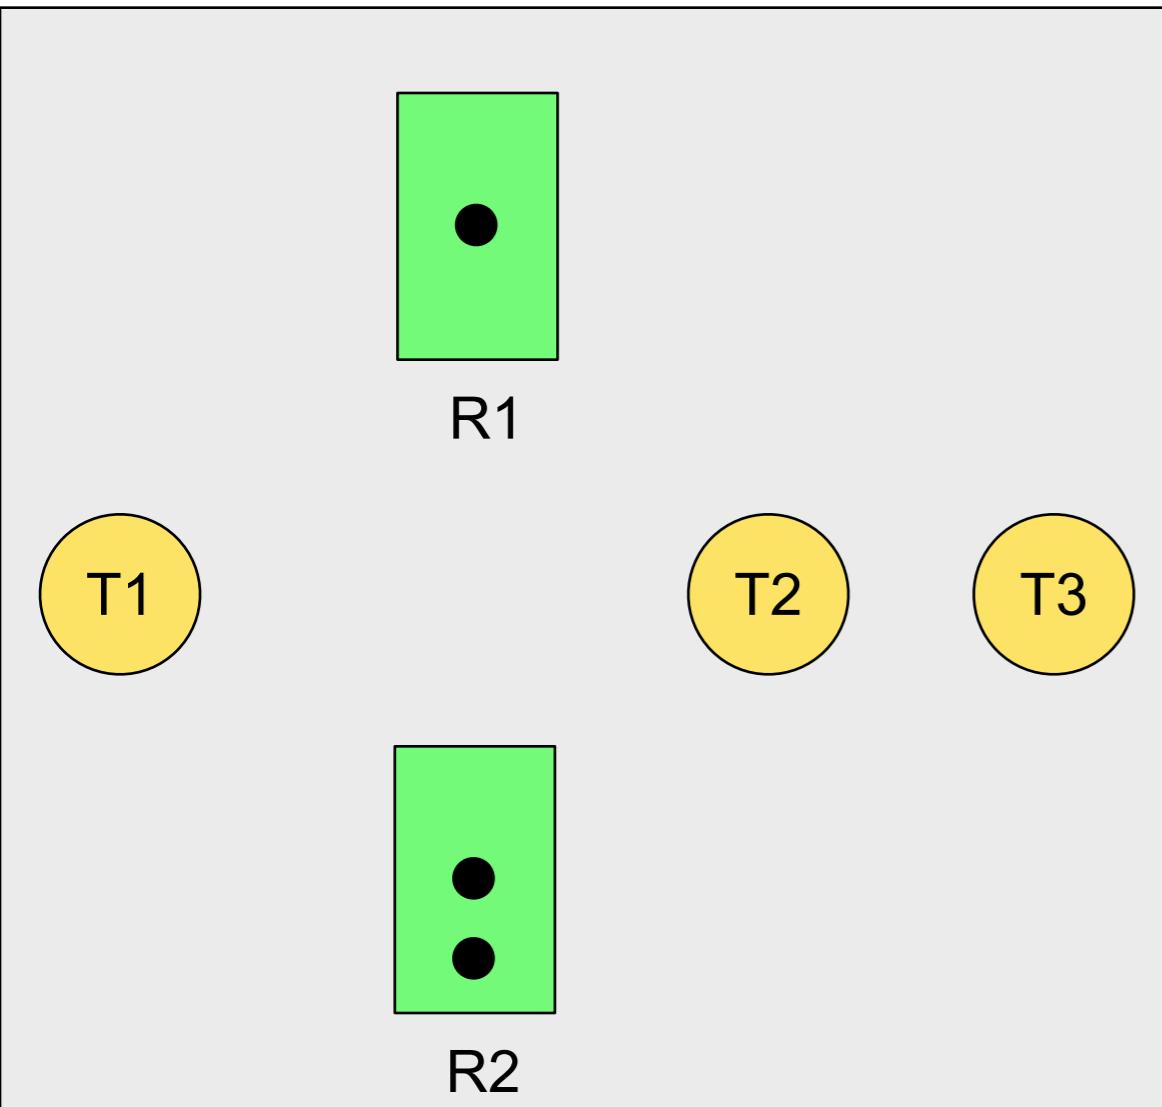
Another mRAG Example (3)

- ▶ The freed instance of R2 gets assigned to T2



- ▶ Request edge transforms into allocation edge

Another mRAG Example (4)

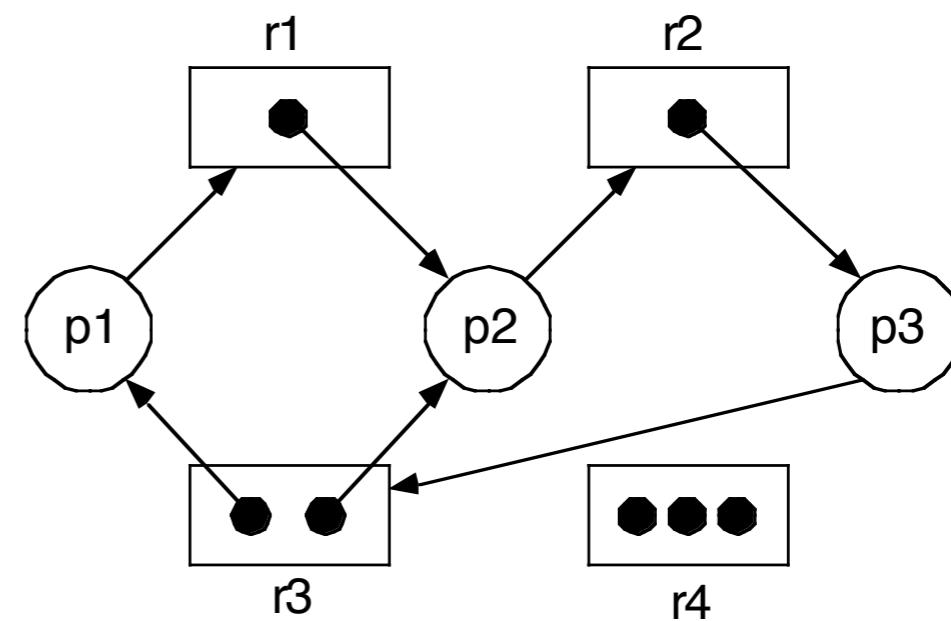
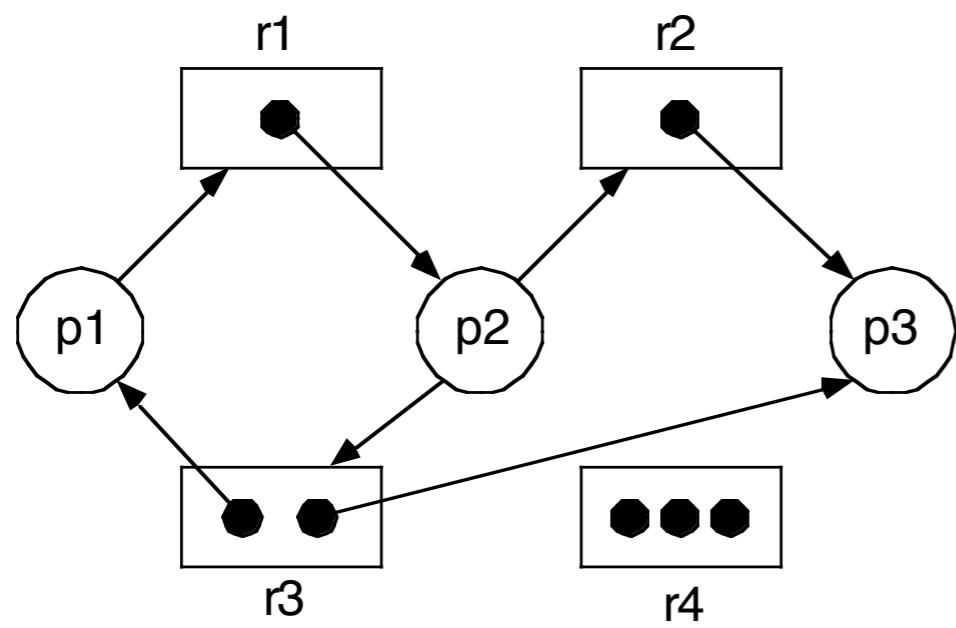


- ▶ Cycle is now removed!
 - T2 can finish, then T1 can finish!
 - **No more deadlock!**
- ▶ What we know: For multi-resource instances,
 - No cycle in RAG => No deadlock
 - Cycle in RAG => No clue!

Towards Deadlock Detection of mRAGs

- ▶ If a cycle is not strong enough for deadlock detection of mRAGs, then what is?
- ▶ **Observe:** One of the following mRAGs has a deadlock.

 - They both have cycles...
 - But what's *different* about the nature of their cycles?



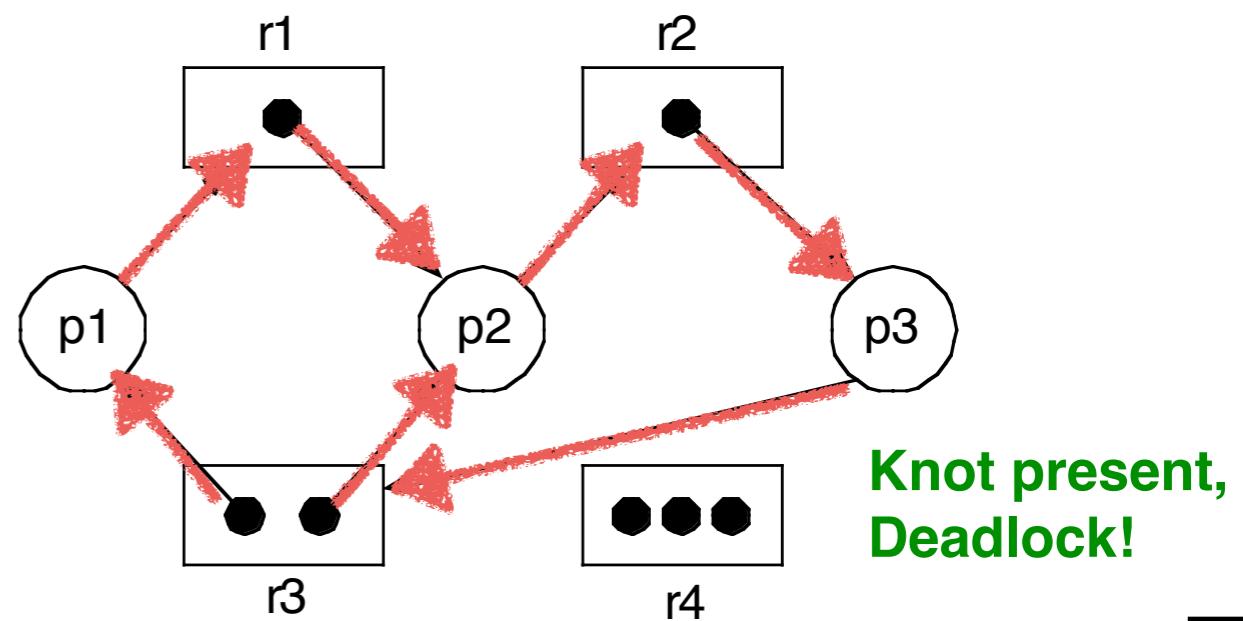
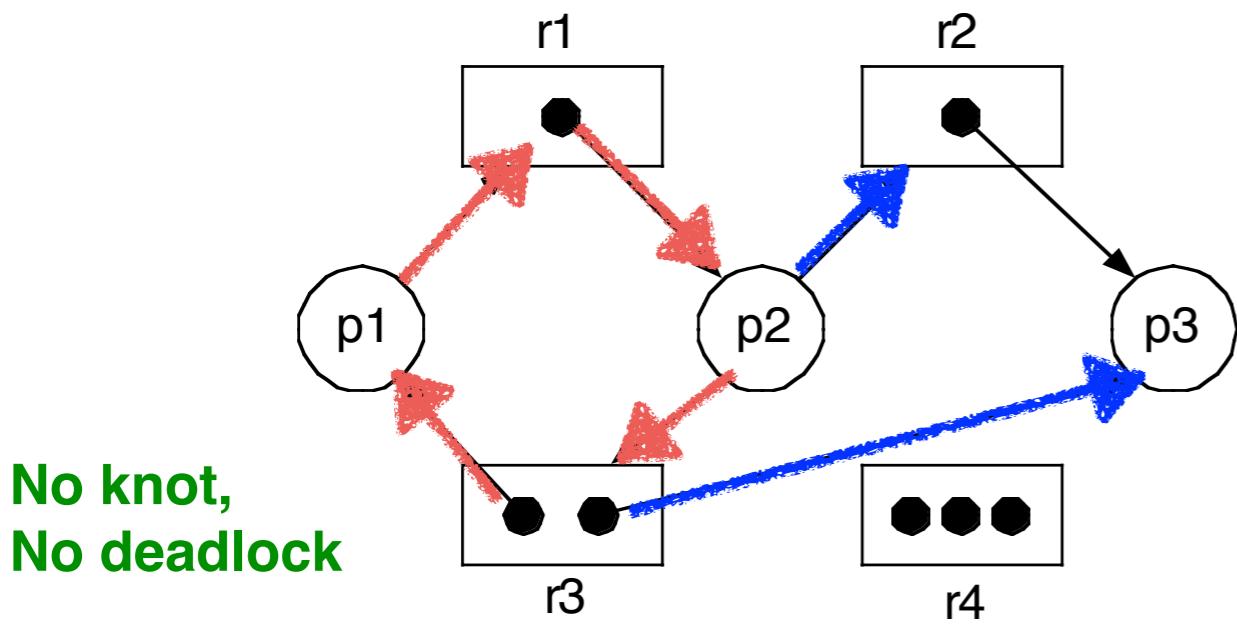
Def'n: Knots!



- ▶ A *Knot* K in a directed graph $G = (V, E)$ is a subgraph such that:
 - K is the *maximal* strongly-connected component (SCC)
 - SCC? Any two nodes $u, v \in K$ must be *reachable* from each other
 - Maximal? No other SCC is a proper subset of K
 - No nodes outside of K is reachable from K
 - (*Essentially, a knot is an SCC you can't escape.*)

Maximal Reachability (SCC)

Escape outside of SCC





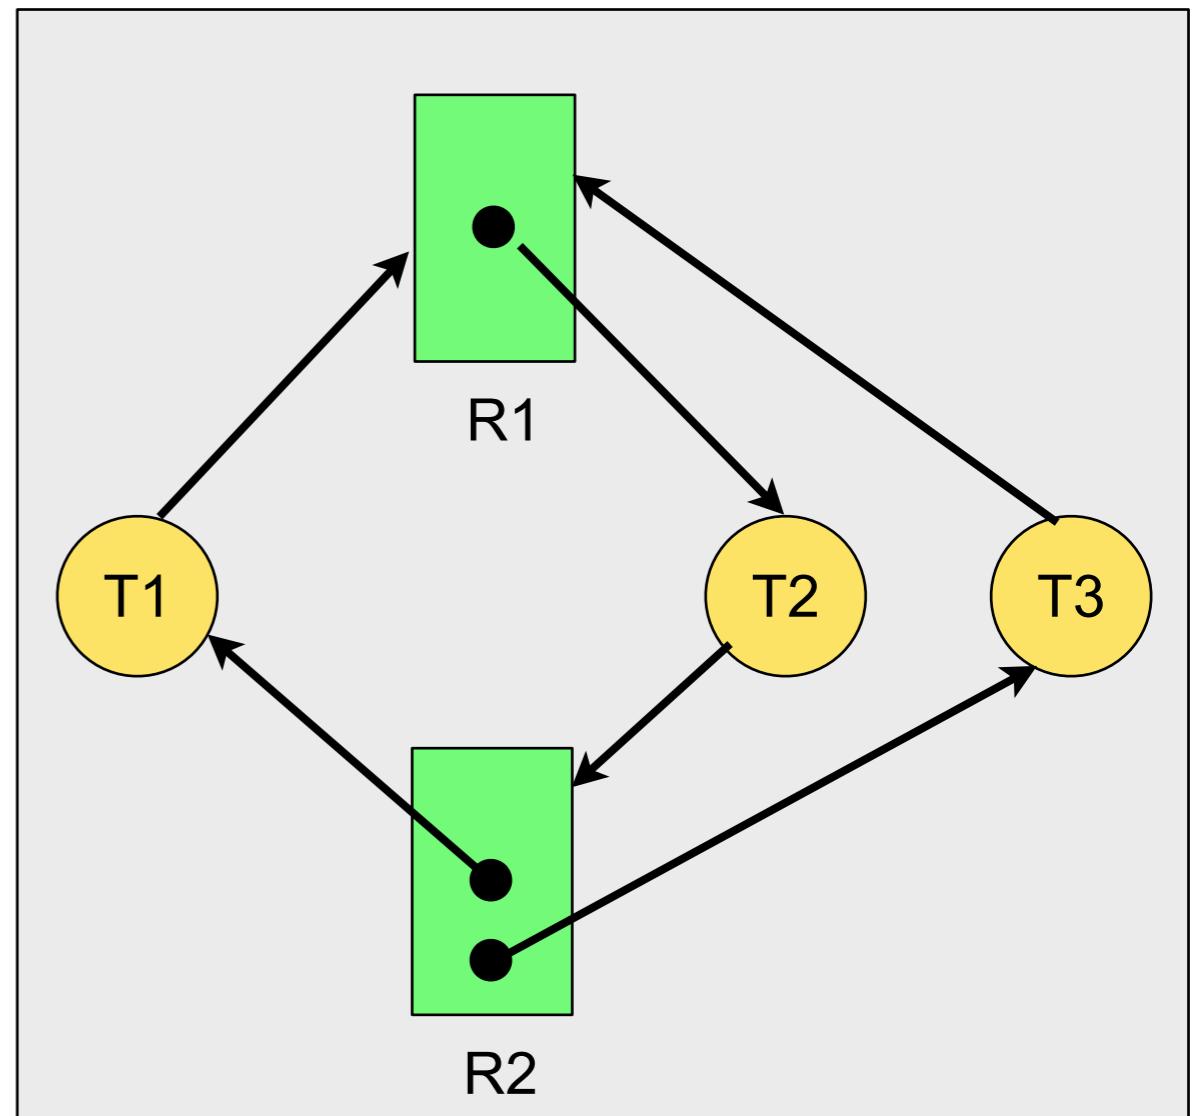
Def'n: Knots!

- ▶ A *Knot* K in a directed graph $G = (V, E)$ is a subgraph such that:
 - K is the *maximal* strongly-connected component (SCC)
 - SCC? Any two nodes $u, v \in K$ must be *reachable* from each other
 - Maximal? No other SCC is a proper subset of K
 - No nodes outside of K is reachable from K
 - Essentially, a knot is a cycle you can't escape.
- ▶ The claim for mRAGs:
 - Knot \implies Deadlock
 - All 4 necessary & sufficient conditions holds in the presence of a knot
 - No Knot \nRightarrow No Deadlock (*Not sure yet, let's check later*)

Back to mRAG Example 1

- ▶ Is there a knot?
 - What are the strongly-connected components (SCCs)?

- ▶ Step 1: List the SCCs
 - $SCC1 = \{T1, R1, T2, R2\}$
 - $SCC2 = \{T2, R1, T3, R2\}$
 - $SCC3 = \{T1, T2, T3, R1, R2\}$
 - ***SCC3 is the maximal SCC since $SCC1 \subset SCC3$, and $SCC2 \subset SCC3$.***

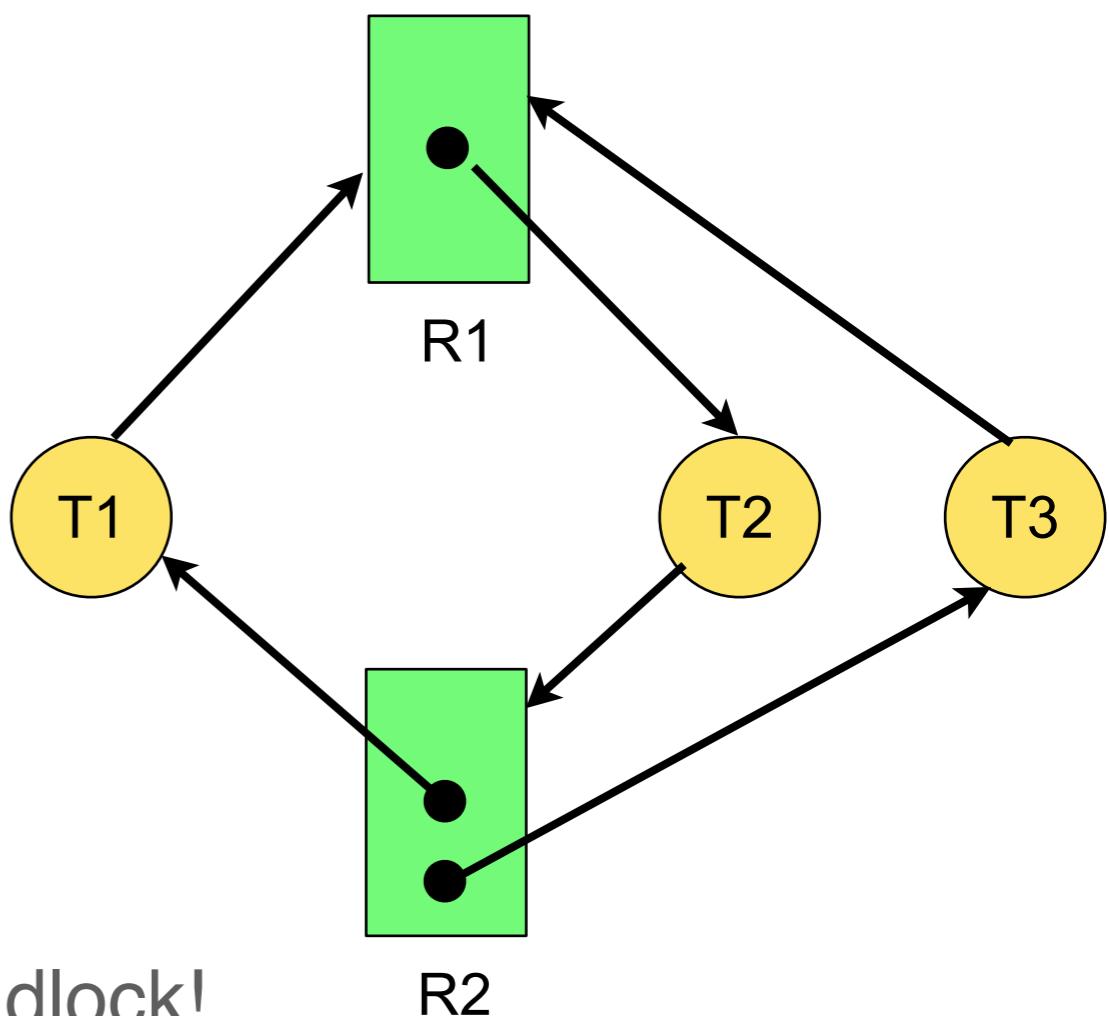


Back to mRAG Example 1

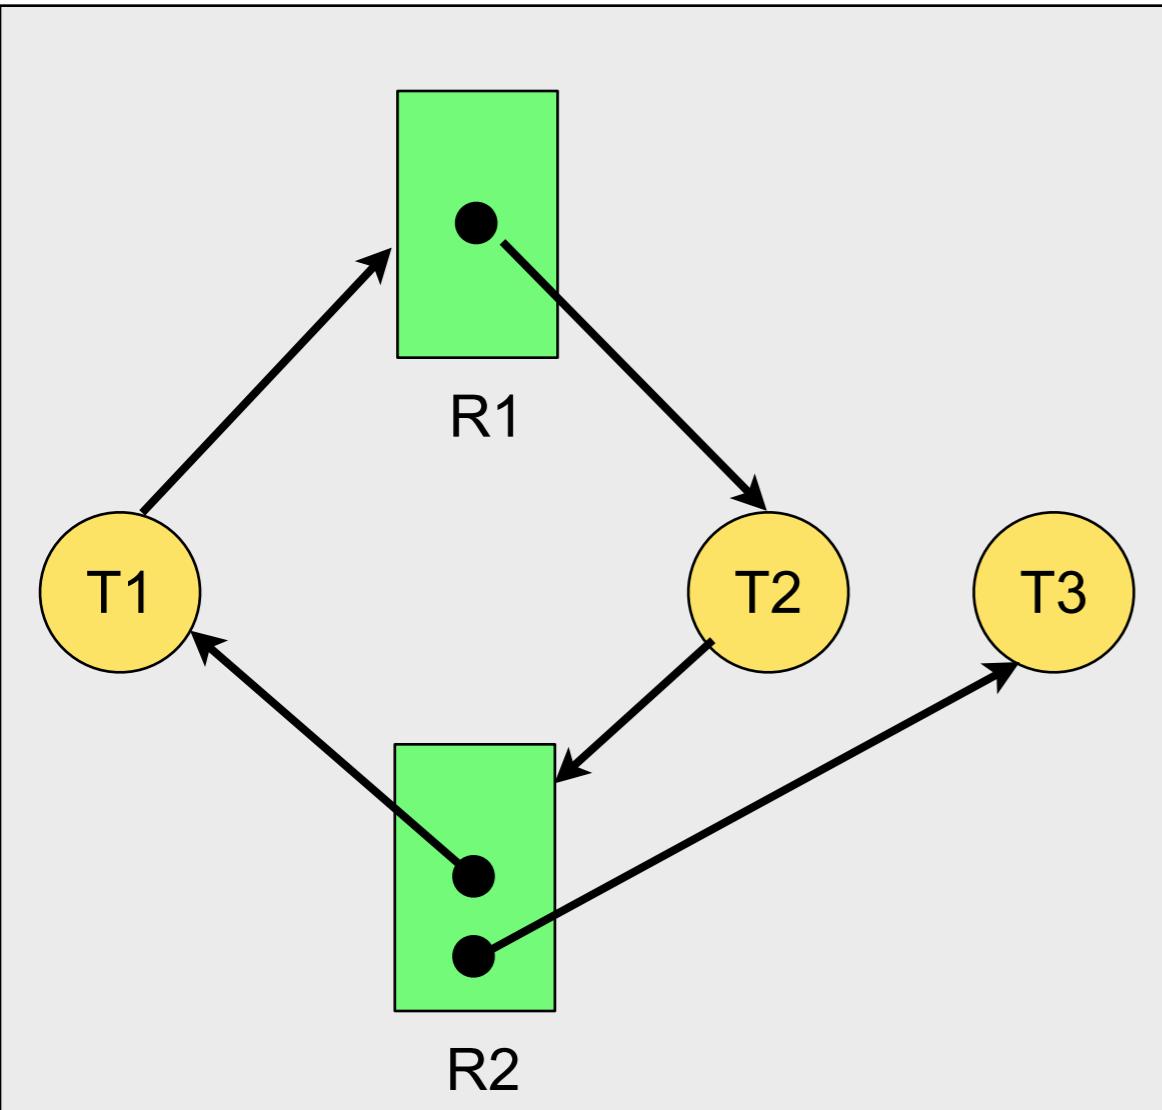
- ▶ Is there a knot?
 - What are the strongly-connected components (SCCs)?

- ▶ Step 1: List the SCCs:
 - $SCC1 = \{T1, R1, T2, R2\}$
 - $SCC2 = \{T2, R1, T3, R2\}$
 - $SCC3 = \{T1, T2, T3, R1, R2\}$

- ▶ Step 2: Is $SCC3$ contained?
 - Yes!
 - So $SCC3$ is a **knot**, and there **is** a deadlock!

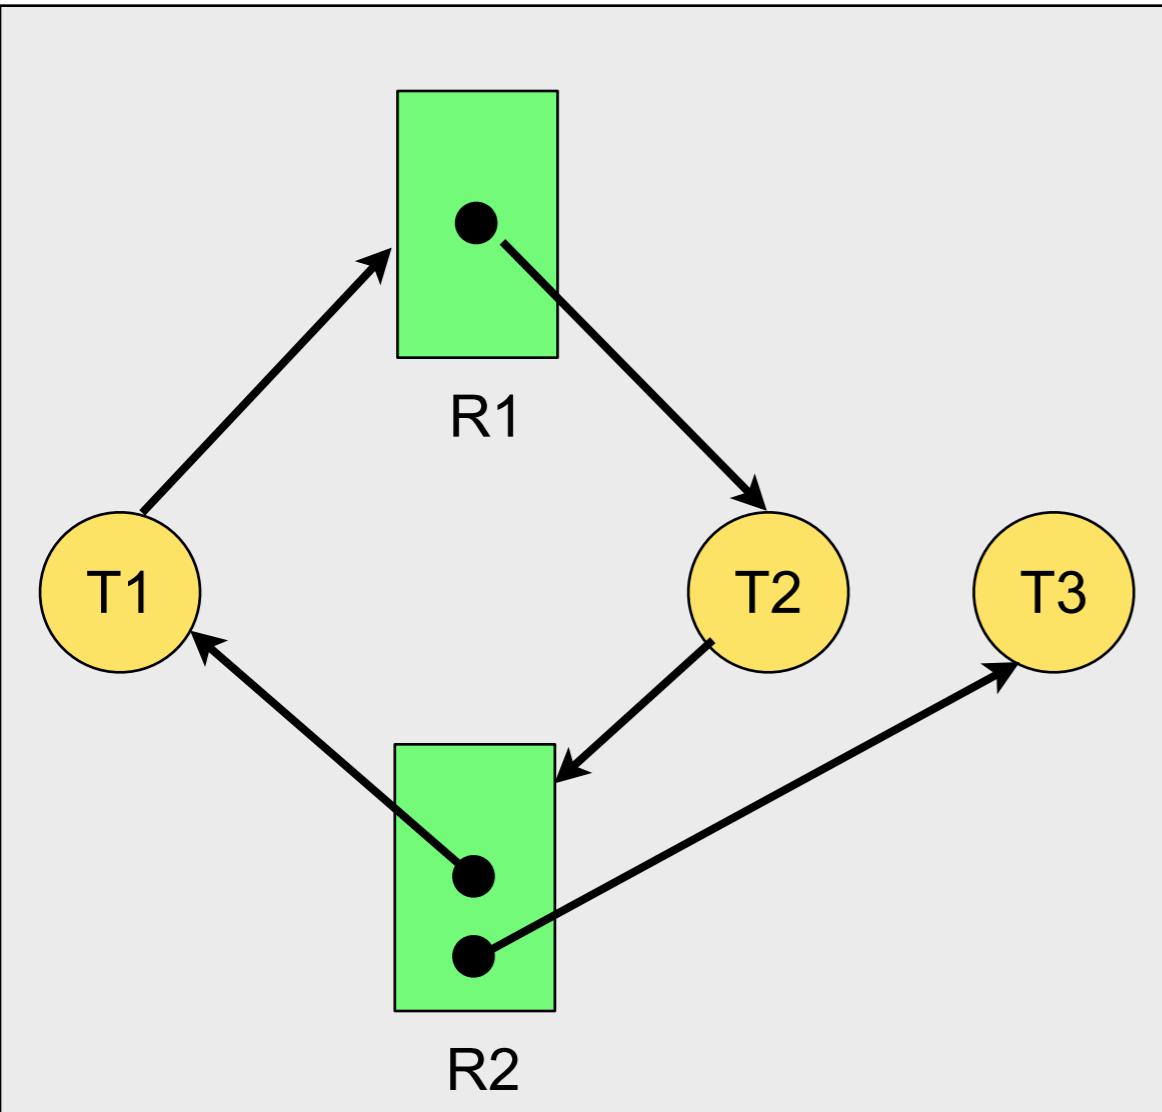


Back to mRAG Example 2



- ▶ How about this example?
 - There's no deadlock here.
- ▶ List the SCCs (just one)
 - $\text{SCC1} = \{\text{T1}, \text{R1}, \text{T2}, \text{R2}\}$
- ▶ Is SCC1 a knot?

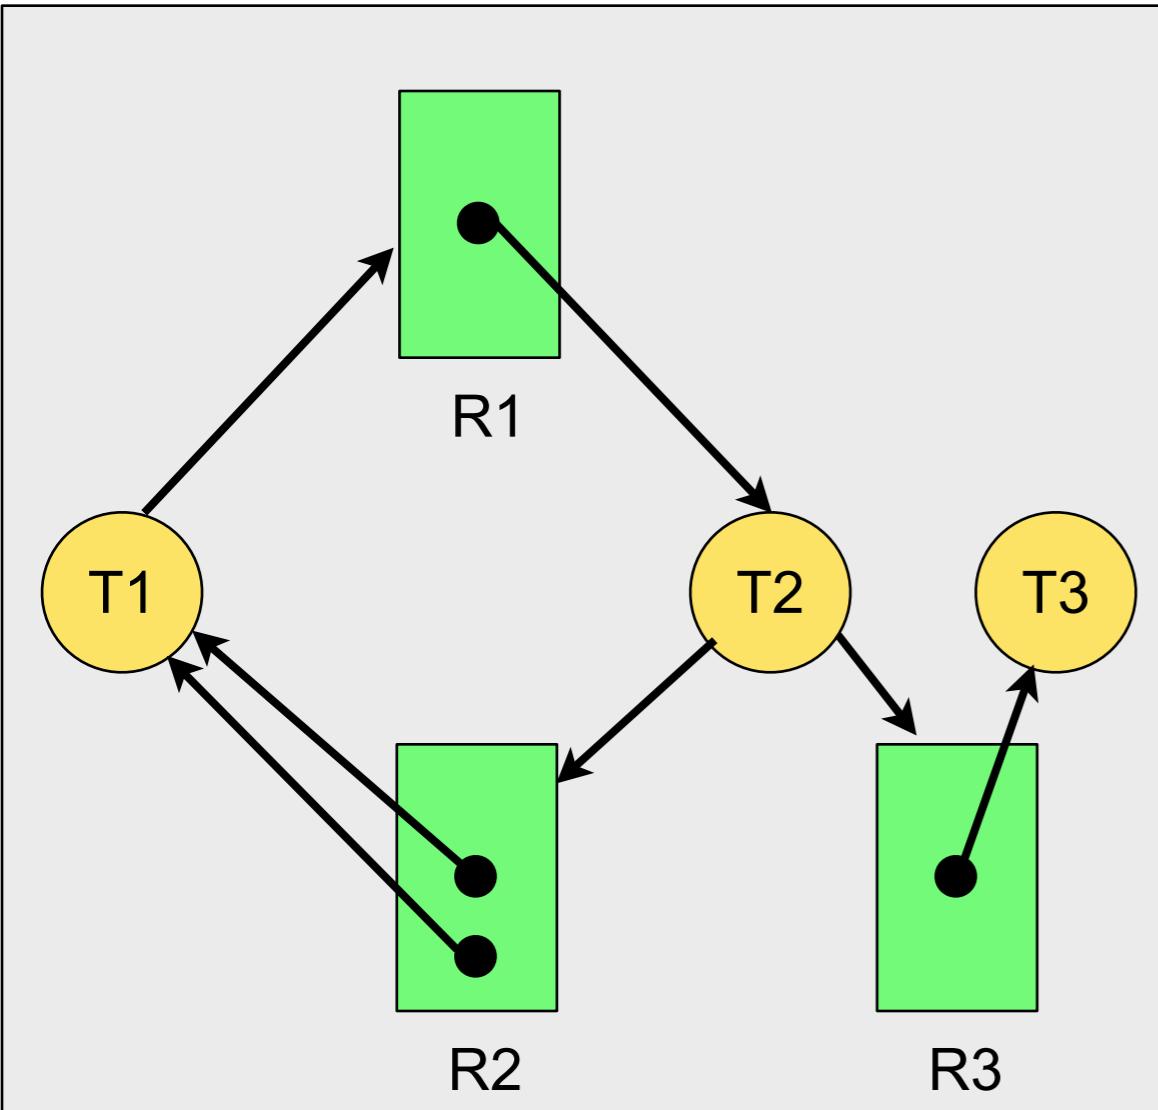
Back to mRAG Example 2



- ▶ How about this example?
 - There's no deadlock here.
- ▶ List the SCCs (just one)
 - $\text{SCC1} = \{\text{T1}, \text{R1}, \text{T2}, \text{R2}\}$
- ▶ Is SCC1 a knot?
 - No! T3 is reachable and it's outside of SCC1. And no deadlock!

One More mRAG Example...

- ▶ T3 is a sink that every other node in the SCC can reach!

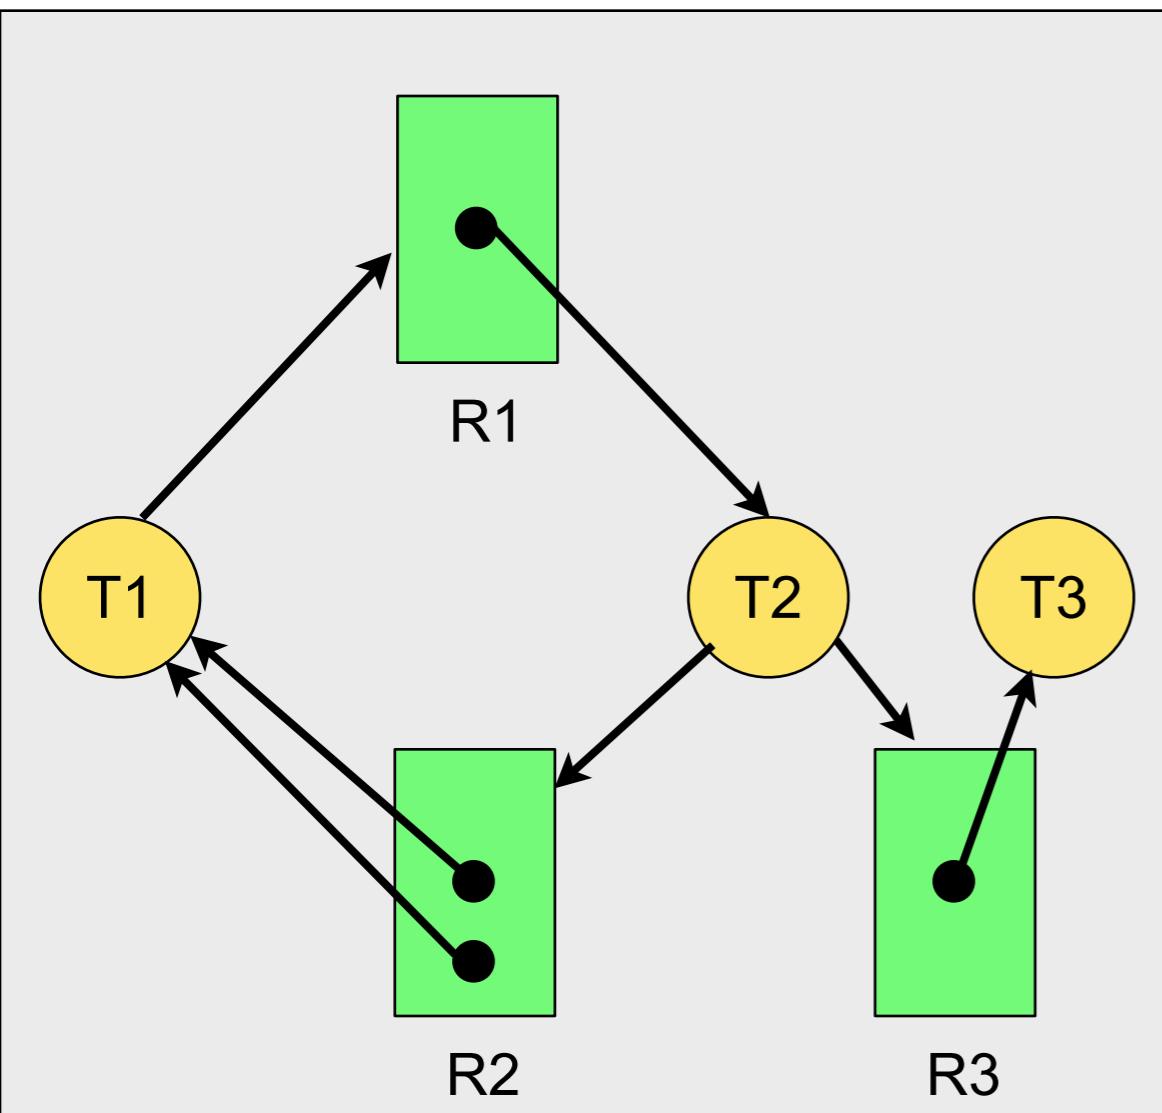


- ▶ So no knot here... but let's confirm by reducing the edges!
 - ▶ (Drat... there's no knot, but there's still a deadlock)

There's no knot! There must not be deadlock?!

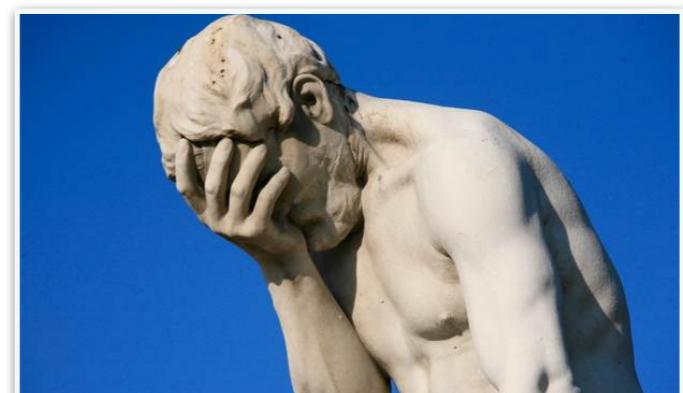
One More mRAG Example... (2)

- ▶ No knot, but a deadlock exists!



- ▶ Conclusion for mRAGS:

- Cycle => No clue
- No Cycle => No Deadlock
- Knot => Deadlock
- No Knot => No clue



There's no knot! We don't know anything.

Summary of RAGs

► Single-Resource Instances (RAGs):

- Existence of cycle is necessary and sufficient for a deadlock
- This can be used in practice
 - *Most resources cannot tolerate multiple simultaneous accesses*
- Example: Linux (lockdep)

► Multi-Resource Instance (mRAGs):

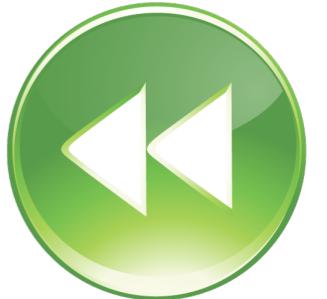
- Existence of a cycle is necessary but insufficient for a deadlock
- Existence of a knot is sufficient for deadlock but unnecessary
- We can't use knot detection reliably: False negatives are possible!

Deadlock Recovery

- ▶ Recovery
 - So we detect a deadlock in a single-instance RAG. Now what?
- ▶ Options:
 - *Do nothing* (This is what Linux does; detect and do nothing)
 - *Terminate* an deadlocked thread to force it to give up resources
 - *Roll back* an offending thread
 - Pick an thread that is involved in the deadlock
 - Rollback execution to a previous state (*how?*), then replay
 - Hope for a deadlock-free schedule to play out

Recovery after Detection (Cont.)

- ▶ How to implement rollbacks?



- ▶ Need to take a *checkpoint* (snapshot of complete system state) periodically.
 - Roll back to last check point, replay all the instructions again.
 - Sounds expensive (*It is!*)
 - Can't things replay exactly the same way? (Yep)
 - What happens to OS performance?

Topics for Today...

- ▶ Deadlock Conditions
- ▶ Ways to deal with deadlocks:
 - Deadlock Detection
 - Single Resource Allocation Graph
 - Multi-Resource Allocation Graph
 - **Deadlock Avoidance**
 - Dijkstra's Bankers Algorithm

Deadlock Avoidance

- ▶ Multi-instance RAGs are a problem.
 - We can show when a deadlock doesn't exist, but...
 - A deadlock can't be reliably detected
- ▶ Another option: Deadlock Avoidance
 - What if we knew all the threads' resource needs in advance?
 - "For me to run, I'll need 3 P's and 2 M's"
 - (Isn't always possible to know in advance)
 - OS keeps track of current system state
 - OS allocates resources in a way that takes OS from one "safe" state to another



Bankers Algorithm (Dijkstra)

- ▶ Risk-averse resource request algorithm
 - Simulate the execution after each resource request
 - Grant the request if the simulation leads to a deadlock-free (safe) state
 - *Works for multi-instance resources!*



What's a "Safe" System State?

There exists an execution sequence T_1, T_2, \dots, T_n such that

T_1 can obtain *all* its resources and finish,

T_2 can obtain *all* its resources and finish, ...

System Model in Banker's Algorithm

n : The number of processes/threads

m : The number of resource types

$$Available = (V_1, \dots, V_m)$$

V_j is the number of still **available** resources of type j

$$Alloc = \begin{pmatrix} A_{1,1} & \dots & A_{1,m} \\ \vdots & \ddots & \vdots \\ A_{n,1} & \dots & A_{n,m} \end{pmatrix}$$

$A_{i,j}$ is the current **allocation** of resource j to thread i

$$Max = \begin{pmatrix} M_{1,1} & \dots & M_{1,m} \\ \vdots & \ddots & \vdots \\ M_{n,1} & \dots & M_{n,m} \end{pmatrix}$$

$M_{i,j}$ is the **maximum demand** of thread i for resource j over the thread's lifetime

Example System State

- Threads: T1, ..., T5 and resource types: Printer, Disk, Semaphore

	P	D	S
<i>Available</i> =	3	3	2

	P	D	S
T1	7	5	3
T2	3	2	2
T3	9	0	2
T4	2	2	2
T5	4	3	3

	P	D	S
T1	0	1	0
T2	2	0	0
T3	3	0	2
T4	2	1	1
T5	0	0	2

- What's the total number of printers we have? **10**
- How many semaphores does T4 still need? **1**
- Say T3 finishes executing.
 - What must change about the state? **Available adds the row at Alloc[3]**
 - Is the OS currently in a safe state? (Yes! But how do we know?)*

Is the Current System State in a Safe State?

m = number of resource types (3)

n = number of threads (5)

```
boolean isSafe(Available, Alloc, Need) {
```

```
    Work[] = Available.clone();
    Finish[] = (0, 0, ..., 0)
```

```
    while (exists unfinished thread i && Need[i] ≤ Work) {
        // pretend that thread i finishes execution
        // then OS can reclaim thread i's allocated resources
        Work += Alloc[i]
        Finish[i] = 1
    }
    // there's an execution order in which all threads
    // can finish. System is safe!
    if (Finish == (1, 1, ..., 1))
        return true
    // there could be a deadlock!
    return false
}
```

Step 1: Initialization

P	D	S
3	3	2

 $Available = (3 \ 3 \ 2)$

P	D	S
0	1	0
2	0	0
3	0	2
2	1	1
0	0	2

 $Alloc = \begin{pmatrix} 0 & 1 & 0 \\ 2 & 0 & 0 \\ 3 & 0 & 2 \\ 2 & 1 & 1 \\ 0 & 0 & 2 \end{pmatrix} \begin{matrix} T1 \\ T2 \\ T3 \\ T4 \\ T5 \end{matrix}$

P	D	S
7	4	3
1	2	2
6	0	0
0	1	1
4	3	1

 $Need = \begin{pmatrix} 7 & 4 & 3 \\ 1 & 2 & 2 \\ 6 & 0 & 0 \\ 0 & 1 & 1 \\ 4 & 3 & 1 \end{pmatrix} \begin{matrix} T1 \\ T2 \\ T3 \\ T4 \\ T5 \end{matrix}$

Is the Current System State in a Safe State?

m = number of resource types (3)

n = number of threads (5)

```
boolean isSafe(Available, Alloc, Need) {
    Work[] = Available.clone();
    Finish[] = (0, 0, ..., 0)

    while (exists unfinished thread i && Need[i] ≤ Work) {
        // pretend that thread i finishes execution
        // then OS can reclaim thread i's allocated resources
        Work += Alloc[i]
        Finish[i] = 1
    }

    // there's an execution order in which all threads
    // can finish. System is safe!
    if (Finish == (1, 1, ..., 1))
        return true
    // there could be a deadlock!
    return false
}
```

Simplified Notation

Let X and Y be vectors of length n .

We say that $X \leq Y$ if and only if $X[i] \leq Y[i]$ for all i

$$Available = \begin{pmatrix} P & D & S \\ 3 & 3 & 2 \end{pmatrix}$$

$$Alloc = \begin{pmatrix} P & D & S \\ 0 & 1 & 0 \\ 2 & 0 & 0 \\ 3 & 0 & 2 \\ 2 & 1 & 1 \\ 0 & 0 & 2 \end{pmatrix} \begin{matrix} T1 \\ T2 \\ T3 \\ T4 \\ T5 \end{matrix}$$

Step 2: Simulation

$$Need = \begin{pmatrix} P & D & S \\ 7 & 4 & 3 \\ 1 & 2 & 2 \\ 6 & 0 & 0 \\ 0 & 1 & 1 \\ 4 & 3 & 1 \end{pmatrix} \begin{matrix} T1 \\ T2 \\ T3 \\ T4 \\ T5 \end{matrix}$$

Yes, we're in a safe state!!!

m = number of resource types (3)

n = number of threads (5)

```

boolean isSafe(Available, Alloc, Need) {
    Work[] = Available.clone();
    Finish[] = (0, 0, ..., 0)

    while (exists unfinished thread i && Need[i] ≤ Work) {
        // pretend that thread i finishes execution
        // then OS can reclaim thread i's allocated resources
        Work += Alloc[i]
        Finish[i] = 1
    }

    // there's an execution order in which all threads
    // can finish. System is safe!
    if (Finish == (1, 1, ..., 1))
        return true
    // there could be a deadlock!
    return false
}
  
```

Step 3: Determine safety

$$Available = \begin{pmatrix} P & D & S \\ 3 & 3 & 2 \end{pmatrix}$$

$$Alloc = \begin{pmatrix} P & D & S \\ 0 & 1 & 0 \\ 2 & 0 & 0 \\ 3 & 0 & 2 \\ 2 & 1 & 1 \\ 0 & 0 & 2 \end{pmatrix} \begin{matrix} T1 \\ T2 \\ T3 \\ T4 \\ T5 \end{matrix}$$

$$Need = \begin{pmatrix} P & D & S \\ 7 & 4 & 3 \\ 1 & 2 & 2 \\ 6 & 0 & 0 \\ 0 & 1 & 1 \\ 4 & 3 & 1 \end{pmatrix} \begin{matrix} T1 \\ T2 \\ T3 \\ T4 \\ T5 \end{matrix}$$

An Unsafe Example: $Available = (0 \ 1 \ 1)$

m = number of resource types (3)
 n = number of threads (5)

```
boolean isSafe(Available, Alloc, Need) {
    Work[] = Available.clone();
    Finish[] = (0, 0, ..., 0)

    while (exists unfinished thread i && Need[i] ≤ Work) {
        // pretend that thread i finishes execution
        // then OS can reclaim thread i's allocated resources
        Work += Alloc[i]
        Finish[i] = 1
    }
    // there's an execution order in which all threads
    // can finish. System is safe!
    if (Finish == (1, 1, ..., 1))
        return true
    // there could be a deadlock!
    return false
}
```

$$Available = \begin{pmatrix} P & D & S \\ 0 & 1 & 1 \end{pmatrix}$$

$$Alloc = \begin{pmatrix} P & D & S \\ 0 & 1 & 0 \\ 2 & 0 & 0 \\ 3 & 0 & 2 \\ 2 & 1 & 1 \\ 0 & 0 & 2 \end{pmatrix} \begin{matrix} T1 \\ T2 \\ T3 \\ T4 \\ T5 \end{matrix}$$

$$Need = \begin{pmatrix} P & D & S \\ 7 & 4 & 3 \\ 1 & 2 & 2 \\ 6 & 0 & 0 \\ 0 & 1 & 1 \\ 4 & 3 & 1 \end{pmatrix} \begin{matrix} T1 \\ T2 \\ T3 \\ T4 \\ T5 \end{matrix}$$

Banker's: What Good Is a Safety Checker?

- ▶ We want the OS to use the safety algo to determine whether a resource request should be granted (safe) or denied (unsafe)
 - *This is Banker's Algorithm*
- ▶ Back to previous example:
 - Let's say **T2** requests 1 printer and 2 semaphores
 - *Should the OS grant this request?*

$$Req[2] = \begin{pmatrix} P & D & S \\ 1 & 0 & 2 \end{pmatrix}$$

$$Available = \begin{pmatrix} P & D & S \\ 3 & 3 & 2 \end{pmatrix}$$

$$Max = \begin{pmatrix} P & D & S \\ 7 & 5 & 3 \\ 3 & 2 & 2 \\ 9 & 0 & 2 \\ 2 & 2 & 2 \\ 4 & 3 & 3 \end{pmatrix} \quad \begin{matrix} T1 \\ T2 \\ T3 \\ T4 \\ T5 \end{matrix}$$

$$Alloc = \begin{pmatrix} P & D & S \\ 0 & 1 & 0 \\ 2 & 0 & 0 \\ 3 & 0 & 2 \\ 2 & 1 & 1 \\ 0 & 0 & 2 \end{pmatrix} \quad \begin{matrix} T1 \\ T2 \\ T3 \\ T4 \\ T5 \end{matrix}$$

Banker's Algorithm

```

requestResources(Req[i], Available, Max, Alloc) {
    // Create and initialize the Need matrix
    Need = Max - Alloc

    if (Req[i] > Need[i])
        perror("Thread %d requested more than it needs\n", i)
    else if (Req[i] > Available)
        // can't fulfill request right now, try again later
        wait(thread[i])
    else {
        // update system state as if the request was granted
        Alloc[i] += Req[i]
        Need[i] -= Req[i]
        Available -= Req[i]

        if (isSafe(Available, Alloc, Need))
            grantRequest(thread[i])
        else {
            // restore system state
            Alloc[i] -= Req[i]
            Need[i] += Req[i]
            Available += Req[i]
            wait(thread[i])
        }
    }
}
  
```

$$P \quad D \quad S$$

$$Req[2] = \begin{pmatrix} 1 & 0 & 2 \end{pmatrix}$$

$$P \quad D \quad S$$

$$Available = \begin{pmatrix} 3 & 3 & 2 \end{pmatrix}$$

$$Alloc = \begin{pmatrix} 0 & 1 & 0 \\ 2 & 0 & 0 \\ 3 & 0 & 2 \\ 2 & 1 & 1 \\ 0 & 0 & 2 \end{pmatrix} \begin{matrix} T1 \\ T2 \\ T3 \\ T4 \\ T5 \end{matrix}$$

$$P \quad D \quad S$$

$$Max = \begin{pmatrix} 7 & 5 & 3 \\ 3 & 2 & 2 \\ 9 & 0 & 2 \\ 2 & 2 & 2 \\ 4 & 3 & 3 \end{pmatrix} \begin{matrix} T1 \\ T2 \\ T3 \\ T4 \\ T5 \end{matrix}$$

Banker's Algorithm

```
requestResources(Req[i], Available, Max, Alloc) {
```

```
  // Create and initialize the Need matrix
```

```
  Need = Max - Alloc
```

Step 1: Initialization

```
  if (Req[i] > Need[i])
    perror("Thread %d requested more than it needs\n", i)
```

```
  else if (Req[i] > Available)
    // can't fulfill request right now, try again later
```

```
    wait(thread[i])
```

```
  else {
```

```
    // update system state as if the request was granted
```

```
    Alloc[i] += Req[i]
```

```
    Need[i] -= Req[i]
```

```
    Available -= Req[i]
```

```
    if (isSafe(Available, Alloc, Need))
      grantRequest(thread[i])
```

```
    else {
```

```
      // restore system state
```

```
      Alloc[i] -= Req[i]
```

```
      Need[i] += Req[i]
```

```
      Available += Req[i]
```

```
      wait(thread[i])
```

```
}
```

```
}
```

P	D	S
1	0	2

Req[2] = (1 0 2)

P	D	S
3	3	2

Available = (3 3 2)

P	D	S
0	1	0
2	0	0
3	0	2
2	1	1
0	0	2

Alloc = (T1 T2 T3 T4 T5)

P	D	S
7	4	3
1	2	2
6	0	0
0	1	1
4	3	1

Need = (T1 T2 T3 T4 T5)

Banker's Algorithm

```

requestResources(Req[i], Available, Max, Alloc) {
    // Create and initialize the Need matrix
    Need = Max - Alloc
    Step 2: Sanity Check

    if (Req[i] > Need[i])
        perror("Thread %d requested more than it needs\n", i)
    else if (Req[i] > Available)
        // can't fulfill request right now, try again later
        wait(thread[i])
    else {
        // update system state as if the request was granted
        Alloc[i] += Req[i]
        Need[i] -= Req[i]
        Available -= Req[i]

        if (isSafe(Available, Alloc, Need))
            grantRequest(thread[i])
        else {
            // restore system state
            Alloc[i] -= Req[i]
            Need[i] += Req[i]
            Available += Req[i]
            wait(thread[i])
        }
    }
}

```

$$P \quad D \quad S$$

$$Req[2] = \begin{pmatrix} 1 & 0 & 2 \end{pmatrix}$$

$$P \quad D \quad S$$

$$Available = \begin{pmatrix} 3 & 3 & 2 \end{pmatrix}$$

$$Alloc = \begin{pmatrix} P & D & S \\ 0 & 1 & 0 \\ 2 & 0 & 0 \\ 3 & 0 & 2 \\ 2 & 1 & 1 \\ 0 & 0 & 2 \end{pmatrix} \begin{matrix} T1 \\ T2 \\ T3 \\ T4 \\ T5 \end{matrix}$$

$$P \quad D \quad S$$

$$Need = \begin{pmatrix} 7 & 4 & 3 \\ 1 & 2 & 2 \\ 6 & 0 & 0 \\ 0 & 1 & 1 \\ 4 & 3 & 1 \end{pmatrix} \begin{matrix} T1 \\ T2 \\ T3 \\ T4 \\ T5 \end{matrix}$$

Banker's Algorithm

```

requestResources(Req[i], Available, Max, Alloc) {
    // Create and initialize the Need matrix
    Need = Max - Alloc

    if (Req[i] > Need[i])
        perror("Thread %d requested more than it needs\n", i)
    else if (Req[i] > Available)
        // can't fulfill request right now, try again later
        wait(thread[i])
    else {
        // update system state as if the request was granted
        Alloc[i] += Req[i]
        Need[i] -= Req[i]
        Available -= Req[i]
    }
}
  
```

Step 3: Assume request is granted...

$$P \quad D \quad S$$

$$Req[2] = \begin{pmatrix} 1 & 0 & 2 \end{pmatrix}$$

$$P \quad D \quad S$$

$$Available = \begin{pmatrix} 2 & 3 & 0 \end{pmatrix}$$

$$Alloc = \begin{pmatrix} 0 & 1 & 0 \\ 3 & 0 & 2 \\ 3 & 0 & 2 \\ 2 & 1 & 1 \\ 0 & 0 & 2 \end{pmatrix} \begin{matrix} T1 \\ T2 \\ T3 \\ T4 \\ T5 \end{matrix}$$

$$P \quad D \quad S$$

$$Need = \begin{pmatrix} 7 & 4 & 3 \\ 0 & 2 & 0 \\ 6 & 0 & 0 \\ 0 & 1 & 1 \\ 4 & 3 & 1 \end{pmatrix} \begin{matrix} T1 \\ T2 \\ T3 \\ T4 \\ T5 \end{matrix}$$

Banker's Algorithm

```

requestResources(Req[i], Available, Max, Alloc) {
    // Create and initialize the Need matrix
    Need = Max - Alloc

    if (Req[i] > Need[i])
        perror("Thread %d requested more than it needs\n", i)
    else if (Req[i] > Available)
        // can't fulfill request right now, try again later
        wait(thread[i])
    else {
        // update system state as if request was granted..
        Alloc[i] += Req[i]
        Need[i] -= Req[i]
        Available -= Req[i]

        if (isSafe(Available, Alloc, Need))
            grantRequest(thread[i])
        else {
            // restore system state
            Alloc[i] -= Req[i]
            Need[i] += Req[i]
            Available += Req[i]
            wait(thread[i])
        }
    }
}
  
```

Step 4: Run the safety algorithm

(In this case, yes, Thread 2's request is granted!)

$$P \quad D \quad S$$

$$Req[2] = \begin{pmatrix} 1 & 0 & 2 \end{pmatrix}$$

$$P \quad D \quad S$$

$$Available = \begin{pmatrix} 2 & 3 & 0 \end{pmatrix}$$

$$Alloc = \begin{pmatrix} 0 & 1 & 0 \\ 3 & 0 & 2 \\ 3 & 0 & 2 \\ 2 & 1 & 1 \\ 0 & 0 & 2 \end{pmatrix} \begin{matrix} T1 \\ T2 \\ T3 \\ T4 \\ T5 \end{matrix}$$

$$P \quad D \quad S$$

$$Need = \begin{pmatrix} 7 & 4 & 3 \\ 0 & 2 & 0 \\ 6 & 0 & 0 \\ 0 & 1 & 1 \\ 4 & 3 & 1 \end{pmatrix} \begin{matrix} T1 \\ T2 \\ T3 \\ T4 \\ T5 \end{matrix}$$

Evaluation of Banker's Algorithm

- ▶ Summary: fulfill resource requests such that deadlocks can't happen
- ▶ Pros
 - Can detect potential deadlocks for multi-instance resources
- ▶ Cons
 - Need a crystal ball: *Max* resource requirements must be known
 - Overhead costs (*run for every resource request*)
 - Resource-Request Algorithm
 - Due in large part to Safety Algorithm
 - Time and space complexities? $O(n^2)$ where $n = \max(\text{number of threads}, \text{number of resource types})$

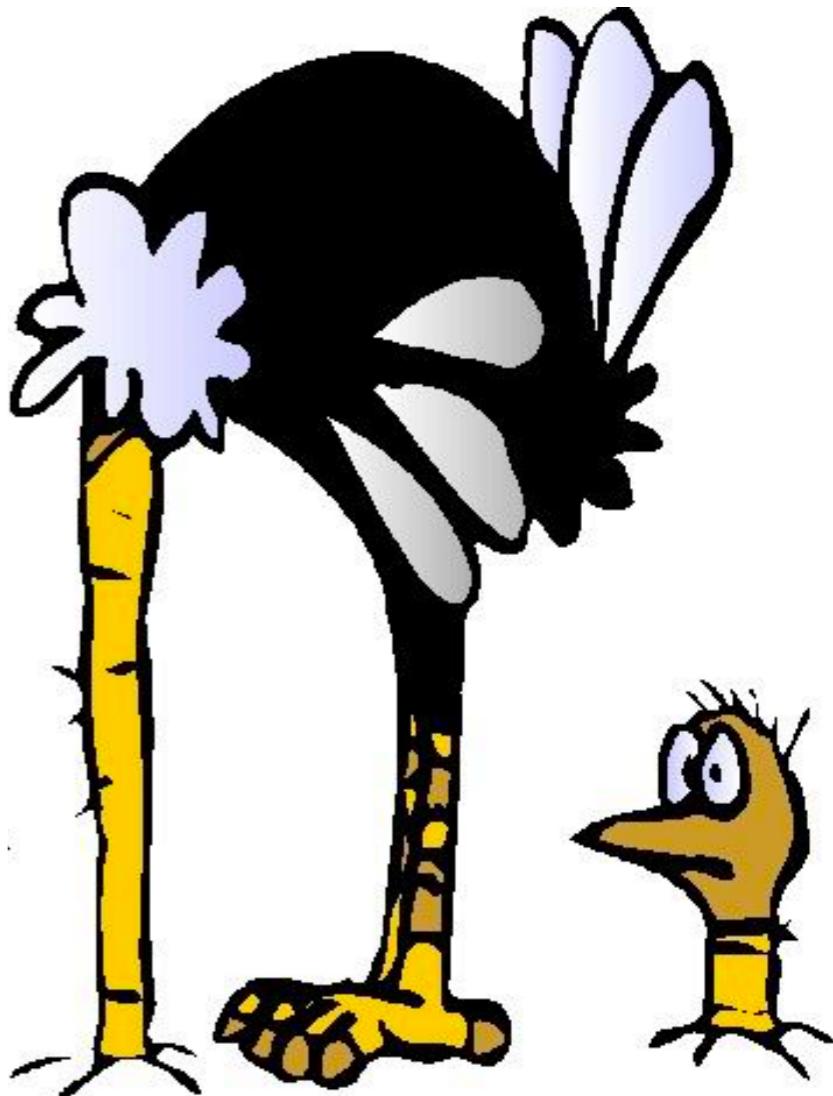


Topics for Today...

- ▶ Deadlock Conditions
- ▶ Ways to deal with deadlocks:
 - Deadlock Detection
 - Single Resource Allocation Graph
 - Multi-Resource Allocation Graph
 - Deadlock Avoidance
 - Banker's Algorithm
 - The Ostrich Approach

Ostrich Approach

- ▶ Key insight: Sometimes more cost-effective to let just deadlocks happen
- ▶ *Ostrich Approach*
 - OS sticks head in sand
 - Leave user to deal with deadlocks
 - Kill the offending thread(s) or process(es) or reboot
- ▶ In Linux, user can detect deadlocks.
- ▶ In Java, you can look at a threaddump: **ctrl+**



Evaluating Deadlock Handling Approaches

► Deadlock Avoidance

- Too much overhead per resource request

► Deadlock Detection

- Run infrequently: could miss deadlock state
- Run too often: high overhead
 - Could still miss deadlock state
- Rollbacks are also costly

In Conclusion...

- ▶ Deadlock \neq Starvation
 - Starvation: waiting indefinitely
 - Deadlock: circular waiting for resources
- ▶ Know what "*necessary and sufficient*" means
- ▶ Know the four conditions for a deadlock
- ▶ Know the different ways to deal with deadlocks

In Conclusion...

- ▶ Many ways to deal with deadlocks, most are impractical
 - Understand implications of prevention schemes
 - Know Banker's Algorithm
 - Know RAGs
 - Single Resource vs Multi-Resource
 - Existence of Cycles and Knots, and what they tell us

▶ Next Topic: Memory Management

