

CS 475

Operating Systems



Department of Mathematics
and Computer Science

Lecture 8
Memory Management

Motivation: Virtualizing Memory

- ▶ Recall the goal of the OS: Timesharing
 - A program must be loaded into main memory before it can be run.
 - Reality: main memory is limited. How to run multiple processes?

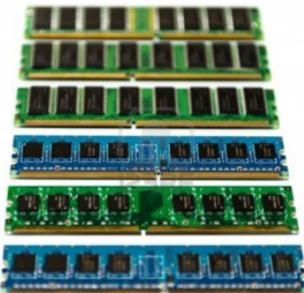
- ▶ Policy: Want *Controlled Multiplexing* of main memory
 - Processes must be able to coexist securely
 - Each process thinks it has access to the entire main memory.



Memory Management Goals

- ▶ Need a new OS subsystem to deal with memory management
 - **Memory Management Unit (MMU)**

- ▶ An MMU must enable the following goals:
 - **Process Relocation:** OS can take process' address space out of memory and give it back later (*possibly* in a different physical location)
 - **Protection of Address Spaces:** Don't want different processes to have access to each other's memory.
 - **Sharing:** There are some cases when processes *do* want access to each other's memory (e.g., share code across multiple processes)

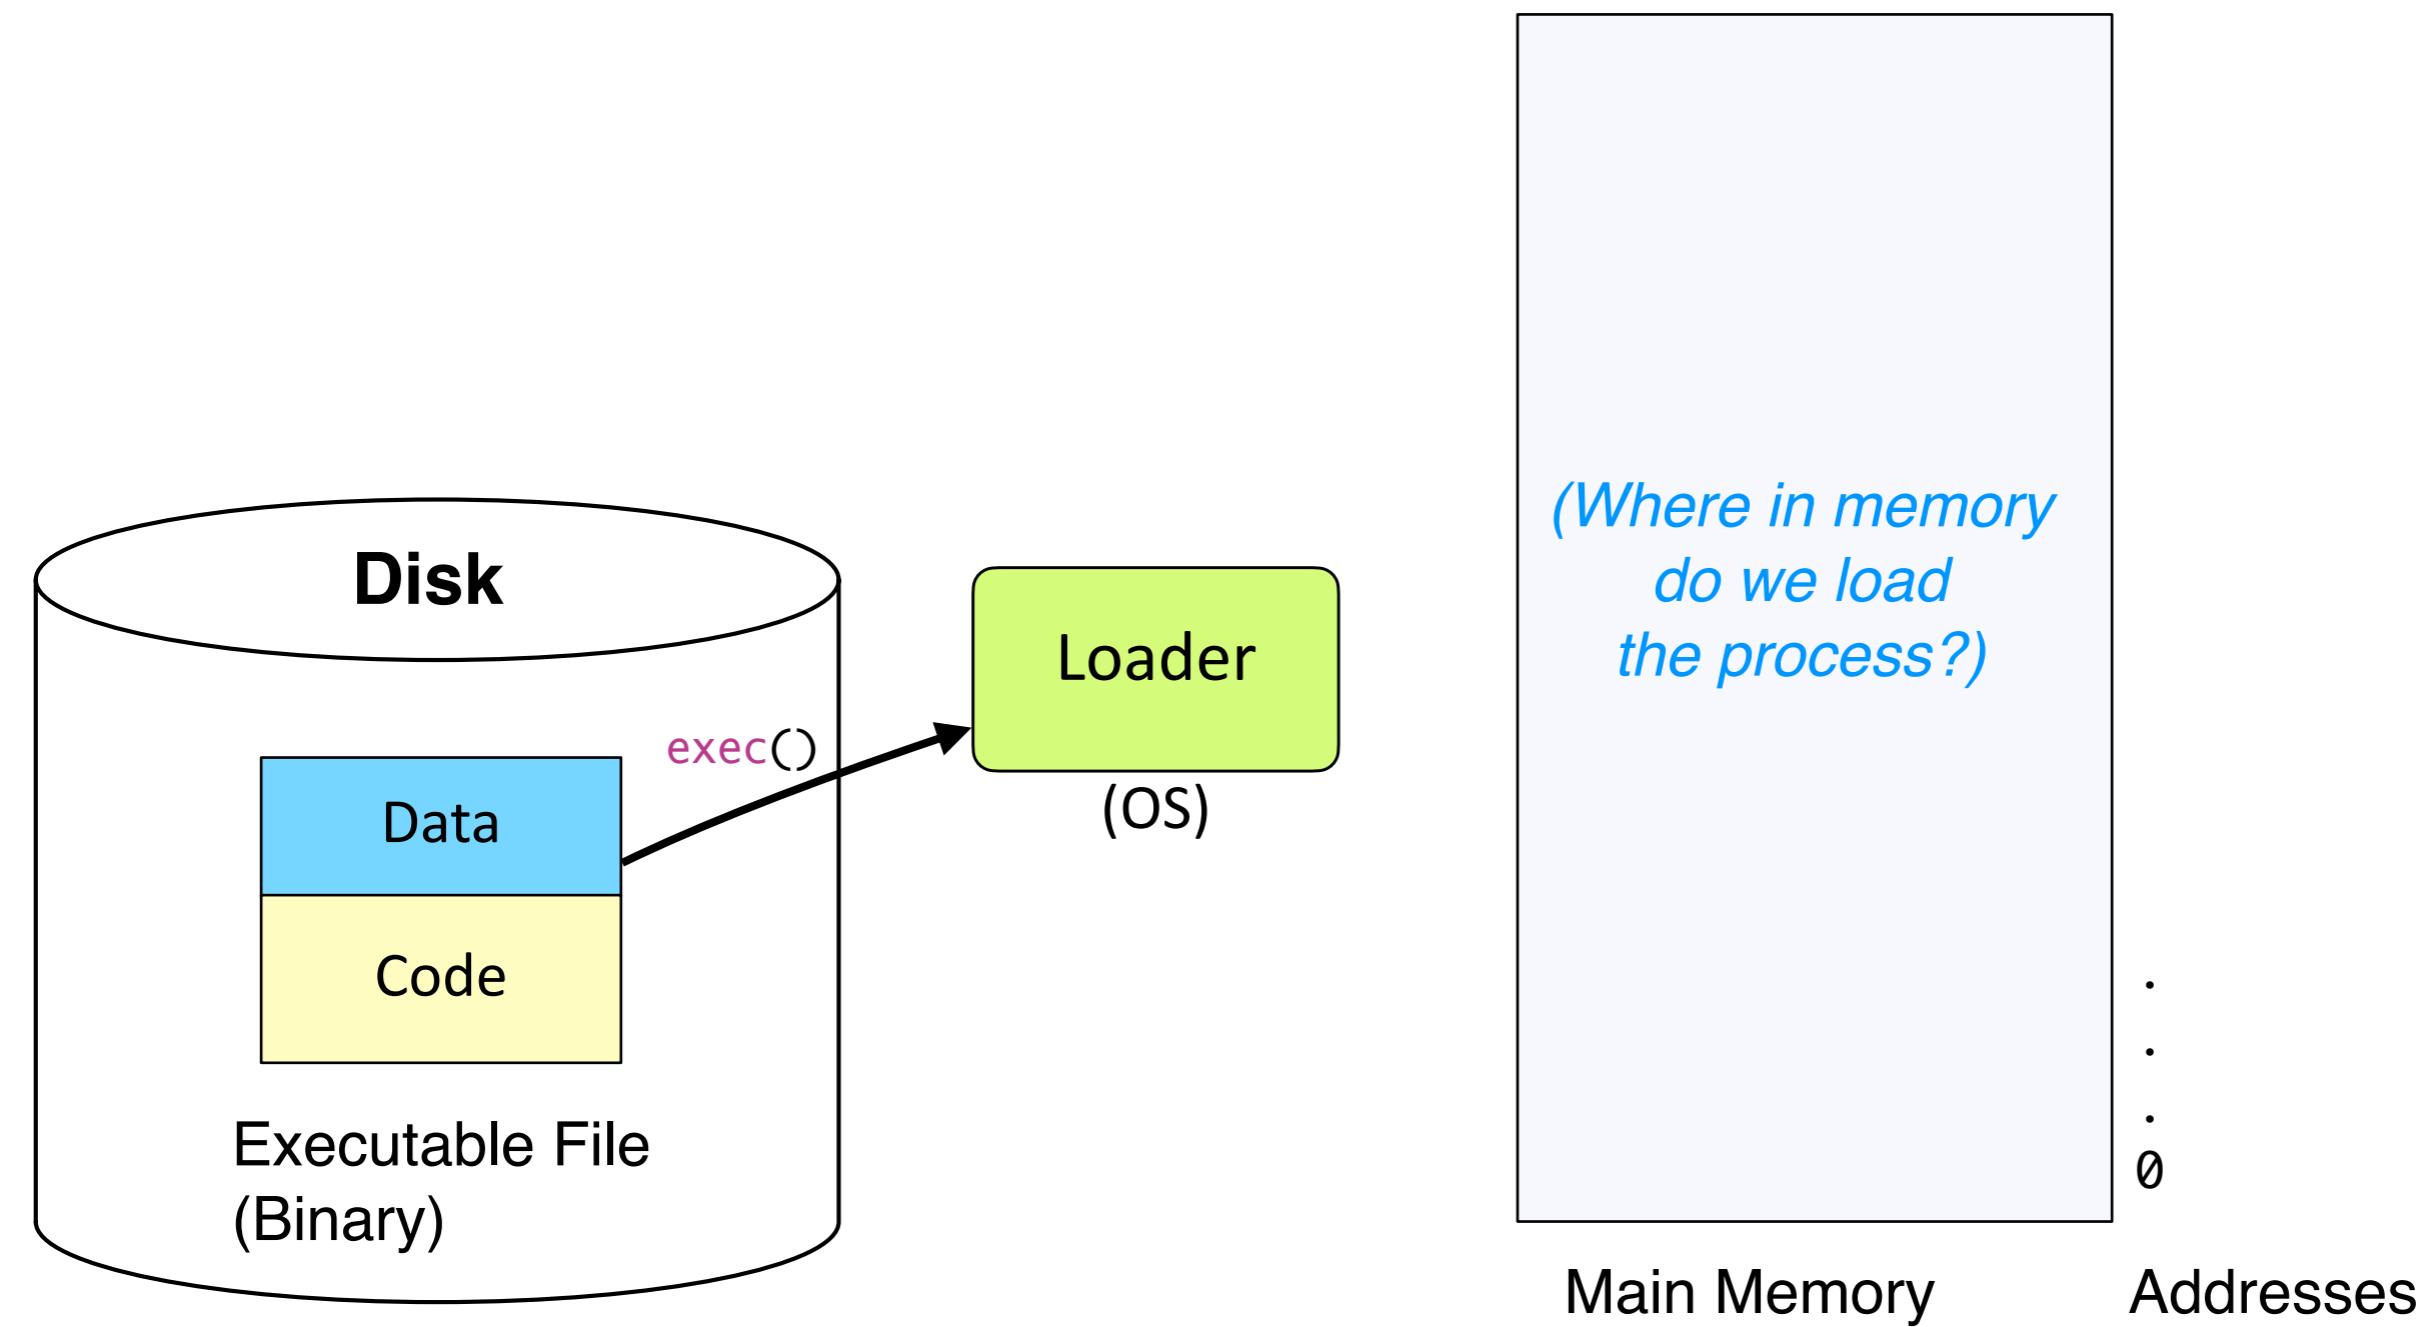


Goals for This Lecture...

- ▶ Motivation and Goals
- ▶ Towards Virtual Addressing
 - Process Relocation
 - Address Translation
- ▶ Partitioning
- ▶ Segmentation
- ▶ Paging

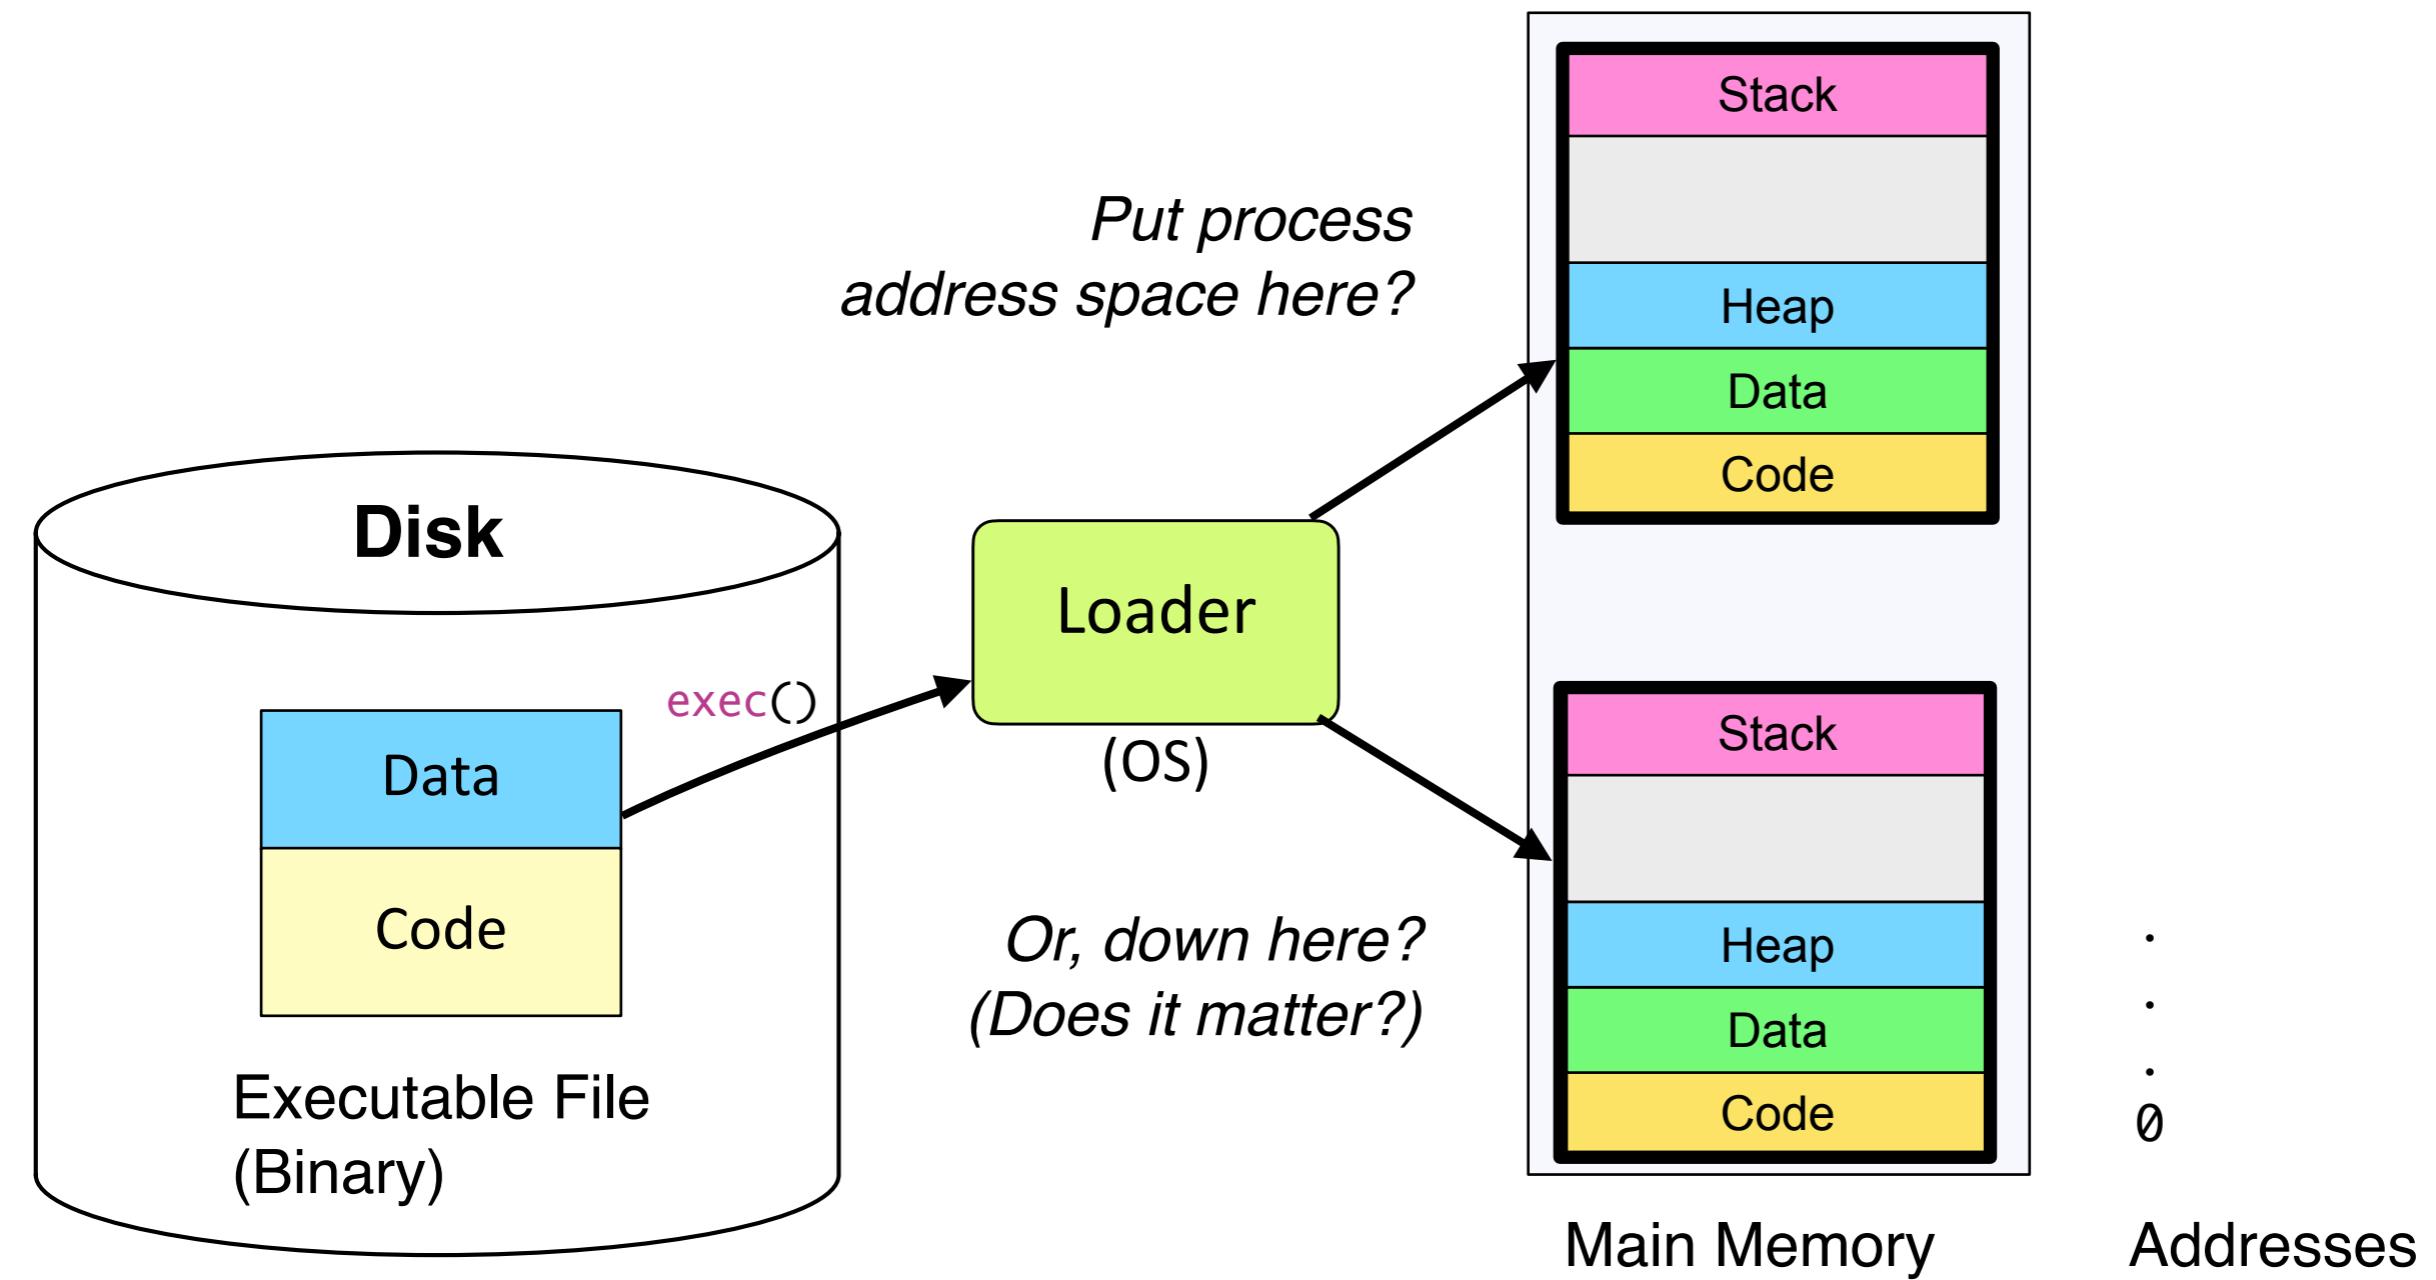
The Loader

- ▶ The `exec("/path/to/file")` system call invokes the OS's *Program Loader* to put an executable file into memory for execution.



The Loader

- ▶ The `exec("/path/to/file")` system call invokes the OS's *Program Loader* to put an executable file into memory for execution.



Consider One of the Compiler's Jobs (1)

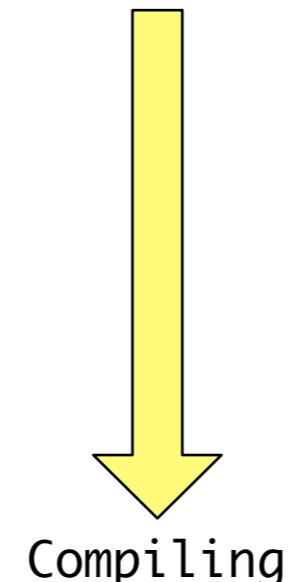
- ▶ Assume each instruction and data takes up 4 bytes of memory.

C code (To be compiled.)

```
int x = 0; //global

void main() {
    x++;
    bar(x);
}

void bar(int x) {
    do other stuff
}
```



Assembly code: To be loaded in Memory to run.

Address in Memory	Data or Instructions
0	main:
4	movl 0(??), %eax ; de-reference x
8	addl \$1, %eax ; increment
16	movl %eax, 0(??) ; write back to x
20	jmp ?? ; call bar(x)
24	.
.	.
.	.

(Still compiling)

- ▶ How does the compiler know what address (??) x will be placed?
 - What address (??) does bar() start at?

Consider One of the Compiler's Jobs (2)

Took a while to resolve, but *now* we know where **bar()** and **x** are located in memory

Ah! There's the start of **bar()**!
It's at address 256.

Ah! There's **x**!
It's at address 1024.

Address	Data or Instruction
0	main:
4	movl 0(<u>????</u>), %eax ;de-reference x
8	addl \$1, %eax ;increment
16	movl %eax, 0(<u>????</u>) ;write back
20	jmp <u>???</u> ;call bar(x)
.	.
.	.
256	bar:
.	.
.	.
1020	.data
1024	x DD 0 ; x in data segment

Address Binding

- ▶ **Address Binding:** Yay variables and functions can now be "*bound*" to addresses in memory.

▶ Key Insight

- Address binding is easy when compiler can assume process address start from **0**
- All other addresses are relative to **0**

Importance of zero as the starting (base) address!

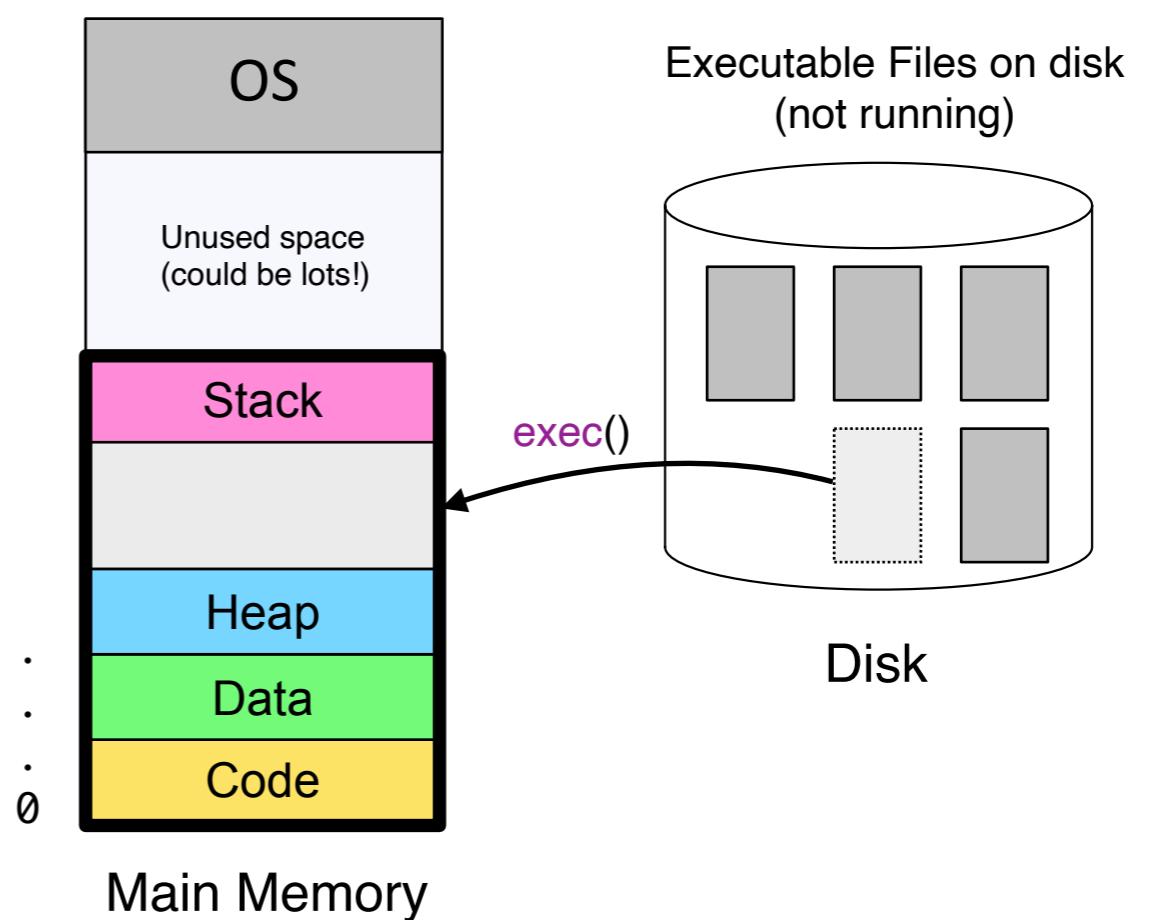
Assembly code: Loaded in Main Memory	
Addr	Data or Instruction
0	main:
4	movl 0(1024), %eax ;de-reference x
8	addl \$1, %eax ;increment
16	movl %eax, 0(1024) ;write back
20	jmp 256 ;call bar(x)
·	
256	bar:
·	
·	
1020	.data
1024	x DD 0 ; x declared as global here

Absolute Loader

Absolute Loading

1. Pin the OS to highest memory location (far far away from user code)
2. **Always** load user process' address space starting from address **0**
(0 is called the "*base address*")

- Only one user process can be in memory (e.g., DOS)
- How does `contextSwitch()` work?
 - Well, *slowly!*
 - (Two disk accesses)*



Absolute Loading (Cont.)

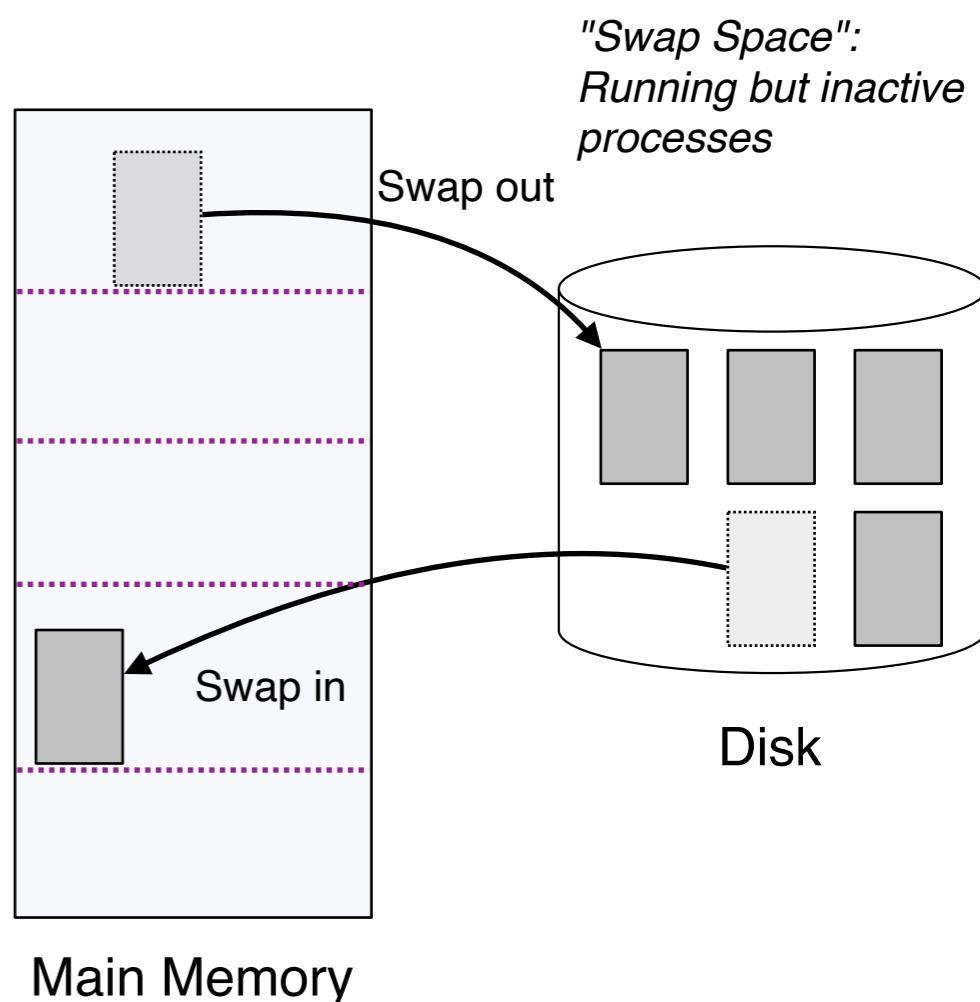
- ▶ Used for batch-processing systems.
- ▶ Compiler binds variables and functions to physical addresses relative to the base address 0
- ▶ Pros:
 - Only need to protect OS code from the user process
 - Raise a protection ("segmentation") fault and terminate user process on encroachment!
- ▶ Cons:
 - Lots of wasted space in main memory (gap between process and OS)
 - High overhead of context switching (disk involved)
 - Can't timeshare effectively if switching overhead is high

Goals for This Lecture...

- ▶ Motivation and Goals
- ▶ Towards Virtual Addressing
 - Process Relocation
 - Address Translation
- ▶ Partitioning
- ▶ Segmentation
- ▶ Paging

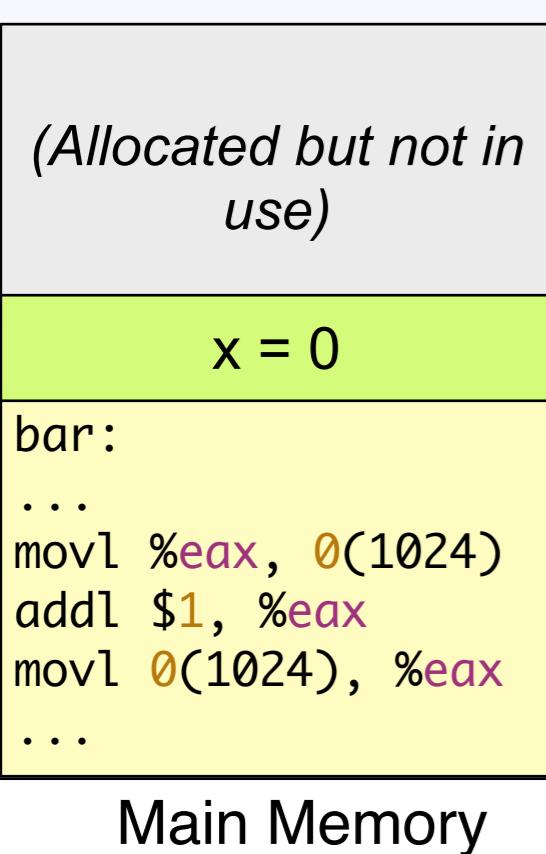
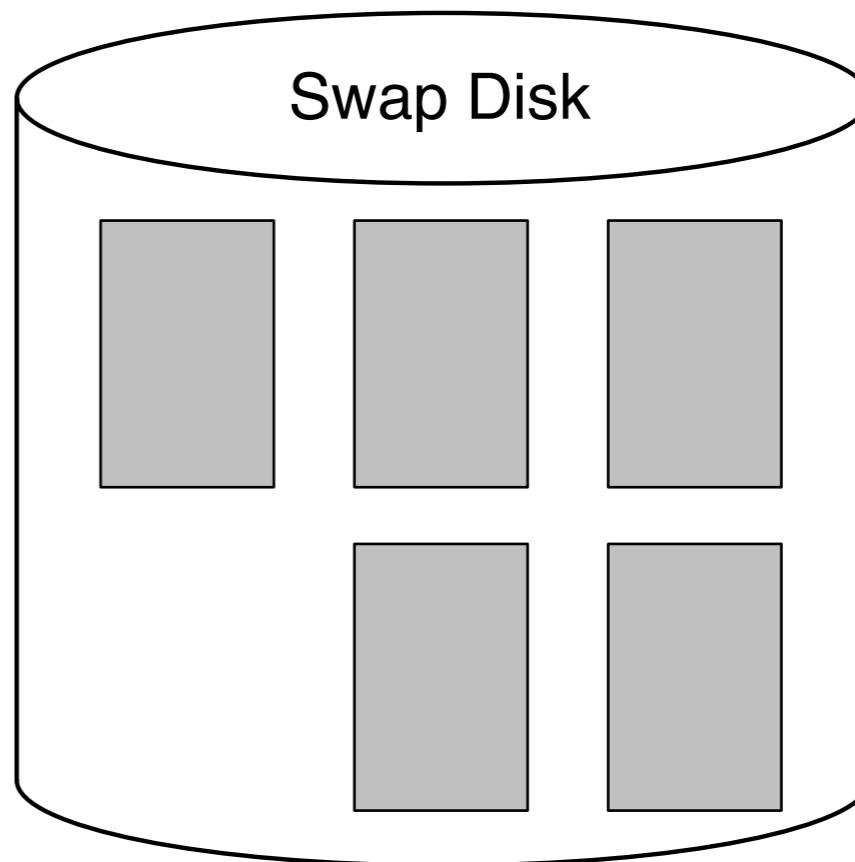
What If We Want Timesharing?

- ▶ But achieve high CPU utilization, we need to timeshare!
 - Memory is shared by *multiple* active process address spaces.
 - Split the main memory up into *partitions* where processes can live.
- But we still need space to put all the inactive processes: "*Swap Space*"
- Reserve some space on disk.
- Worst-case context switch?
 - Swap out to disk → Swap in from disk
 - Best case? Both already in memory



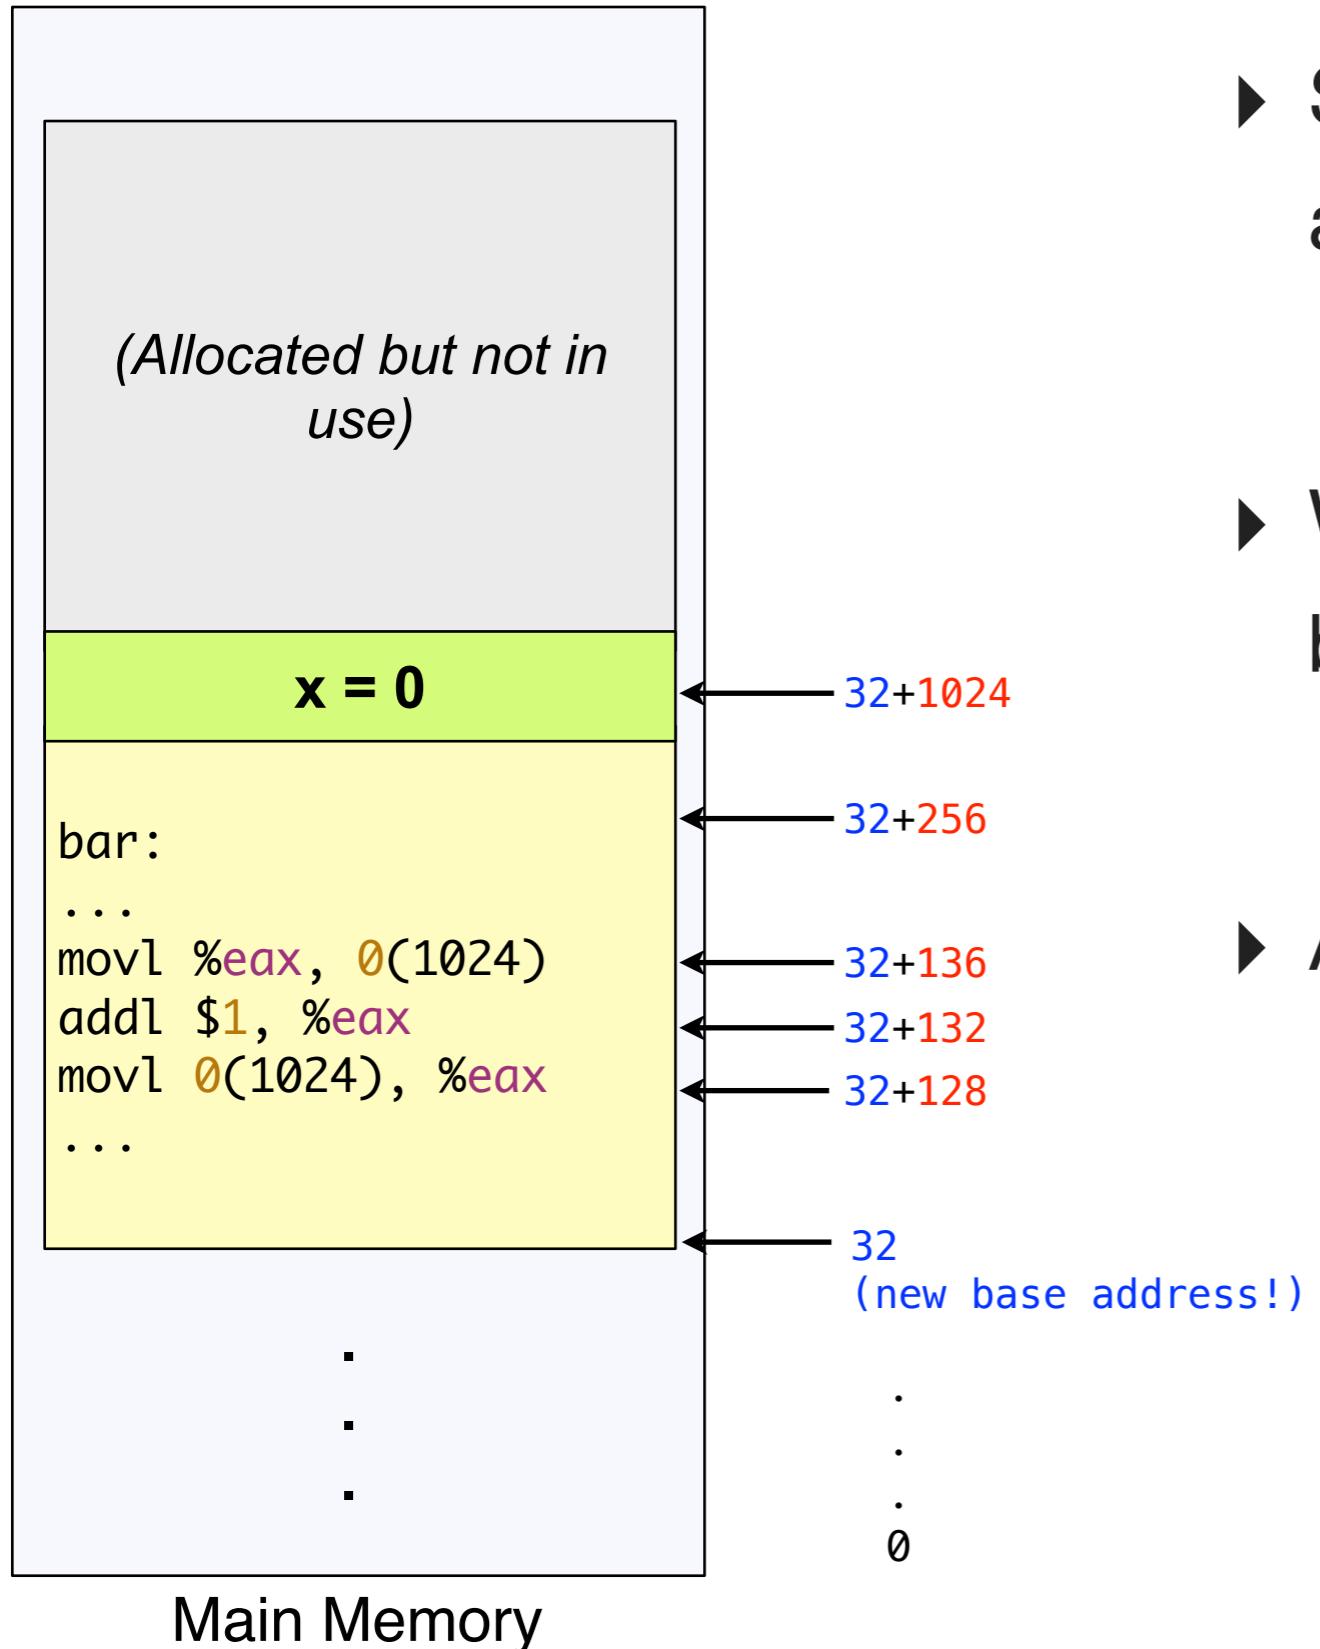
Process Relocation

- Let processes swap into a potentially to a **different** location (i.e., "relocate")



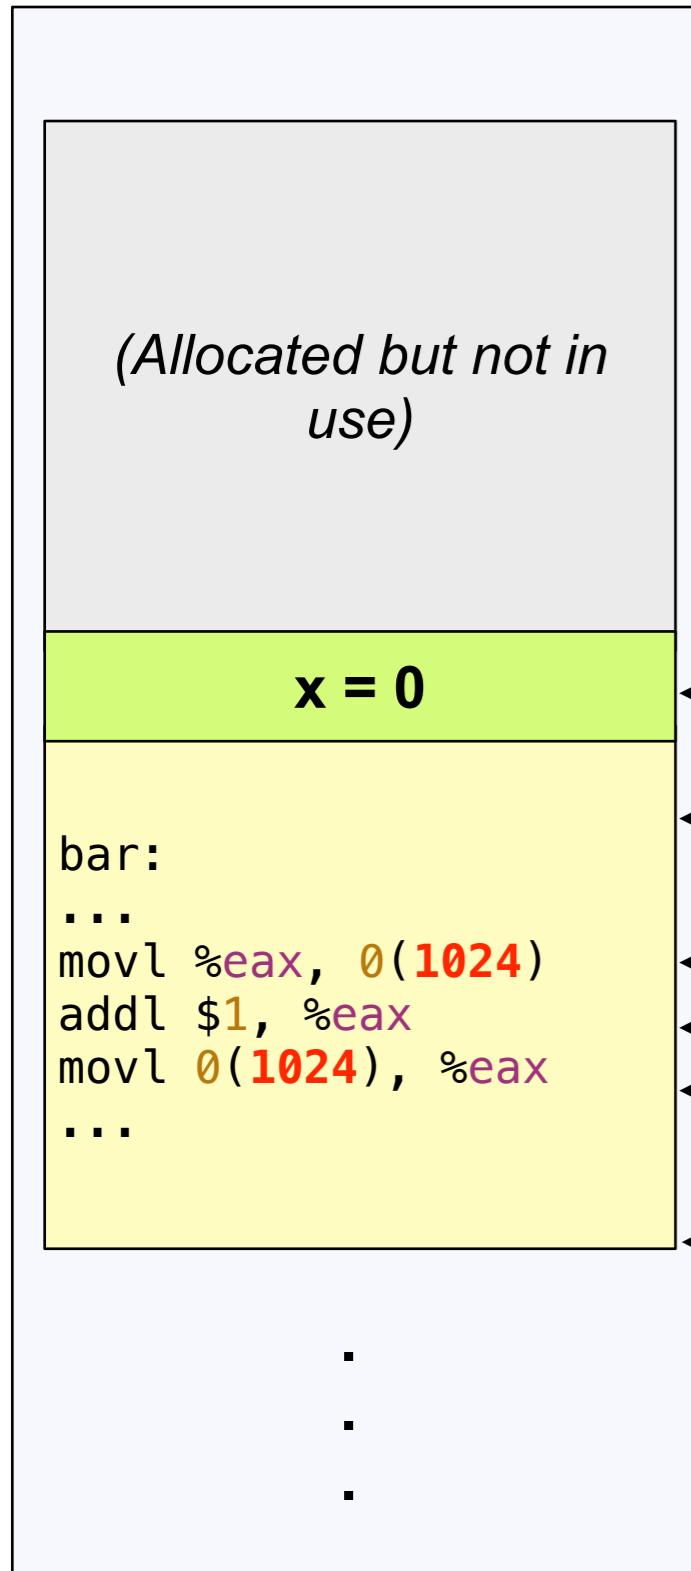
- But why is relocation a desirable policy to achieve higher CPU %? (The alternative is to always swap back into same location in RAM)*

How to Relocate? Need to Rebind Addresses!



- ▶ Suppose this process began at address **0**, but gets swapped out.
- ▶ When it returns, it gets swapped back in at address **32**.
- ▶ After swapping back into memory,
 - Execution will raise a segmentation fault! (*Why?*)

How to Relocate? Need to Rebind Addresses!



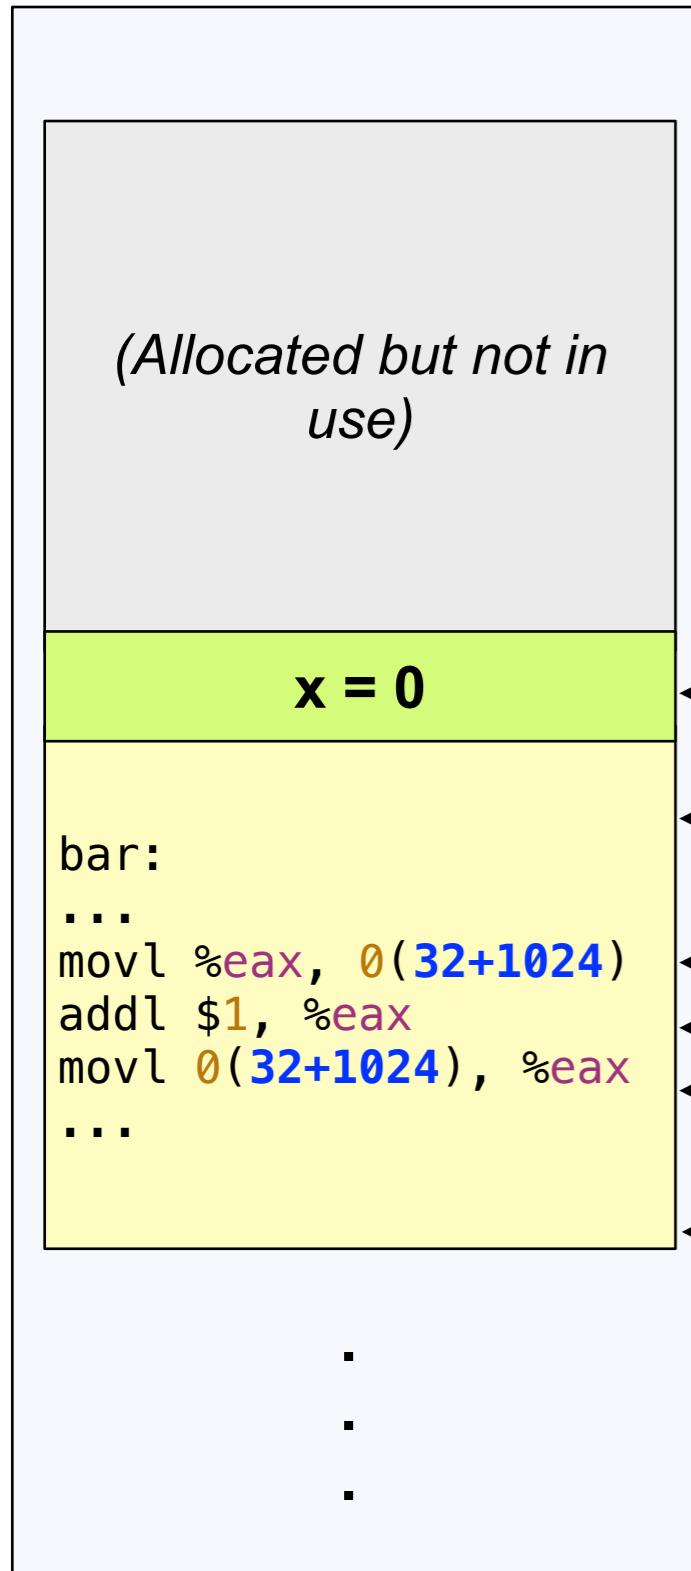
Main Memory

► Why does it segfault? All the compiler-bound addresses are now **wrong**!

- All address references in the code now off by an offset of: **-32**

► How do we fix the code's addresses?

How to Relocate? Need to Rebind Addresses!



Main Memory

- ▶ *Solution: "Address Translation"*

- ▶ When loading a process to run, OS has to **re-bind** the addresses!

- Add the **New Base Address (32)** to all the old address references in the code!
- Oh no we have to retranslate (recompile) all references before swapping a process back in!

Static Relocation (Load-Time Binding)

Static Relocation

1. Split physical memory into partitions

Compiler-generated addresses aren't fixed (or "real") anymore

2. Put OS in one of these partitions

3. Compiler ***still binds*** user code's addresses relative to address 0

4. But at *Load Time*, the OS:

(a) Assigns the process to an appropriate partition in memory

(b) Loader binds the compiler-generated *addresses* to the OS-assigned base address to reflect assigned location

These addresses are real!

Loading is really slow (Need to translate every address in the program first!)

Evaluation of Static Relocation

▶ Pros:

- Physical memory can be multiplexed across different processes
- Totally software-based
 - Translate all addresses directly within the binary executable file at load time
- No changes to compiler technology
 - Compiler still thinks that all user program's addresses start at 0 unless told otherwise

▶ Cons:

- No protection of address spaces
 - A process could still refer to an illegal address (un)intentionally.
- Very poor performance
 - Swap-in times suffer since the *loader* must re-bind every address reference to a potentially new **base** address!

Goals for This Lecture...

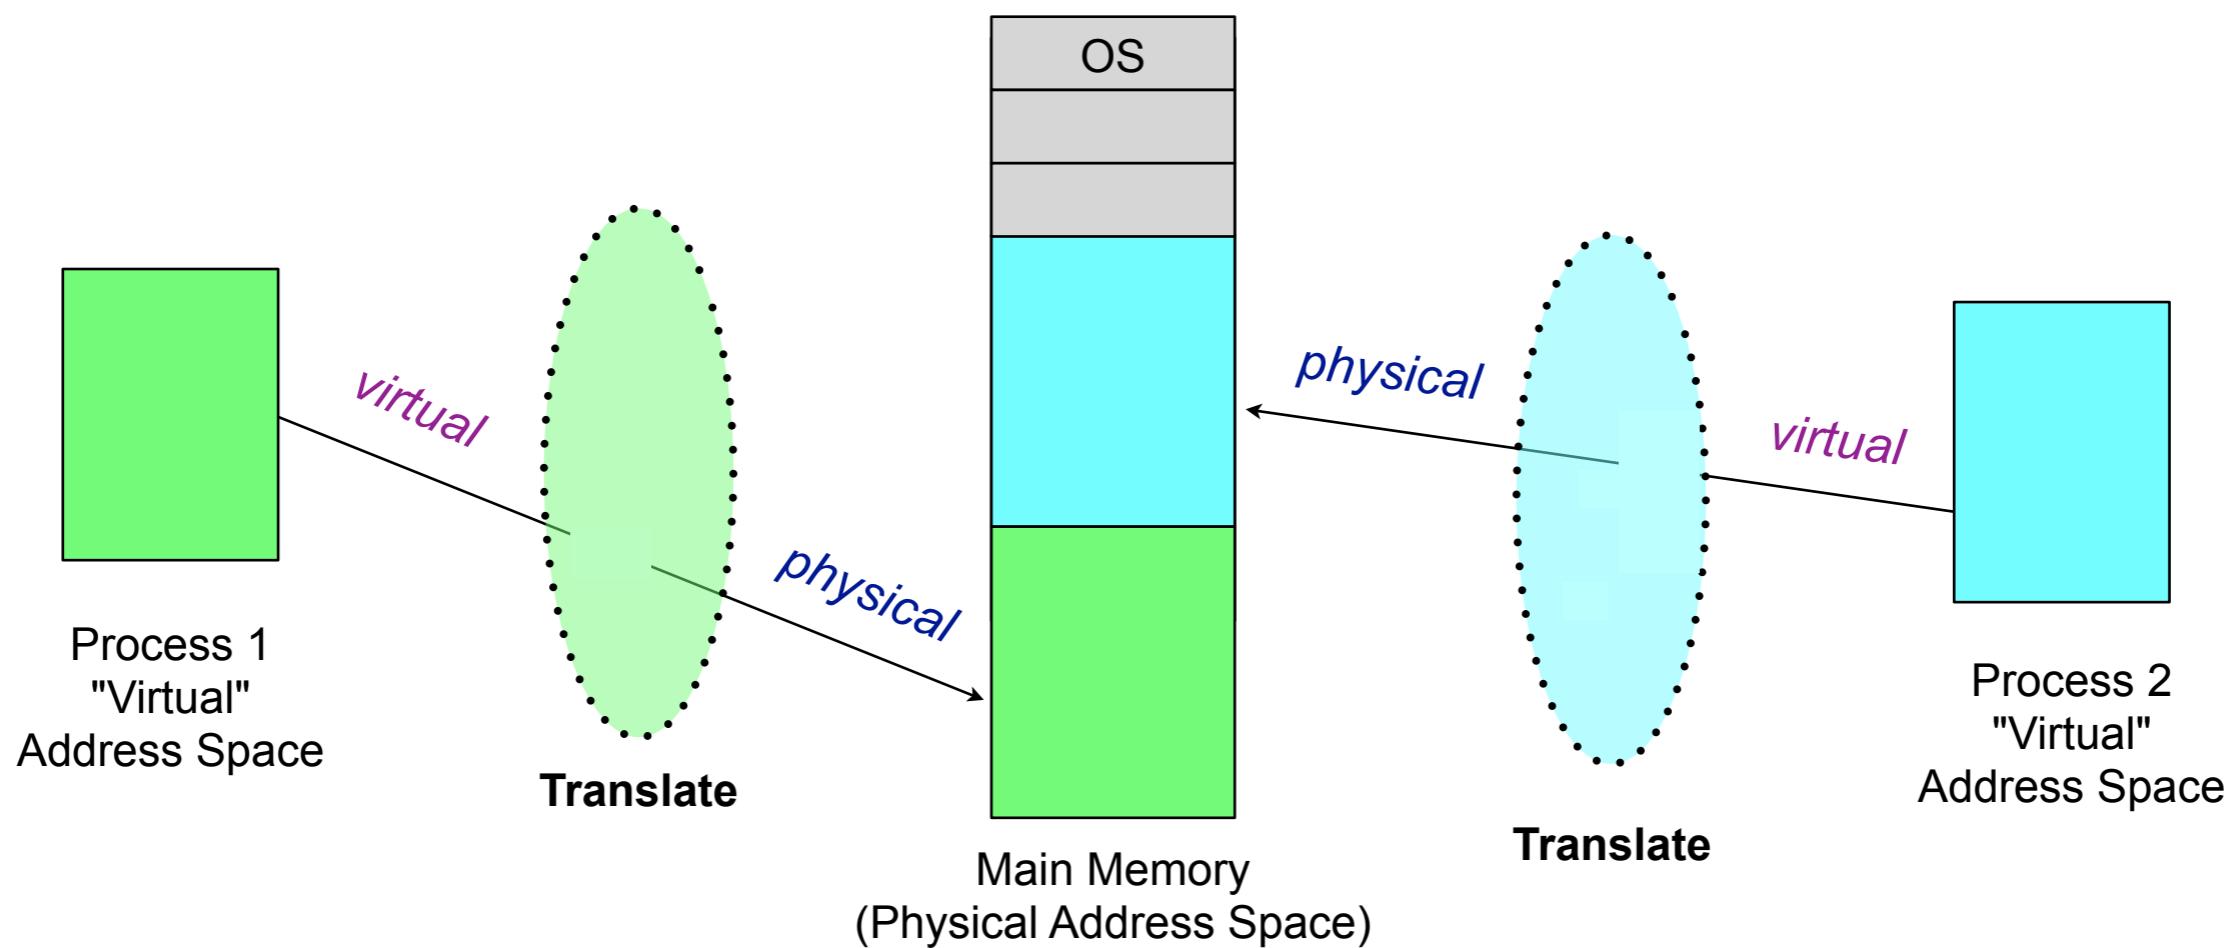
- ▶ Motivation and Goals
- ▶ Towards Virtual Addressing
 - Process Relocation
 - Address Translation
- ▶ Partitioning
- ▶ Segmentation
- ▶ Paging

Virtual Addressing

- ▶ *Virtual Addressing* is considered one of the biggest breakthroughs in computer science. (First appeared in Multics)
- ▶ Ren and Stimpy go to the post office.
 - Post office has many physical mailboxes to rent out.
 - Adam wants PO Box 1, 2, and 3.
 - America wants PO Box 1, 2, 3, and 4.
 - Can the post office satisfy both customers?

Virtual Addressing & Address Translation

- ▶ **Virtual Address:** an address that a process uses to access its own address space (Processes know nothing about real addresses!)
- ▶ **Physical Address:** the address in physical memory
 - Address translation done by the loader (for now...)



Dynamic Relocation (Runtime Binding)

Dynamic Relocation (Hardware-Supported)

1. Split physical memory up into partitions
2. Put the OS in one of these partitions.
3. Compiler *still binds* user code's addresses relative to base 0
4. Translate each virtual address to a physical address as *the process runs.*

- ▶ Implementing this mechanism will be our focus
 - *How does address translation at runtime (vs. at load-time) help us multiprogram?*

Goals for This Lecture...

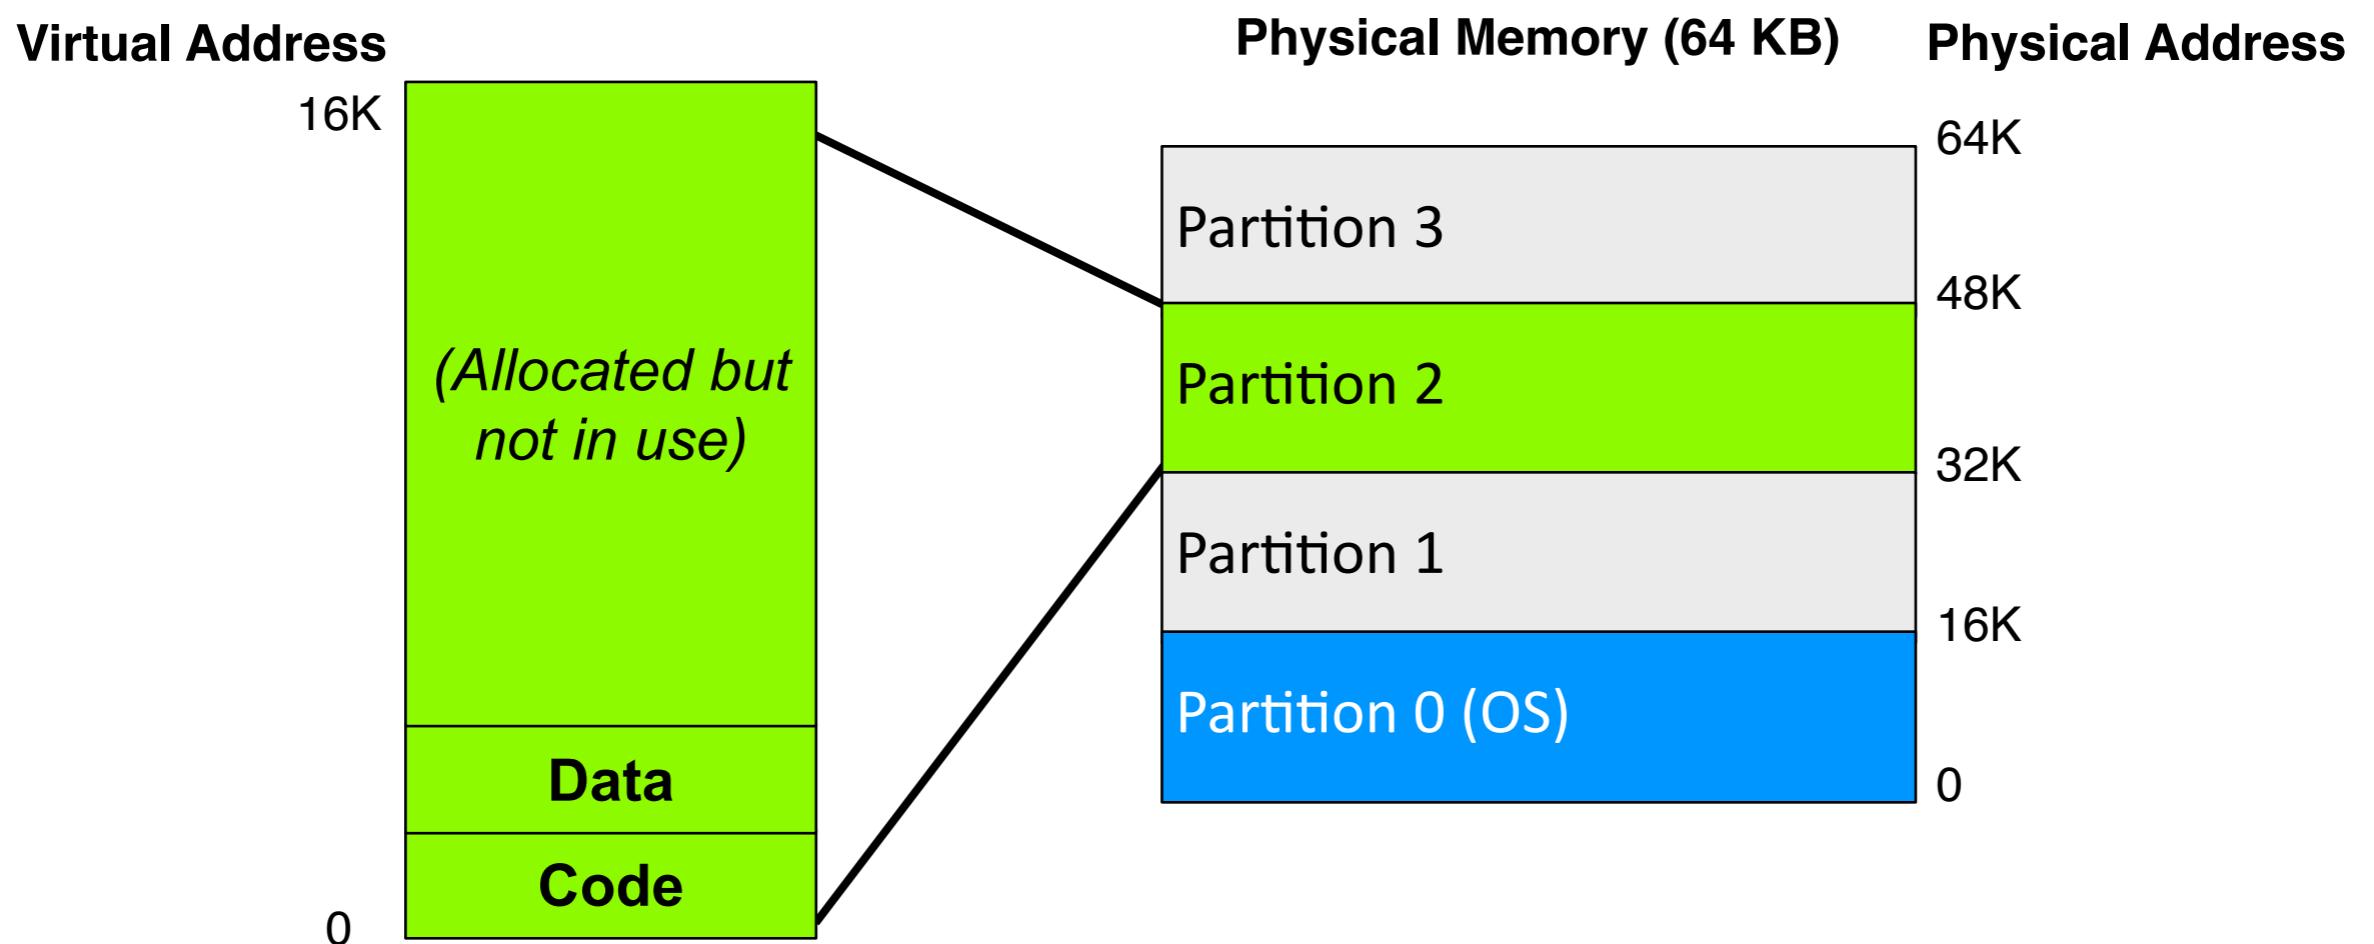
- ▶ Motivation and Goals
- ▶ Towards Virtual Addressing
- ▶ **Memory Partitioning**
 - Fixed Partitioning
 - Dynamic Partitioning (Base and Bounds)
- ▶ Segmentation
- ▶ Paging

Some Assumptions

- ▶ Let's make some simplifying assumptions:
 1. Size of process address space \leq available physical memory
 2. The address space is monolithic and must be placed *contiguously* in physical memory.
 - (This *will* change later...)
 3. Specifically for this lecture:
 - 64 KB available physical memory
 - Process' address spaces are allowed a maximum of 16 KB
 - Therefore, $64\text{ KB} / 16\text{ KB} = 4$ processes can be loaded in memory simultaneously

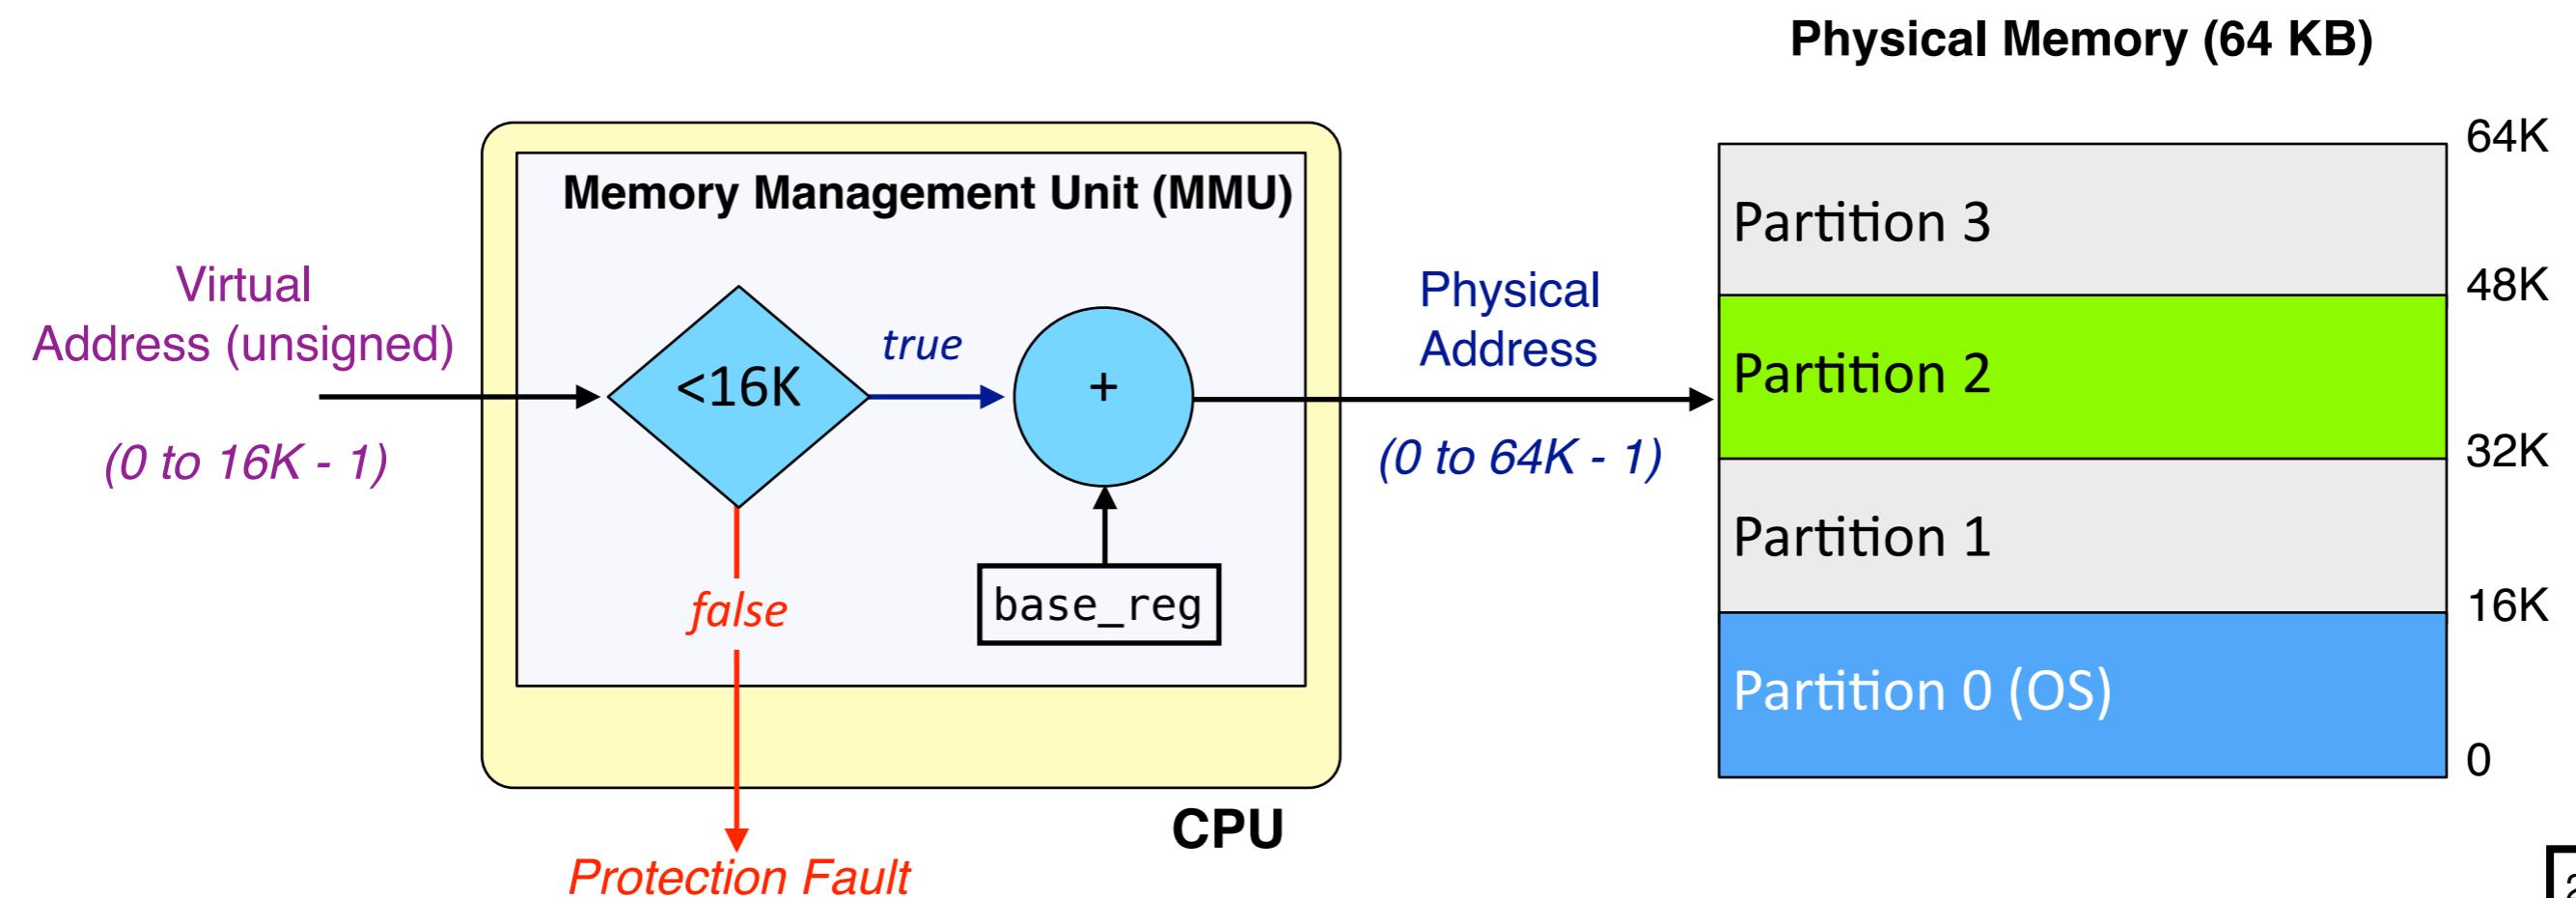
Fixed Partitioning

- ▶ *Fixed Partitioning*: The original dynamic-relocation technique
 - Break physical memory into equal-sized partitions
- ▶ A process' entire address space is loaded into a partition



Implementing Fixed Partitioning

- ▶ Address translation (base address of process stored in process' PCB)
 - $\text{physical addr} = \text{virtual addr} + \text{base addr}$
- ▶ Memory protection is now possible!
 - If $\text{virtual addr} \geq 16K$, generate a segmentation fault and terminate
- ▶ Updated hardware (MMU):



Implementing Fixed Partitioning (cont.)

► How to context switch now?

- Store a process' **base address** in its PCB
 - Holds the *base address* of the partition to which a process is assigned
- Add a **base register** to the CPU hardware
 - Holds the *base address* of the currently running process
- A chance that the scheduled process was out on disk (or was new)
 - Loader looks for a free partition and loads the process from disk
 - Update the *base address* in the process PCB
 - Update the **base register** on the CPU
 - (If **no empty partitions** exist, then OS first swaps out another process)

Evaluation of Fixed Partitioning

▶ Pros:

- Enables code relocation and basic memory protection!
- Address translation is fast and with minimal hardware changes
- Switching doesn't slow badly (better than static relocation)
- Partitions are of equal size: Don't need a complicated algorithm to find a "gap" big enough to fit the process' entire address space.

Evaluation of Fixed Partitioning

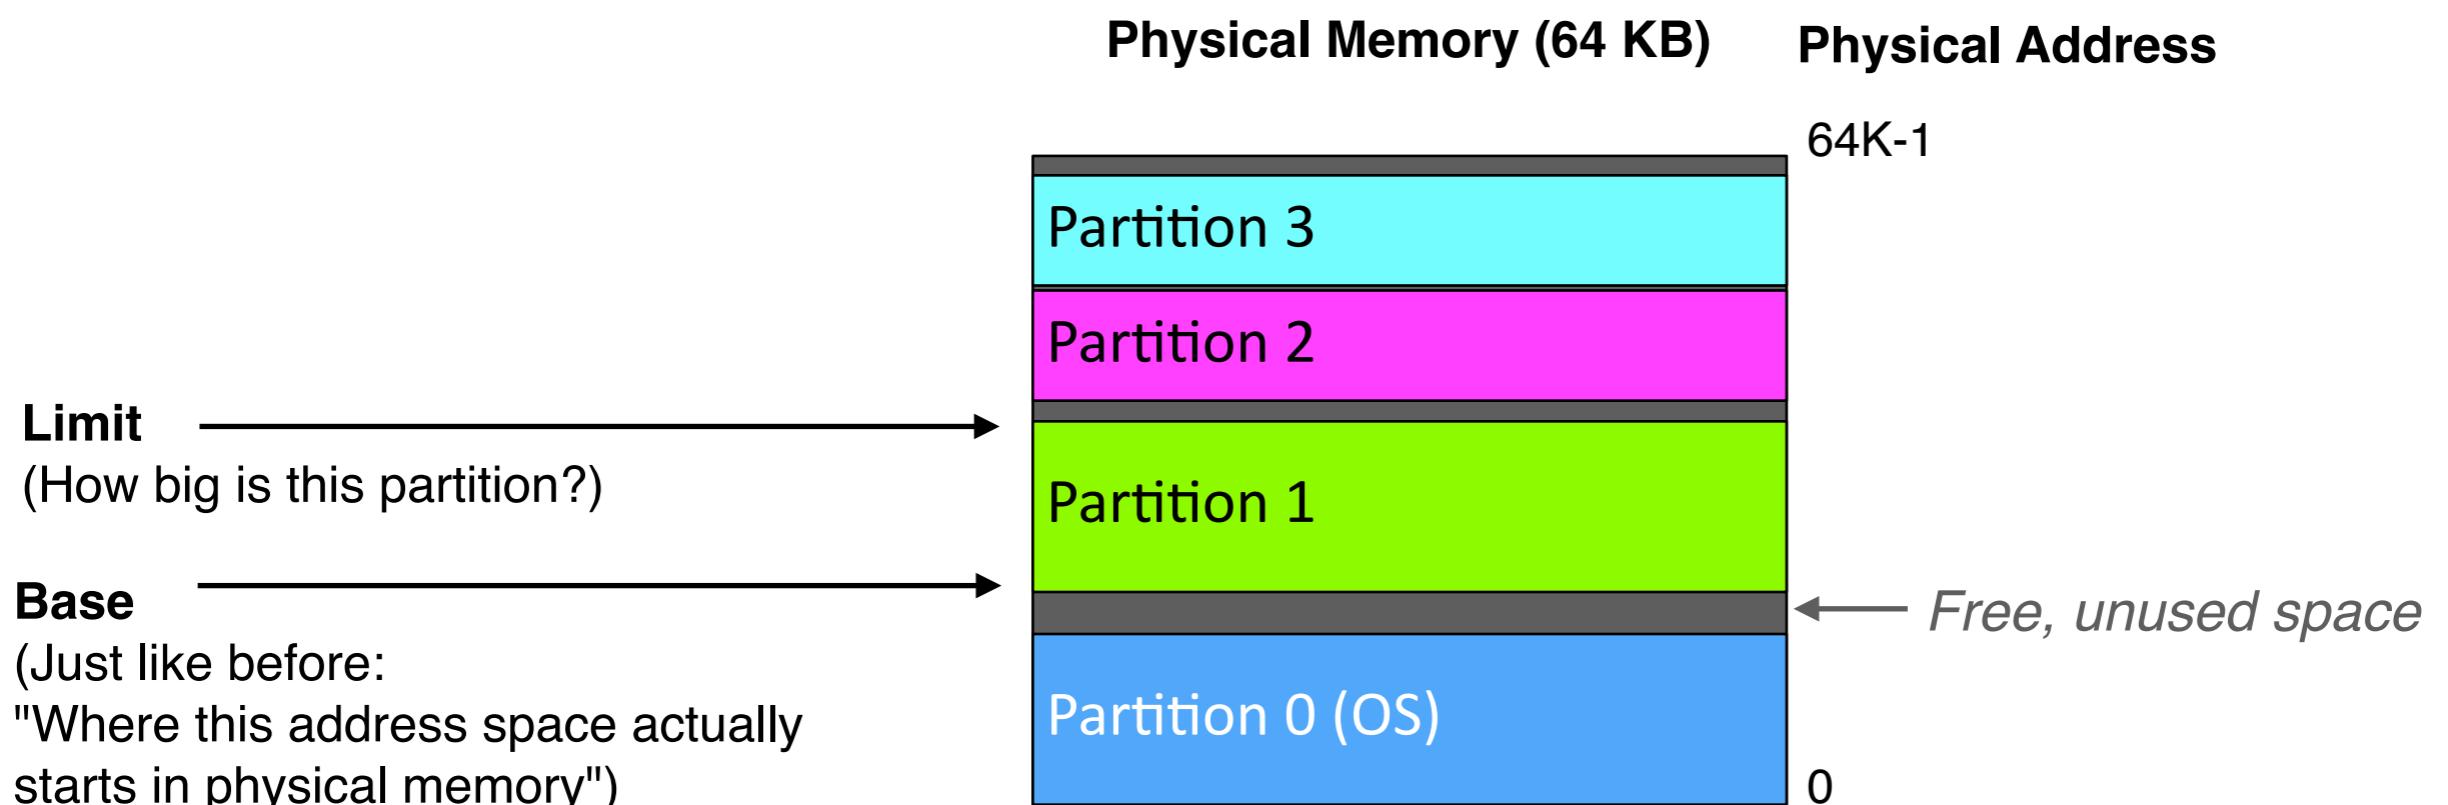
▶ Cons:

- The fixed number of partitions set by the OS **limits** the degree of multiprogramming
 - Memory partitions are 1-size-fits-all, but processes vary wildly in actual memory usage
 - Processes *can't grow* beyond the OS-defined partition size
-
- ## ▶ Susceptible to "*internal fragmentation*"
- That is, unused space *within* memory partitions

Variable Partitioning (Base and Bounds)

▶ Variable Partitioning

- Seeks to fix internal fragmentation and to remove limit on number of partitions OS can have.
- When a process gets brought into memory, it gets a partition of exactly the right size.

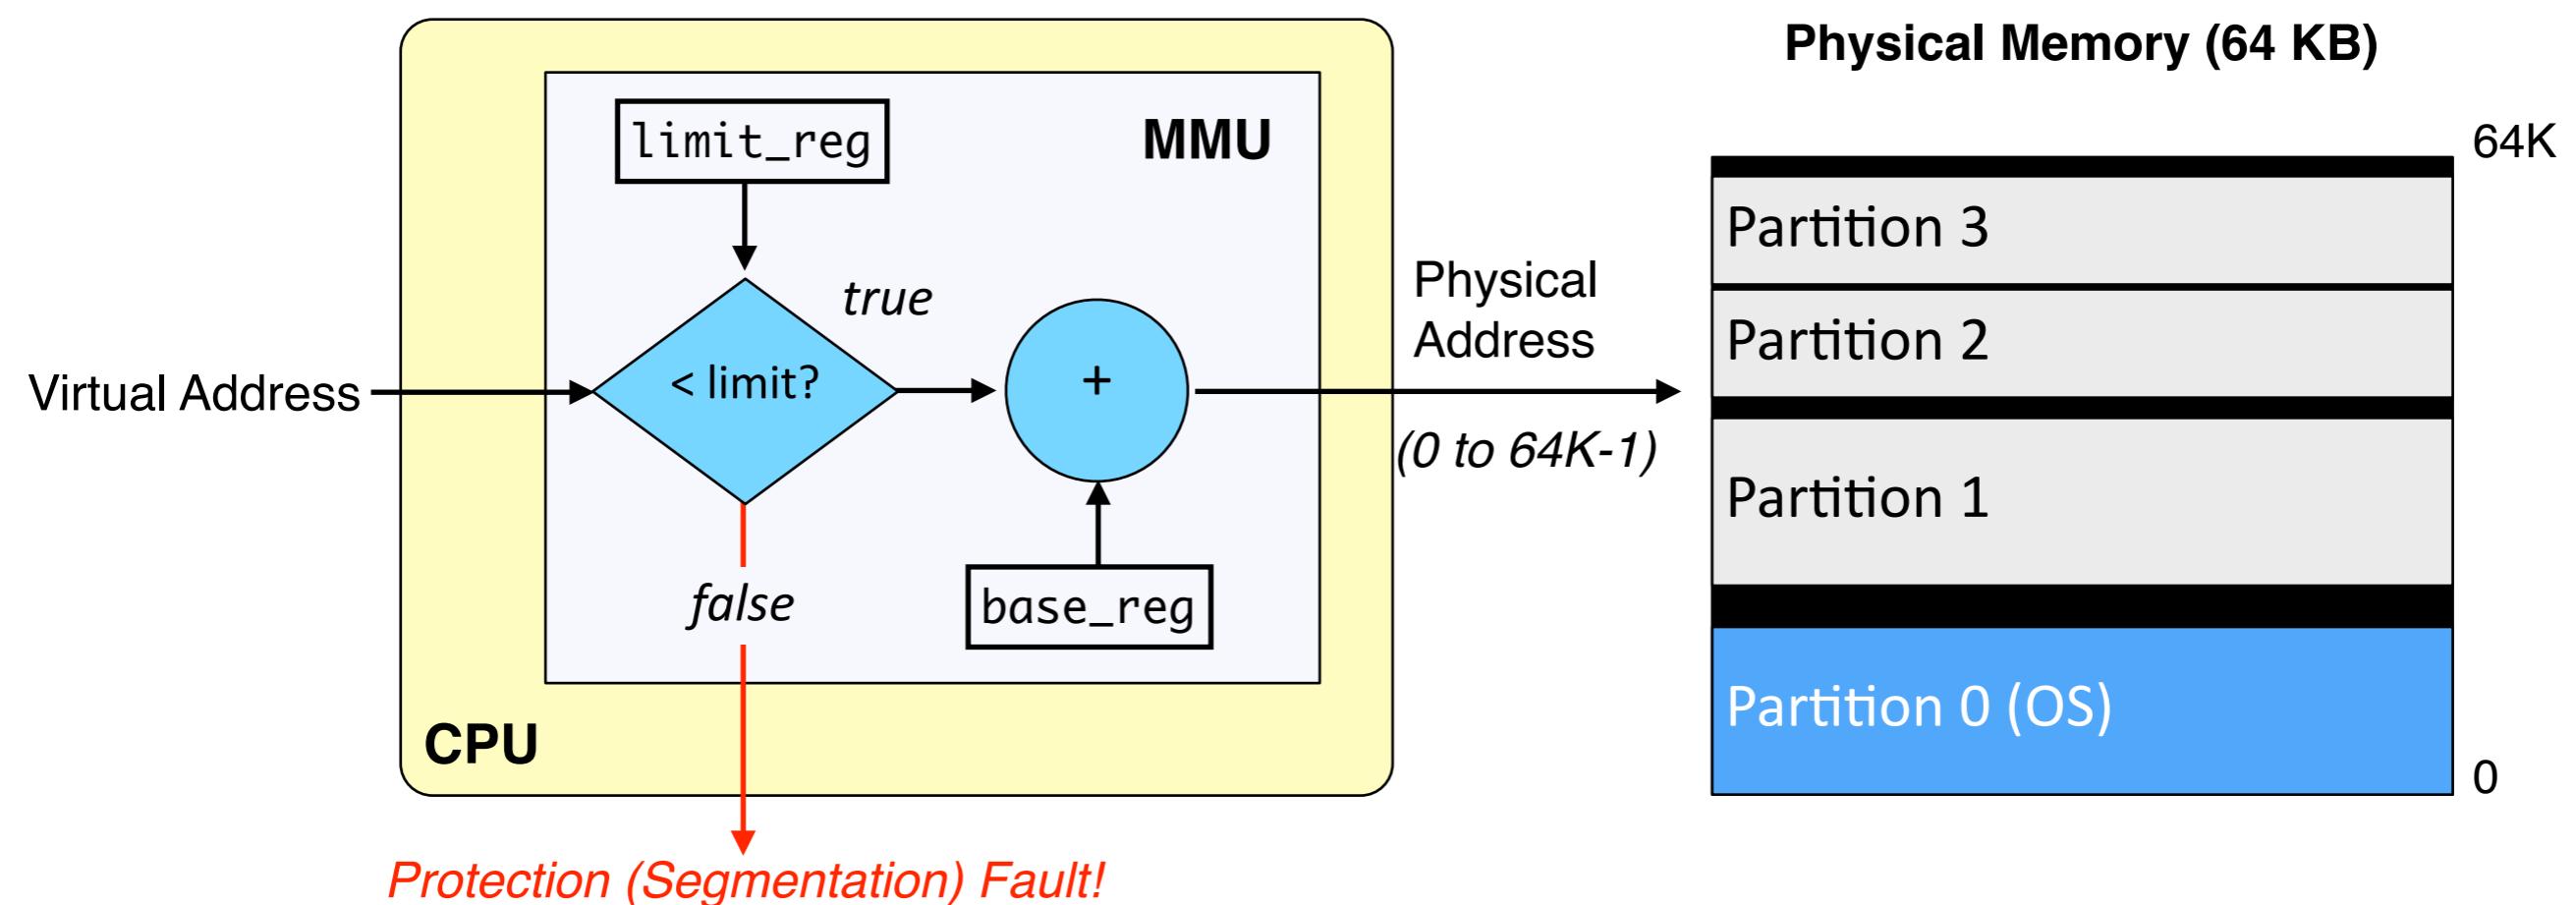


Variable Partitioning (Cont.)

- ▶ Each process (PCB) now has to remember:
 - **base_addr** - Where does process address space actually start in physical memory?
 - **limit** - How big is your partition?
- ▶ Hardware requirements:
 - **base_reg** and **limit_reg** added to the CPU that will get populated when the process PCB is loaded or switched in.

Variable Partitioning (Cont.)

- ▶ Translation logic (same as fixed partitioning)
 - $\text{physical addr} = \text{virtual addr} + \text{base addr}$
- ▶ Memory Protection:
 - If $\text{virtual addr} \geq \text{limit}$, generate segfault and terminate process



Evaluation of Variable Partitioning

▶ Pros:

- No more internal fragmentation!
- Degree of multiprogramming no longer confined by no. of partitions
- Processes can even grow in size during runtime (i.e. malloc)
 - Swap address space out to disk, increase *limit*, swap back in.
- Minimal hardware changes

▶ Cons:

- External fragmentation is now possible!
- Need a memory allocation algorithm (next slide)
 - Adds overhead to context switching

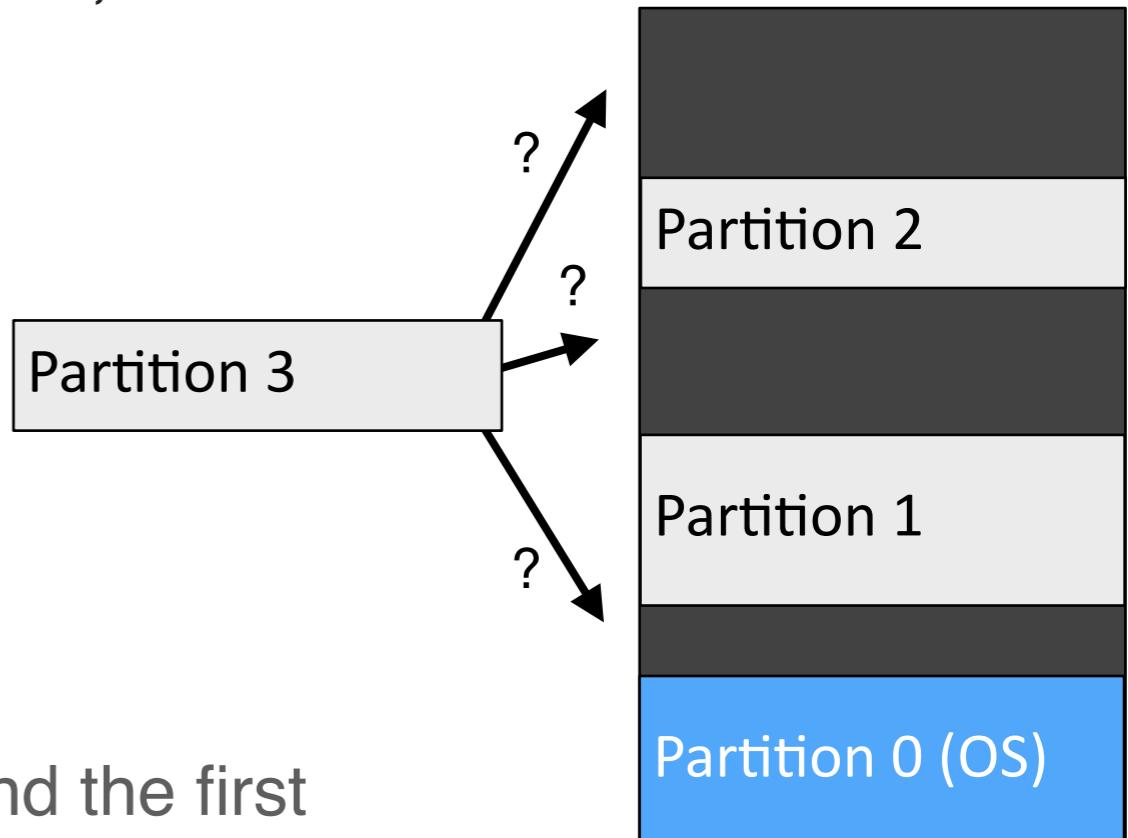
Memory Allocation Policy Needed!

- ▶ **Memory allocation:** With variable partitions, there are differently sized gaps in memory.

- What's the right *gap* to assign a new process?

- ▶ Policies:

- **First Fit:** Starting from physical address 0 , find the first gap that is wide enough to fit the process.
 - **Best Fit:** Find the best-fitting gap.
 - Best-Fit usually the worst.
 - Gives rise to **Worst-Fit** usually being the best -- *wait what?*

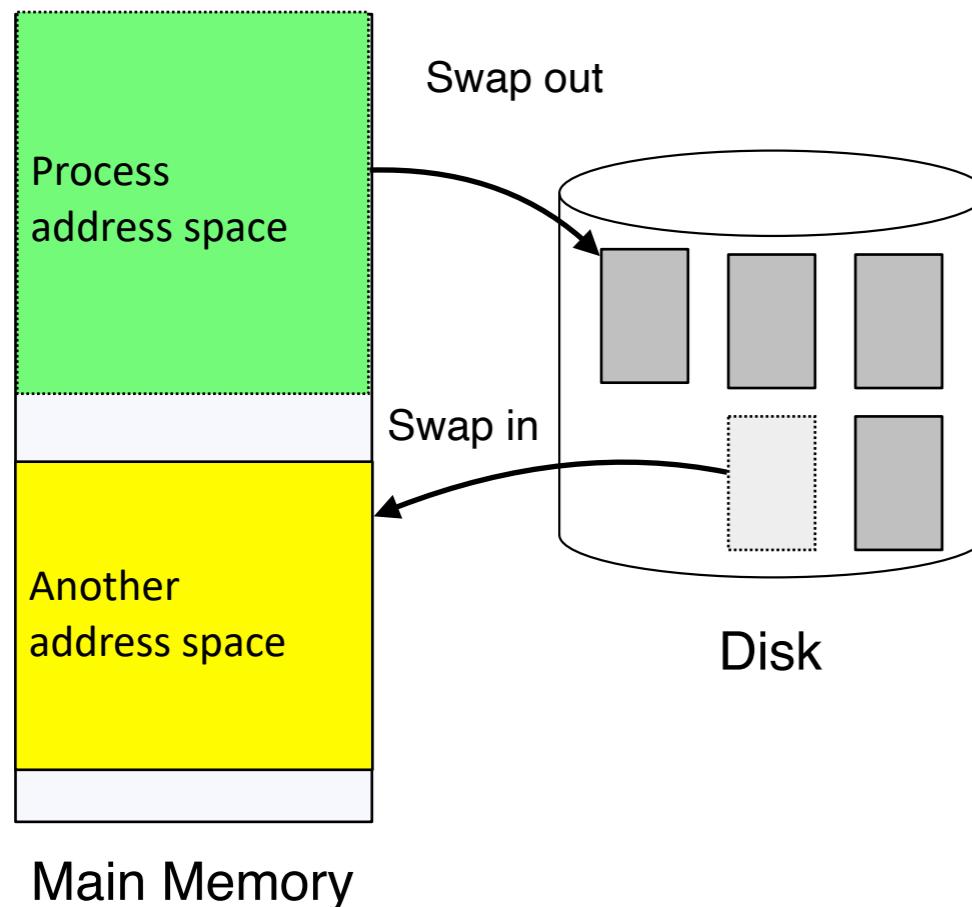


Goals for This Lecture...

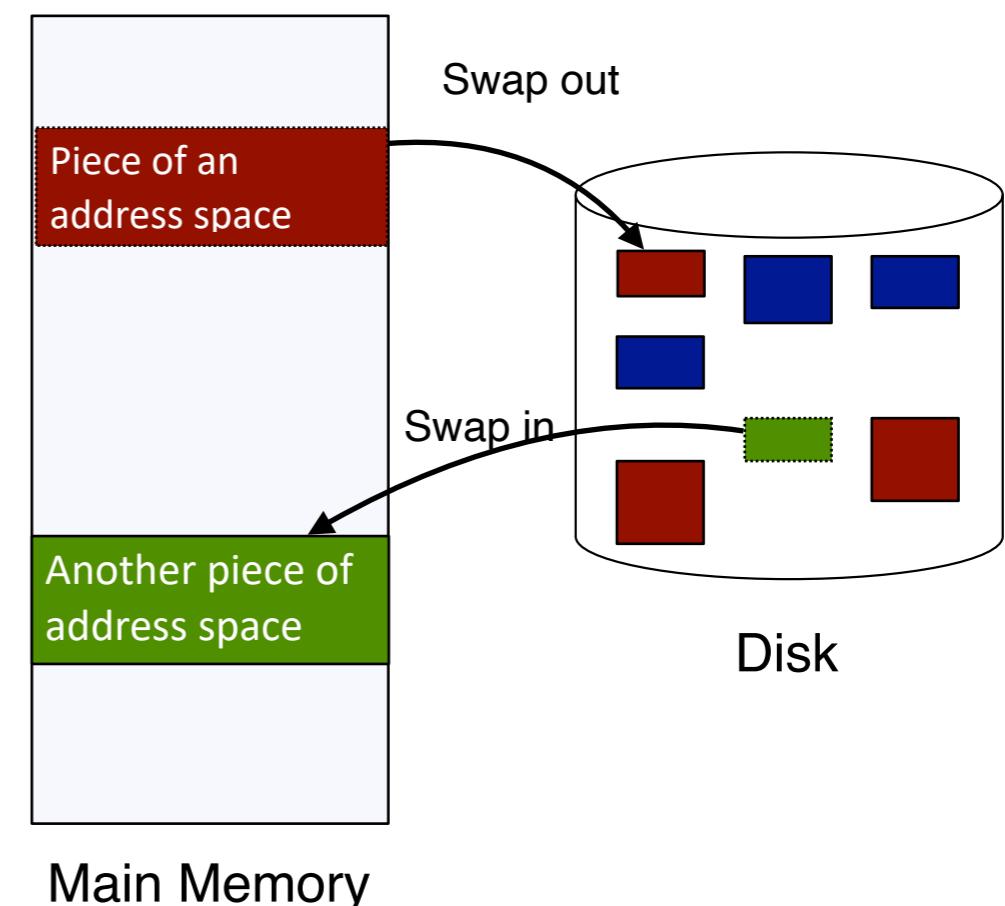
- ▶ Motivation and Goals
- ▶ Towards Virtual Addressing
- ▶ Partitioning
 - Fixed Partitioning
 - Variable Partitioning (Base and Bounds)
- ▶ Segmentation
- ▶ Paging

Partitions Are Monolithic and Restrictive

- ▶ Currently, the **whole** address space must fit into one partition.
 - Context switching can be a pain if a partition isn't available.



Old (Partitioned) Memory Model



New (Segmented) Memory Model

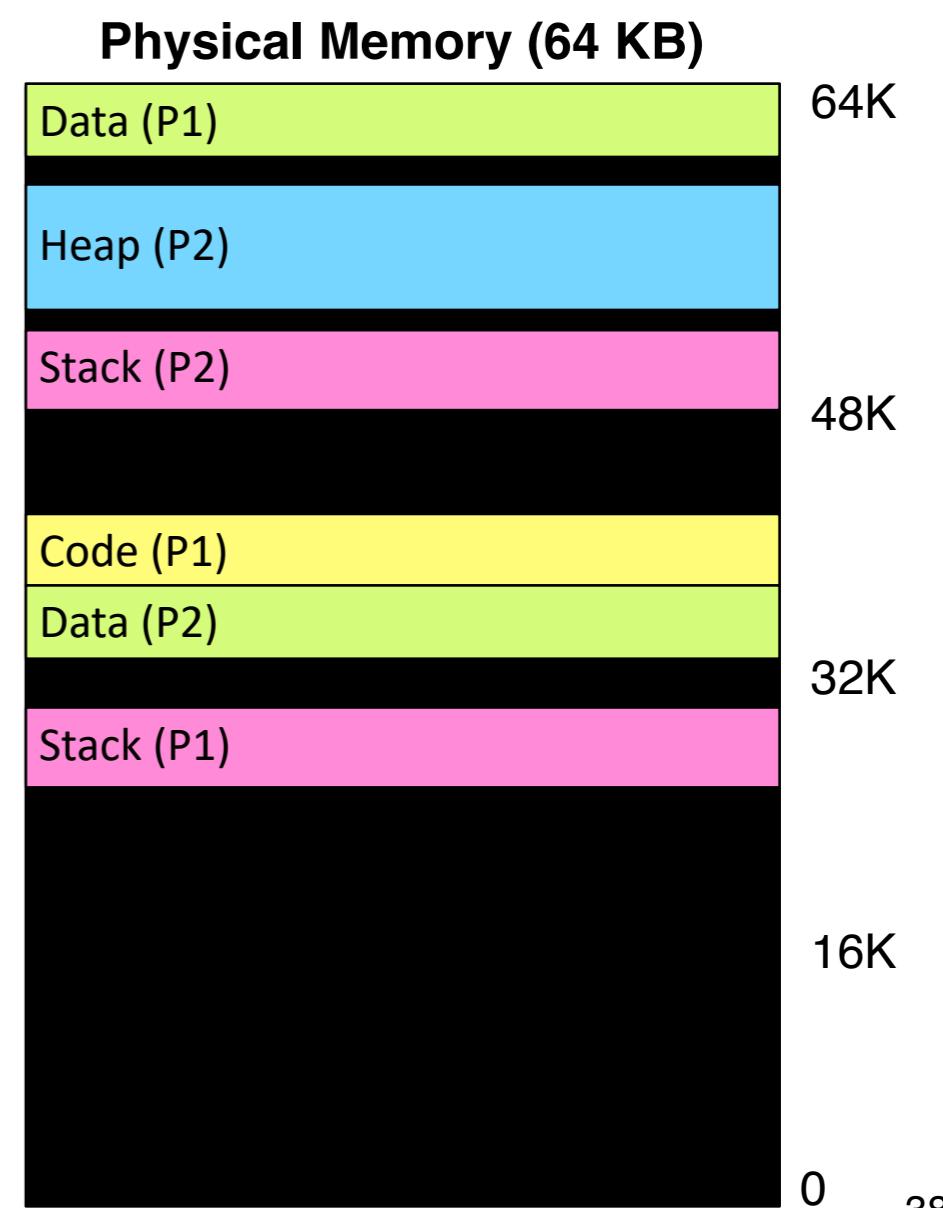
Memory Segmentation

► New Idea: **Segmentation**

- Break each address space down into:
code, data, heap, stack
- Each segment gets its own "base" and "limit"

Key difference from Variable Partitioning:

- Each segment is independently mapped to physical memory.
- Flexibility for multiprogramming: Not *all* segments need to be loaded in order for process to run.
- Requires a "segment table" for each process.



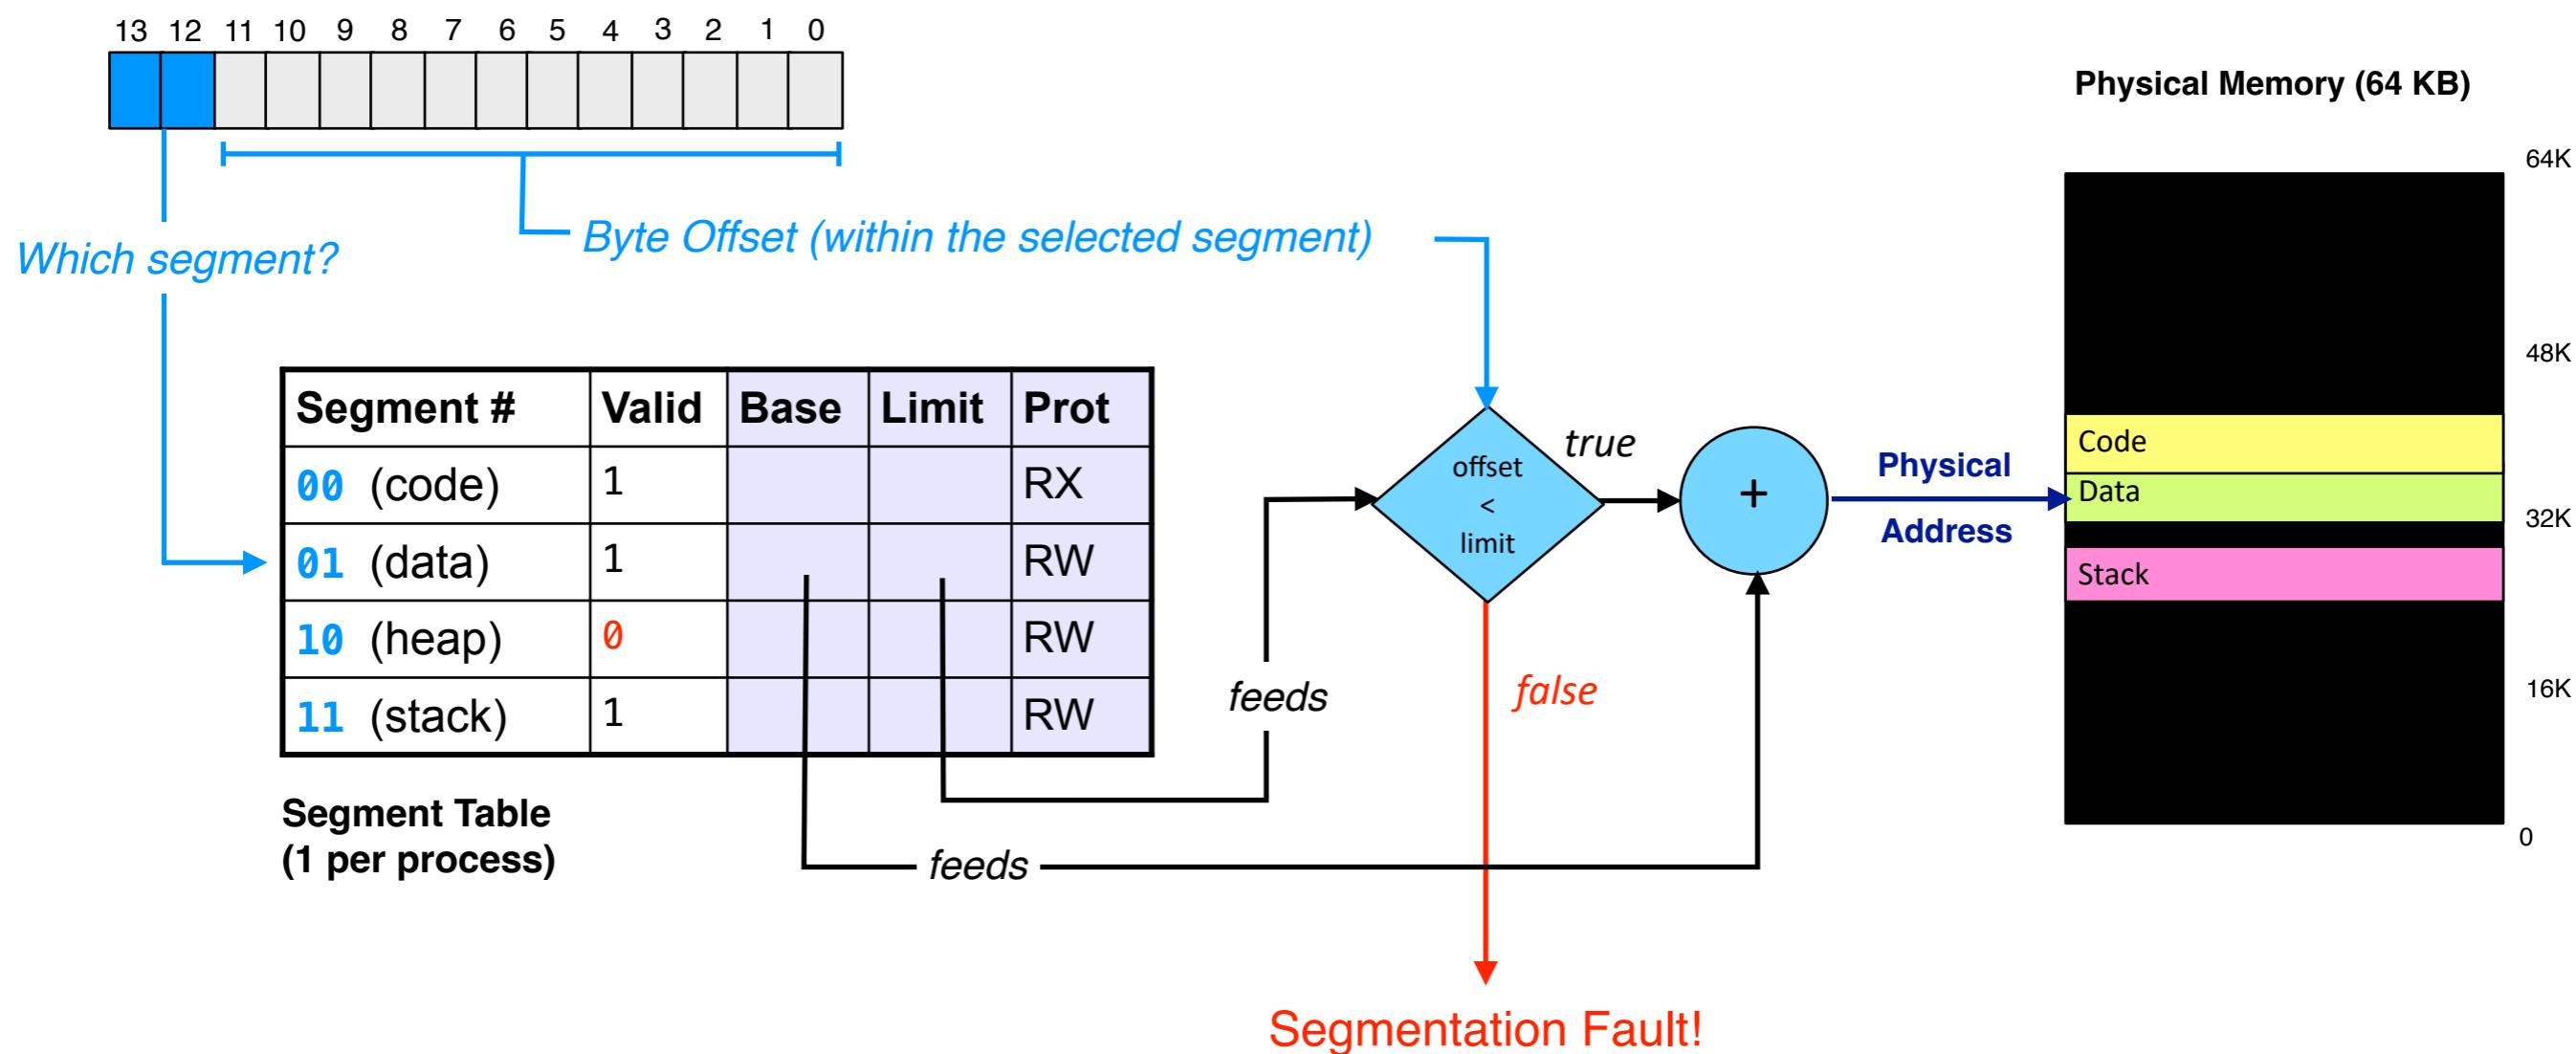
Segmentation: Address Translation

- ▶ Each process now has a *Segment Table* (or "Segment Map")

- What's the max size (limit) of each segment in this scheme below?

Virtual Address for a byte in process space:

(Assuming 16 KB address space for processes, the virtual address requires $\log_2 16K = 14$ bits to address each byte)



Sanity Check: Segmentation

- ▶ Consider a system that supports 32-bit addressing.
 - How much physical memory (RAM) can this system handle?
 - What is the possible range of addresses per process?
 - (I.e., how many bytes can each process have access to?)
 - What is the maximum segment size?

Sanity Check: Segmentation (Sol'n)

- ▶ Consider a system that supports 32-bit addressing.
 - How much physical memory (RAM) can this system handle?
 - Solution: 1 address per byte, so there are 2^{32} possible addresses.
 - So, 2^{32} bytes = 4GB. Anything over that would be un-addressable with only 32 bits.
 - What is the possible range of addresses per process?
 - (I.e., how many bytes can each process have access to?)
 - Solution: addressable range: 0 to $2^{32} - 1$... so yes one process address space could take up all the physical memory!
 - What is the maximum segment size?
 - Solution: You need the left-most two bits to address each segment, leaving 30 bits for the byte offset.
 - Therefore, 2^{30} bytes can be addressed per segment, so the segment size is 2^{30} bytes = 1GB

Segmentation Summary

▶ Pros:

- Better space utilization and flexibility than partitions since segments can be placed (re-located) anywhere in physical memory.
- Segments can be swapped out independently of each other
 - "Valid Bit" in the segment table in tells us whether a segment is loaded in memory or is out on disk
- Segment table allows for easy memory sharing
 - Example: Two processes can use the same (`base, limit`) for the code segment
- Segments can grow + shrink during runtime
 - Swap segment out to disk, change its "limit" and swap back in

Segmentation Summary (Cont.)

▶ Cons:

- Overhead in OS
 - A segment table needs to be set up for each process
 - (Process creation time is now slower)
- Address translation is much slower than memory partitioning
 - Each memory access needs to first consult segment table instead of just adding the offset to a base register
- External fragmentation of segments *still* possible
 - (Holes *between* segments get smaller and smaller and will add up and make it hard to swap segments in)

Goals for This Lecture...

- ▶ Motivation and Goals
- ▶ Towards Virtual Addressing
- ▶ Partitioning
- ▶ Segmentation
- ▶ **Paging**
 - Page Tables
 - Translation Lookaside Buffer (TLB)

Segmentation Review

► Segmentation offers...

- Dynamic relocation and protection ✓
 - Virtual addressing and translation
- Code sharing ✓
 - Via segment tables
 - Just map a segment from different processes to the same base/limit

► One major problem remains:

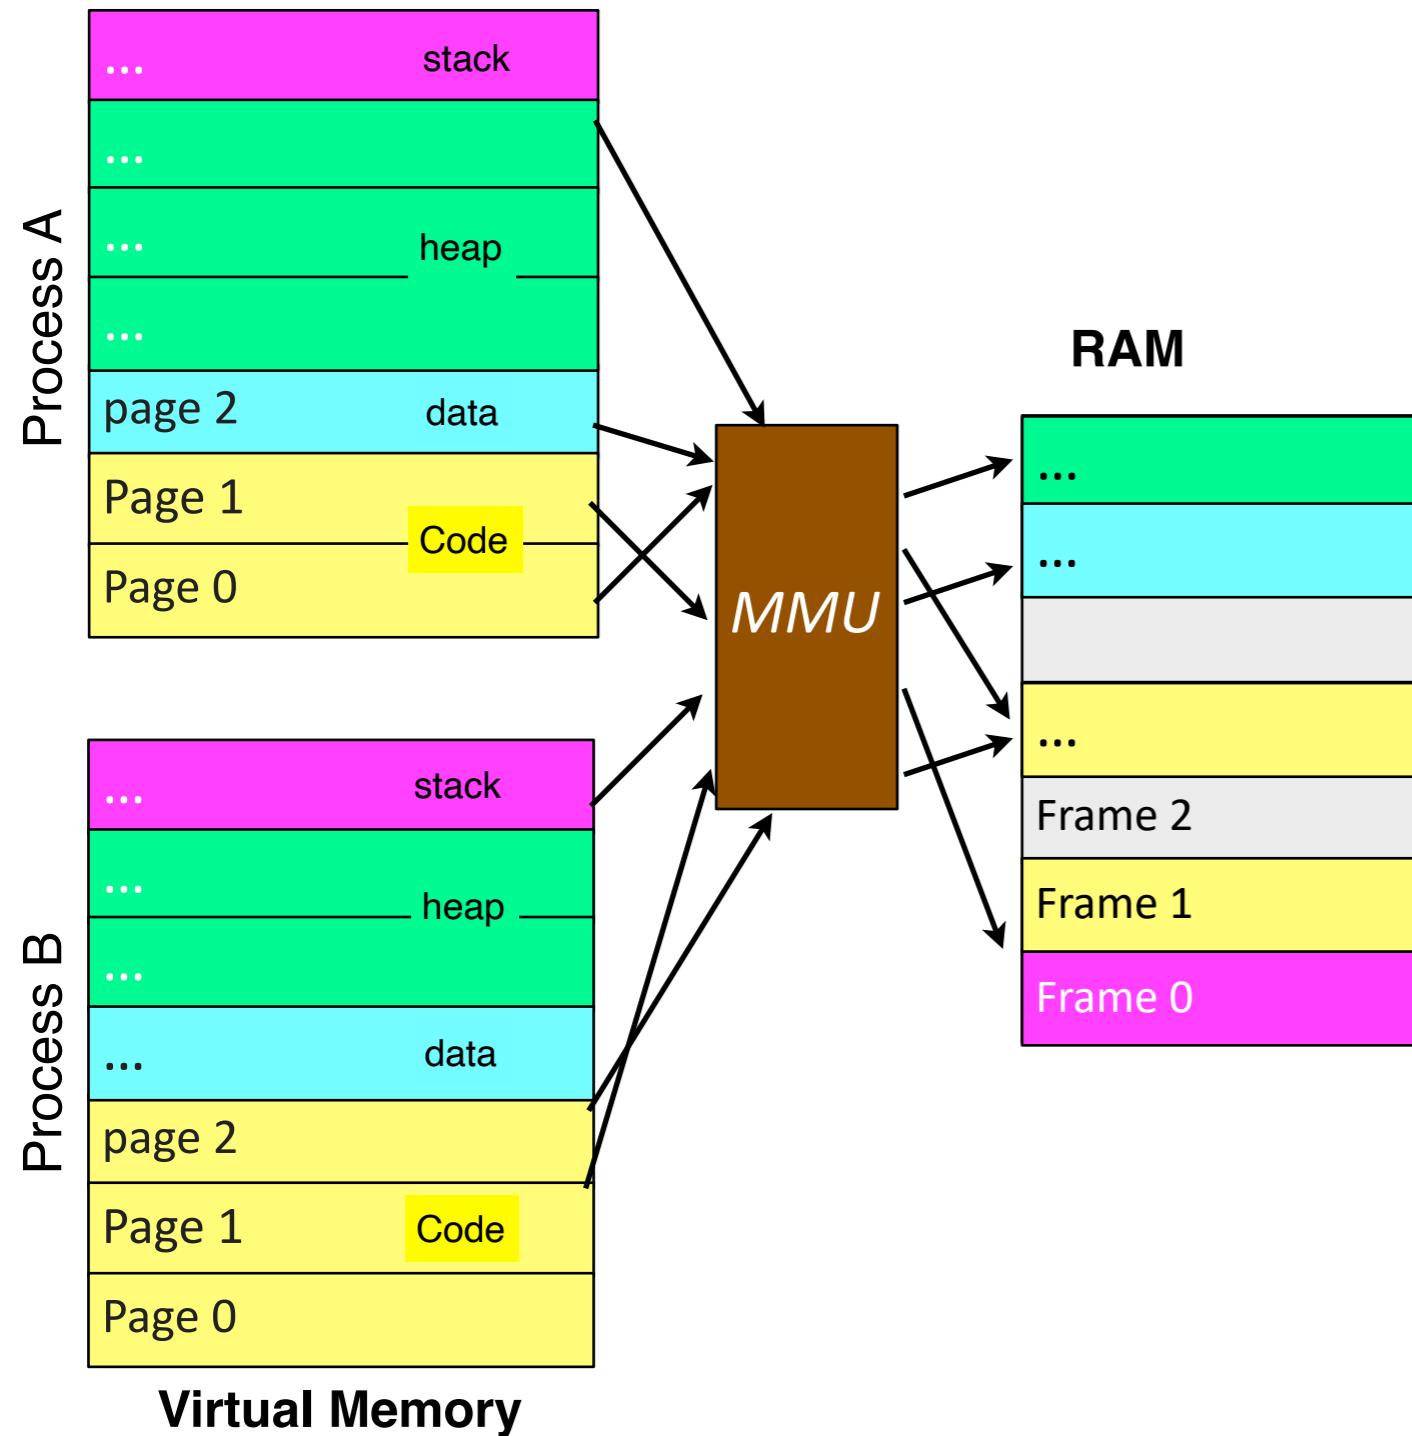
- **External Fragmentation of Segments**

We need to focus on eliminating **External Fragmentation**.

- * They arise from being variable length.
- * We should return to fixed-length chunks!
- * Which means, we will *allow* internal fragmentation!

Modern Memory Management: Paging

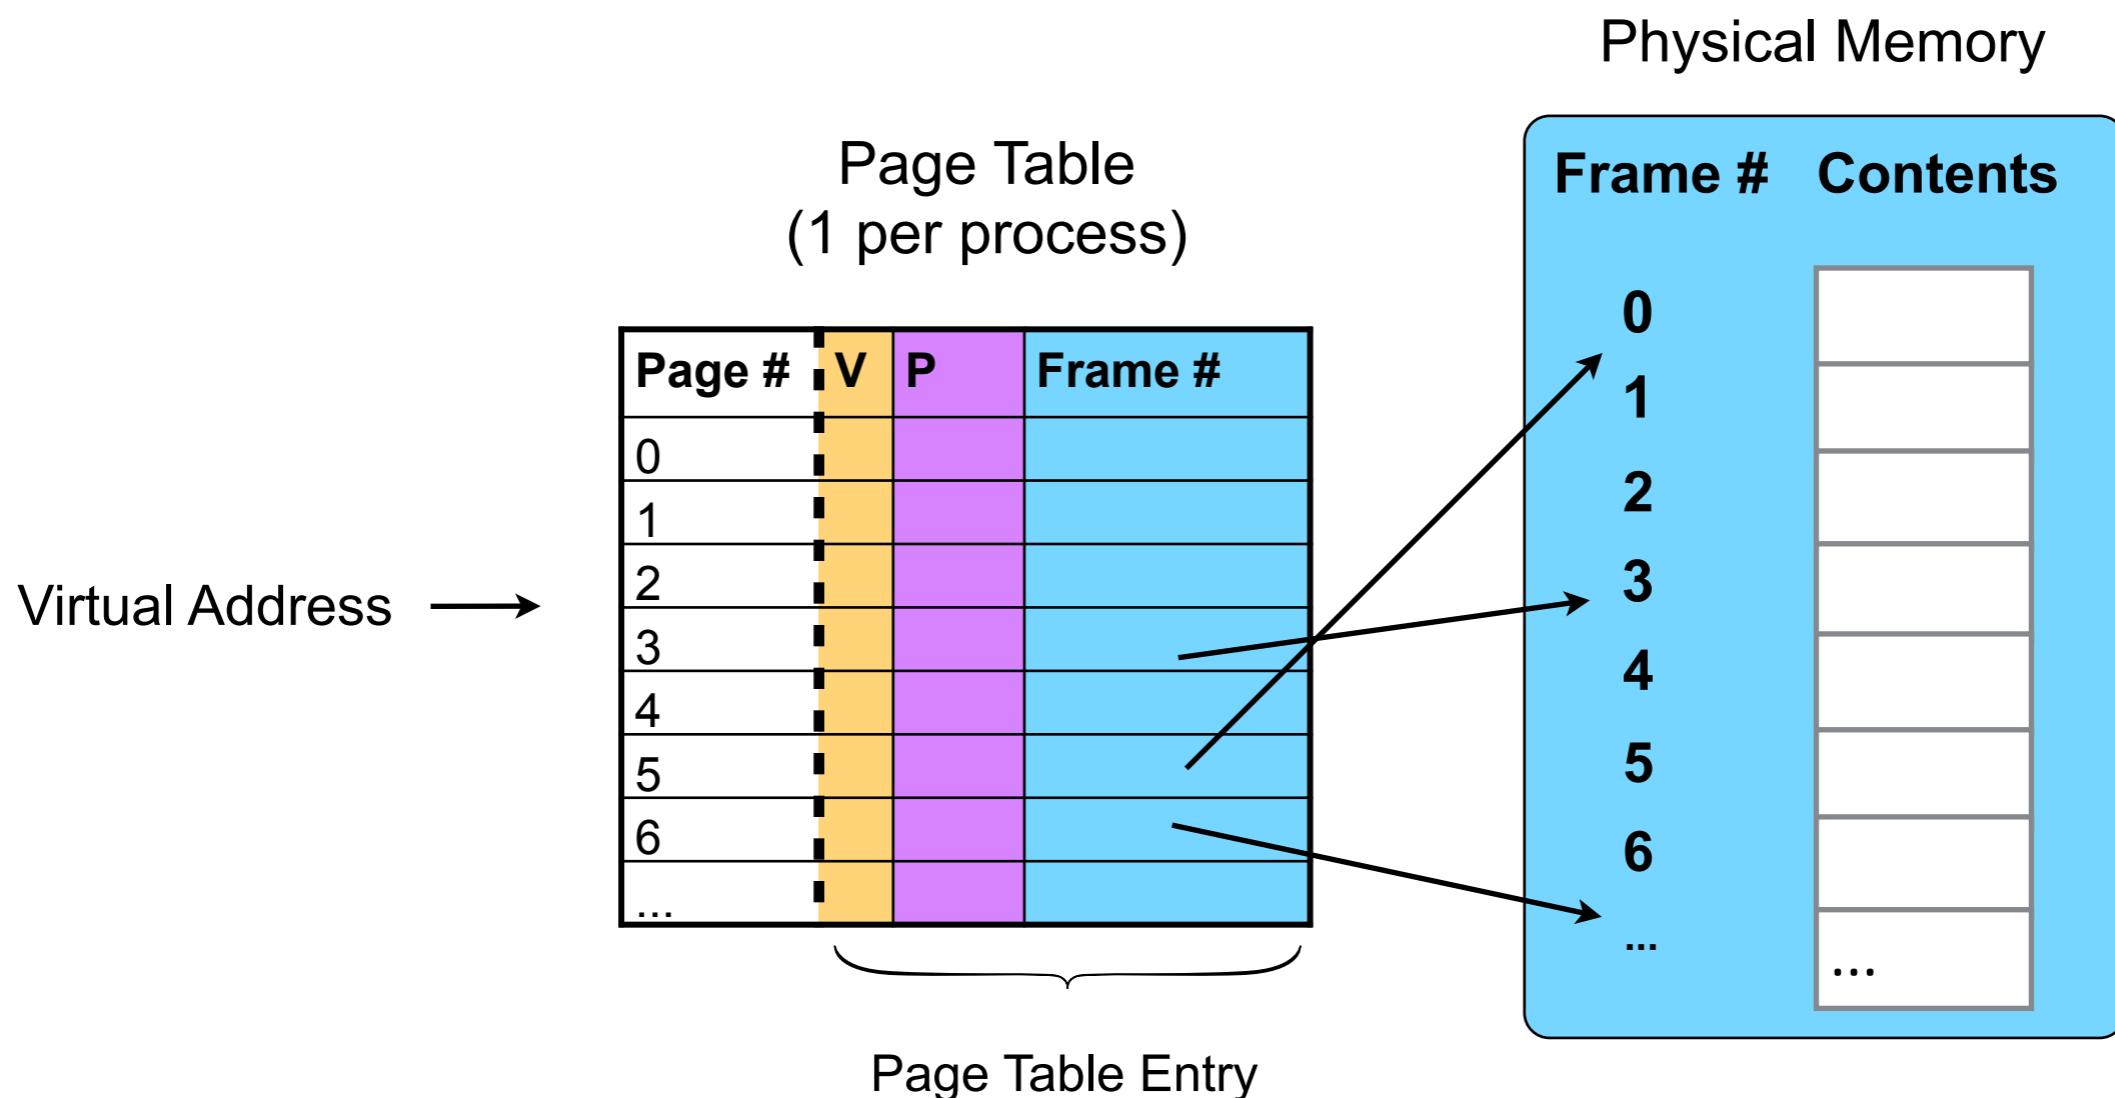
Lesson learned from segmentation: Processes don't need a contiguous, monolithic view of its own address space. **We can break it up!**



- ▶ Use **small, fixed-sized** chunks of virtual and physical memory
 - *Page*: a virtual memory chunk
 - *Frame*: a physical memory chunk
 - ***Page Size = Frame Size***
 - This constraint *must* be satisfied!
 - Anywhere from 512 bytes to 16 KB in practice.

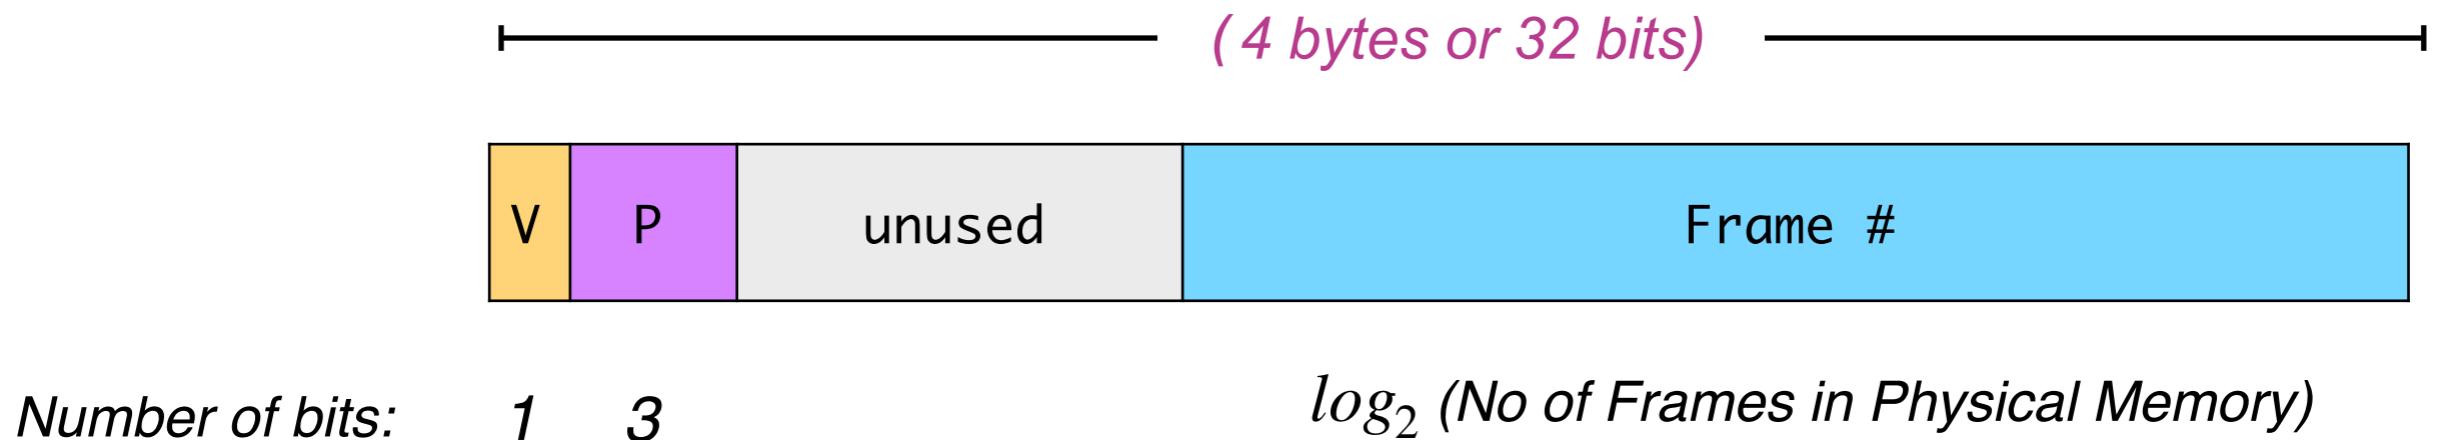
Page Tables

- ▶ How does the OS keep track of pages for a process?



Page Table Entries

- ▶ A "page table entry" is really just 32-bit sequence (i.e., an int)

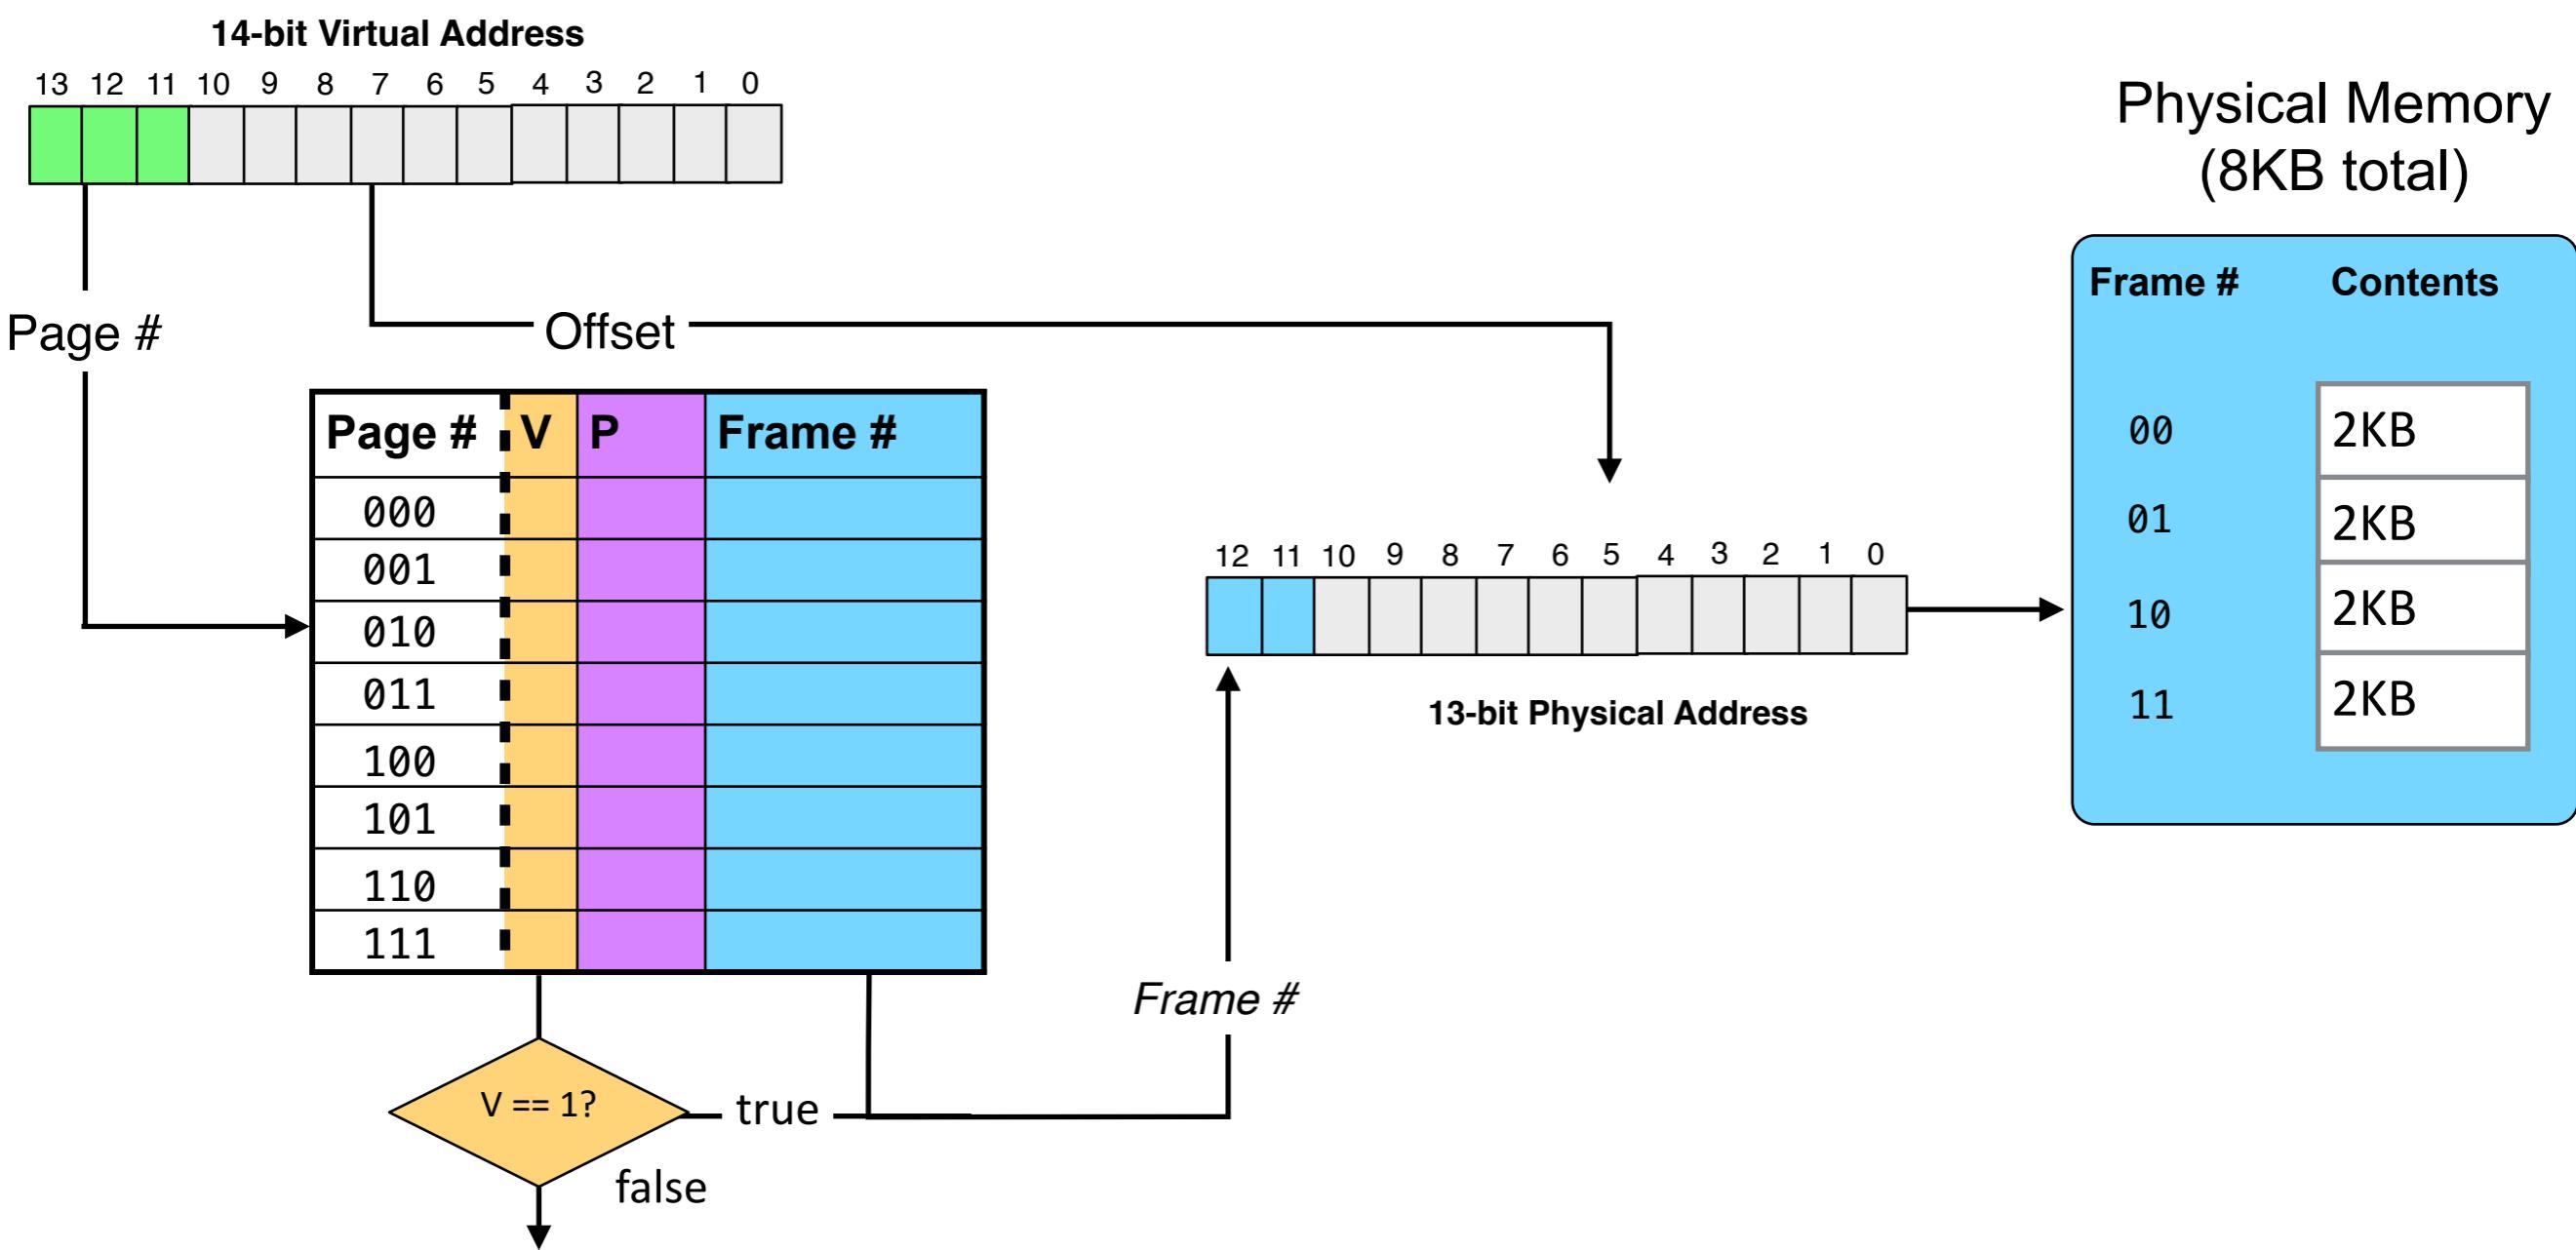


- ▶ What do these bits mean?

- **Valid Bit (V):** Is the page present in memory (1) or out on disk (0)?
- **Protection Bits (P):** Read-Only? Read-Write? Read-Execute? ...
- **Frame #:** Which frame in physical memory does this page map to?
- **Unused:** any remaining bits might be useful in the future.

Address Translation Logic (Paging)

- ▶ Example: 8 pages; 4 frames in physical memory; page size = 2KB
 - (In paging, processes can address more memory than available physical memory!)



Trap a Page Fault! (Need to load a page/frame into memory from disk)

Paging Example

- ▶ Consider the following scheme:
 - 8 frames and 8 pages (Size = 8 KB)
 - Addresses need at least 16 bits (*How do I know that? What's the address-format?*)
 - Four processes: A, B, C, and D.
 - Show the page tables after this schedule is run:
 - Load(A); Load(B); Load(C); SwapOut(B); Load(D); Load(B)
 - For now, assume all pages of a process get swapped in and out
- ▶ Say D's address space is currently 20,000 bytes.
 - What actually happens when D accesses its own (virtual) address,
9217 (= 0010 0100 0000 0001)?

Paging Example

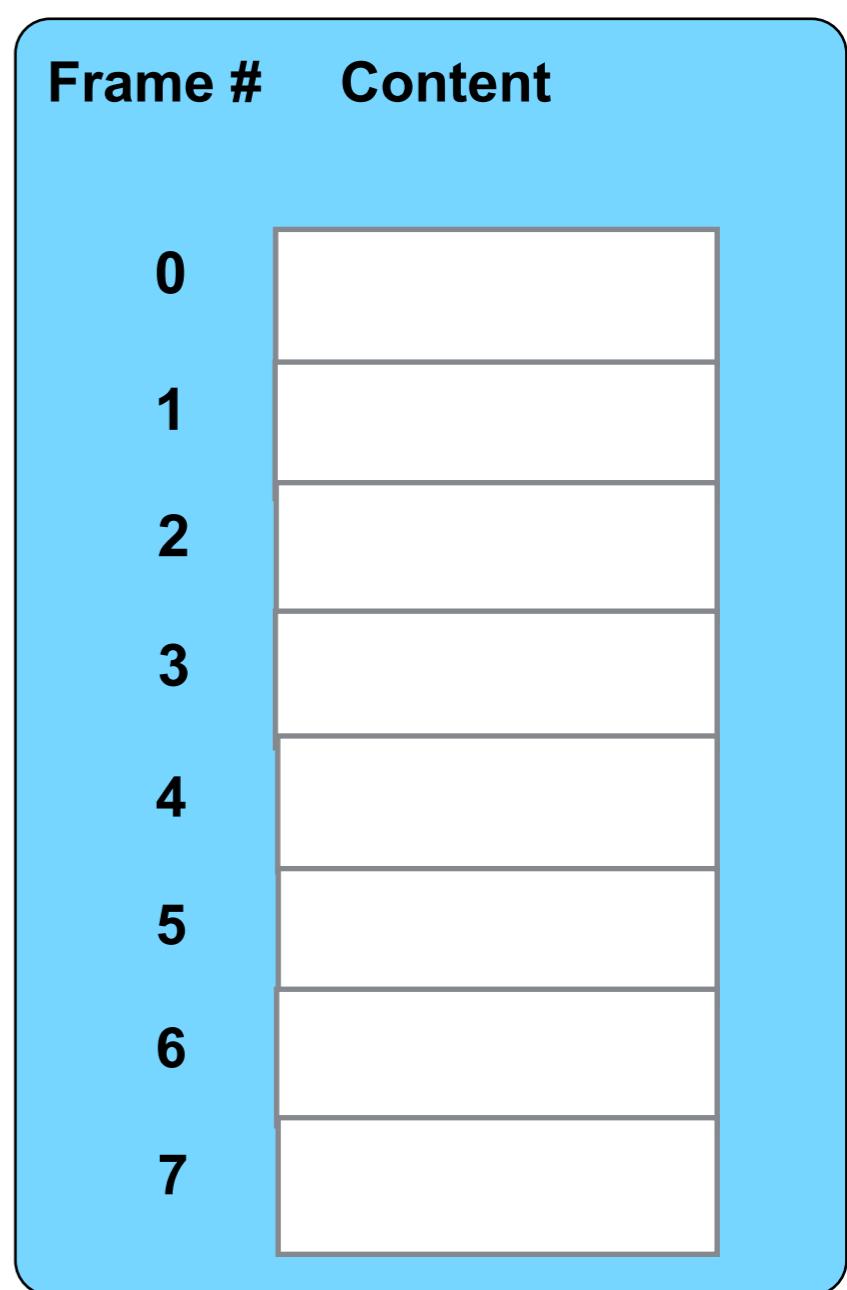
PageTable: Process A

Page #	V	P	Frame #
0 (000)			
1 (001)			
2 (010)			
3 (011)			
4 (100)			
5 (101)			
6 (110)			
7 (111)			

PageTable: Process B

Page #	V	P	Frame #
0 (000)			
1 (001)			
2 (010)			
3 (011)			
4 (100)			
5 (101)			
6 (110)			
7 (111)			

Physical Memory



PageTable: Process C

Page #	V	P	Frame #
0 (000)			
1 (001)			
2 (010)			
3 (011)			
4 (100)			
5 (101)			
6 (110)			
7 (111)			

PageTable: Process D

Page #	V	P	Frame #
0 (000)			
1 (001)			
2 (010)			
3 (011)			
4 (100)			
5 (101)			
6 (110)			
7 (111)			

Solution: After Running the Schedule

PageTable: Process A

Page #	V	P	Frame #
0 (000)	1		0
1 (001)	1		1
2 (010)	0		-
3 (011)	0		-
4 (100)	0		-
5 (101)	0		-
6 (110)	0		-
7 (111)	0		-

PageTable: Process B

Page #	V	P	Frame #
0 (000)	1		6
1 (001)	1		7
2 (010)	0		-
3 (011)	0		-
4 (100)	0		-
5 (101)	0		-
6 (110)	0		-
7 (111)	0		-

Physical Memory

Frame #	Content
0	A[0]
1	A[1]
2	D[0]
3	D[1]
4	C[0]
5	D[2]
6	B[0]
7	B[1]

PageTable: Process C

Page #	V	P	Frame #
0 (000)	1		4
1 (001)	0		-
2 (010)	0		-
3 (011)	0		-
4 (100)	0		-
5 (101)	0		-
6 (110)	0		-
7 (111)	0		-

PageTable: Process D

Page #	V	P	Frame #
0 (000)	1		2
1 (001)	1		3
2 (010)	1		5
3 (011)	0		-
4 (100)	0		-
5 (101)	0		-
6 (110)	0		-
7 (111)	0		-

Evaluation of Paging

► Pros

- No External Fragmentation!
- Separation of address spaces: Virtual vs. Physical
 - Relocatability!
 - Processes can address **more memory** than is *physically* available!
 - (How? Multiple pages can map to the same frame, and only one page is **valid**)
- Memory Protection
 - Page tables enable controlled multiplexing of processes in RAM
- Memory Sharing
 - Map pages from different processes to the same frame in RAM

Paging Makes Memory Sharing Easy

Virtual Address (Proc A)



Page Table (Proc A)

Page #	V	M	R	P	frame #
0	1	1	1	RX	3
1					
...					

Physical Memory

Frame #	Contents
0	
1	
2	
3	Shared Page
4	
5	

Page Table (Proc B)

Page #	V	M	R	P	frame #
0					
...					
N	1	1	1	RX	3



Virtual Address (Proc B)

Evaluation of Paging (Cont.)

► Cons

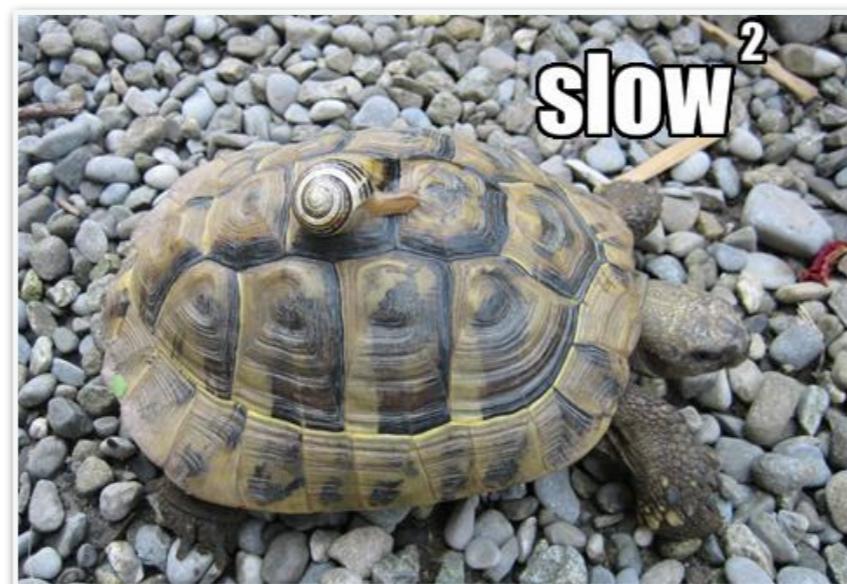
- Still some internal fragmentation
 - Worse for larger sized pages
 - Not a huge deal... if we just keep pages small (4 KB - 16 KB in practice)
 - Page size is an OS configuration parameter
- In the best case, paging slows down memory access by a factor of 2!
 - Page tables are too large to fit into CPU registers.
 - *You effectively need **two** memory accesses to resolve each memory reference:*
 - *(1 to lookup the page in the page table + 1 to reference physical memory)*

Goals for This Lecture...

- ▶ Motivation and Goals
- ▶ Towards Virtual Addressing
- ▶ Partitioning
- ▶ Segmentation
- ▶ **Paging**
 - Page Tables
 - Translation Lookaside Buffer (TLB)

Problem: Address Translation Is Slow!

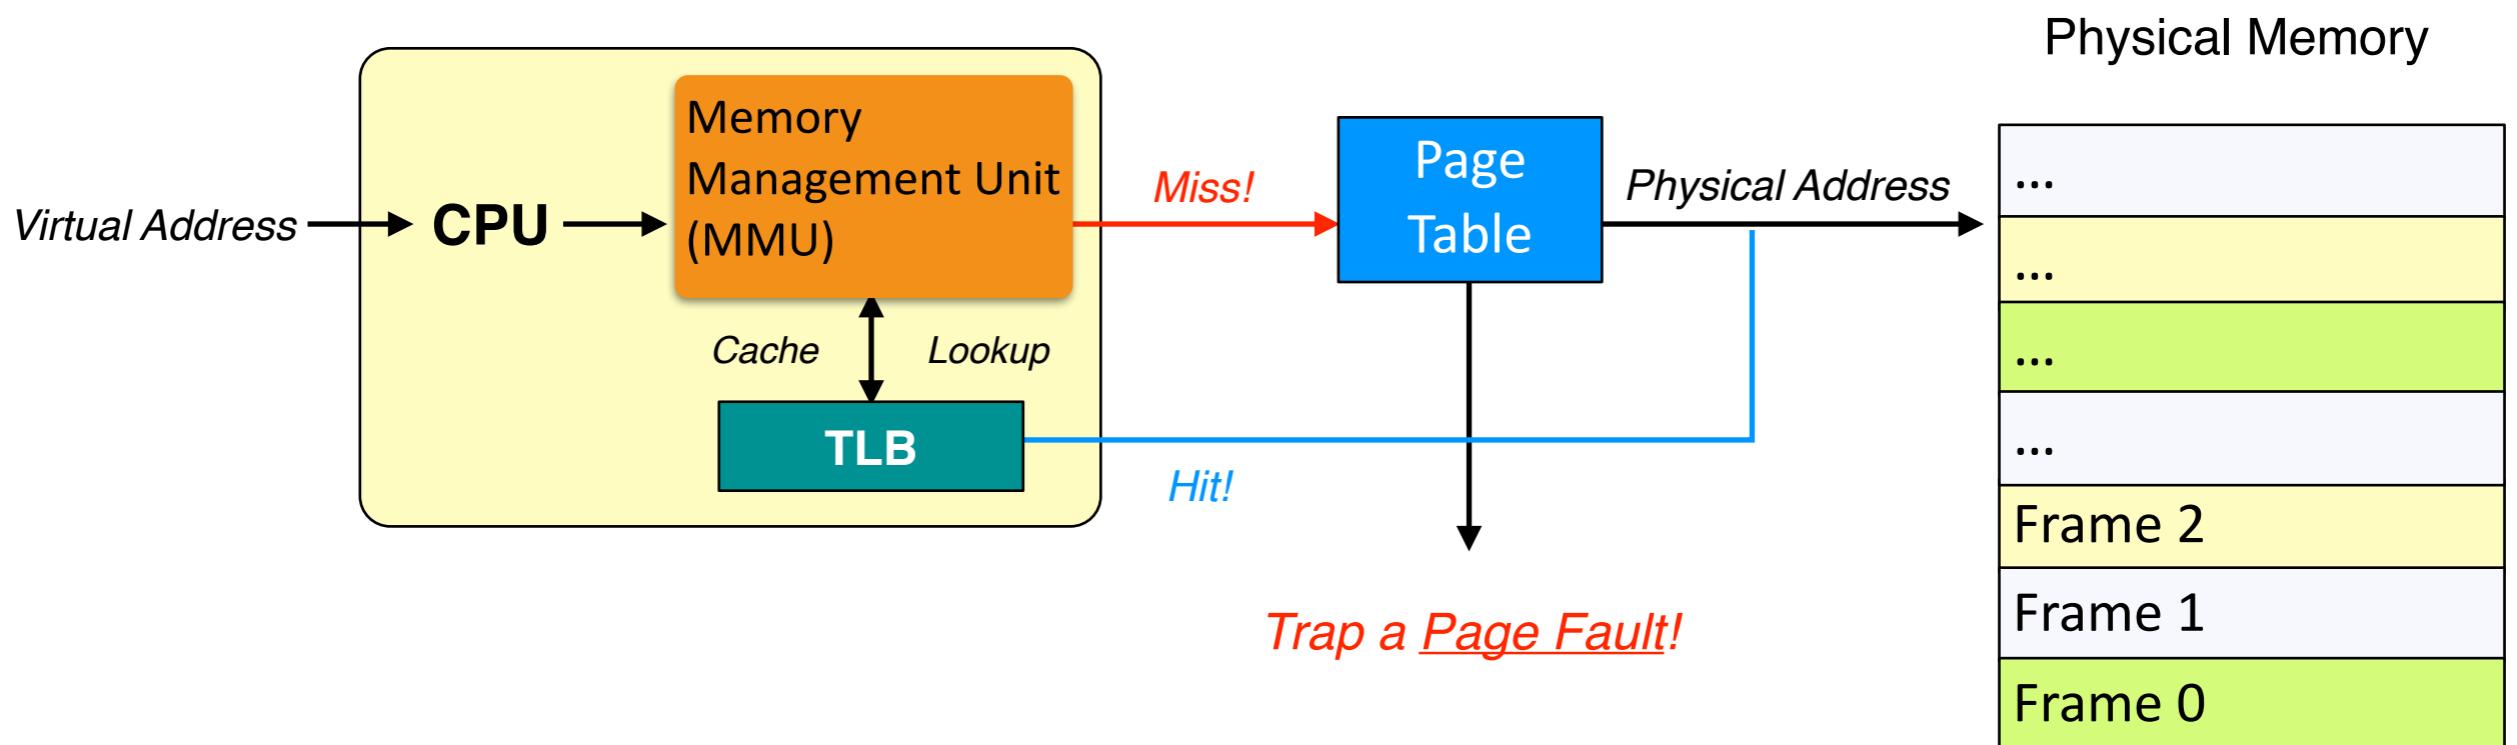
- ▶ New Problem: **S-L-O-W** Address Translations
- ▶ Every memory reference is at least twice as slow as before:
 - Given a virtual address:
 - Page Table lookup (1 memory access)
 - Retrieve the byte in physical memory (1 memory access)
 - If the page-table entry was invalid, then you have to first swap the page in from disk (even slower!)



Updated Translation Hardware

- ▶ Cache the address translations!

- The *Translation Lookaside Buffer (TLB)* first appeared in IBM System 370 (just a fancy name for an "translation cache")



Translation Lookaside Buffer (TLB)

- ▶ TLBs are used to "cache" the address translations!
 - It lives on the CPU (the fastest cache on chip ~ faster than L1 cache!)
 - Its sole purpose is to map virtual addresses to physical addresses
 - A typical TLB stores 64 to 512 entries

Translation Lookaside Buffer (TLB)

Page #	Frame #
3	2032
68	309
...	...

A blue arrow labeled "Lookup Page #" points to the "Page #" column header. A red arrow labeled "Access Frame #" points to the "Frame #" column header.

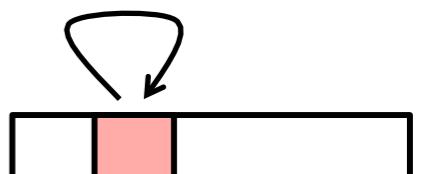
Would a TLB Really Work?

▶ Principle of Locality?

- Do programs tend to use data and instructions with addresses near or equal to those used recently?

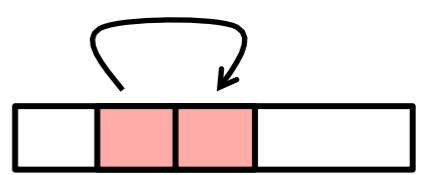
▶ "Temporal Locality"

- Recently-used elements will likely be used again in the near future.



▶ "Spatial Locality"

- Elements with nearby addresses will tend to be referenced in the near future.



Principle of Locality Example

- ▶ For translation caching to work, both **code** and **data** references need to observe the principle of locality.
- ▶ Both **temporal** and **spatial** locality of **code (instructions)** and **data** exhibited in this simple example.

```
double sum_vec(double a[M] ) {  
    double sum = 0;  
  
    for (int i = 0; i < M; i++) {  
        sum += a[i];  
    }  
  
    return sum;  
}
```

Locality Example: Code

► Instruction (Code) Locality

- When loaded, the compiled binary code gets placed inside pages
- Repeated references of instructions in the *same page* is very likely.
 - e.g., Loops tend to dominate code
 - e.g., Calls to popular functions too

```
double sum_vec(double a[M] ) {  
    double sum = 0;  
  
    for (int i = 0; i < M; i++) {  
        sum += a[i];  
    }  
  
    return sum;  
}
```

Locality Example: Data

► Data Locality

- We know `sizeof(double)` is 8 bytes
- If page size = 4KB, then a page can store $4K/8 = 512$ doubles:
 - $a[0] \dots a[511]$ fit in one page
 - $a[512] \dots a[1023]$ fit in another page, and so on.
 - *Repeated references to sequential array elements are to the same page!*

```
double sum_vec(double a[M]) {
    double sum = 0;

    for (int i = 0; i < M; i++) {
        sum += a[i];
    }

    return sum;
}
```

Handling Context Switches

- Remember, due to virtual addressing,
 - Data at address **X** for **Process A** is different than for **Process B**
 - (Unless the page containing address **X** is shared)
 - **Problem:** TLB entries would be invalid after a context switch!
- ▶ Two options:
 - Small TLB: Just clear it out on a context switch
 - Larger TLB can store more entries. May not want to flush all.
 - Add address-space identifiers (*e.g.*, PID) to the TLB.

Measuring Memory Access Performance

- ▶ We have two access paths for an address reference:
 1. Address translation is found in TLB (*TLB hit*)
 2. Or, not found in TLB, need a Page-Table Walk (*TLB miss*)
- ▶ Definition: *Effective Memory Access Time (EAT)*
 - $EAT = hT_{hit} + (1 - h)T_{miss} + mem_access_time$
 - h = TLB hit rate
 - T_{hit} = TLB lookup time
 - T_{miss} = TLB lookup time + page-table lookup time

(*How much can a TLB improve EAT?*)

Example

- ▶ Assume:
 - Memory Access Time = 200 ns
 - Page Table lookup requires 1 additional memory access per translation
- ▶ What's the EAT without a TLB?
- ▶ What's the EAT with a TLB with following specs?
 - TLB hit rate (h) = 0.75
 - TLB lookup time (T_{hit}) = 25 ns
- ▶ If we can only absorb a 20% slowdown of memory access, what does the TLB hit rate have to be?

In Summary... (Cont.)

- ▶ Limitation: We have less physical memory than virtual address space
 - Currently: *All* segments or pages of a process must be in physical memory for the process to execute
 - Context switching and swapping slow because of it!
- ▶ Do all pages need to be in physical memory for process to run?
 - At any point in time, only part of the code is actually executing
 - *Next time: Virtual Memory*

Administrivia 4/14

- ▶ Reminders:
 - Hwk 7 (Due 4/18 Friday)
- ▶ Last time ...
 - Banker's algorithm
- ▶ Today
 - Start memory management (Chap 9)
 - Desirable features of memory management
 - Loaders: Absolute loading; Static Relocation; Virtual Addresses

Administrivia 4/18

- ▶ Reminders:
 - Hwk 7 (extension - due Monday)
- ▶ Last time...
 - Virtual addressing (*Why do we even need this?*)
 - Fixed partitioning
 - Variable partitioning ("Base & Bounds")
 - Segmentation
- ▶ Today
 - The modern approach: Paging

Administrivia 4/21

- ▶ Reminders
 - Homework 8 posted (due 5/9)
- ▶ Last time...
 - Paging: the modern memory management solution
 - Split address space up into fixed size chunks (called "pages")
- ▶ Today:
 - Review of paging
 - Updated address translation scheme
 - Translation lookaside buffer (TLB)