

CS 475

Operating Systems



Department of Mathematics
and Computer Science

Lecture 9
Virtual Memory (Part II)


Brad and Adam Go to the Post Office

- ▶ Post Office keeps mailboxes for customers out in the front.



*Fast, but limited
(RAM)*

- ▶ Mail that "spills over" go in the warehouse in the back end.



*Slow, but large
(Swap disk)*

Brad and Adam Go to the Post Office

- ▶ The Post Office assigns 3 mailboxes each to Adam & Brad
 - Brad only uses one of his mailboxes as a secondary storage for ham.
 - Adam owns a t-shirt company called "*Purrrrfectly Adam*" that is surprisingly successful
 - Needs 9 (virtual) PO Boxes

Brad's Mapping

PO Box	V	Real Box
1	1	F
--	--	E
--	--	A

Adam's Mapping

PO Box	V	Real Box
1	1	B
2	1	D
3	0	
4	0	
5	1	C
6	0	
7	0	
8	0	
9	0	



Motivation: Memory Allocation

- ▶ If a process is given *too few* frames...
 - Could lead to many page faults during the runtime of a process
- ▶ If a process is given *too many* frames...
 - Frames could go unused
 - Wasteful
 - May lead to page faults in *other* processes
- ▶ Why does memory allocation matter?
 - Minimize page faults!
 - Avoid *thrashing*: process spends more time swapping than executing

Goals for This Lecture...

- ▶ Demand Paging
 - Motivation
 - Implementation
- ▶ Page Replacement Policies
- ▶ Memory Allocation
 - Equal
 - Proportional
 - Working Set Model
 - How Do malloc() and free() Work?

Equal Allocation (This is the P.O. Problem)

► Equal Allocation

- Suppose we have m frames and n processes ($m > n$)
- Give every process an equal share of m/n frames
- Share of frames doesn't change over time

► Example: $m = 93$ frames, $n = 10$ processes

- Each process gets 9 frames
- Remaining 3 frames used as *page buffers*

► Pros: Easy, static algorithm. Fair.

► Cons?

- Larger processes starve for frames (swap with disk more often)

Aside: Page Buffers

- ▶ *Page (Frame) Buffering is an OS optimization.*
- ▶ When a new page gets allocated and an "dirty page" gets replaced:
 - Without page buffering:
 - Write out old frame to disk
 - Give old frame to new page
 - With page buffering:
 - Select a new frame from page-buffer
 - Give to new frame while writing old page out to disk

Proportional Allocation

► *Proportional Allocation Policy*

- Suppose we have m frames in memory and n processes
- Continuously monitor memory usage for all processes
 - Let s_i denote the current virtual memory size of process P_i
 - Allocate a_i frames for P_i where $a_i = m \frac{s_i}{\sum_{i=1}^n s_i}$

► Pros:

- Allows memory allocation to change over time

► Cons:

- Unfair to existing processes when adding a new process
- Existing processes all have to contribute some frames


Important Side Note: "Thrashing"

- ▶ Page faults are expensive
 - Suppose a process wasn't allocated enough frames
 - During execution: page-in, page-out, page-in, page-out.
 - Not doing anything effective!
- ▶ We say that a process "*thrashes*" when it is spending more time paging than executing effective instructions.



Why Does Thrashing Occur?

- ▶ OS wants to keep CPU fully utilized
 - If CPU utilization % is low, OS dispatches another process to run
 - New processes → More memory usage
 - When memory runs out, we have increased page-fault rate
 - CPU utilization decreases due to page replacements
 - Causing OS to **add even more processes...** CPU utilization % eventually goes to 0




Working Set Model (1969, Multics)

- ▶ Process execution observes temporal and spatial locality
 - This implies there is a "**working set**" of pages that a process commonly uses at any time instant.
- ▶ To prevent thrashing we need to:
 - Track the **working set** across all processes over time
 - Try to give a process as many frames as it needs
 - Suspend a process if the system's **page demand** would exceed number of available frames

The "Working Set" Model


- ▶ For each process, track its page reference history
 - *Working Set*: the set of the Δ most recently-used pages
 - Δ is the size of the "sliding window"
 - τ (tau) is the time interval between checks
 - Consider the following page references across time:



- ▶ What's the working set at time u ? At time $u + \tau$?

The "Working Set" Model (2)

- ▶ Example: Consider the following page references



- ▶ Working sets for process P_i

$$WS_i(u) = \{B, A, F, E, G\}$$

$$WS_i(u + \tau) = \{C, D\}$$

- ▶ *(What do they tell us?)*

Working Set Model (Cont.)

- ▶ If $WS_i(t)$ denotes the *working set* of process P_i at time t , then the OS's total demand for pages at time t is:

$$pageDemand(t) = \sum_{i=1}^n |WS_i(t)|$$

- ▶ Suppose our system supports m frames of memory, **thrashing** occurs when $pageDemand(t) > m$.

Working-Set Model (Cont.)

► *Working-Set Allocation Policy*

- Select a value for Δ , the sliding window size
- Monitor working set $WS_i(t)$ of all existing processes
- Allocate $|WS_i(t)|$ frames for process P_i
 - If $pageDemand(t) > m$ // *OS will thrash!*
 - Choose a process to suspend (swap it out to reduce degree of multiprogramming)
 - Redistribute its frames to other processes

Working Set: Sliding Window Size

- ▶ The sliding window size Δ is an important parameter.
- ▶ Too Large:
 - Captures too many references
 - May overestimate page demand
 - Causes processes to swap out when they do not need to → low CPU utilization
- ▶ Too Small:
 - Could underestimate page demand
 - Causes OS to aggressively dispatch processes to run → thrashing

Goals for This Lecture...

- ▶ Demand Paging
 - Motivation
 - Implementation
- ▶ Page Replacement Policies
- ▶ Memory Allocation
 - Equal
 - Proportional
 - Working Set Model
 - How Do malloc() and free() Work?

How Does C Implement malloc() and free()

- ▶ Recall C's dynamic (de)allocation we can't live without!

```
/**  
 * Attempts to allocate size bytes on the heap and returns pointer to  
 * first byte  
 */  
void* malloc(size_t size);
```

@param size: the number of bytes to allocate on the heap.

@return pointer to the first byte if successful, otherwise **NULL**

```
/**  
 * Frees the memory pointed to by ptr  
 */  
void free(void* ptr);
```

@param ptr: the pointer to the allocated memory.


A Review: Process' Address Space

- ▶ For each process, its PCB stores a pointer called the *program break* that points to top of the heap
 - How do we move it?

`malloc()` calls this system call: *Program Break (brk)*


```
#include <stdlib.h>

/**
 * A system call to ask the kernel for numbytes bytes to
 * be allocated to the calling process. Moves brk
 * pointer accordingly.
 */
void *sbrk(intptr_t numbytes);
```



Step 1: Process Starts, No Heap

Address Space



```

int main() {
    double *dptr = (double*) malloc(sizeof(double));
    //...
    free(dptr);


    int *iptr = (int*) malloc(10*sizeof(int));
    //...
    free(iptr);

    int *iptr2 = (int*) malloc(5*sizeof(int));

    char *cptr = (char*) malloc(6*sizeof(char));
}
  
```

Step 2: Allocate 8 bytes on heap

Address Space



```

int main() {
    double *dptr = (double*) malloc(sizeof(double));
    //...
    free(dptr);


    int *iptr = (int*) malloc(10*sizeof(int));
    //...
    free(iptr);

    int *iptr2 = (int*) malloc(5*sizeof(int));

    char *cptr = (char*) malloc(6*sizeof(char));
}
  
```


Step 3: Freeing the Space

Address Space



Step 4: Process Starts, No Heap

Address Space



```

int main() {
    double *dptr = (double*) malloc(sizeof(double));
    //...
    free(dptr);


    int *iptr = (int*) malloc(10*sizeof(int)); ←
    //...
    free(iptr);

    int *iptr2 = (int*) malloc(5*sizeof(int));

    char *cptr = (char*) malloc(6*sizeof(char));
}
  
```

Step 4: Process Starts, No Heap

Address Space



```

int main() {
    double *dptr = (double*) malloc(sizeof(double));
    //...
    free(dptr);


    int *iptr = (int*) malloc(10*sizeof(int));
    //...
    free(iptr); ←

    int *iptr2 = (int*) malloc(5*sizeof(int));

    char *cptr = (char*) malloc(6*sizeof(char));
}
  
```

Step 4: Process Starts, No Heap

Address Space



```


int main() {
    double *dptr = (double*) malloc(sizeof(double));
    //...
    free(dptr);

    int *iptr = (int*) malloc(10*sizeof(int));
    //...
    free(iptr);

    int *iptr2 = (int*) malloc(5*sizeof(int)); ←
    char *cptr = (char*) malloc(6*sizeof(char));
}
  
```

Step 4: Process Starts, No Heap

Address Space



```

int main() {
    double *dptr = (double*) malloc(sizeof(double));
    //...
    free(dptr);

    int *iptr = (int*) malloc(10*sizeof(int));
    //...
    free(iptr);

    int *iptr2 = (int*) malloc(5*sizeof(int));


    char *cptr = (char*) malloc(6*sizeof(char)); ←
}
  
```

Let's Focus on the Heap

- ▶ How is the *heap* segment managed?
 - Just a list of "chunks": each is the size `malloc()` requested
- ▶ A chunk has a header, called the *memory control block (mcb)*:

```
struct mem_control_blk {
    unsigned int avail; //is this chunk available?
    size_t size;        //size of this chunk
};
```

- ▶ Heap content:



Example Showing Reuse of Freed Memory

```
#include <stdio.h>
#include <stdlib.h>

int main() {
    double *a = (double*) malloc(sizeof(double));
    *a = 123.4;
    printf("a (before free) = %0.6f\n", *a);
    free(a);

    double *b = (double*) malloc(sizeof(double));
    *b = 5.678;
    printf("b = %0.6f\n", *b);
    printf("a (after free) = %0.6f\n", *a); // What's printed?

    return 0;
}
```

Example Showing Reuse of Freed Memory

```
#include <stdio.h>
#include <stdlib.h>

int main() {
    double *a = (double*) malloc(sizeof(double));
    *a = 123.4;
    printf("a (before free) = %0.6f\n", *a);
    free(a);

    double *b = (double*) malloc(sizeof(double));
    *b = 5.678;
    printf("b = %0.6f\n", *b);
    printf("a (after free) = %0.6f\n", *a); // What's printed?

    return 0;
}
```

Program Output

```
-----
a (before free) = 123.400000
b = 5.678000
a (after free) = 5.678000
```

Importance of setting "loose pointers" to NULL after freeing!

free()

- ▶ Things to note:
 - Accessing the MCB
 - No changes to heap contents

```
void free(void *ptr) {  
    // remember to move pointer in reverse to account for the MCB header  
    struct mem_control_blk *mcb = (struct mem_control_blk*)  
        (ptr - sizeof(struct mem_control_blk));  
    mcb->avail = 1;  
}
```

Naive malloc()

```

void* malloc(size_t size) {
    void* current_addr = start_addr; // get start addr of heap
    void* end_addr = sbrk(0);           // get end addr of heap (top of heap)

    while (current_addr < end_addr) {
        //get the MCB of current chunk
        struct mem_control_blk *mcb = (struct mem_control_blk*) (current_addr);
        if (mcb->avail && size <= mcb->size) {
            //found a freed chunk! Use it!
            mcb->avail = 0;
            return current_addr + sizeof(struct mem_control_blk);
        }

        //skip to next chunk. Need to account for header
        current_addr + mcb->size + sizeof(struct mem_control_blk);
    }

    //must mean there were no freed chunks that are available or large enough
    current_addr = sbrk(size); //request the chunk from kernel
    struct mem_control_blk *mcb = (struct mem_control_blk*) (current_addr);
    mcb->avail = 0;
    mcb->size = size;
    return current_addr + sizeof(struct mem_control_blk);
}


```

What Made It Naive?

- ▶ May have to traverse entire list of chunks

- Better: use a *free list!*

```
struct mem_control_blk {
    unsigned int avail; // is this chunk available?
    size_t size;        // size of this chunk
    void* next;         // <-- ADDED THIS: pointer to the next free chunk
};
```



Goals for This Lecture...

- ▶ Demand Paging
 - Motivation
 - Implementation
- ▶ Page Replacement Policies
- ▶ Memory Allocation
- ▶ Conclusion

In Conclusion...

- ▶ **Virtual Memory:** Allows processes to operate when their virtual address space is **LARGER** than physical address space!
- ▶ Demand paging treats RAM as a cache for pages
 - Program stored on disk to begin with
 - Swap pages in and out on as-needed basis
 - Page fault => Heavy penalty (disk access)

In Conclusion... (Cont.)

- ▶ Good replacement policies:
 - Minimize page faults
 - Are fast to make decisions on which page to toss out
 - There are many loads and stores in a program
 - That's why LRU is good, but impractical
 - Also why RAND is actually used in practice!